

Computational Études

A Spectral Approach

Dr. Denys Dutykh

*Mathematics Department
College of Computing and Mathematical Sciences
Khalifa University of Science and Technology
Abu Dhabi, UAE*



Contents

Preface	2
Acknowledgements	5
1 Introduction	8
1.1 The Spectral Promise	9
1.2 A Brief History	10
1.3 Limitations and Trade-offs	10
1.4 The Philosophy of “Études”	11
1.5 Collocation: Computing in Physical Space	11
1.6 A Modern Workflow	12
1.7 A non-exhaustive literature overview	12
1.8 Summary	13
2 Classical Second Order PDEs and Separation of Variables	15
2.1 Heat Equation with Periodic Boundary Conditions	16
2.2 Numerical Illustration	20
2.3 Wave Equation with Dirichlet Boundary Conditions	23
2.4 Numerical Illustration	27
2.5 Laplace Equation in a Periodic Strip	29
2.6 Numerical Illustration	33
2.7 A non-exhaustive literature overview	35
2.8 Summary	36
3 Mise en Bouche	39
3.1 The Method of Weighted Residuals	40
3.2 A First Collocation Example	41
3.3 Collocation versus Galerkin	44
3.4 Conclusions and Questions	49
3.5 A Broader Perspective	50
3.6 A non-exhaustive literature overview	52
3.7 Summary	53
4 The Geometry of Nodes	55
4.1 The Problem: Polynomial Interpolation	56
4.2 The Runge Phenomenon	58
4.3 Theoretical Explanation: Potential Theory	60
4.4 The Solution: Chebyshev Points	62
4.5 Lagrange Basis Functions and Lebesgue Constants	64

4.6	Barycentric Interpolation	69
4.7	Convergence Analysis	70
4.8	Computational Étude: Random Nodes	72
4.9	Computational Étude: Random Angles on the Circle	75
4.10	Practical Guidelines and Outlook	77
4.11	A non-exhaustive literature overview	78
4.12	Summary	79
5	Differentiation Matrices	82
5.1	From Interpolation to Differentiation	83
5.2	Finite Difference Matrices	84
5.3	The Periodic Spectral Differentiation Matrix	86
5.4	Fornberg's Recursive Algorithm	94
5.5	Computational Étude: The Rational Trigonometric Test	98
5.6	Higher-Order Derivatives	100
5.7	Computational Étude: The Quantum Harmonic Oscillator	104
5.8	Looking Ahead: The Non-Periodic Case	107
5.9	A non-exhaustive literature overview	107
5.10	Summary	109
6	Smoothness and Spectral Accuracy	112
6.1	The Two Steps of Accuracy	113
6.2	The Decay of Fourier Coefficients	115
6.3	The Aliasing Phenomenon	118
6.4	Accuracy of Spectral Differentiation	119
6.5	A non-exhaustive literature overview	122
6.6	Summary	123
7	Chebyshev Differentiation Matrices	125
7.1	The Non-Periodic Challenge	126
7.2	The Chebyshev Differentiation Matrix	128
7.3	Small- N Examples	130
7.4	Matrix Structure and Properties	131
7.5	Demonstration: The Witch of Agnesi	133
7.6	Spectral Convergence	136
7.7	A non-exhaustive literature overview	138
7.8	Summary	139

8	Boundary Value Problems	141
8.1	Second Derivatives and Matrix Squaring	142
8.2	Imposing Boundary Conditions	142
8.3	Linear BVP: The Poisson Equation	143
8.4	Variable Coefficient Problems	146
8.5	Mixed Boundary Conditions	148
8.6	Nonlinear BVP: The Bratu Equation	151
8.7	Eigenvalue Problems	155
8.8	Two-Dimensional Problems	158
8.9	The Helmholtz Equation	163
8.10	Computational Étude: The Quantum Harmonic Oscillator Revisited	165
8.11	A non-exhaustive literature overview	168
8.12	Summary	170
9	Physical and Fourier Space on Grids	172
9.1	Motivation: Two Views of the Same Function	173
9.2	The Continuous Fourier Transform on $\mathbb{R}$	175
9.3	Sampling on an Infinite Grid: The Semidiscrete Fourier Transform	176
9.4	Band-Limited Interpolation and the Sinc Kernel	178
9.5	Periodic Grids: The Discrete Fourier Transform and FFT	182
9.6	Aliasing and Spectra on Periodic Grids	184
9.7	Spectra and Smoothness	186
9.8	Periodic Band-Limited Interpolation	187
9.9	A non-exhaustive literature overview	189
9.10	Summary	192
10	Spectral PDE Solvers with Chebyshev and Fourier Grids	195
10.1	Chebyshev–Fourier Recap	196
10.2	Chebyshev Differentiation via FFT	197
10.3	Method of Lines and Time Stepping	200
10.4	One-Dimensional Wave Equation	202
10.5	Two-Dimensional Wave Equation	205
10.6	One-Dimensional Heat Equation	209
10.7	Two-Dimensional Heat Equation	213
10.8	Two-Dimensional Poisson Equation	216
10.9	Additional Physics Études	221

10.10 Efficiency: Matrices versus FFT	223
10.11 A non-exhaustive literature overview	224
10.12 Summary	225
Bibliography	228



Preface

The purpose of this book is twofold: to explain to students why spectral methods work and why they are so remarkably efficient, and to teach them how to implement these methods using modern computational tools.

This text is conceived as a collection of **Computational Études**. In musical education, an étude is a composition designed to perfect a specific technique (be it rapid scales, complex arpeggios, or delicate phrasing) while remaining a pleasing piece of music in its own right. Similarly, each chapter here presents a focused mathematical concept paired with its computational realization: a study that is both instructive and complete.

Who Is This Book For?

We have written primarily for graduate students in applied mathematics, physics, and engineering who seek a hands-on understanding of spectral methods. The reader should be comfortable with calculus, linear algebra, and basic programming. Prior exposure to numerical methods is helpful but not essential; we build the necessary foundations as we proceed.

How Is the Book Organized?

Each chapter is designed to be largely self-contained. We begin with fundamental concepts (interpolation and differentiation) before advancing to time-stepping schemes and applications. The mathematical exposition is deliberately concise, favoring clarity over exhaustive rigor. Proofs are included when they illuminate; otherwise, we direct the reader to authoritative references.

Reproducible Science

This book is also an experiment in **reproducible science**. Every figure, every table, every numerical result you see in these pages was generated by code available in the accompanying repository. We provide implementations in both Python and MATLAB, allowing readers to choose their preferred environment. The Python code emphasizes accessibility and integration with the open-source scientific ecosystem; the MATLAB code leverages its historical significance in numerical computing and, where appropriate, the Advanpix Multiprecision Computing Toolbox for extended precision arithmetic.

How to Use This Book

We invite you to treat this book not as a static reference, but as a workshop. Clone the repository, run the scripts, modify the parameters, break the code, and fix it. That is the only way to truly master the art of spectral methods.

The complete source code and the Typst manuscript are available at:

<https://github.com/dutykh/computational-etudes/>

Dr. Denys Dutykh

Abu Dhabi, UAE
February 2026



Acknowledgements

I would like to express my sincere gratitude to the students who took this course, carefully reread the manuscript, and helped me present it in a clean and polished form. Their diligent attention to detail and thoughtful feedback have been invaluable.

In particular, I thank:

- **Farrell Adriano**
- **Aziz Sharofov**
- **Mark Essa Sukaiti**

Their contributions have significantly improved the quality of this book.



CHAPTER 1

Introduction

Differential equations serve as the fundamental language of the physical sciences, describing phenomena ranging from the propagation of sound waves to the flow of heat and the dynamics of fluids. Finding exact analytical solutions to these equations is a luxury rarely afforded in practical applications. Consequently, the scientist and the engineer must turn to numerical approximation.

Broadly speaking, numerical methods for differential equations fall into two categories: local methods and global methods. The former, including Finite Difference and Finite Element Methods, approximate the unknown solution using functions that are non-zero only on small sub-domains (elements or grid stencils). These methods are robust and flexible, handling complex geometries with grace. However, their accuracy is typically algebraic; refining the grid by a factor of two might improve the error by a factor of four or eight, but rarely more. From a computational perspective, local methods are *myopic*: to compute a derivative at a grid point, they look only at immediate neighbors.

Spectral methods represent the global approach. They approximate the solution as a linear combination of continuous, global basis functions, typically trigonometric polynomials (Fourier series) for periodic problems or Chebyshev polynomials for non-periodic ones. In stark contrast to local schemes, spectral methods are *holistic*: the derivative at any single point depends on the function values at *every other point* in the domain. Mathematically, this is equivalent to fitting a single high-degree polynomial through all data points. This global coupling is what allows information to propagate instantly across the grid, granting us the remarkable convergence that we call “spectral accuracy.”

There is, however, a unifying viewpoint that bridges these apparently distinct philosophies. Pseudospectral methods can be understood as the natural limit of finite difference methods as the order of accuracy increases without bound. Imagine a sequence of finite difference stencils: first using two neighbors, then four, then eight, and so on. As the stencil width grows to encompass the entire grid, the local method transforms smoothly into a global one. This perspective, developed by Kreiss and Oliger and elaborated by Fornberg, offers both intuition and practical computational strategies. The theory and practice of spectral methods are comprehensively developed in the classical texts by [47], [120], and [21].

1.1 The Spectral Promise

The fundamental argument for spectral methods is one of efficiency. If the solution to a problem is smooth, the coefficients of its expansion in a proper global basis decay exponentially fast. This phenomenon is known as spectral accuracy.

In practical terms, this means that spectral methods can achieve a level of precision with a few dozen degrees of freedom that a finite difference scheme might require thousands of grid points to match. While a fourth-order finite difference method implies that the error $\varepsilon \sim O(N^{-4})$, a spectral method boasts $\varepsilon \sim O(c^{-N})$ for some constant $c > 1$ [73]. When the solution is analytic, the convergence is explosive; the error drops into the “spectral valley” until it hits the floor of machine precision.

Beyond raw accuracy, spectral methods possess another virtue that is often decisive in physical applications: they are virtually free of both dissipative and dispersive numerical errors. Finite difference schemes, by their very nature, introduce artificial dissipation that can overwhelm the true physical dissipation in problems such as high-Reynolds number fluid flows. They also suffer from dispersive errors that cause different frequency components to propagate at slightly different speeds, turning sharp gradients into spurious wavetrains. Spectral methods avoid both pathologies. This fidelity to the underlying physics explains their dominance in demanding applications: turbulence modeling, global weather and climate prediction, nonlinear wave dynamics, and seismic analysis all rely heavily on spectral techniques.

This global dependence has a computational consequence: spectral differentiation matrices are *dense*, not sparse. Where a finite difference scheme produces banded matrices that are cheap to store and invert, spectral methods fill in every entry. However, the extraordinary accuracy means we need so few points (often just dozens where finite differences would require thousands) that we can afford this density. The cost per degree of freedom is higher, but the total cost for a given accuracy is dramatically lower.

However, this power is not without its price. Spectral methods are unforgiving regarding grid placement. We cannot simply choose points where we please; for non-periodic problems, the mathematics dictates that points must cluster at boundaries (the celebrated Chebyshev points) to prevent the interpolation from diverging. Attempting high-degree polynomial interpolation on an equispaced grid leads to the notorious Runge phenomenon [103], where oscillations grow without bound near the boundaries. This sensitivity to geometry is what restricts spectral methods primarily to simple domains, but within those domains, they reign supreme.

This book aims to demystify this “spectral magic.” We will see that it is not magic at all, but a direct consequence of the smoothness of the underlying functions and the careful choice of basis and grid.

1.2 A Brief History

Spectral representations have been used for analytic studies of differential equations since the days of Fourier in the early nineteenth century. The idea of employing them for *numerical* computation, however, emerged much later. Lanczos, in the 1930s, pioneered the use of Chebyshev expansions for solving ordinary differential equations numerically [74]. This approach remained somewhat academic until the early 1970s, when Kreiss and Olinger introduced the pseudospectral method for partial differential equations [73]. Their innovation was to work with function values at grid points rather than expansion coefficients, dramatically simplifying the treatment of nonlinear terms. The practical viability of these methods was ensured by the fast Fourier transform algorithm, rediscovered by Cooley and Tukey in 1965 [30], which reduced the cost of transforming between physical and spectral space from $O(N^2)$ to $O(N \log N)$ operations. This confluence of theoretical insight and algorithmic efficiency launched spectral methods into the mainstream of computational science.

1.3 Limitations and Trade-offs

Honesty compels us to acknowledge that spectral methods are not a panacea. Several factors can limit their applicability or efficiency:

- **Boundary conditions** can be awkward to impose, particularly for problems with complex constraints or time-dependent boundaries.
- **Irregular domains** resist the tensor-product structure that makes spectral methods efficient. While domain decomposition and mapping techniques exist, they sacrifice some of the method’s elegance.
- **Strong shocks and discontinuities** violate the smoothness assumptions that underlie spectral accuracy. The Gibbs phenomenon produces persistent oscillations near discontinuities, requiring filtering or other remediation.

- **Variable resolution requirements** across a large domain are difficult to accommodate. Unlike adaptive mesh refinement in finite element methods, spectral grids are inherently uniform in their polynomial degree.

These limitations explain why finite element and finite difference methods continue to thrive in many applications. The art lies in recognizing when spectral methods are the right tool. When the geometry is simple, the solution is smooth, and high accuracy is paramount, no other approach comes close.

1.4 The Philosophy of “Études”

The title of this volume, Computational Études, reflects a specific pedagogical philosophy. In musical education, an étude is a composition designed to practice a particular technical skill (be it rapid scales or complex arpeggios) while remaining a pleasing piece of music in its own right.

In this text, our “technical skills” are not rapid scales or arpeggios, but rather handling stiffness in time-stepping, managing aliasing in nonlinear products, enforcing boundary conditions through tau methods or lifting functions, and filtering spurious oscillations. Just as a Chopin Étude transforms a technical exercise into art, a well-written spectral code transforms a mathematical formula into a robust simulation. The études collected here are designed to cultivate this virtuosity.

We approach spectral methods not through dry, abstract theorems, but through concrete, self-contained studies. Each chapter focuses on a specific mathematical concept (interpolation, differentiation, aliasing, or time-stepping) and explores it through a compact, runnable implementation.

We deliberately restrict our focus primarily to one-dimensional problems. This choice is strategic. The mathematical essence of spectral methods (the treatment of boundaries, the distribution of collocation points, and the structure of differentiation matrices) is fully present in one dimension. Extending these ideas to two or three dimensions usually involves tensor products, which add significant programming overhead without necessarily adding new conceptual depth. By staying in 1D, we keep our code short, readable, and focused on the physics and mathematics.

1.5 Collocation: Computing in Physical Space

While the theory of spectral methods relies on orthogonal expansions (Fourier series for periodic problems, Chebyshev series otherwise), the actual computation often proceeds differently. Rather than manipulating expansion coefficients directly (the *modal* or *Galerkin* approach), we typically work with function values at carefully chosen grid points (the *nodal* or *collocation* approach, also called *pseudospectral*) [94].

The collocation philosophy dominates this book for a practical reason: it handles nonlinear terms with ease. When the governing equation contains products like $u \cdot u_x$, the modal approach requires computing convolutions of coefficient sequences, a tedious operation. Collocation simply evaluates the product pointwise on the grid. This directness makes pseudospectral methods the tool of choice for most computational applications, and it is the approach we shall master through these études.

The reader should be aware that both viewpoints (modal and nodal) illuminate the same underlying mathematics. The Fast Fourier Transform provides the bridge, allowing us to move efficiently between coefficient space and physical space as needed.

1.6 A Modern Workflow

Finally, this book is an experiment in reproducible science [33]. The days of presenting numerical results as static, unverifiable images are passing. The results you see in these pages were generated by the code available in the accompanying repository. We utilize a dual-language approach:

- **Python:** For accessibility and integration with the vast open-source scientific ecosystem.
- **Matlab:** For its historical significance in this field and its concise matrix syntax, often utilizing the Advanpix Multiprecision Computing Toolbox [2] to explore phenomena that lie beyond standard double precision.

We invite you to treat this book not as a static reference, but as a workshop. Run the scripts, change the parameters, break the code, and fix it. That is the only way to truly learn the art of spectral methods.

1.7 A non-exhaustive literature overview

The ascent of spectral methods from analytical curiosity to the gold standard of high-accuracy computing is a narrative defined by the convergence of approximation theory and algorithmic innovation. The theoretical bedrock of the field lies in the 19th-century analysis of Weierstrass [130], who established the density of polynomials in the space of continuous functions, thereby guaranteeing the existence of polynomial approximants. However, the translation of this existence into construction was fraught with instability, most famously demonstrated by Runge [103]. His discovery that high-degree polynomial interpolation on equispaced grids diverges near boundaries — the “Runge phenomenon” — necessitated the adoption of non-uniform node distributions, cementing the importance of Chebyshev and Legendre points in non-periodic spectral methods.

The transition to numerical practicality began in the 1930s with Lanczos [74], who pioneered the use of Chebyshev polynomials for solving differential equations and introduced the “Tau method” to satisfy boundary conditions. Lanczos’s insight that Chebyshev expansions minimize maximum error provided the “spectral” alternative to the local Taylor-series approximations of finite difference methods. Yet, the computational cost of these global methods remained prohibitive until the rediscovery of the Fast Fourier Transform (FFT) by Cooley and Tukey [30]. By reducing the complexity of transformations from $O(N^2)$ to $O(N \log N)$, the FFT rendered global approximation techniques competitive for large-scale problems.

In the 1970s, the field crystallized through the foundational work of Kreiss and Oliger [73] and Orszag [94]. Kreiss and Oliger formalized the “pseudospectral” approach, demonstrating its superior phase accuracy for wave propagation compared to finite difference schemes. Simultaneously, Orszag applied these techniques to fluid dynamics, solving the Orr–Sommerfeld stability equation with unprecedented precision and establishing spectral methods as essential tools for turbulence simulation. The conceptual unification of these approaches was later articulated by Fornberg [45], who demonstrated that pseudospectral methods can be viewed as the limit of finite difference schemes as the stencil width extends to the full domain.

Modern developments continue to push the boundaries of the field. The **Chebfun** project, led by Trefethen [120], has automated the application of spectral methods, treating functions as continuous objects discretized adaptively to machine precision. Furthermore, the increasing complexity of computational workflows has elevated the importance of reproducible research. The philosophies of Donoho [33] and LeVeque [77] underscore that code and data are integral to the scientific record, a principle that guides the “computational étude” approach of this text. Recent advancements in sparse spectral

methods and high-dimensional approximation suggest that the “spectral promise” continues to evolve, offering new solutions for the curse of dimensionality in complex physical systems.

1.8 Summary

This introductory chapter has set the stage for spectral methods by contrasting them with local approaches:

1. **Spectral accuracy:** When the solution is smooth, global polynomial approximation yields exponential convergence — far surpassing the algebraic rates of finite difference and finite element methods.
2. **Global coupling:** Spectral methods compute derivatives using information from *every* grid point, making them holistic rather than myopic. This global dependence produces dense matrices, but the small number of degrees of freedom required keeps the total cost low.
3. **Limitations:** Spectral methods are most effective on simple geometries with smooth solutions. Discontinuities trigger the Gibbs phenomenon, and irregular domains resist the tensor-product structure that underpins efficient spectral computation.
4. **Pseudospectral philosophy:** Working with function values at grid points (collocation) rather than expansion coefficients handles nonlinear terms naturally and is the approach developed throughout this book.
5. **Reproducibility:** All numerical results are generated by code provided in the accompanying repository, in Python, MATLAB, and Julia, embodying the principle that code and data are integral to the scientific record.

The chapters that follow translate these ideas into algorithms: separation of variables, interpolation theory, differentiation matrices, and ultimately the spectral solution of partial differential equations.



CHAPTER 2

Classical Second Order PDEs and Separation of Variables

In this opening chapter we derive exact solutions for three classical linear partial differential equations: the *heat equation* (parabolic), the *wave equation* (hyperbolic), and the *Laplace equation* (elliptic). These solutions are found by the *method of separation of variables*, which expresses the solution as an infinite series of eigenfunctions.

Why begin a book on *numerical* methods with *analytical* solutions? Because separation of variables is the theoretical ancestor of spectral methods [55, 94, 120]. When we later truncate these infinite series at some finite N and compute with only the first N modes, we are doing exactly what a spectral solver does, but with pen and paper first. This chapter thus serves as the conceptual bridge between classical analysis and modern computation. The methods presented here are classical and thoroughly developed in standard references such as [119], [21], and [27].

We treat three model problems:

- heat equation with periodic boundary conditions in one spatial dimension,
- wave equation on a bounded interval,
- Laplace equation in a simple domain.

We begin with a complete analytic solution of the heat equation. The other two examples will follow the same pattern.

2.1 Heat Equation with Periodic Boundary Conditions

We consider the one dimensional heat equation on the interval $[0, 2\pi]$ with periodic boundary conditions. The unknown $u(x, t)$ represents, for example, the temperature at point x and time t .

The problem is

$$\frac{\partial u}{\partial t}(x, t) = \frac{\partial^2 u}{\partial x^2}(x, t), \quad x \in [0, 2\pi], \quad t > 0,$$

with periodic boundary conditions

$$u(x + 2\pi, t) = u(x, t)$$

for all real x and all $t > 0$, and initial condition

$$u(x, 0) = f(x), \quad x \in [0, 2\pi].$$

We assume that f is smooth and 2π periodic:

$$f(x + 2\pi) = f(x).$$

Our goal is to obtain an explicit representation of $u(x, t)$ as an infinite series and to see how separation of variables leads naturally to a Fourier series in space. The technique we employ here follows the classical approach presented in [119].

2.1.1 Step 1: Separation Ansatz

We look for nontrivial solutions of the form

$$u(x, t) = X(x) \cdot T(t),$$

where X depends only on x and T depends only on t .

Substituting into the PDE gives

$$X(x) \cdot T'(t) = X''(x) \cdot T(t).$$

We assume X and T are not identically zero, so we can divide both sides by $X(x) \cdot T(t)$:

$$\frac{T'(t)}{T(t)} = \frac{X''(x)}{X(x)}.$$

The left side depends only on t , the right side only on x . Therefore both sides must be equal to the same constant, which we denote by $-\lambda$:

$$\frac{T'(t)}{T(t)} = \frac{X''(x)}{X(x)} = -\lambda.$$

We obtain two ordinary differential equations:

$$T'(t) + \lambda T(t) = 0,$$

$$X''(x) + \lambda X(x) = 0.$$

The periodic boundary conditions for u imply periodic conditions for X :

$$X(0) = X(2\pi), \quad X'(0) = X'(2\pi).$$

We have arrived at a spatial eigenvalue problem for X . The systematic treatment of such eigenvalue problems is a cornerstone of the theory of partial differential equations, with roots in Fourier's original treatise [48]; see [119] for a comprehensive exposition.

2.1.2 Step 2: Spatial Eigenvalue Problem with Periodic Boundary Conditions

We now solve

$$X''(x) + \lambda X(x) = 0,$$

with

$$X(0) = X(2\pi), \quad X'(0) = X'(2\pi).$$

We consider three cases: $\lambda < 0$, $\lambda = 0$, and $\lambda > 0$.

Case 1: $\lambda < 0$

Write $\lambda = -\mu^2$ with $\mu > 0$. The equation becomes

$$X''(x) - \mu^2 X(x) = 0.$$

The general solution is

$$X(x) = Ae^{\mu x} + Be^{-\mu x}.$$

Imposing periodicity $X(0) = X(2\pi)$ gives

$$A + B = Ae^{2\mu\pi} + Be^{-2\mu\pi}.$$

Imposing $X'(0) = X'(2\pi)$ gives

$$\mu(A - B) = \mu(Ae^{2\mu\pi} - Be^{-2\mu\pi}).$$

Since $\mu > 0$, we can divide by μ and rewrite both conditions as a homogeneous linear system:

$$A(1 - e^{2\mu\pi}) + B(1 - e^{-2\mu\pi}) = 0,$$

$$A(1 - e^{2\mu\pi}) - B(1 - e^{-2\mu\pi}) = 0.$$

Adding these two equations gives

$$2A(1 - e^{2\mu\pi}) = 0.$$

Since $\mu > 0$, we have $e^{2\mu\pi} > 1$, so $1 - e^{2\mu\pi} \neq 0$. Therefore $A = 0$.

Subtracting the second equation from the first gives

$$2B(1 - e^{-2\mu\pi}) = 0.$$

Since $\mu > 0$, we have $e^{-2\mu\pi} < 1$, so $1 - e^{-2\mu\pi} \neq 0$. Therefore $B = 0$.

Hence the only solution is $A = B = 0$, the trivial solution. There are no nontrivial periodic eigenfunctions for $\lambda < 0$, and we discard this case.

Case 2: $\lambda = 0$

The equation reduces to

$$X''(x) = 0.$$

Its general solution is

$$X(x) = A + Bx.$$

Periodicity $X(0) = X(2\pi)$ gives

$$A = A + 2\pi B$$

so $B = 0$. Then $X(x) = A$ is constant.

The derivative is $X'(x) = 0$, so $X'(0) = X'(2\pi)$ is automatically satisfied.

Thus $\lambda = 0$ gives one eigenfunction

$$X_0(x) = 1$$

(up to a multiplicative constant).

Case 3: $\lambda > 0$

Write $\lambda = k^2$ with $k > 0$. The equation becomes

$$X''(x) + k^2 X(x) = 0.$$

The general solution is

$$X(x) = A \cos(kx) + B \sin(kx).$$

We now impose periodicity. First,

$$X(0) = A, \quad X(2\pi) = A \cos(2\pi k) + B \sin(2\pi k).$$

The condition $X(0) = X(2\pi)$ gives

$$A = A \cos(2\pi k) + B \sin(2\pi k).$$

Next,

$$X'(x) = -Ak \sin(kx) + Bk \cos(kx),$$

so

$$X'(0) = Bk, \quad X'(2\pi) = -Ak \sin(2\pi k) + Bk \cos(2\pi k).$$

The condition $X'(0) = X'(2\pi)$ gives

$$Bk = -Ak \sin(2\pi k) + Bk \cos(2\pi k).$$

We can divide by k the last identity (for $k \neq 0$) and write the system as

$$A(1 - \cos(2\pi k)) - B \sin(2\pi k) = 0,$$

$$A \sin(2\pi k) + B(1 - \cos(2\pi k)) = 0.$$

For a nontrivial pair (A, B) the determinant of the system must vanish:

$$(1 - \cos(2\pi k))^2 + (\sin(2\pi k))^2 = 0.$$

The left side is a sum of squares, so it is zero if and only if

$$1 - \cos(2\pi k) = 0, \quad \sin(2\pi k) = 0.$$

Hence

$$\cos(2\pi k) = 1, \quad \sin(2\pi k) = 0.$$

This happens exactly when k is an integer:

$$k = n, \quad n \in \mathbb{Z}.$$

The case $n = 0$ corresponds to $\lambda = 0$, which we have already treated. For $n \geq 1$ we obtain eigenvalues

$$\lambda_n = n^2, \quad n = 1, 2, 3, \dots$$

For each $n \geq 1$ the corresponding eigenfunctions can be chosen as

$$X_n^{(c)}(x) = \cos(nx), \quad X_n^{(s)}(x) = \sin(nx).$$

These functions are 2π periodic, and their derivatives are also 2π periodic, so the boundary conditions are satisfied.

We have therefore found a complete set of spatial eigenfunctions for the heat equation with periodic boundary conditions:

- a constant mode $X_0(x) = 1$ (eigenvalue $\lambda_0 = 0$),
- cosine modes $X_n^{(c)}(x) = \cos(nx)$,
- sine modes $X_n^{(s)}(x) = \sin(nx)$,

with eigenvalues $\lambda_n = n^2$ for $n \geq 1$.

2.1.3 Step 3: Time Dependent Factors

For each eigenvalue λ the corresponding time factor satisfies

$$T'(t) + \lambda T(t) = 0.$$

The solution is

$$T(t) = Ce^{-\lambda t}.$$

For the constant mode $\lambda_0 = 0$ we obtain

$$T_0(t) = C_0$$

(constant in time).

For $n \geq 1$ we have

$$T_n(t) = C_n e^{-n^2 t}.$$

Combining space and time, we form products $u(x, t) = X(x) \cdot T(t)$ to obtain separated solutions. For each mode, the arbitrary constant from the time factor can be absorbed into a single new constant. Specifically, we write

$$\begin{aligned} u_0(x, t) &= A_0, \\ u_n^{(c)}(x, t) &= A_n \cos(nx) \cdot e^{-n^2 t}, \\ u_n^{(s)}(x, t) &= B_n \sin(nx) \cdot e^{-n^2 t}, \end{aligned}$$

where we have renamed the constants: $A_0 = C_0$, and for $n \geq 1$, the constant A_n absorbs the coefficient from the cosine mode while B_n absorbs the coefficient from the sine mode. Note that each separated solution carries its own independent arbitrary constant.

Since the heat equation is linear, any linear combination of these separated solutions is again a solution. Therefore the general solution that satisfies the periodic boundary conditions can be written as an infinite series

$$u(x, t) = A_0 + \sum_{n=1}^{\infty} (A_n \cos(nx) + B_n \sin(nx)) e^{-n^2 t}.$$

The coefficients A_0, A_n, B_n remain to be determined from the initial condition.

2.1.4 Step 4: Imposing the Initial Condition and Fourier Series

We impose the initial condition

$$u(x, 0) = f(x).$$

Setting $t = 0$ in the general solution we obtain

$$u(x, 0) = A_0 + \sum_{n=1}^{\infty} (A_n \cos(nx) + B_n \sin(nx)) = f(x).$$

This is exactly the Fourier series expansion of the 2π periodic function f . Under mild regularity assumptions, f has a Fourier series

$$f(x) = a_0 + \sum_{n=1}^{\infty} (a_n \cos(nx) + b_n \sin(nx))$$

with Fourier coefficients

$$\begin{aligned} a_0 &= \frac{1}{2\pi} \int_0^{2\pi} f(x) dx, \\ a_n &= \frac{1}{\pi} \int_0^{2\pi} f(x) \cos(nx) dx, \quad n \geq 1, \\ b_n &= \frac{1}{\pi} \int_0^{2\pi} f(x) \sin(nx) dx, \quad n \geq 1. \end{aligned}$$

By uniqueness of the Fourier expansion, we must have

$$A_0 = a_0, \quad A_n = a_n, \quad B_n = b_n.$$

Thus the coefficients in the heat equation solution are exactly the Fourier coefficients of the initial data.

2.1.5 Step 5: Final Explicit Solution and Interpretation

Substituting these coefficients into the general solution we obtain the explicit formula

$$u(x, t) = a_0 + \sum_{n=1}^{\infty} (a_n \cos(nx) + b_n \sin(nx)) e^{-n^2 t}, \quad t > 0.$$

This infinite sum solves the heat equation with periodic boundary conditions and initial data f . Each Fourier mode decays exponentially in time at a rate proportional to its eigenvalue n^2 . High frequency modes (large n) decay faster, which expresses the smoothing effect of the heat equation.

The constant term a_0 does not decay. It represents the average value of f on $[0, 2\pi]$, which is preserved by the heat flow.

From a spectral viewpoint, the functions

$$1, \cos(x), \sin(x), \cos(2x), \sin(2x), \dots$$

form an eigenbasis of the spatial operator

$$Lu = \frac{\partial^2 u}{\partial x^2}$$

with periodic boundary conditions. The evolution of each eigenmode is independent and given simply by multiplication by $e^{-n^2 t}$ in time.

Later, in the numerical part of this book, we will approximate $u(x, t)$ by truncating the sum to finitely many modes:

$$u_{N(x,t)} = a_0 + \sum_{n=1}^N (a_n \cos(nx) + b_n \sin(nx)) e^{-n^2 t}.$$

This truncation is the essence of a Fourier spectral method; recent extensions treat both space and time spectrally [71]. The analytic solution derived here is the infinite dimensional limit of that numerical procedure.

2.2 Numerical Illustration

To visualize the smoothing effect of heat diffusion, we compute the truncated Fourier series solution for a triangle wave initial condition:

$$f(x) = \pi - |x - \pi|, \quad x \in [0, 2\pi].$$

This function is continuous but has a corner (non-differentiable point) at $x = \pi$. Its Fourier series contains only cosine terms with coefficients decaying as $1/n^2$:

$$f(x) = \frac{\pi}{2} + \frac{4}{\pi} \sum_{k=1}^{\infty} \frac{\cos((2k-1)x)}{(2k-1)^2}.$$

The key portion of the implementation evaluates the truncated Fourier series at any point in space and time. In Python:

```
1 def heat_solution(x, t, a0, a_n, b_n):
2     u = np.full_like(x, a0, dtype=float)
3     n_modes = len(a_n) - 1
4     for n in range(1, n_modes + 1):
5         decay = np.exp(-n**2 * t)
6         u += (a_n[n] * np.cos(n * x) + b_n[n] * np.sin(n * x)) * decay
```

 Python

```
7 return u
```

The equivalent MATLAB implementation:

```
1 u = a0 * ones(size(x));
2 for n = 1:N_MODES
3     decay = exp(-n^2 * t);
4     u = u + (a_n(n+1) * cos(n*x) + b_n(n+1) * sin(n*x)) * decay;
5 end
```

Matlab

The Julia implementation:

```
1 function heat_solution(x, t, a0, a_n, b_n)
2     u = fill(a0, length(x))
3     n_modes = length(a_n) - 1
4     for n in 1:n_modes
5         decay = exp(-n^2 * t)
6         @. u += (a_n[n+1] * cos(n * x) + b_n[n+1] * sin(n * x)) * decay
7     end
8     return u
9 end
```

Julia

Figure 1 shows the evolution of $u_N(x, t)$ with $N = 50$ modes at several time values. At $t = 0$ the triangle wave is faithfully reproduced. As time increases, the higher frequency modes decay exponentially faster than the lower ones (the n -th mode decays as $e^{-n^2 t}$), and the solution rapidly smooths toward the constant equilibrium $u = \pi/2$.




Figure 1: Evolution of the heat equation solution with a triangle wave initial condition. The higher frequency modes decay rapidly, smoothing the initial corner at $x = \pi$.

The code generating Figure 1 is available in:

- codes/python/ch02/heat_equation_evolution.py

- `codes/matlab/ch02/heat_equation_evolution.m`
- `codes/julia/ch02/heat_equation_evolution.jl`

A complementary view of the solution is provided by the waterfall plot in Figure 2, which displays the entire space-time evolution as a three-dimensional surface. The smoothing effect of the heat equation is clearly visible: the initial sharp triangle wave rapidly flattens as time progresses, with the solution approaching the constant equilibrium state $u = \pi/2$.




Figure 2: Waterfall plot showing the complete space-time evolution of the heat equation solution. The initial triangle wave smooths rapidly as higher frequency modes decay exponentially.

The code generating Figure 2 is available in:

- `codes/python/ch02/heat_equation_waterfall.py`
- `codes/matlab/ch02/heat_equation_waterfall.m`
- `codes/julia/ch02/heat_equation_waterfall.jl`

2.3 Wave Equation with Dirichlet Boundary Conditions

We now consider the one dimensional wave equation on a finite interval with homogeneous Dirichlet boundary conditions. This is the classical model for a vibrating string of length L with both ends fixed, first studied by d'Alembert [102] and Euler [39].

Let $u(x, t)$ denote the vertical displacement of the string at position $x \in [0, L]$ and time $t > 0$. The equation of motion is

$$\frac{\partial^2 u}{\partial t^2}(x, t) = c^2 \frac{\partial^2 u}{\partial x^2}(x, t), \quad 0 < x < L, \quad t > 0,$$

where $c > 0$ is the wave speed.

The boundary conditions express that the endpoints of the string are clamped:

$$u(0, t) = 0, \quad u(L, t) = 0, \quad t > 0.$$

We prescribe the initial displacement and initial velocity:

$$u(x, 0) = f(x), \quad \frac{\partial u}{\partial t}(x, 0) = g(x), \quad 0 < x < L,$$

with suitable functions f and g that vanish at $x = 0$ and $x = L$.

As in the heat equation example, we use separation of variables and obtain a representation of the solution as an infinite series in spatial eigenfunctions. This time the temporal factors are oscillatory instead of decaying. The mathematical theory of the vibrating string is a classical topic covered extensively in [119].

2.3.1 Step 1: Separation Ansatz

We look for nontrivial solutions of the form

$$u(x, t) = X(x) \cdot T(t).$$

Substituting into the wave equation gives

$$X(x) \cdot T''(t) = c^2 X''(x) \cdot T(t).$$

Assuming X and T are not identically zero, we divide both sides by $c^2 X(x) \cdot T(t)$:

$$\frac{T''(t)}{c^2 T(t)} = \frac{X''(x)}{X(x)}.$$

The left side depends only on t , the right side only on x . Therefore both sides must be equal to a constant, which we denote by $-\lambda$:

$$\frac{T''(t)}{c^2 T(t)} = \frac{X''(x)}{X(x)} = -\lambda.$$

We obtain the pair of ordinary differential equations

$$\begin{aligned} T''(t) + c^2 \lambda T(t) &= 0, \\ X''(x) + \lambda X(x) &= 0, \end{aligned}$$

with boundary conditions

$$X(0) = 0, \quad X(L) = 0.$$

As in the heat equation case, we have a spatial eigenvalue problem for X .

2.3.2 Step 2: Spatial Eigenvalue Problem with Dirichlet Boundary Conditions

We must solve

$$X''(x) + \lambda X(x) = 0, \quad 0 < x < L,$$

with

$$X(0) = 0, \quad X(L) = 0.$$

We again consider three cases: $\lambda < 0$, $\lambda = 0$, and $\lambda > 0$.

Case 1: $\lambda < 0$

Write $\lambda = -\mu^2$ with $\mu > 0$. The equation becomes

$$X''(x) - \mu^2 X(x) = 0.$$

The general solution is

$$X(x) = Ae^{\mu x} + Be^{-\mu x}.$$

The boundary condition at $x = 0$ gives

$$X(0) = A + B = 0 \Rightarrow B = -A.$$

Then

$$X(L) = Ae^{\mu L} - Ae^{-\mu L} = A(e^{\mu L} - e^{-\mu L}).$$

The condition $X(L) = 0$ implies

$$A(e^{\mu L} - e^{-\mu L}) = 0.$$

Since $e^{\mu L} \neq e^{-\mu L}$ for $\mu > 0$, we must have $A = 0$. Then $B = 0$ and the solution is trivial. Therefore there are no nontrivial eigenfunctions for $\lambda < 0$.

Case 2: $\lambda = 0$

The equation reduces to

$$X''(x) = 0,$$

whose general solution is

$$X(x) = A + Bx.$$

The boundary conditions give

$$\begin{aligned} X(0) &= A = 0, \\ X(L) &= A + BL = BL = 0. \end{aligned}$$

Hence $B = 0$ and X is again trivial. There is no nontrivial eigenfunction for $\lambda = 0$.

Case 3: $\lambda > 0$

Write $\lambda = k^2$ with $k > 0$. The equation becomes

$$X''(x) + k^2 X(x) = 0.$$

The general solution is

$$X(x) = A \cos(kx) + B \sin(kx).$$

The boundary condition at $x = 0$ gives

$$X(0) = A = 0.$$

So $X(x) = B \sin(kx)$. The boundary condition at $x = L$ gives

$$X(L) = B \sin(kL) = 0.$$

For a nontrivial solution we need $B \neq 0$, so we must have

$$\sin(kL) = 0.$$

Therefore

$$kL = n\pi, \quad n = 1, 2, 3, \dots$$

The corresponding values of k are

$$k_n = \frac{n\pi}{L}, \quad n = 1, 2, 3, \dots$$

We conclude that the eigenvalues and eigenfunctions are

$$\lambda_n = k_n^2 = \left(\frac{n\pi}{L}\right)^2,$$

$$X_n(x) = \sin\left(\frac{n\pi x}{L}\right), \quad n = 1, 2, 3, \dots$$

Each X_n vanishes at $x = 0$ and $x = L$, as required by the Dirichlet boundary conditions. The family $\{X_n\}_{n \geq 1}$ is orthogonal in $L^2(0, L)$:

$$\int_0^L \sin\left(\frac{n\pi x}{L}\right) \sin\left(\frac{m\pi x}{L}\right) dx = \begin{cases} 0 & \text{if } n \neq m \\ L/2 & \text{if } n = m. \end{cases}$$

These eigenfunctions will form the spatial basis in our series solution.

2.3.3 Step 3: Time Dependent Factors

For each eigenvalue λ_n the time factor T_n satisfies

$$T_n''(t) + c^2 \lambda_n T_n(t) = 0.$$

Using $\lambda_n = (n\pi/L)^2$ we can write

$$T_n''(t) + \omega_n^2 T_n(t) = 0,$$

where

$$\omega_n = c \frac{n\pi}{L}, \quad n = 1, 2, 3, \dots$$

The general solution of this second order linear ODE is

$$T_n(t) = A_n \cos(\omega_n t) + B_n \sin(\omega_n t),$$

where A_n and B_n are constants.

Combining the space and time factors, we obtain separated solutions

$$u_n(x, t) = (A_n \cos(\omega_n t) + B_n \sin(\omega_n t)) \sin\left(\frac{n\pi x}{L}\right), \quad n = 1, 2, 3, \dots$$

Because the wave equation is linear, any linear combination of these separated solutions is again a solution. Therefore the general solution satisfying the Dirichlet boundary conditions can be written as an infinite series

$$u(x, t) = \sum_{n=1}^{\infty} (a_n \cos(\omega_n t) + b_n \sin(\omega_n t)) \sin\left(\frac{n\pi x}{L}\right),$$

for suitable coefficients a_n and b_n .

These coefficients will be determined from the initial conditions.

2.3.4 Step 4: Imposing the Initial Conditions and Sine Series

We now use the initial displacement and velocity.

At $t = 0$ we have

$$u(x, 0) = \sum_{n=1}^{\infty} (a_n \cos(0) + b_n \sin(0)) \sin\left(\frac{n\pi x}{L}\right) = \sum_{n=1}^{\infty} a_n \sin\left(\frac{n\pi x}{L}\right).$$

The initial condition $u(x, 0) = f(x)$ becomes

$$f(x) = \sum_{n=1}^{\infty} a_n \sin\left(\frac{n\pi x}{L}\right).$$

This is the Fourier sine series of f on the interval $(0, L)$.

Similarly, we differentiate u with respect to t :

$$\frac{\partial u}{\partial t}(x, t) = \sum_{n=1}^{\infty} (-a_n \omega_n \sin(\omega_n t) + b_n \omega_n \cos(\omega_n t)) \sin\left(\frac{n\pi x}{L}\right).$$

Evaluating at $t = 0$ gives

$$\frac{\partial u}{\partial t}(x, 0) = \sum_{n=1}^{\infty} b_n \omega_n \sin\left(\frac{n\pi x}{L}\right).$$

The initial condition $\frac{\partial u}{\partial t}(x, 0) = g(x)$ becomes

$$g(x) = \sum_{n=1}^{\infty} b_n \omega_n \sin\left(\frac{n\pi x}{L}\right).$$

So the sequence $\{a_n\}$ consists of the Fourier sine coefficients of f , and the sequence $\{b_n\omega_n\}$ consists of the Fourier sine coefficients of g .

Using the orthogonality relations, we obtain explicit formulas for the coefficients. For $n \geq 1$,

$$a_n = \frac{2}{L} \int_0^L f(x) \sin\left(\frac{n\pi x}{L}\right) dx,$$

and

$$b_n\omega_n = \frac{2}{L} \int_0^L g(x) \sin\left(\frac{n\pi x}{L}\right) dx.$$

Therefore

$$b_n = \frac{2}{L\omega_n} \int_0^L g(x) \sin\left(\frac{n\pi x}{L}\right) dx = \frac{2}{L} \frac{1}{cn\pi/L} \int_0^L g(x) \sin\left(\frac{n\pi x}{L}\right) dx.$$

In summary,

$$\begin{aligned} a_n &= \frac{2}{L} \int_0^L f(x) \sin\left(\frac{n\pi x}{L}\right) dx, \\ b_n &= \frac{2}{L\omega_n} \int_0^L g(x) \sin\left(\frac{n\pi x}{L}\right) dx, \quad \omega_n = c \frac{n\pi}{L}. \end{aligned}$$

2.3.5 Step 5: Final Explicit Solution and Interpretation

Substituting these coefficients into the series, we obtain the explicit solution of the wave equation with Dirichlet boundary conditions:

$$u(x, t) = \sum_{n=1}^{\infty} [a_n \cos(\omega_n t) + b_n \sin(\omega_n t)] \sin\left(\frac{n\pi x}{L}\right),$$

where $\omega_n = cn\pi/L$ and the Fourier sine coefficients are

$$a_n = \frac{2}{L} \int_0^L f(y) \sin\left(\frac{n\pi y}{L}\right) dy, \quad b_n = \frac{2}{n\pi c} \int_0^L g(y) \sin\left(\frac{n\pi y}{L}\right) dy.$$

Each term in the sum is a normal mode of vibration: a standing wave with spatial shape $\sin(n\pi x/L)$ and temporal oscillation at frequency ω_n . The coefficients of $\cos(\omega_n t)$ and $\sin(\omega_n t)$ are determined by the initial displacement f and initial velocity g through their Fourier sine coefficients.

Comparing with the heat equation:

- For the heat equation, each mode decayed like $e^{-n^2 t}$ and the solution became smoother in time.
- For the wave equation, each mode oscillates periodically in time with constant amplitude, reflecting conservation of energy in the undamped string.

From a spectral viewpoint, the functions

$$\sin\left(\frac{\pi x}{L}\right), \sin\left(\frac{2\pi x}{L}\right), \sin\left(\frac{3\pi x}{L}\right), \dots$$

form an eigenbasis of the spatial operator

$$Lu = \frac{\partial^2 u}{\partial x^2}$$

with Dirichlet boundary conditions. In this basis, the evolution is diagonal: each mode evolves independently according to a simple harmonic oscillator in time.

As in the heat equation example, a spectral method will approximate $u(x, t)$ by truncating the infinite sum. For some integer $N \geq 1$ we consider the finite approximation

$$u_N(x, t) = \sum_{n=1}^N (a_n \cos(\omega_n t) + b_n \sin(\omega_n t)) \sin\left(\frac{n\pi x}{L}\right).$$

The analytic series above is the infinite dimensional limit of this spectral representation.

2.4 Numerical Illustration

To visualize the oscillatory behavior of the vibrating string, we compute the truncated Fourier sine series solution for a plucked string initial condition. The string is plucked at its center, forming a triangular initial displacement:

$$f(x) = \begin{cases} \frac{2h}{L}x & \text{for } 0 \leq x \leq L/2 \\ 2h(1 - x/L) & \text{for } L/2 \leq x \leq L \end{cases}$$

with zero initial velocity $g(x) = 0$. Here h denotes the height of the pluck at the center.

The Fourier sine coefficients of this triangular shape are

$$a_n = \frac{8h}{n^2\pi^2} \sin\left(\frac{n\pi}{2}\right),$$

which gives nonzero values only for odd n , with alternating signs. The slow algebraic decay of these coefficients reflects the limited regularity of the triangular initial shape; see [21] for a thorough discussion of convergence rates for non-smooth functions.

The key portion of the implementation computes the solution at any point in space and time. In Python:

```
1 def wave_solution(x, t, a_n, b_n, L, c):
2     u = np.zeros_like(x, dtype=float)
3     n_modes = len(a_n) - 1
4     for n in range(1, n_modes + 1):
5         omega_n = c * n * np.pi / L
6         spatial = np.sin(n * np.pi * x / L)
7         temporal = a_n[n] * np.cos(omega_n * t) + b_n[n] * np.sin(omega_n * t)
8         u += temporal * spatial
9     return u
```

Python

The equivalent MATLAB implementation:

```
1 u = zeros(size(x));
2 for n = 1:N_MODES
3     omega_n = C * n * pi / L;
4     spatial = sin(n * pi * x / L);
5     temporal = a_n(n+1) * cos(omega_n * t) + b_n(n+1) * sin(omega_n * t);
6     u = u + temporal * spatial;
7 end
```

Matlab

The Julia implementation:

```
1 function wave_solution(x, t, a_n, b_n, L, c)
2     u = zeros(length(x))
3     n_modes = length(a_n) - 1
4     for n in 1:n_modes
5         omega_n = c * n * pi / L
6         spatial = @. sin(n * pi * x / L)
7         temporal = a_n[n+1] * cos(omega_n * t) + b_n[n+1] * sin(omega_n * t)
8         @. u += temporal * spatial
```

Julia

```

9     end
10    return u
11 end
    
```

Figure 3 shows the evolution of $u_N(x, t)$ with $N = 50$ modes at several time values within half a period $T = 2L/c$. The string oscillates back and forth, with the triangular shape inverting at $t = T/2$. Unlike the heat equation, the wave equation preserves energy and the solution does not decay; it continues oscillating indefinitely.




Figure 3: Evolution of the wave equation solution with a plucked string initial condition. The string oscillates with period $T = 2\pi$, inverting at $t = T/2$.

The code generating Figure 3 is available in:

- codes/python/ch02/wave_equation_evolution.py
- codes/matlab/ch02/wave_equation_evolution.m
- codes/julia/ch02/wave_equation_evolution.jl

The waterfall plot in Figure 4 provides a complete view of the oscillatory dynamics over one full period. Unlike the heat equation, the wave equation conserves energy: the solution oscillates indefinitely without decay, and the periodic nature of the motion is clearly visible in the three-dimensional representation.




Figure 4: Waterfall plot showing the complete space-time evolution of the wave equation solution over one period T . The plucked string oscillates between its initial shape and its mirror image.

The code generating Figure 4 is available in:

- `codes/python/ch02/wave_equation_waterfall.py`
- `codes/matlab/ch02/wave_equation_waterfall.m`
- `codes/julia/ch02/wave_equation_waterfall.jl`

2.5 Laplace Equation in a Periodic Strip

For the elliptic case we consider the Laplace equation in a simple two dimensional domain that is periodic in one direction and bounded in the other. This setting connects naturally with the periodic heat equation example and again leads to a Fourier series representation in the periodic direction.

Let

$$D = \{(x, y) \in \mathbb{R}^2 : 0 < x < 2\pi, 0 < y < 1\}.$$

We seek a harmonic function $u(x, y)$ solving

$$u_{xx}(x, y) + u_{yy}(x, y) = 0, \quad (x, y) \in D,$$

with periodic boundary conditions in x

$$u(x + 2\pi, y) = u(x, y), \quad \text{for all real } x, \quad 0 < y < 1,$$

and Dirichlet conditions in y

$$u(x, 0) = f(x), \quad u(x, 1) = 0, \quad 0 \leq x \leq 2\pi.$$

We assume that f is 2π periodic and smooth:

$$f(x + 2\pi) = f(x).$$

As in the parabolic and hyperbolic examples, we apply separation of variables and obtain a representation of u as an infinite Fourier series in x with y dependent coefficients. The theory of harmonic functions and Laplace's equation is presented in detail in [119]. Modern spectral approaches for elliptic problems in complex domains include multi-domain methods [133] and fictitious domain techniques with circular embedding [57].

2.5.1 Step 1: Separation Ansatz

We look for nontrivial separated solutions of the form

$$u(x, y) = X(x) \cdot Y(y).$$

Substituting into the Laplace equation gives

$$X''(x) \cdot Y(y) + X(x) \cdot Y''(y) = 0.$$

Assuming X and Y are not identically zero, we divide by $X(x) \cdot Y(y)$:

$$\frac{X''(x)}{X(x)} + \frac{Y''(y)}{Y(y)} = 0.$$

The first term depends only on x , the second only on y . Therefore both must be equal to constants whose sum is zero. We introduce a separation constant λ and write

$$\frac{X''(x)}{X(x)} = -\lambda, \quad \frac{Y''(y)}{Y(y)} = \lambda.$$

This leads to the two ordinary differential equations

$$\begin{aligned} X''(x) + \lambda X(x) &= 0, \\ Y''(y) - \lambda Y(y) &= 0. \end{aligned}$$

The periodic boundary conditions for u in the x direction imply the periodic conditions

$$X(x + 2\pi) = X(x), \quad \text{for all real } x.$$

The Dirichlet conditions in y will be enforced later on the full series. For the separated functions Y we will impose appropriate conditions at $y = 1$, while the condition at $y = 0$ will be handled via the Fourier coefficients of f .

We have once more an eigenvalue problem in the periodic direction.

2.5.2 Step 2: Eigenfunctions in the Periodic Direction

The equation for X with periodic boundary conditions is exactly the same as in the heat equation example:

$$X''(x) + \lambda X(x) = 0, \quad X(x + 2\pi) = X(x).$$

We recall the result: there is a constant mode with eigenvalue $\lambda_0 = 0$,

$$X_0(x) = 1,$$

and for each integer $n \geq 1$ there are cosine and sine modes with eigenvalues

$$\lambda_n = n^2,$$

$$X_n^{(c)}(x) = \cos(nx), \quad X_n^{(s)}(x) = \sin(nx).$$

The family

$$1, \cos(x), \sin(x), \cos(2x), \sin(2x), \dots$$

forms an orthogonal basis in $L^2(0, 2\pi)$ for 2π periodic functions.

For each such eigenvalue we now solve the corresponding equation for Y .

2.5.3 Step 3: Equations for the y Dependent Factors

For each eigenvalue λ we have

$$Y''(y) - \lambda Y(y) = 0.$$

We treat separately the constant mode $\lambda_0 = 0$ and the nonzero modes $\lambda_n = n^2$ for $n \geq 1$.

Constant mode $\lambda_0 = 0$

For $\lambda_0 = 0$ the equation reduces to

$$Y_0''(y) = 0.$$

The general solution is

$$Y_0(y) = A_0 + B_0 y.$$

We want separated solutions that satisfy the homogeneous boundary condition at $y = 1$:

$$u(x, 1) = 0 \quad \text{for all } x.$$

For the constant mode this means

$$X_0(x) \cdot Y_0(1) = Y_0(1) = 0,$$

hence

$$Y_0(1) = A_0 + B_0 = 0.$$

We choose a convenient normalization so that $Y_0(0) = 1$. Then $A_0 = 1$ and the relation $A_0 + B_0 = 0$ gives $B_0 = -1$. Thus

$$Y_0(y) = 1 - y.$$

This separated mode

$$u_0(x, y) = X_0(x) \cdot Y_0(y) = 1 - y$$

is harmonic, periodic in x , and vanishes at $y = 1$, with value 1 at $y = 0$.

Higher modes $\lambda_n = n^2$ for $n \geq 1$

For $\lambda_n = n^2$ the equation is

$$Y_n''(y) - n^2 Y_n(y) = 0.$$

The general solution can be written in hyperbolic form

$$Y_n(y) = \alpha_n \cosh(ny) + \beta_n \sinh(ny).$$

We require that each separated mode vanish at $y = 1$:

$$Y_n(1) = 0.$$

To match later the Fourier coefficients of f at $y = 0$, it is convenient to normalize so that

$$Y_n(0) = 1.$$

Imposing $Y_n(0) = 1$ gives

$$\alpha_n = 1.$$

Then

$$Y_n(1) = \cosh(n) + \beta_n \sinh(n) = 0$$

so

$$\beta_n = -\frac{\cosh(n)}{\sinh(n)} = -\coth(n).$$

Thus

$$Y_n(y) = \cosh(ny) - \coth(n) \cdot \sinh(ny).$$

An alternative and simpler expression uses the hyperbolic sine function of the distance to the boundary $y = 1$. One checks that

$$Y_n(y) = \frac{\sinh(n(1-y))}{\sinh(n)}$$

satisfies

$$\begin{aligned} Y_n''(y) - n^2 Y_n(y) &= 0, \\ Y_n(1) &= 0, \\ Y_n(0) &= 1. \end{aligned}$$

Indeed, $Y_n(1) = \sinh(0)/\sinh(n) = 0$, $Y_n(0) = \sinh(n)/\sinh(n) = 1$, and

$$Y_n''(y) = n^2 \frac{\sinh(n(1-y))}{\sinh(n)} = n^2 Y_n(y).$$

We will use the form

$$Y_n(y) = \frac{\sinh(n(1-y))}{\sinh(n)}, \quad n \geq 1.$$

2.5.4 Step 4: Building the Series Solution

Each separated solution corresponding to the eigenfunctions in x and the functions Y_n in y has the form

$$u_0(x, y) = C_0 \cdot Y_0(y) = C_0(1-y),$$

$$u_n^{(c)}(x, y) = C_n \cdot \cos(nx) \cdot Y_n(y),$$

$$u_n^{(s)}(x, y) = D_n \cdot \sin(nx) \cdot Y_n(y), \quad n \geq 1,$$

for some constants C_0, C_n, D_n .

Since the Laplace equation is linear and the boundary condition at $y = 1$ is homogeneous, any linear combination of these separated solutions is again a solution that vanishes at $y = 1$. Therefore a general solution satisfying the periodic condition in x and the Dirichlet condition $u(x, 1) = 0$ can be written as

$$u(x, y) = C_0(1-y) + \sum_{n=1}^{\infty} [C_n \cos(nx) + D_n \sin(nx)] \frac{\sinh(n(1-y))}{\sinh(n)}.$$

It remains to impose the boundary condition at $y = 0$,

$$u(x, 0) = f(x).$$

At $y = 0$ we obtain

$$u(x, 0) = C_0 + \sum_{n=1}^{\infty} [C_n \cos(nx) + D_n \sin(nx)],$$

because $Y_0(0) = 1$ and $Y_n(0) = 1$ for $n \geq 1$.

Hence the boundary condition $u(x, 0) = f(x)$ becomes

$$f(x) = C_0 + \sum_{n=1}^{\infty} [C_n \cos(nx) + D_n \sin(nx)].$$

This is exactly the Fourier series expansion of f . For a 2π periodic f we have

$$f(x) = a_0 + \sum_{n=1}^{\infty} [a_n \cos(nx) + b_n \sin(nx)],$$

with coefficients

$$a_0 = \frac{1}{2\pi} \int_0^{2\pi} f(x) dx,$$

$$a_n = \frac{1}{\pi} \int_0^{2\pi} f(x) \cos(nx) dx, \quad n \geq 1,$$

$$b_n = \frac{1}{\pi} \int_0^{2\pi} f(x) \sin(nx) dx, \quad n \geq 1.$$

By uniqueness of the Fourier expansion, we must have

$$C_0 = a_0, \quad C_n = a_n, \quad D_n = b_n.$$

2.5.5 Step 5: Final Explicit Solution and Interpretation

Substituting these coefficients into the series we obtain the explicit representation

$$u(x, y) = a_0(1 - y) + \sum_{n=1}^{\infty} [a_n \cos(nx) + b_n \sin(nx)] \frac{\sinh(n(1 - y))}{\sinh(n)}, \quad 0 < y < 1.$$

Here a_0 , a_n , and b_n are the Fourier coefficients of the boundary data f as defined above.

This series converges (under mild assumptions on f) to the unique harmonic function that is periodic in x , equal to f on $y = 0$, and zero on $y = 1$.

From a spectral viewpoint:

- The functions $1, \cos(nx), \sin(nx)$ are eigenfunctions of the one dimensional Laplacian $X \mapsto X''$ with periodic boundary conditions in x , with eigenvalues $\lambda_0 = 0$ and $\lambda_n = n^2$.
- For each spatial frequency n in the periodic direction, the dependence in the transverse direction y is determined by the simple ordinary differential equation $Y'' - \lambda_n Y = 0$ with boundary condition $Y(1) = 0$ and normalization $Y(0) = 1$. This gives the hyperbolic profiles

$$Y_0(y) = 1 - y, \\ Y_n(y) = \frac{\sinh(n(1 - y))}{\sinh(n)}, \quad n \geq 1.$$

- The boundary data f at $y = 0$ is expanded in the eigenbasis $\{1, \cos(nx), \sin(nx)\}$ and each Fourier mode is propagated into the interior of the strip with its own y dependent factor $Y_n(y)$.

Analytically, the solution is an infinite sum of separated solutions. In a spectral method we will truncate this sum to finitely many modes in x ,

$$u_N(x, y) = a_0(1 - y) + \sum_{n=1}^N [a_n \cos(nx) + b_n \sin(nx)] \frac{\sinh(n(1 - y))}{\sinh(n)},$$

and approximate the harmonic function inside the strip by this finite Fourier representation. Related spectral techniques for Stokes and Laplace problems in confined periodic domains are developed in [99].

2.6 Numerical Illustration

To visualize the structure of harmonic functions in the strip, we compute the truncated Fourier series solution for a boundary condition that contains two modes:

$$f(x) = \sin(x) + \frac{1}{2} \sin(3x).$$

For this particular boundary data, the Fourier coefficients are simply $b_1 = 1$ and $b_3 = 1/2$, with all other coefficients zero. The solution can be written explicitly as

$$u(x, y) = \sin(x) \frac{\sinh(1 - y)}{\sinh(1)} + \frac{1}{2} \sin(3x) \frac{\sinh(3(1 - y))}{\sinh(3)}.$$

The key portion of the implementation evaluates the truncated series on a two-dimensional grid. In Python:

```
1 def laplace_solution(x, y, a0, a_n, b_n):
2     u = a0 * (1 - y)
3     n_modes = len(a_n) - 1
4     for n in range(1, n_modes + 1):
5         if abs(a_n[n]) < 1e-15 and abs(b_n[n]) < 1e-15:
6             continue
7         y_factor = np.sinh(n * (1 - y)) / np.sinh(n)
8         u += (a_n[n] * np.cos(n * x) + b_n[n] * np.sin(n * x)) * y_factor
9     return u
```

 Python

The equivalent MATLAB implementation:

```
1 U = a0 * (1 - Y);
2 for n = 1:N_MODES
3     if abs(a_n(n+1)) < 1e-15 && abs(b_n(n+1)) < 1e-15
4         continue;
5     end
6     y_factor = sinh(n * (1 - Y)) / sinh(n);
7     U = U + (a_n(n+1) * cos(n*X) + b_n(n+1) * sin(n*X)) .* y_factor;
8 end
```

 Matlab

The Julia implementation:

```
1 function laplace_solution(X, Y, a0, a_n, b_n)
2     U = a0 .* (1.0 .- Y)
3     n_modes = length(a_n) - 1
4     for n in 1:n_modes
5         if abs(a_n[n+1]) < 1e-15 && abs(b_n[n+1]) < 1e-15
6             continue
7         end
8         y_factor = sinh.(n .* (1.0 .- Y)) ./ sinh(n)
9         @. U += (a_n[n+1] * cos(n * X) + b_n[n+1] * sin(n * X)) * y_factor
10    end
11    return U
12 end
```

 Julia

Figure 5 shows the solution $u(x, y)$ in the strip $[0, 2\pi] \times [0, 1]$. At the bottom boundary $y = 0$, the solution matches the prescribed boundary data $f(x)$. As y increases toward the top boundary, the solution decays to zero. Crucially, the higher frequency mode ($n = 3$) decays much faster than the lower frequency mode ($n = 1$), as the hyperbolic factor $\sinh(n(1 - y)) / \sinh(n)$ decreases more rapidly for larger n . This illustrates the smoothing effect of harmonic extension into the interior.




Figure 5: Solution of the Laplace equation in the periodic strip with boundary data $f(x) = \sin(x) + \frac{1}{2}\sin(3x)$ at $y = 0$ and $u = 0$ at $y = 1$. Higher frequency modes decay faster toward the interior.

The code generating Figure 5 is available in:

- `codes/python/ch02/laplace_equation_2d.py`
- `codes/matlab/ch02/laplace_equation_2d.m`
- `codes/julia/ch02/laplace_equation_2d.jl`

2.7 A non-exhaustive literature overview

The method of separation of variables and the spectral methods derived from it share a common intellectual lineage that spans three centuries. This overview traces that arc from the foundational debates of the Enlightenment to the high-performance computing algorithms of the present day.

The Historical Foundations. The techniques presented in this chapter originated in the “vibrating string controversy” of the mid-18th century, a debate that defined the very concept of a mathematical function. In 1747, Jean le Rond d’Alembert derived the wave equation and proposed a solution based on traveling waves, $u(x, t) = f(x - ct) + g(x + ct)$, but argued that the functions must be analytically given [102]. Leonhard Euler expanded this view in 1748, arguing for “discontinuous” functions (in the sense of having corners, like the plucked string) [39]. It was Daniel Bernoulli who, in 1753, proposed that the solution could be represented as an infinite sum of trigonometric modes — the series solutions used in this chapter.

The legitimacy of these infinite series was firmly established by Joseph Fourier in his 1822 treatise *Théorie Analytique de la Chaleur* [48]. Fourier demonstrated that even discontinuous functions could be represented by trigonometric series, applying the method to the heat equation. This work laid the groundwork for Sturm–Liouville theory [114, 115], which generalized the concept of eigenfunctions to a broad class of second-order operators.

The Computational Era. The transition from analytical separation of variables to numerical spectral methods began with Cornelius Lanczos in 1938, who utilized Chebyshev polynomials for differential equations (the Tau method) [74]. The modern era of spectral methods was launched by the rediscovery of the Fast Fourier Transform (FFT) by Cooley and Tukey in 1965 [30], which reduced the computational cost of transformations from $O(N^2)$ to $O(N \log N)$. In the 1970s, Steven Orszag and David Gottlieb codified the theory of spectral methods, demonstrating their superior convergence properties for smooth flow problems compared to finite difference methods [55, 94].

State of the Art (2024–2025). Today, spectral methods remain at the forefront of computational research. Recent developments have focused on overcoming historical limitations regarding domain geometry and temporal discretization.

- **Complex Geometries:** Yao, Wang, and Wang [133] have developed multi-domain Legendre spectral methods capable of solving wave equations in exterior domains with complex obstacles. Similarly, fictitious domain methods with circular embedding have been refined to solve elliptic PDEs in irregular geometries [57].
- **Space-Time Methods:** Moving beyond the traditional method of lines, Kaur *et al.* [71] have introduced space-time spectral methods that treat time as a fourth spectral dimension, achieving high-order accuracy in both space and time for linear and nonlinear PDEs.
- **Non-Local Operators:** Mustapha *et al.* [90] have extended Fourier spectral methods to non-local equations and fractional Laplacians on bounded domains, addressing the unique challenges of boundary conditions in non-local mechanics.
- **Fluid Dynamics:** In the *Journal of Fluid Mechanics*, Peláez *et al.* [99] recently applied spectral solvers to oscillatory Stokes flow in doubly periodic confined domains, demonstrating the method's continued vitality in microfluidics research.

These recent works illustrate that the fundamental principle of this chapter — expanding a solution in a global basis of eigenfunctions — remains a powerful and evolving tool in modern scientific computing.

2.8 Summary

The three examples presented in this chapter (the heat equation, the wave equation, and the Laplace equation) share a common mathematical structure that will guide us throughout the rest of this book. In each case, separation of variables reduces a partial differential equation to a family of ordinary differential equations: an eigenvalue problem in the spatial variable and a simpler equation governing the behavior in the remaining variable (time or the transverse coordinate). The eigenfunctions form an orthogonal basis, and the solution is expressed as an infinite series

$$u(x, t) = \sum_k \hat{u}_k(t) \varphi_k(x)$$

whose coefficients are determined by the initial or boundary data through Fourier projections. This is the DNA of spectral methods.

Let us distill the key ideas:

1. **Separation.** We decompose the solution into modes that evolve independently (or nearly so). Each mode satisfies a simpler equation than the original PDE.
2. **Basis.** The spatial part of each mode is an eigenfunction of a differential operator: Fourier exponentials for periodic problems, trigonometric functions for Dirichlet conditions, Chebyshev or Legendre polynomials for more general settings.

3. **Truncation.** In practice we cannot sum infinitely many terms due to the apparent finitude of our Universe. We retain only the first N modes, and the accuracy of this approximation depends critically on how fast the coefficients $\hat{u}_k$ decay.

When the solution is smooth, the coefficients decay *exponentially*, and a modest N suffices for high accuracy. This is the source of spectral methods' legendary efficiency. For a comprehensive treatment of these classical methods and their mathematical foundations, the reader is referred to [119].

The analytical solutions derived in this chapter are beautiful, but they are also fragile. They apply only to linear equations on simple domains with special boundary conditions. The moment we encounter a nonlinearity, a complicated geometry, or variable coefficients, we must turn to computation. The chapters that follow will develop the computational machinery needed to turn these analytical insights into practical algorithms:

- **FFT and its cousins:** how to move efficiently between physical space and coefficient space.
- **Differentiation matrices:** how to compute derivatives spectrally.
- **Quadrature rules:** how to compute inner products and projections.
- **Time stepping:** how to advance the ODE system for the coefficients.

Armed with these tools, we will be able to solve problems far beyond the reach of pen-and-paper analysis.



CHAPTER 3

Mise en Bouche

In French cuisine, a *mise en bouche* is a small appetizer offered by the chef to stimulate the palate before the main courses arrive. In this chapter we offer a similar intellectual appetizer: a compact, self-contained taste of spectral methods that illuminates the essential mechanics before we develop the full machinery of Fourier and Chebyshev approximation.

Rather than jumping immediately to high-degree polynomials with $N = 100$, we perform hand calculations with just $N = 2$ or $N = 3$ unknowns. This low-dimensional setting makes every step transparent. We can verify each formula by direct computation and gain intuition that will guide us through the more sophisticated developments to come.

The techniques presented here follow the classical exposition in [21], adapted to our pedagogical goals. The unifying framework of the Method of Weighted Residuals was first systematized in the seminal review by [44] and later expanded in [43]. This framework connects the collocation (pseudospectral) approach we favor in this book with the Galerkin methods that dominate finite element analysis.

3.1 The Method of Weighted Residuals

3.1.1 Series Expansions and the Residual Function

The central idea of spectral methods is to approximate the unknown function $u(x)$ by a finite sum of basis functions:

$$u(x) \approx u_N(x) = \sum_{n=0}^N a_n \varphi_n(x),$$

where $\{\varphi_n(x)\}$ are known basis functions and $\{a_n\}$ are unknown coefficients to be determined.

When we substitute this approximation into a differential equation

$$\mathcal{L}u = f(x),$$

where $\mathcal{L}$ is a linear differential operator, the result is generally not zero. The *residual function* measures this discrepancy:

$$R(x; a_0, a_1, \dots, a_N) = \mathcal{L}u_N - f.$$

For the exact solution, $R(x) \equiv 0$. The challenge is to choose the coefficients $\{a_n\}$ so that the residual is as small as possible. Different spectral methods correspond to different strategies for minimizing this residual.

3.1.2 Two Minimization Strategies

The two most important strategies are:

1. **Collocation (Pseudospectral) Method:** Force the residual to be exactly zero at a set of carefully chosen points $\{x_j\}$, called collocation points:

$$R(x_j; a_0, \dots, a_N) = 0, \quad j = 1, 2, \dots, N + 1.$$

This gives $N + 1$ equations for $N + 1$ unknowns.

2. **Galerkin Method:** Require the residual to be orthogonal to each basis function (a concept originating with [51]) in the sense of a weighted inner product:

$$\int_{-1}^1 R(x) \varphi_k(x) w(x) dx = 0, \quad k = 0, 1, \dots, N,$$

where $w(x)$ is a weight function (often $w(x) = 1$ for polynomial bases).

Both methods convert the differential equation into a system of algebraic equations for the unknown coefficients. The collocation approach is simpler to implement and handles nonlinear terms easily, while the Galerkin approach often provides better global accuracy in weighted norms.

3.2 A First Collocation Example

We illustrate the collocation method with a complete worked example that can be verified by hand calculation.

3.2.1 Problem Statement

Consider the boundary value problem on $[-1, 1]$:

$$u''(x) - (4x^2 + 2)u(x) = 0, \quad -1 \leq x \leq 1,$$

with boundary conditions

$$u(-1) = 1, \quad u(1) = 1.$$

3.2.2 The Exact Solution

The exact solution is

$$u_{\text{exact}}(x) = e^{x^2-1}.$$

Let us verify this claim. The first derivative is

$$u'_{\text{exact}}(x) = 2xe^{x^2-1}.$$

The second derivative is

$$u''_{\text{exact}}(x) = (2 + 4x^2)e^{x^2-1} = (4x^2 + 2)u_{\text{exact}}(x).$$

Substituting into the ODE:

$$u''_{\text{exact}} - (4x^2 + 2)u_{\text{exact}} = (4x^2 + 2)u_{\text{exact}} - (4x^2 + 2)u_{\text{exact}} = 0. \checkmark$$

The boundary conditions are satisfied:

$$u_{\text{exact}}(\pm 1) = e^{1-1} = e^0 = 1. \checkmark$$

3.2.3 Trial Function

To satisfy the boundary conditions automatically, we write the approximation in a form that equals 1 at $x = \pm 1$ regardless of the coefficient values. A convenient choice is:

$$u_2(x) = 1 + (1 - x^2)(a_0 + a_1x + a_2x^2).$$

The factor $(1 - x^2)$ vanishes at the endpoints, so

$$u_2(\pm 1) = 1 + 0 \cdot (\dots) = 1$$

for any values of a_0, a_1, a_2 . We have three undetermined coefficients.

Expanding the trial function:

$$\begin{aligned} u_2(x) &= 1 + a_0 + a_1x + a_2x^2 - a_0x^2 - a_1x^3 - a_2x^4 \\ &= (1 + a_0) + a_1x + (a_2 - a_0)x^2 - a_1x^3 - a_2x^4. \end{aligned}$$

3.2.4 Computing the Residual

The residual is

$$R(x; a_0, a_1, a_2) = u_2''(x) - (4x^2 + 2)u_2(x).$$

Computing the second derivative of u_2 :

$$u_2'(x) = a_1 + 2(a_2 - a_0)x - 3a_1x^2 - 4a_2x^3,$$

$$u_2''(x) = 2(a_2 - a_0) - 6a_1x - 12a_2x^2.$$

Substituting into the residual and simplifying (a calculation best verified with computer algebra), the residual is a polynomial of degree six in x with coefficients that depend linearly on a_0, a_1, a_2 :

$$\begin{aligned} R(x) = & (-2 - 4a_0 + 2a_2) - 8a_1x + (-4 - 2a_0 - 14a_2)x^2 - 2a_1x^3 \\ & + (4a_0 - 6a_2)x^4 + 4a_1x^5 + 4a_2x^6. \end{aligned}$$

3.2.5 Collocation Conditions

We have three unknowns, so we choose three collocation points in the interior of the interval. A simple choice is

$$x_1 = -\frac{1}{2}, \quad x_2 = 0, \quad x_3 = \frac{1}{2}.$$

Setting the residual to zero at these points gives three linear equations:

At $x = -1/2$:

$$R\left(-\frac{1}{2}\right) = -\frac{17}{4}a_0 + \frac{33}{8}a_1 - \frac{25}{16}a_2 - 3 = 0.$$

At $x = 0$:

$$R(0) = -4a_0 + 2a_2 - 2 = 0.$$

At $x = 1/2$:

$$R\left(\frac{1}{2}\right) = -\frac{17}{4}a_0 - \frac{33}{8}a_1 - \frac{25}{16}a_2 - 3 = 0.$$

3.2.6 Solving the System

From the second equation:

$$-4a_0 + 2a_2 = 2 \Rightarrow a_2 = 1 + 2a_0.$$

Adding the first and third equations (the a_1 terms cancel):

$$-\frac{17}{2}a_0 - \frac{25}{8}a_2 = 6.$$

Substituting $a_2 = 1 + 2a_0$:

$$-\frac{17}{2}a_0 - \frac{25}{8}(1 + 2a_0) = 6$$

$$-\frac{17}{2}a_0 - \frac{25}{8} - \frac{25}{4}a_0 = 6$$

$$-\left(\frac{34}{4} + \frac{25}{4}\right)a_0 = 6 + \frac{25}{8}$$

$$-\frac{59}{4}a_0 = \frac{73}{8}$$

$$a_0 = -\frac{73}{118}.$$

Then

$$a_2 = 1 + 2 \cdot \left(-\frac{73}{118}\right) = 1 - \frac{146}{118} = -\frac{28}{118} = -\frac{14}{59}.$$

Subtracting the first equation from the third:

$$-\frac{33}{4}a_1 = 0 \Rightarrow a_1 = 0.$$

The vanishing of a_1 reflects the symmetry of the problem: both the differential equation and the boundary conditions are symmetric about $x = 0$, so the solution must be an even function. An odd coefficient like a_1 would break this symmetry.

3.2.7 The Approximate Solution

Substituting the coefficients back:

$$u_2(x) = 1 + (1 - x^2) \left(-\frac{73}{118} - \frac{14}{59}x^2 \right).$$

After simplification, this becomes the even polynomial:

$$u_2(x) = \frac{14}{59}x^4 + \frac{45}{118}x^2 + \frac{45}{118}.$$

The boundary conditions are satisfied:

$$u_2(\pm 1) = \frac{14}{59} + \frac{45}{118} + \frac{45}{118} = \frac{28}{118} + \frac{90}{118} = \frac{118}{118} = 1. \checkmark$$

The implementation of this approximation is straightforward. In Python:

```
1 def u_approx(x):
2     """Evaluate the collocation approximation."""
3     a0, a1, a2 = -73/118, 0, -14/59
4     return 1 + (1 - x**2) * (a0 + a1*x + a2*x**2)
```

 Python

The equivalent MATLAB implementation:

```
1 % Collocation coefficients
2 a0 = -73/118; a1 = 0; a2 = -14/59;
3
4 % Approximate solution (anonymous function)
5 u_approx = @(x) 1 + (1 - x.^2) .* (a0 + a1*x + a2*x.^2);
```

 Matlab

The Julia implementation:

```
1 function u_approx(x)
2     # Coefficients from solving the collocation system
3     a0 = -73 / 118
4     a1 = 0.0
5     a2 = -14 / 59
6     return 1 + (1 - x^2) * (a0 + a1 * x + a2 * x^2)
7 end
```

 Julia

3.2.8 Error Analysis

The following table compares the exact and approximate solutions at several points:

x	$u_{\text{exact}}(x)$	$u_{\text{approx}}(x)$	Error
-1	1.000 00	1.000 00	0.000 00
-0.5	0.472 37	0.491 53	-0.019 16
0	0.367 88	0.381 36	-0.013 48
0.5	0.472 37	0.491 53	-0.019 16
1	1.000 00	1.000 00	0.000 00

Table 1: Comparison of exact and three-coefficient collocation approximation.

The maximum error is approximately 2×10^{-2} , which is remarkably good for such a low-order approximation. Figure 6 shows the solutions graphically.




Figure 6: Left: exact solution $u(x) = e^{x^2-1}$ compared with the three-coefficient collocation approximation. The collocation points $x = -1/2, 0, 1/2$ are marked with squares. Right: the error $u_{\text{exact}} - u_{\text{approx}}$.

The code generating Figure 6 is available in:

- codes/python/ch03/collocation_example1.py
- codes/matlab/ch03/collocation_example1.m
- codes/julia/ch03/collocation_example1.jl

3.3 Collocation versus Galerkin

To compare the two main approaches to the Method of Weighted Residuals, a comparison famously analyzed by [125], we consider a second example where both methods can be applied with explicit hand calculations.

3.3.1 Problem Statement

Consider the reaction-diffusion equation on $[-1, 1]$:

$$\frac{d^2u}{dx^2} - 4u = -1, \quad -1 \leq x \leq 1,$$

with homogeneous Dirichlet boundary conditions:

$$u(-1) = 0, \quad u(1) = 0.$$

3.3.2 The Exact Solution

The homogeneous equation $u'' - 4u = 0$ has the general solution $u_h = A \cosh(2x) + B \sinh(2x)$. A particular solution for the constant forcing -1 is $u_p = 1/4$. The general solution is therefore

$$u(x) = A \cosh(2x) + B \sinh(2x) + \frac{1}{4}.$$

The boundary condition $u(-1) = 0$ gives $A \cosh(2) - B \sinh(2) + 1/4 = 0$. The boundary condition $u(1) = 0$ gives $A \cosh(2) + B \sinh(2) + 1/4 = 0$.

Adding these equations: $2A \cosh(2) + 1/2 = 0$, so $A = -1/(4 \cosh(2))$. Subtracting: $2B \sinh(2) = 0$, so $B = 0$.

The exact solution is:

$$u_{\text{exact}}(x) = \frac{1}{4} \left(1 - \frac{\cosh(2x)}{\cosh(2)} \right).$$

The maximum value occurs at $x = 0$:

$$u_{\text{exact}}(0) = \frac{1}{4} \left(1 - \frac{1}{\cosh(2)} \right) \approx 0.1834.$$

3.3.3 Trial Function and Basis

Since the boundary conditions are homogeneous, we choose basis functions that automatically vanish at $x = \pm 1$. For a symmetric problem like this one, we use even functions:

$$\varphi_0(x) = 1 - x^2, \quad \varphi_1(x) = (1 - x^2)x^2 = x^2 - x^4.$$

Both functions vanish at $x = \pm 1$ and are even in x . Our trial function is:

$$u_1(x) = a_0 \varphi_0(x) + a_1 \varphi_1(x) = a_0(1 - x^2) + a_1(x^2 - x^4).$$

3.3.4 The Residual

The operator is $\mathcal{L}u = u'' - 4u$. We compute:

$$\varphi_0'' = -2, \quad \varphi_1'' = 2 - 12x^2.$$

Applying the operator to each basis function:

$$\mathcal{L}\varphi_0 = -2 - 4(1 - x^2) = -6 + 4x^2,$$

$$\mathcal{L}\varphi_1 = (2 - 12x^2) - 4(x^2 - x^4) = 2 - 16x^2 + 4x^4.$$

The residual is:

$$\begin{aligned} R(x) &= a_0 \mathcal{L}\varphi_0 + a_1 \mathcal{L}\varphi_1 - (-1) \\ &= a_0(-6 + 4x^2) + a_1(2 - 16x^2 + 4x^4) + 1. \end{aligned}$$

3.3.5 Collocation Method

With two unknowns, we need two collocation points. Due to symmetry, we can use points in $[0, 1)$. We choose $x_1 = 0$ and $x_2 = 0.5$.

At $x = 0$:

$$\mathcal{L}\varphi_0(0) = -6, \quad \mathcal{L}\varphi_1(0) = 2.$$

$$R(0) = -6a_0 + 2a_1 + 1 = 0 \Rightarrow 6a_0 - 2a_1 = 1.$$

At $x = 0.5$:

$$\mathcal{L}\varphi_0(0.5) = -6 + 4 \cdot 0.25 = -5,$$

$$\mathcal{L}\varphi_1(0.5) = 2 - 16 \cdot 0.25 + 4 \cdot 0.0625 = 2 - 4 + 0.25 = -1.75.$$

$$R(0.5) = -5a_0 - 1.75a_1 + 1 = 0 \Rightarrow 5a_0 + 1.75a_1 = 1.$$

Solving the system:

$$\begin{cases} 6a_0 - 2a_1 = 1 \\ 5a_0 + 1.75a_1 = 1 \end{cases}$$

Multiply the first equation by 5 and the second by 6:

$$30a_0 - 10a_1 = 5, \quad 30a_0 + 10.5a_1 = 6.$$

Subtracting: $20.5a_1 = 1$, so $a_1 \approx 0.04878$.

Substituting back: $6a_0 = 1 + 2 \cdot 0.04878 = 1.09756$, so $a_0 \approx 0.1829$.

The collocation solution is:

$$u_{\text{coll}}(x) \approx 0.1829(1 - x^2) + 0.0488(x^2 - x^4).$$

The following Python code assembles and solves the collocation system:

```
1 def solve_collocation():
2     """Solve the collocation system at x = 0 and x = 0.5."""
3     # Operator L[φ] = φ'' - 4φ applied to basis functions
4     L_phi0 = lambda x: -2 - 4*(1 - x**2)
5     L_phi1 = lambda x: (2 - 12*x**2) - 4*(x**2 - x**4)
6
7     # Build system matrix at collocation points
8     A = np.array([[L_phi0(0.0), L_phi1(0.0)],
9                  [L_phi0(0.5), L_phi1(0.5)]])
10    b = np.array([-1.0, -1.0]) # RHS from f = -1
11    return np.linalg.solve(A, b)
```

 Python

The equivalent MATLAB implementation:

```
1 % Operator L = d^2/dx^2 - 4 applied to basis functions
2 L_phi0 = @(x) -2 - 4*(1 - x.^2);
3 L_phi1 = @(x) (2 - 12*x.^2) - 4*(x.^2 - x.^4);
4
5 % Build and solve collocation system
6 A_coll = [L_phi0(0.0), L_phi1(0.0);
7           L_phi0(0.5), L_phi1(0.5)];
8 coeffs = A_coll \ [-1; -1];
```

 Matlab

The Julia implementation:

```
1 # Operator L = d^2/dx^2 - 4 applied to basis functions
2 L_phi0(x) = -2 - 4 * (1 - x^2)
3 L_phi1(x) = (2 - 12x^2) - 4 * (x^2 - x^4)
4
5 # Build and solve collocation system
6 A = [L_phi0(0.0) L_phi1(0.0);
7      L_phi0(0.5) L_phi1(0.5)]
8 coeffs = A \ [-1.0, -1.0]
```

 Julia

3.3.6 Galerkin Method

The Galerkin conditions require the residual to be orthogonal to each basis function:

$$\int_{-1}^1 R(x) \varphi_k(x) dx = 0, \quad k = 0, 1.$$

This gives a symmetric matrix system $\mathbf{A}\mathbf{a} = \mathbf{b}$ where:

$$A_{ij} = \int_{-1}^1 \mathcal{L}\varphi_j(x) \cdot \varphi_i(x) dx,$$

$$b_i = \int_{-1}^1 (-1) \cdot \varphi_i(x) dx = - \int_{-1}^1 \varphi_i(x) dx.$$

Computing the integrals (using standard formulas for powers of x):

For $\varphi_0 = 1 - x^2$:

$$\int_{-1}^1 \varphi_0(x) dx = \int_{-1}^1 (1 - x^2) dx = \left[x - \frac{x^3}{3} \right]_{-1}^1 = 2 - \frac{2}{3} = \frac{4}{3}.$$

For $\varphi_1 = x^2 - x^4$:

$$\int_{-1}^1 \varphi_1(x) dx = \left[\frac{x^3}{3} - \frac{x^5}{5} \right]_{-1}^1 = \frac{2}{3} - \frac{2}{5} = \frac{4}{15}.$$

The right-hand side is:

$$b_0 = -\frac{4}{3}, \quad b_1 = -\frac{4}{15}.$$

The matrix entries require more computation. Using computer algebra or careful integration:

$$A_{00} = -\frac{104}{15}, \quad A_{01} = A_{10} = -\frac{8}{7}, \quad A_{11} = -\frac{328}{315}.$$

Solving the system yields:

$$a_0 \approx 0.1832, \quad a_1 \approx 0.0550.$$

The Galerkin method requires numerical integration to assemble the system. In Python:

```
1 def solve_galerkin():
2     """Solve using Galerkin:  $\langle R, \varphi_k \rangle = 0$  for  $k = 0, 1$ ."""
3     from scipy import integrate
4
5     # Matrix entries:  $A_{ij} = \int L[\varphi_j] \varphi_i dx$ 
6     A00, _ = integrate.quad(lambda x: L_phi0(x) * phi0(x), -1, 1)
7     A01, _ = integrate.quad(lambda x: L_phi1(x) * phi0(x), -1, 1)
8     A10, _ = integrate.quad(lambda x: L_phi0(x) * phi1(x), -1, 1)
9     A11, _ = integrate.quad(lambda x: L_phi1(x) * phi1(x), -1, 1)
10
11     A = np.array([[A00, A01], [A10, A11]])
12     b0, _ = integrate.quad(lambda x: -phi0(x), -1, 1)
13     b1, _ = integrate.quad(lambda x: -phi1(x), -1, 1)
14     return np.linalg.solve(A, [b0, b1])
```

 Python

The equivalent MATLAB implementation uses the built-in integral function:

```
1 % Matrix entries:  $A_{ij} = \int L[\varphi_j] \varphi_i dx$ 
2 A00 = integral(@(x) L_phi0(x) .* phi0(x), -1, 1);
3 A01 = integral(@(x) L_phi1(x) .* phi0(x), -1, 1);
4 A10 = integral(@(x) L_phi0(x) .* phi1(x), -1, 1);
```

 Matlab

```

5  A11 = integral(@(x) L_phi1(x) .* phi1(x), -1, 1);
6
7  A_gal = [A00, A01; A10, A11];
8
9  % RHS: b_i = ∫ f φ_i dx where f = -1
10 b0 = integral(@(x) -phi0(x), -1, 1);
11 b1 = integral(@(x) -phi1(x), -1, 1);
12 coeffs = A_gal \ [b0; b1];

```

The Julia implementation uses QuadGK for numerical integration:

```

1  using QuadGK
2
3  # Matrix entries: A_{ij} = ∫ L[φ_j] φ_i dx
4  A00, _ = quadgk(x -> L_phi0(x) * phi0(x), -1, 1)
5  A01, _ = quadgk(x -> L_phi1(x) * phi0(x), -1, 1)
6  A10, _ = quadgk(x -> L_phi0(x) * phi1(x), -1, 1)
7  A11, _ = quadgk(x -> L_phi1(x) * phi1(x), -1, 1)
8
9  A = [A00 A01; A10 A11]
10
11 # RHS: b_i = ∫ f φ_i dx where f = -1
12 b0, _ = quadgk(x -> -phi0(x), -1, 1)
13 b1, _ = quadgk(x -> -phi1(x), -1, 1)
14 coeffs = A \ [b0, b1]

```

 Julia

3.3.7 Comparison

The following table compares the two methods at the central point $x = 0$:

Method	$u(0)$	Absolute Error
Exact	0.1835	0
Collocation ($N = 2$)	0.1829	0.0006
Galerkin ($N = 2$)	0.1832	0.0003

Table 2: Comparison of spectral approximations at the central maximum.

For this problem, the Galerkin method is more accurate both at the central point and in a global sense. This is consistent with the error plot in Figure 7, which shows the Galerkin error (green) remaining closer to zero across the entire interval. The Galerkin method minimizes the error in a root-mean-square sense, which typically leads to better overall accuracy for smooth problems.

Figure 7 shows both approximate solutions compared to the exact solution, along with the error profiles.




Figure 7: Left: exact solution compared with collocation and Galerkin approximations. The collocation points $x = 0$ and $x = 0.5$ are marked. Right: error profiles for both methods.

The code generating Figure 7 is available in:

- `codes/python/ch03/collocation_vs_galerkin.py`
- `codes/matlab/ch03/collocation_vs_galerkin.m`
- `codes/julia/ch03/collocation_vs_galerkin.jl`

3.4 Conclusions and Questions

These simple examples illuminate several important features of spectral methods:

1. **Automatic satisfaction of boundary conditions.** By choosing trial functions that vanish at the boundaries (or equal the prescribed boundary values), we eliminate the boundary conditions from the algebraic system.
2. **Symmetry exploitation.** When the problem has symmetry, the solution inherits that symmetry. In our examples, the vanishing of odd coefficients ($a_1 = 0$) reflects the even symmetry of the exact solution.
3. **Simplicity of implementation.** Even with hand calculations, we can obtain remarkably accurate approximations with just a few coefficients.
4. **Trade-offs between methods.** Collocation is simpler to implement (no integrals to compute), while Galerkin typically provides better accuracy by minimizing a global error measure.

These examples raise important questions that we will address in subsequent chapters:

- **What is the optimal choice of basis functions?** Using simple powers of x works for small N , but becomes numerically unstable for large N . Chebyshev and Fourier bases are far superior.
- **What are the optimal collocation points?** Our ad hoc choices $x = -1/2, 0, 1/2$ worked well, but there exist optimal point distributions (Gauss and Gauss–Lobatto points) derived from orthogonal polynomial theory.
- **How fast does the error decrease as N increases?** For smooth solutions, spectral methods achieve exponential convergence (the error decreases like c^{-N} for some $c > 1$), which is dramatically faster than the algebraic convergence $O(N^{-p})$ of finite difference and finite element methods.

The following chapters will develop the theory and algorithms needed to answer these questions and to apply spectral methods to a wide range of problems.

3.5 A Broader Perspective

For demanding students who wish to understand how spectral methods fit into the wider landscape of numerical analysis, this section provides a high-level comparison with other approaches, particularly finite element methods. The discussion follows [21].

3.5.1 Local versus Global Basis Functions

The fundamental distinction between spectral methods and finite element methods lies in the *support* of the basis functions. In finite element methods, the computational domain is divided into many small sub-intervals (or triangles, tetrahedra in higher dimensions), and the basis functions $\varphi_n(x)$ are *local*: they are polynomials of fixed, low degree (typically linear or quadratic) that are non-zero only over one or two adjacent elements.

In contrast, spectral methods use *global* basis functions. Each $\varphi_n(x)$ is a polynomial (or trigonometric polynomial) of potentially high degree that is non-zero (except at isolated points) over the entire computational domain. This global character is both the source of spectral methods' power and the reason for some of their limitations.

Figure 8 illustrates this distinction schematically. The finite element basis function (left) has compact support and contributes to the solution only locally. The spectral basis function (right) influences the solution everywhere.




Figure 8: Schematic comparison of local (finite element) and global (spectral) basis functions. The finite element “hat function” is non-zero only over two adjacent elements, while the spectral basis function oscillates across the entire domain.

3.5.2 Refinement Strategies

When a numerical approximation is insufficiently accurate, there are three fundamentally different strategies to improve it, illustrated schematically in Figure 9:

1. ***h*-refinement**: Subdivide each element into smaller pieces, reducing the mesh spacing h uniformly throughout the domain. This increases the number of elements while keeping the polynomial degree fixed.
2. ***r*-refinement** (adaptive): Redistribute the mesh points, clustering them in regions where the solution has steep gradients or other features requiring high resolution. The total number of degrees of freedom remains roughly constant.

3. ***p*-refinement:** Keep the mesh fixed while increasing p , the polynomial degree within each element. For a single-element domain, this is precisely what spectral methods do. This approach is mathematically equivalent to the p -version of the finite element method established by [7].




Figure 9: Three strategies for improving accuracy in numerical methods. Spectral methods employ p -refinement: increasing the polynomial degree while using a single element (or few elements) spanning the entire domain.

Spectral methods can be viewed as the extreme form of p -refinement: a single element spans the entire domain, and accuracy is improved solely by increasing the polynomial degree. This strategy is devastatingly effective when the solution is smooth, but struggles when the solution has discontinuities or sharp gradients.

3.5.3 Trade-offs: Sparse versus Full Matrices

The choice between local and global basis functions entails fundamental trade-offs:

Finite element advantages:

- *Sparse matrices:* Since each basis function is non-zero over only a few elements, the stiffness matrix has mostly zero entries. Sparse matrix solvers can exploit this structure, dramatically reducing computational cost for large systems.
- *Geometric flexibility:* The small elements (triangles, tetrahedra) can be fitted to irregularly shaped domains like automobile bodies or aircraft wings.

Finite element disadvantages:

- *Low accuracy per degree of freedom:* Each basis function is a polynomial of low degree, so many elements are needed for high accuracy.

Spectral method advantages:

- *High accuracy for smooth problems:* The high-degree global polynomials capture smooth solutions with far fewer degrees of freedom.
- *Efficiency with iterative solvers:* When fast iterative methods are used, spectral methods can be much more efficient than low-order methods for many problem classes.

Spectral method disadvantages:

- *Full matrices:* The global basis functions create dense matrices where most entries are non-zero.
- *Geometric limitations:* Spectral methods are most natural on simple domains (intervals, rectangles, disks) and require more sophisticated techniques for irregular geometries.

For problems with smooth solutions on regular domains (many important problems in fluid dynamics, quantum mechanics, and wave propagation fall into this category), the accuracy advantage of spectral methods often outweighs the matrix structure disadvantage.

3.5.4 Spectral Element Methods

A natural question arises: can we combine the geometric flexibility of finite elements with the high accuracy of spectral methods? The answer is yes, through *spectral element methods*, as introduced by [98].

In spectral element methods, the domain is subdivided into elements (as in finite elements), but within each element, the polynomial degree p is chosen to be moderately high, typically $p = 6$ to 8 . This hybrid approach inherits several advantages:

- The element subdivision provides geometric flexibility and matrix sparsity.
- The high polynomial degree within each element provides spectral-like accuracy.
- The theoretical framework is essentially the same as for global spectral methods.

Spectral element codes are typically written so that p is a user-adjustable parameter, allowing practitioners to balance accuracy and cost for their specific application. We will not develop spectral element methods in detail in this book, but the reader should be aware that much of the theory developed for global spectral methods transfers directly to the spectral element context.

3.5.5 The Convergence of Methods at High Order

Perhaps the most profound insight from the comparison between finite element and spectral methods is this: *for sufficiently high polynomial degree, the two approaches become essentially equivalent*.

Low-order finite elements (linear, quadratic) can be derived, justified, and implemented without knowledge of Fourier or Chebyshev convergence theory. However, as the polynomial degree increases, ad hoc approaches become increasingly ill-conditioned and numerically unstable, a phenomenon rigorously analyzed in [55]. The only practical way to implement well-behaved high-order finite elements (say, sixth order or higher) is to use the technology of spectral methods: Chebyshev or Legendre basis functions, Gaussian quadrature, and the convergence theory we will develop in subsequent chapters.

Thus, the question “Are finite elements or spectral methods better?” becomes somewhat artificial for high-order approximations. The real question is: *Does the problem at hand require high-order accuracy, or is second or fourth order sufficient?*

When the solution is smooth and high accuracy is needed, the spectral/high-order approach is clearly superior. When the solution has discontinuities, shocks, or boundary layers, or when the geometry is highly irregular, low-order methods with adaptive mesh refinement may be more practical. The wise practitioner chooses the tool appropriate to the problem.

3.6 A non-exhaustive literature overview

The Method of Weighted Residuals (MWR), which forms the backbone of the techniques explored in this chapter, is not a single invention but a synthesis of ideas that spanned the early 20th century. The unifying framework was formalized in the landmark review by [44], who demonstrated that diverse approximation schemes—including the method of moments, the subdomain method, and the Galerkin method—could all be rigorously classified by their choice of weight function. This taxonomy transformed numerical analysis from a collection of ad hoc recipes into a coherent mathematical discipline.

The specific tension between **collocation** and **Galerkin** methods discussed in this chapter was a central debate in the computational community for decades. While Galerkin methods offered theoretical optimality in energy norms, their implementation was hampered by the high cost of numerical integration. This bottleneck was resolved by the seminal work of [125] in chemical engineering. They introduced **orthogonal collocation**, proving that by selecting collocation points as the roots of

orthogonal polynomials, one could achieve the accuracy of Galerkin methods with the computational efficiency of point-wise evaluation. This insight is the direct mathematical ancestor of the pseudospectral methods we utilize in this text.

The rigorous analysis of these methods was cemented in the 1970s. [94] and [55] established the error estimates for “pseudospectral” methods, demonstrating that the aliasing errors introduced by collocation are generally of the same order as the truncation errors, thus validating the use of the Fast Fourier Transform for nonlinear problems in fluid dynamics.

The “broader perspective” of local versus global bases mirrors the historical development of the **Finite Element Method** (FEM). While classical FEM relies on mesh refinement (h -version), [7] pioneered the p -version of the finite element method, proving that increasing the polynomial degree p on a fixed mesh yields exponential convergence for smooth solutions. This philosophy evolved into the **Spectral Element Method** (SEM) introduced by [98], which combines the geometric flexibility of finite elements with the high-order accuracy of spectral expansions—a synthesis that remains the gold standard for high-fidelity simulation in complex geometries.

In the contemporary era (2020–2026), spectral methods are undergoing a renaissance at the intersection of machine learning and nonlocal physics. [87] and [91] review how Deep Operator Networks (DeepONet) and Physics-Informed Neural Networks (PINNs) are being used to “learn” optimal spectral bases, overcoming the curse of dimensionality in high-dimensional PDEs. Furthermore, the extension of spectral methods to **fractional differential equations**, as detailed in the recent monograph by [134], and the development of spectral solvers for **nonlocal diffusion on bounded domains** [90], illustrate the enduring adaptability of the spectral approach to the frontiers of mathematical physics.

3.7 Summary

This chapter has provided a first taste of spectral methods through low-dimensional, hand-computable examples:

1. **Method of Weighted Residuals:** Spectral methods approximate the solution as a finite sum of basis functions and determine the unknown coefficients by minimizing the residual — either pointwise (collocation) or in a weighted integral sense (Galerkin).
2. **Collocation simplicity:** Forcing the residual to vanish at selected points converts a differential equation into a small algebraic system, with no integrals to evaluate.
3. **Galerkin accuracy:** Requiring orthogonality of the residual against basis functions typically yields better global accuracy, at the cost of computing integrals.
4. **Boundary conditions:** Choosing trial functions that automatically satisfy the boundary data eliminates boundary conditions from the algebraic system.
5. **Symmetry exploitation:** When the problem possesses symmetry, the solution inherits that symmetry, and odd coefficients vanish automatically — reducing the effective size of the algebraic system.
6. **Local versus global bases:** Finite element methods use low-degree, compactly supported basis functions on many small elements, producing sparse matrices. Spectral methods use high-degree global polynomials on the entire domain, producing dense but small matrices. For smooth solutions, the spectral approach requires far fewer degrees of freedom.

The examples in this chapter were intentionally simple. The chapters that follow will replace ad hoc polynomial bases with Chebyshev and Fourier expansions, and hand-chosen collocation points with optimal distributions derived from orthogonal polynomial theory.



CHAPTER 4

The Geometry of Nodes

Before we can differentiate functions numerically using spectral methods, we must first understand how to represent them. Polynomial interpolation is the process of constructing a polynomial that passes through a given set of data points. It is the foundation upon which pseudospectral methods are built. In this chapter, we explore a fascinating paradox: while polynomial interpolation seems entirely straightforward, the choice of interpolation nodes determines whether the method succeeds brilliantly or fails catastrophically.

The story begins with a surprising discovery by the German mathematician Carl Runge in 1901 [103]. Attempting to approximate a simple, smooth function by interpolating polynomials, Runge found that increasing the polynomial degree made the approximation *worse*, not better. This counterintuitive phenomenon, now bearing his name, reveals deep connections between numerical analysis, complex analysis, and potential theory.

Our journey through this “étude in grid geometry” will explain the Runge phenomenon, develop the theoretical framework of potential theory that predicts where interpolation succeeds or fails, and introduce the Chebyshev points that form the cornerstone of practical spectral methods.

4.1 The Problem: Polynomial Interpolation

4.1.1 Lagrange Interpolation

Given $N + 1$ distinct points $\{x_0, x_1, \dots, x_N\}$ on an interval $[a, b]$ and corresponding function values $\{f_0, f_1, \dots, f_N\}$ where $f_j = f(x_j)$, there exists a unique polynomial $p_N(x)$ of degree at most N such that

$$p_N(x_j) = f_j, \quad j = 0, 1, \dots, N.$$

This interpolating polynomial can be written explicitly using the *Lagrange formula*:

$$p_N(x) = \sum_{k=0}^N f_k L_k(x),$$

where the *Lagrange basis polynomials* are

$$L_k(x) = \prod_{j=0, j \neq k}^N \frac{x - x_j}{x_k - x_j}.$$

Each basis polynomial $L_k(x)$ has the cardinal property: it equals 1 at x_k and 0 at all other nodes. This property ensures that substituting any node x_j into the Lagrange formula yields f_j .

4.1.2 Interpolation versus Best Approximation

It is important to distinguish interpolation from best approximation. The *Weierstrass Approximation Theorem*, proved by Karl Weierstrass in 1885 [130] when he was 70 years old, guarantees that any continuous function on $[a, b]$ can be uniformly approximated to arbitrary accuracy by polynomials: if f is continuous on $[-1, 1]$ and $\varepsilon > 0$ is arbitrary, then there exists a polynomial p such that $\|f - p\|_\infty < \varepsilon$.

The theorem was independently discovered at about the same time by Carl Runge, the same mathematician whose name we attached to the phenomenon of divergent equispaced interpolation. This coincidence is not accidental: Runge’s deep investigations into polynomial approximation led him to both the positive existence result and the negative convergence phenomenon.

Weierstrass's original proof is remarkably elegant, connecting approximation theory to the heat equation. The idea is to extend $f(x)$ to a function $\tilde{f}$ with compact support on the real line, then convolve it with the Gaussian kernel

$$\varphi(x) = \frac{e^{-x^2/4t}}{\sqrt{4\pi t}}.$$

This convolution solves the diffusion equation $\partial u / \partial t = \partial^2 u / \partial x^2$ and converges uniformly to f as $t \rightarrow 0$. Since the convolution is an entire function (analytic throughout the complex plane), it has a uniformly convergent Taylor series that can be truncated to yield polynomial approximations of arbitrary accuracy.

The theorem stimulated an extraordinary amount of mathematics in the early twentieth century, with many alternative proofs discovered in rapid succession: Picard [100], Lebesgue [76] (in his first paper at age 23), Fejér [42] (at age 20), Landau [75], Jackson [67], and Bernstein [15], among others.

Bernstein Polynomials

The most constructive proof was given by Sergei Bernstein in 1912 [15]. For $f \in C([0, 1])$, the *Bernstein polynomial* of degree n is defined by

$$B_n(x) = \sum_{k=0}^n f(k/n) \binom{n}{k} x^k (1-x)^{n-k}.$$

Bernstein proved that $B_n(x) \rightarrow f(x)$ uniformly as $n \rightarrow \infty$. The convergence can be understood through a beautiful probabilistic interpretation: $B_n(x)$ represents the expected value of f evaluated at the proportion of heads in n independent coin flips, where each coin has probability x of landing heads. As $n \rightarrow \infty$, the law of large numbers ensures this proportion concentrates near x , and the expected value converges to $f(x)$.

Bernstein polynomials provide explicit approximations that converge for *any* continuous function, though the convergence is typically slow: $O(1/\sqrt{n})$ for Lipschitz continuous functions. This is far inferior to the geometric convergence achieved by Chebyshev interpolation for analytic functions.

The Limits of Interpolation

Despite the Weierstrass theorem's guarantee that approximating polynomials exist, a fundamental negative result constrains how we can find them. The *Faber–Bernstein theorem* [41] asserts that there is no fixed array of interpolation grids with 1, 2, 3, ... points that achieves convergence as $n \rightarrow \infty$ for *all* continuous functions.

This result might seem discouraging, but it has led to an unfortunate overemphasis on pathological functions in the numerical analysis literature. The practical reality is far more favorable: for functions with even modest smoothness, Chebyshev interpolation converges beautifully. The Weierstrass theorem applies to highly irregular functions like $\sin(1/x) \sin(1/\sin(1/x))$, which oscillates infinitely often near infinitely many points. Such functions are mathematical curiosities, not the smooth solutions to differential equations that arise in scientific computing.

Interpolation constructs a specific polynomial by enforcing exact agreement at the nodes. The hope is that as $N \rightarrow \infty$, the interpolating polynomials p_N converge uniformly to f . As we shall see, this hope is fulfilled for some node distributions but dramatically fails for others. The key insight of this chapter is that for smooth functions and well-chosen nodes, interpolation not only succeeds but achieves the optimal approximation rate up to a small constant factor.

4.1.3 Equispaced Nodes

The most natural choice of interpolation nodes is equally spaced points:

$$x_j = -1 + \frac{2j}{N}, \quad j = 0, 1, \dots, N.$$

These nodes divide the interval $[-1, 1]$ into N equal subintervals. For low-degree polynomials and well-behaved functions, equispaced interpolation works reasonably well. However, a fundamental problem emerges as N increases.

4.2 The Runge Phenomenon

4.2.1 A Smooth but Troublesome Function

In 1901, Carl Runge [103] studied the interpolation of a deceptively simple function:

$$f(x) = \frac{1}{1 + 25x^2}.$$

This *Runge function* is infinitely differentiable on the entire real line. Its graph is a smooth bell curve centered at the origin with maximum value $f(0) = 1$ and asymptotic decay to zero as $|x| \rightarrow \infty$.

Runge discovered that polynomial interpolation on equispaced nodes *diverges* for this function. Rather than improving with increasing N , the interpolating polynomials develop wild oscillations near the endpoints $x = \pm 1$.

4.2.2 Numerical Demonstration

Figure 10 illustrates this phenomenon. For $N = 6$, the interpolant reasonably approximates the function. But for $N = 10$ and $N = 14$, the approximation deteriorates dramatically at the boundaries, with oscillations that grow unboundedly as N increases.




Figure 10: The Runge phenomenon: polynomial interpolation of $f(x) = 1/(1 + 25x^2)$ on equispaced nodes. As the polynomial degree increases, oscillations near the boundaries $x = \pm 1$ grow unboundedly. The exact function (solid) and interpolants (dashed) are shown for $N = 6, 10$, and 14 .

The following Python code computes the Lagrange interpolant for a given set of nodes:

```
1 def lagrange_interpolate(x_nodes, f_nodes, x_eval):
2     """Evaluate the Lagrange interpolating polynomial."""
3     n = len(x_nodes)
4     p_eval = np.zeros_like(x_eval)
5
6     for k in range(n):
7         L_k = np.ones_like(x_eval)
8         for j in range(n):
9             if j != k:
10                 L_k *= (x_eval - x_nodes[j]) / (x_nodes[k] - x_nodes[j])
11         p_eval += f_nodes[k] * L_k
12
13     return p_eval
```

Python

The equivalent MATLAB implementation:

```
1 function p = lagrange_interp(x_nodes, f_nodes, x_eval)
2     n = length(x_nodes);
3     p = zeros(size(x_eval));
4
5     for k = 1:n
6         L_k = ones(size(x_eval));
7         for j = 1:n
```

Matlab

```

8         if j ~= k
9             L_k = L_k .* (x_eval - x_nodes(j)) / (x_nodes(k) - x_nodes(j));
10        end
11    end
12    p = p + f_nodes(k) * L_k;
13 end
14 end

```

The Julia implementation:

```

1  function lagrange_interpolate(x_nodes, f_nodes, x_eval)
2      n = length(x_nodes)
3      p_eval = zeros(length(x_eval))
4
5      for k in 1:n
6          L_k = ones(length(x_eval))
7          for j in 1:n
8              if j != k
9                  L_k .*= (x_eval .- x_nodes[j]) ./ (x_nodes[k] - x_nodes[j])
10             end
11         end
12         p_eval .+= f_nodes[k] .* L_k
13     end
14
15     return p_eval
16 end

```

 Julia

The code generating Figure 10 is available in:

- codes/python/ch04/runge_phenomenon.py
- codes/matlab/ch04/runge_phenomenon.m
- codes/julia/ch04/runge_phenomenon.jl

4.2.3 Why Does This Happen?

The Runge phenomenon seems paradoxical: how can adding more information (more interpolation points) make the approximation worse? The answer lies in the *Lebesgue constant*, which we will explore in Section 4.5. But first, let us develop a deeper understanding through potential theory.

4.3 Theoretical Explanation: Potential Theory

4.3.1 Singularities in the Complex Plane

The key to understanding the Runge phenomenon lies in extending our view from the real line to the complex plane. The Runge function has singularities (poles) where its denominator vanishes:

$$1 + 25z^2 = 0 \Rightarrow z = \pm \frac{i}{5} = \pm 0.2i.$$

These poles lie on the imaginary axis, at a distance of only 0.2 from the real interval $[-1, 1]$. This proximity is responsible for the divergence of equispaced interpolation.

4.3.2 The Potential Function

The convergence of polynomial interpolation is governed by a *potential function* associated with the node distribution [104]. For a distribution with density $\mu(x)$ on $[-1, 1]$, the potential at a point z in the complex plane is:

$$\varphi(z) = - \int_{-1}^1 \mu(x) \ln|z - x| dx.$$

The *equipotential curves* $\varphi(z) = \text{const}$ form closed contours around the interval $[-1, 1]$. The largest equipotential curve that does not enclose any singularity of f determines the region of convergence.

4.3.3 Uniform versus Chebyshev Density

For equispaced nodes, the density is asymptotically uniform: $\mu(x) = 1/2$. The resulting equipotential curves are roughly elliptical, but they extend only a short distance from the interval before reaching the Runge function's poles at $\pm 0.2i$.

For Chebyshev nodes (which we introduce in the next section), the density is:

$$\mu(x) = \frac{1}{\pi\sqrt{1-x^2}}.$$

This density diverges at the endpoints, concentrating nodes near $x = \pm 1$. The corresponding potential simplifies to:

$$\varphi(z) = \ln|z + \sqrt{z^2 - 1}| - \ln 2 = \ln \rho - \ln 2,$$

where $\rho = |z + \sqrt{z^2 - 1}|$. The equipotential curves are *Bernstein ellipses* with foci at ± 1 .

Figure 11 compares the equipotential curves for both distributions, showing why Chebyshev interpolation succeeds where equispaced interpolation fails.




Figure 11: Equipotential curves in the complex plane. Left: uniform node density (equispaced nodes). Right: Chebyshev density, where equipotentials are Bernstein ellipses. The poles of the Runge function at $z = \pm 0.2i$ are marked with crosses. For Chebyshev interpolation, the critical ellipse passing through the poles determines the convergence rate.

The code generating Figure 11 is available in:

- codes/python/ch04/equipotential_curves.py
- codes/matlab/ch04/equipotential_curves.m
- codes/julia/ch04/equipotential_curves.jl

4.4 The Solution: Chebyshev Points

4.4.1 Definition and Geometric Construction

The *Chebyshev-Gauss-Lobatto points* (often simply called Chebyshev points) are defined by:

$$x_j = \cos\left(\frac{j\pi}{N}\right), \quad j = 0, 1, \dots, N.$$

These points have a beautiful geometric interpretation: they are the projections onto the x -axis of $N + 1$ equally spaced points on the upper semicircle of the unit circle. Figure 12 illustrates this construction.




Figure 12: Geometric construction of Chebyshev points. Points equally spaced on the upper semicircle (by angle $\theta_j = j\pi/N$) project vertically onto the Chebyshev nodes on the x -axis. Equal spacing on the circle maps to clustering near the endpoints $x = \pm 1$.

The code generating Figure 12 is available in:

- codes/python/ch04/chebyshev_points_circle.py
- codes/matlab/ch04/chebyshev_points_circle.m
- codes/julia/ch04/chebyshev_points_circle.jl

4.4.2 Node Density and Clustering

The geometric construction reveals why Chebyshev points cluster near the endpoints. Equal angular spacing on the circle corresponds to denser horizontal spacing near $x = \pm 1$. The node density is approximately:

$$\rho(x) \approx \frac{N}{\pi\sqrt{1-x^2}},$$

which diverges as $x \rightarrow \pm 1$. This clustering counteracts the growth of interpolation error at the boundaries.

4.4.3 Chebyshev Interpolation of the Runge Function

Figure 13 demonstrates the dramatic improvement when using Chebyshev points instead of equispaced nodes for the Runge function. Where equispaced interpolation diverges, Chebyshev interpolation converges rapidly and uniformly.




Figure 13: Chebyshev interpolation of the Runge function. Left: the function and interpolants for $N = 6, 10$, and 14 . Right: interpolation error on a logarithmic scale. Unlike equispaced interpolation, the error decreases rapidly with increasing N .

The Chebyshev nodes in Python and MATLAB:

```
1 def chebyshev_nodes(N):
2     """Chebyshev-Gauss-Lobatto points on [-1, 1]."""
3     j = np.arange(N + 1)
4     return np.cos(j * np.pi / N)
```

Python

```
1 % Chebyshev-Gauss-Lobatto points
2 j = 0:N;
3 x_cheb = cos(j * pi / N);
```

Matlab

The Julia implementation:

```
1 chebyshev_nodes(N) = [cos(j * pi / N) for j in 0:N]
```

Julia

Note that this formula produces nodes ordered from right to left: $x_0 = \cos(0) = +1$ down to $x_N = \cos(\pi) = -1$. This ordering is natural for the cosine function and is the standard convention in spectral methods.

The code generating Figure 13 is available in:

- codes/python/ch04/chebyshev_success.py
- codes/matlab/ch04/chebyshev_success.m
- codes/julia/ch04/chebyshev_success.jl

4.5 Lagrange Basis Functions and Lebesgue Constants

4.5.1 The Interpolation Operator

To understand why node choice matters so dramatically, we formalize the interpolation process. For any continuous function $u \in C([-1, 1])$, there exists a unique interpolating polynomial $\mathcal{P}_N(x) \in \mathbb{P}_N[\mathbb{R}]$ of degree N satisfying:

$$\mathcal{P}_N(x_j) = u_j \equiv u(x_j), \quad j = 0, 1, \dots, N.$$

We denote this interpolating polynomial as $\mathcal{J}_N[u]$, emphasizing its role as a *linear operator* acting on continuous functions.

The *best approximating polynomial* $\mathcal{P}^* \in \mathbb{P}_N[\mathbb{R}]$ is defined by:

$$\|u - \mathcal{P}^*\|_\infty = \inf_{\mathcal{P} \in \mathbb{P}_N[\mathbb{R}]} \|u - \mathcal{P}\|_\infty.$$

While it is not obliged that $\mathcal{P}^* \equiv \mathcal{J}_N[u]$, since $\mathcal{P}^* \in \mathbb{P}_N[\mathbb{R}]$, we necessarily have $\mathcal{P}^* \equiv \mathcal{J}_N[u]$ when interpolating a polynomial of degree at most N .

4.5.2 The Lebesgue Function

Examining the Lagrange basis functions more closely, we define the *Lebesgue function* as the sum of absolute values of all basis polynomials:

$$\Lambda_N(x) = \sum_{k=0}^N |L_k(x)|.$$

The *Lebesgue constant* is its maximum over the interval:

$$\Lambda_N = \max_{x \in [-1, 1]} \Lambda_N(x).$$

4.5.3 Error Bounds via Operator Norm

The Lebesgue constant bounds the interpolation error through a beautiful argument. For any continuous u with best approximation $\mathcal{P}^*$:

$$\begin{aligned} \|u - \mathcal{J}_N[u]\|_\infty &= \|u - \mathcal{P}^* + \mathcal{P}^* - \mathcal{J}_N[u]\|_\infty \\ &\equiv \|u - \mathcal{P}^* + \mathcal{J}_N[\mathcal{P}^*] - \mathcal{J}_N[u]\|_\infty \\ &\leq \|u - \mathcal{P}^*\|_\infty + \|\mathcal{J}_N[\mathcal{P}^*] - \mathcal{J}_N[u]\|_\infty \\ &\leq \|u - \mathcal{P}^*\|_\infty + \|\mathcal{J}_N\| \cdot \|\mathcal{P}^* - u\|_\infty \\ &\leq (1 + \|\mathcal{J}_N\|) \cdot \|u - \mathcal{P}^*\|_\infty. \end{aligned}$$

The norm of the interpolation operator $\mathcal{J}_N[\cdot]$ is defined as:

$$\|\mathcal{J}_N\| := \sup_{\|u\|_\infty=1} \|\mathcal{J}_N[u]\|_\infty.$$

This operator norm is precisely the Lebesgue constant Λ_N for the node set $\{x_j\}_{j=0}^N$.

Thus, if p_N^* is the best polynomial approximation of degree N to f , and p_N is the Lagrange interpolant, then:

$$\|f - p_N\|_\infty \leq (1 + \Lambda_N) \|f - p_N^*\|_\infty = (1 + \Lambda_N) E_N(f),$$

where $E_N(f)$ denotes the best approximation error.

This bound reveals the problem with equispaced nodes: even if $E_N(f)$ decreases with N (as guaranteed by the Weierstrass theorem), the factor $(1 + \Lambda_N)$ may grow so fast that the product increases.

4.5.4 Asymptotic Growth Rates

The growth rate of Λ_N depends critically on the node distribution:

- **Chebyshev nodes:** $\Lambda_N^{\text{Ch}} = \frac{2}{\pi} \ln N + O(1)$ (logarithmic growth) [22]
- **Legendre nodes:** $\Lambda_N^{\text{Leg}} = O(\sqrt{N})$ (slow algebraic growth)
- **Equispaced nodes:** $\Lambda_N^{\text{eq}} \approx \frac{2^{N+1}}{eN \ln N}$ (exponential growth!) [124]

The exponential growth of the Lebesgue constant for equispaced nodes explains the Runge phenomenon: even though the best approximation error decreases geometrically for smooth functions, the exponentially growing factor $(1 + \Lambda_N)$ eventually dominates.

Remark. As shown by Vértési [126], the Lebesgue constant for Chebyshev nodes is remarkably close to the smallest possible Lebesgue constant for *any* node distribution:

$$\Lambda_N^{\text{Ch}} = \frac{2}{\pi} \left(\ln N + \gamma + \ln \frac{8}{\pi} \right) + o(1),$$

$$\Lambda_N^{\text{min}} = \frac{2}{\pi} \left(\ln N + \gamma + \ln \frac{4}{\pi} \right) + o(1),$$

where $\gamma \approx 0.5772156649\dots$ is the Euler–Mascheroni constant. The difference between these two expressions is only $\frac{2}{\pi} \ln 2 \approx 0.44$, demonstrating that Chebyshev nodes are essentially optimal for polynomial interpolation.

4.5.5 Derivation of the Lebesgue Constant for Equi-Spaced Interpolation

We now derive the asymptotic formula for the Lebesgue constant of equi-spaced nodes. Let $L_k(x)$ denote the fundamental Lagrange polynomials and define the equidistant Lebesgue constant by

$$1 + \Lambda_N^{\text{eq}} = \max_{x \in [-1, 1]} \sum_{k=0}^N |L_k(x)|.$$

The interpretation is the usual one: if the interpolation error in the max norm satisfies $\|f - p_N\|_{\infty} \leq \varepsilon$, then the data values fed into the interpolation formula may differ from the exact polynomial values $p_N(x_j)$ by at most ε at each node x_j , and the factor $1 + \Lambda_N^{\text{eq}}$ is the largest possible amplification of these perturbations by Lagrange interpolation.

The size of the interval does not affect the value of the Lebesgue constant, so for convenience we work with equi-spaced nodes on $[0, N]$ instead of $[-1, 1]$:

$$x_j = j, \quad j = 0, 1, \dots, N.$$

For these nodes the k -th fundamental Lagrange polynomial can be written as

$$L_k(x) = \prod_{j=0, j \neq k}^N \frac{x - j}{k - j} = \frac{x(x-1)\cdots(x-k+1)(x-k-1)\cdots(x-N)}{k(k-1)\cdots 1 \cdot (-1)\cdots(k-N)}.$$

Hence

$$1 + \Lambda_N^{\text{eq}} = \max_{x \in [0, 1]} \sum_{k=0}^N \left| \frac{x(x-1)\cdots(x-k+1)(x-k-1)\cdots(x-N)}{k(k-1)\cdots 1 \cdot (-1)\cdots(k-N)} \right|.$$

We now multiply numerator and denominator of each term by $x - k$. This yields

$$x(x-1)\cdots(x-k+1)(x-k)(x-k-1)\cdots(x-N) = x(1-x)(2-x)\cdots(N-x),$$

so that

$$|L_k(x)| = \frac{x(1-x)(2-x)\cdots(N-x)}{|k-x|k!(N-k)!}.$$

Consequently,

$$1 + \Lambda_N^{\text{eq}} = \max_{x \in [0, 1]} \sum_{k=0}^N \frac{x(1-x)(2-x)\cdots(N-x)}{|k-x|k!(N-k)!}.$$

Using the Gamma function and repeated application of $\Gamma(z+1) = z\Gamma(z)$, we obtain

$$(N-x)(N-1-x)\cdots(1-x) = \frac{\Gamma(N+1-x)}{\Gamma(1-x)},$$

and therefore

$$x(1-x)(2-x)\cdots(N-x) = x \frac{\Gamma(N+1-x)}{\Gamma(1-x)}.$$

Thus the exact expression becomes

$$1 + \Lambda_N^{\text{eq}} = \max_{x \in [0,1]} x \frac{\Gamma(N+1-x)}{\Gamma(1-x)} \sum_{k=0}^N \frac{1}{|k-x|k!(N-k)!}.$$

Up to this point everything is exact. To estimate Λ_N^{eq} for large N , we introduce several standard approximations.

We first use the ratio asymptotics for the Gamma function:

$$\frac{\Gamma(N+1-x)}{\Gamma(N+1)} \sim N^{-x} \quad (N \rightarrow \infty),$$

which gives

$$\Gamma(N+1-x) \sim N!N^{-x}.$$

Next observe that, as N grows, the binomial weights $1/(k!(N-k)!)$ are sharply peaked near $k = N/2$. In this region $|k-x| \approx N/2$ (the optimal x will turn out to be small), so we approximate

$$\frac{1}{|k-x|} \approx \frac{2}{N}$$

and factor this out of the sum:

$$\sum_{k=0}^N \frac{1}{|k-x|k!(N-k)!} \approx \frac{2}{N} \sum_{k=0}^N \frac{1}{k!(N-k)!}.$$

Using the binomial identity

$$\sum_{k=0}^N \frac{N!}{k!(N-k)!} = 2^N,$$

we obtain

$$\sum_{k=0}^N \frac{1}{k!(N-k)!} = \frac{1}{N!} \sum_{k=0}^N \binom{N}{k} = \frac{2^N}{N!},$$

and hence

$$1 + \Lambda_N^{\text{eq}} \approx \max_{x \in [0,1]} x \frac{N!N^{-x}}{\Gamma(1-x)} \cdot \frac{2}{N} \cdot \frac{2^N}{N!}.$$

After cancelling $N!$ we arrive at

$$1 + \Lambda_N^{\text{eq}} \approx \max_{x \in [0,1]} \frac{2^{N+1}}{N} \frac{xN^{-x}}{\Gamma(1-x)}.$$

To understand the dominant dependence on N we analyse the factor xN^{-x} for large N . Consider

$$\varphi(x) = xN^{-x} \quad \text{for } x > 0.$$

Then

$$\ln \varphi(x) = \ln x - x \ln N, \quad \frac{d}{dx} \ln \varphi(x) = \frac{1}{x} - \ln N.$$

The maximum occurs where this derivative vanishes:

$$\frac{1}{x_*} - \ln N = 0 \quad \Rightarrow \quad x_* = \frac{1}{\ln N}.$$

Evaluating φ at x_* gives

$$\varphi(x_*) = x_* N^{-x_*} = \frac{1}{\ln N} \exp(-x_* \ln N) = \frac{1}{\ln N} e^{-1} = \frac{1}{e \ln N}.$$

For large N this maximiser lies in $(0, 1)$, so it is admissible. Moreover $\Gamma(1 - x) \rightarrow 1$ as $x \rightarrow 0$, hence $\Gamma(1 - x_*) \approx 1$,

and the prefactor 1 in $1 + \Lambda_N^{\text{eq}}$ is negligible compared with the rapidly growing right-hand side. Thus

$$\Lambda_N^{\text{eq}} \approx \frac{2^{N+1}}{N} \cdot \frac{1}{e \ln N} = \frac{2^{N+1}}{e N \ln N} = O\left(\frac{2^N}{N \ln N}\right).$$

A slightly sharper estimate follows from the expansion

$$\Gamma(1 - x) = 1 + \gamma x + O(x^2),$$

where γ is the Euler–Mascheroni constant. Substituting $x_* = 1/\ln N$ yields

$$\Gamma(1 - x_*) = 1 + \frac{\gamma}{\ln N} + O\left(\frac{1}{(\ln N)^2}\right),$$

so at leading order in $1/\ln N$ we obtain

$$\Lambda_N^{\text{eq}} \approx \frac{2^{N+1}}{e N (\gamma + \ln N)}.$$

This recovers the classical asymptotic growth of the Lebesgue constant for equi-spaced interpolation nodes: it increases essentially like $2^N/(N \ln N)$ as $N \rightarrow \infty$.

4.5.6 Lebesgue Points and Multi-Dimensional Considerations

The nodes that *minimize* the Lebesgue constant $\Lambda_N \rightarrow \min$ for a given polynomial space $\mathbb{P}_N[\mathbb{R}]$ are called *Lebesgue points*. Finding these optimal nodes is a challenging optimization problem that has been solved only for moderate values of N in one dimension.

In multiple dimensions, the situation becomes considerably more complex. The Lebesgue constant depends not only on the node distribution $\{x_j\}_{j=0}^N$ but also on the *shape of the domain* $\mathcal{U}$. The estimation of the Lebesgue constant in multi-dimensional non-Cartesian domains is a problem that remains essentially open today. The most precious information would be to find the node distribution over a triangle that minimizes the Lebesgue constant—knowledge that would be crucial for the design of new spectral elements [63].

On *quadrilateral* domains, the best known point distributions are tensor products of Gauss–Lobatto¹ or Chebyshev nodes. These tensor product constructions inherit the favorable properties of their one-dimensional counterparts.

On *triangular* domains, the situation is more subtle. The *Fekete points*² currently represent the best known choice. Fekete points are defined as the nodes that maximize the determinant of the Vandermonde matrix—equivalently, they maximize the volume of the interpolation simplex in function space. While not provably optimal in the Lebesgue sense, they exhibit excellent numerical properties and remain the standard choice for triangular spectral elements.

4.5.7 Visualization of Basis Functions

Figure 14 compares the Lagrange basis polynomials for equispaced and Chebyshev nodes. For equispaced nodes, the basis functions near the endpoints develop large oscillations. For Chebyshev nodes, all basis functions remain well-behaved.

¹Rehuel Lobatto (1797–1866) was a Dutch mathematician born into a Portuguese family.

²Michael Fekete (1886–1957) was a Hungarian mathematician who completed his PhD under the supervision of Lipót Fejér. Fekete also provided private tutoring to János Neumann, known today as John von Neumann.




Figure 14: Lagrange basis polynomials $L_{k(x)}$ for $N = 10$. Left: equispaced nodes show large oscillations near the boundaries. Right: Chebyshev nodes yield bounded basis functions throughout the interval.

Figure 15 shows the Lebesgue functions and the growth of Lebesgue constants for the three node distributions.




Figure 15: Left: Lebesgue functions $\Lambda_N(x)$ for $N = 10$ with equispaced, Legendre, and Chebyshev nodes. The equispaced case has large peaks near the boundaries. Right: Growth of Lebesgue constants with N . The equispaced constant grows exponentially, while Chebyshev grows only logarithmically.

The exponential growth of the equispaced Lebesgue constant dominates Figure 15, making it difficult to distinguish the Legendre and Chebyshev curves. Figure 16 provides a zoomed comparison of these two well-behaved distributions, clearly showing the $O(\sqrt{N})$ growth of Legendre versus the superior $O(\ln N)$ growth of Chebyshev.




Figure 16: Lebesgue constants for Legendre and Chebyshev nodes (equispaced omitted for clarity). Chebyshev nodes achieve the slowest possible growth rate $O(\ln N)$, while Legendre nodes grow as $O(\sqrt{N})$. Both are far superior to equispaced nodes, but Chebyshev remains the optimal choice.

The code generating these figures is available in:

- codes/python/ch04/lagrange_basis.py
- codes/python/ch04/lebesgue_functions.py
- codes/python/ch04/lebesgue_constants_zoom.py
- codes/matlab/ch04/lagrange_basis.m
- codes/matlab/ch04/lebesgue_functions.m
- codes/julia/ch04/lagrange_basis.jl
- codes/julia/ch04/lebesgue_functions.jl
- codes/julia/ch04/lebesgue_constants_zoom.jl

4.6 Barycentric Interpolation

4.6.1 Numerical Stability Issues

The Lagrange formula, while mathematically elegant, is numerically problematic. Computing each basis polynomial $L_k(x)$ requires N multiplications and divisions, and the resulting values can be very large or very small, leading to catastrophic cancellation when summed.

4.6.2 The Barycentric Formula

A more stable and efficient approach is the *barycentric interpolation formula* [16], whose numerical stability was rigorously established by [64]:

$$p_N(x) = \frac{\sum_{j=0}^N \frac{w_j}{x-x_j} f_j}{\sum_{j=0}^N \frac{w_j}{x-x_j}},$$

where the *barycentric weights* are:

$$w_j = \frac{1}{\prod_{k \neq j} (x_j - x_k)}.$$

This formula requires only $O(N)$ operations to evaluate $p_{N(x)}$ once the weights are precomputed.

4.6.3 Weights for Chebyshev Points

For Chebyshev-Gauss-Lobatto points, the weights have a particularly simple form:

$$w_j = (-1)^j \delta_j, \quad \text{where } \delta_j = \begin{cases} 1/2 & \text{if } j = 0 \text{ or } j = N \\ 1 & \text{otherwise.} \end{cases}$$

This simplicity, combined with the excellent approximation properties, makes Chebyshev points the standard choice for spectral methods.

4.7 Convergence Analysis

4.7.1 Geometric Convergence

For functions analytic in a region containing $[-1, 1]$, Chebyshev interpolation converges geometrically. If f is analytic inside the Bernstein ellipse $\mathcal{E}_\rho$ with parameter $\rho > 1$, then:

$$\|f - p_N\|_\infty \leq \frac{C}{\rho^N},$$

where C depends on f and ρ but not on N .

The parameter ρ is determined by the location of the nearest singularity: if the closest singularity to $[-1, 1]$ lies on the ellipse $\mathcal{E}_\rho$, then convergence is at rate $O(\rho^{-N})$.

4.7.2 The Runge Function Revisited

For the Runge function with poles at $z = \pm 0.2i$, we can compute:

$$\rho = |0.2i + \sqrt{(0.2i)^2 - 1}| = |0.2i + i\sqrt{1.04}| \approx 1.22.$$

Thus Chebyshev interpolation converges at rate $O(1.22^{-N})$; convergence is geometric but modest due to the poles being close to the real axis. In contrast, equispaced interpolation diverges because its effective $\rho < 1$.

The divergence of equispaced interpolation can be understood through the error bound $\|f - p_N\|_\infty \leq (1 + \Lambda_N) E_N(f)$. For the Runge function, the best approximation error $E_N(f)$ still decreases geometrically—the function is analytic, after all. However, recall from Section 4.5 that the equispaced Lebesgue constant grows as $\Lambda_N^{\text{eq}} \approx 2^{N+1}/(eN \ln N)$. This exponential growth eventually overwhelms the geometric decay of $E_N(f)$, causing the interpolation error $(1 + \Lambda_N) E_N(f)$ to grow without bound.

4.7.3 Numerical Verification

Figure 17 compares the convergence behavior for both node distributions. The Chebyshev error decreases geometrically, while the equispaced error grows without bound.




Figure 17: Convergence comparison for the Runge function. The maximum interpolation error is plotted against polynomial degree N on a semilogarithmic scale. Chebyshev interpolation (green) converges at rate $O(\rho^{-N})$ with $\rho \approx 1.22$. Equispaced interpolation (red) diverges exponentially.

Figure 18 shows the Chebyshev convergence alone, without the divergent equispaced curve that dominates the vertical scale. This reveals the beautiful geometric convergence: the error decreases by a factor of approximately $\rho \approx 1.22$ with each increase in polynomial degree. The theoretical rate matches the computed errors almost perfectly.




Figure 18: Chebyshev interpolation convergence for the Runge function (equispaced omitted for clarity). The error decreases geometrically at rate $O(\rho^{-N})$ with $\rho \approx 1.22$, matching the theoretical prediction based on the pole locations at $z = \pm 0.2i$.

The code generating these figures is available in:

- codes/python/ch04/convergence_comparison.py
- codes/python/ch04/convergence_zoom.py
- codes/matlab/ch04/convergence_comparison.m
- codes/julia/ch04/convergence_comparison.jl
- codes/julia/ch04/convergence_zoom.jl

4.8 Computational Étude: Random Nodes

Having studied the optimal Chebyshev distribution and the problematic equispaced distribution, a natural question arises: what happens if we choose interpolation nodes *randomly*? This question leads us into the realm of *experimental mathematics*, where we use computation to discover and conjecture mathematical relationships.

4.8.1 Motivation

The dramatically different behaviors of equispaced and Chebyshev nodes might suggest that the key factor is *regularity* or *structure* in node placement. Perhaps any “reasonable” arrangement would work? To test this hypothesis, we investigate the simplest unstructured choice: nodes drawn uniformly at random from $[-1, 1]$.

This étude serves two purposes. First, it tests whether the special clustering of Chebyshev points near the endpoints is truly essential, or whether it is merely one of many acceptable distributions. Second, it demonstrates the methodology of computational investigation, using numerical evidence to formulate conjectures about asymptotic behavior.

4.8.2 Experimental Setup

For each polynomial degree N from 2 to 30, we generate $N + 1$ random points uniformly distributed on $[-1, 1]$, sort them, and compute the Lebesgue constant. We repeat this process $M = 200$ times to obtain statistical estimates of the mean, standard deviation, and range of Λ_N for random nodes.

The Python implementation uses NumPy's random number generation:

```

1  def random_nodes(N, rng=None):
2      """Generate N+1 random nodes, sorted, on [-1, 1]."""
3      if rng is None:
4          rng = np.random.default_rng()
5      return np.sort(rng.uniform(-1, 1, N + 1))
6
7  def monte_carlo_lebesgue(N, M=200):
8      """Compute M samples of Lebesgue constant for random nodes."""
9      samples = np.zeros(M)
10     for m in range(M):
11         x_rand = random_nodes(N)
12         samples[m] = lebesgue_constant(x_rand)
13     return samples

```

Python

The equivalent MATLAB code:

```

1  % Generate M samples for polynomial degree N
2  samples = zeros(M, 1);
3  for m = 1:M
4      x_rand = sort(2 * rand(N+1, 1) - 1); % Uniform on [-1, 1]
5      samples(m) = max(lebesgue_function(x_rand, x_fine));
6  end

```

Matlab

The Julia implementation:

```

1  random_nodes(N, rng) = sort(2.0 .* rand(rng, N + 1) .- 1.0)
2
3  function monte_carlo_lebesgue(N, M, rng; n_eval=2000)
4      samples = zeros(M)
5      for m in 1:M
6          x_rand = random_nodes(N, rng)
7          samples[m] = lebesgue_constant(x_rand, n_eval=n_eval)
8      end
9      return samples
10 end

```

Julia

4.8.3 Results

Figure 19 shows the results of our Monte Carlo experiment. The left panel displays the growth of the Lebesgue constant with polynomial degree, including shaded regions showing the statistical variability. The right panel shows the distribution of Λ_N values for a fixed degree $N = 15$.




Figure 19: Lebesgue constant for random nodes. Left: growth with polynomial degree N , showing mean (solid line) and min-max range (light shading). Right: distribution of $\log_{10}(\Lambda_{15})$ over 500 random trials on a logarithmic scale, revealing the enormous spread spanning many orders of magnitude. Vertical lines mark the mean, median, equispaced, and Chebyshev values.

The key observations are striking:

1. **Explosive exponential growth:** Random nodes exhibit exponential growth of the Lebesgue constant, but surprisingly *faster* than equispaced nodes. The mean Λ_N grows roughly as e^{bN} with $b \approx 1.0$, compared to $b \approx 0.6$ for equispaced.
2. **Extreme variability:** There is enormous variability between different random samples. For $N = 15$, Λ_N can range from around 10^2 to 10^{10} or more, spanning many orders of magnitude. This variability increases dramatically with N .
3. **Worse than equispaced:** Counter to naive intuition, random nodes perform *worse* on average than the structured equispaced distribution. The reason is that random sampling occasionally produces clusters of nearby nodes, creating near-singular interpolation conditions. These worst-case realizations dominate the mean.

4.8.4 Empirical Asymptotic Formula

Fitting an exponential model to the mean Lebesgue constants yields an empirical formula:

$$\Lambda_N^{\text{rand}} \approx a \cdot e^{bN},$$

where the fitted constants are approximately $a \approx 0.7$ and $b \approx 1.0$, giving:

$$\Lambda_N^{\text{rand}} \approx 0.7 \cdot 2.7^N.$$

Remarkably, the growth rate for random nodes is *faster* than for equispaced nodes (which grow as $\approx 1.8^N$). This counterintuitive result arises because random sampling occasionally produces clusters of nodes that are extremely close together, causing the Lebesgue constant to explode. The mean is dominated by these worst-case realizations.

The fundamental problem is twofold: (1) the lack of *clustering near the endpoints*, which only carefully designed distributions like Chebyshev points possess, and (2) the risk of *random clustering* in the interior, which creates severely ill-conditioned interpolation problems.

4.8.5 Discussion

This computational étude provides strong numerical evidence for a fundamental principle: **the clustering of Chebyshev points near the endpoints is not merely convenient but essential.**

Random nodes, despite their apparent “fairness” in covering the interval, fail in two ways: they do not provide the boundary resolution needed to control Lagrange basis oscillations, and they risk interior clustering that creates catastrophically ill-conditioned systems. These findings are consistent with the theoretical analysis of [111], who established lower bounds for the expected Lebesgue constant of random point sets.

The enormous variability in Λ_N for random nodes is perhaps the most striking finding. While equispaced nodes are suboptimal, they at least provide *predictable* (if exponentially growing) behavior. Random nodes introduce an additional layer of uncertainty—any given random realization might be acceptable or catastrophic.

This étude illustrates the power of computational mathematics. By systematic numerical investigation, we have:

- Tested a natural hypothesis (random nodes might work)
- Discovered unexpected behavior (random is *worse* than equispaced)
- Quantified the asymptotic growth through data fitting
- Reinforced our understanding of why structured distributions matter

The code generating Figure 19 is available in:

- `codes/python/ch04/lebesgue_random_nodes.py`
- `codes/matlab/ch04/lebesgue_random_nodes.m`
- `codes/julia/ch04/lebesgue_random_nodes.jl`

4.9 Computational Étude: Random Angles on the Circle

The previous étude revealed that uniform random nodes on $[-1, 1]$ perform poorly because they lack the endpoint clustering of Chebyshev points. This raises a natural follow-up question: what if we generate random points that *do* cluster near the endpoints? Specifically, what happens if we generate random angles uniformly on the semicircle and project them via cosine, mimicking the construction of Chebyshev nodes?

4.9.1 The Arcsine Distribution

Recall that Chebyshev–Gauss–Lobatto nodes are defined as $x_j = \cos(j\frac{\pi}{N})$ for $j = 0, 1, \dots, N$. The angles $\theta_j = j\frac{\pi}{N}$ are *equispaced* on $[0, \pi]$, and the cosine projection creates the characteristic endpoint clustering.

Our experiment generates *random* angles uniformly on $[0, \pi]$ and applies the same cosine projection:

$$\theta_k \sim \text{Uniform}(0, \pi), \quad x_k = \cos(\theta_k).$$

If θ is uniformly distributed on $[0, \pi]$, then $x = \cos(\theta)$ has the *arcsine distribution* with probability density function:

$$f(x) = \frac{1}{\pi\sqrt{1-x^2}}, \quad x \in [-1, 1].$$

This is precisely the Chebyshev weight function! The arcsine distribution naturally concentrates probability mass near ± 1 , providing the endpoint clustering we seek.

4.9.2 Implementation

The implementation is straightforward. In Python:

```
1 def random_chebyshev_nodes(N, rng=None):
2     """Generate N+1 nodes by projecting random angles via cosine."""
3     if rng is None:
4         rng = np.random.default_rng()
5     theta = rng.uniform(0, np.pi, N + 1)
6     return np.sort(np.cos(theta))
```

Python

The equivalent MATLAB code:

```
1 theta = pi * rand(N+1, 1); % Uniform on [0, pi]
2 x_rand_cheb = sort(cos(theta)); % Project via cosine and sort
```

Matlab

The Julia implementation:

```
1 function random_chebyshev_nodes(N, rng)
2     theta = π .* rand(rng, N + 1)
3     return sort(cos.(theta))
4 end
```

Julia

We conduct the same Monte Carlo experiment as before: for each N from 2 to 30, we generate 200 random realizations and compute the Lebesgue constant for each.

4.9.3 Surprising Results

Figure 20 displays the results, which are profoundly counterintuitive.




Figure 20: Lebesgue constant for “random Chebyshev” nodes (uniform angles projected via cosine). Left: growth with polynomial degree, showing explosive exponential behavior with enormous variability. Right: distribution of $\log_{10}(\Lambda_{15})$ over 500 trials. Despite having the “correct” arcsine marginal distribution, these nodes perform catastrophically worse than any other distribution studied.

The key observations are striking and instructive:

1. **Catastrophic exponential growth:** Random Chebyshev nodes exhibit the *worst* performance of any distribution we have studied. The mean Lebesgue constant grows explosively, reaching astronomical values (10^{20} or higher) by $N = 30$.

2. **Extreme variability:** The spread across realizations is staggering, spanning 10 to 20 orders of magnitude. Some realizations produce reasonable Lebesgue constants (the minimum values grow polynomially), while others are catastrophically large.
3. **Worse than uniform random:** Counter to our hypothesis, having the “correct” marginal distribution (arcsine) actually makes things *worse*, not better.

4.9.4 Understanding the Failure

Why does this happen? The answer lies in the distinction between *marginal distributions* and *joint distributions*.

When we generate $N + 1$ random angles uniformly on $[0, \pi]$, there is no guarantee of minimum separation. Two angles θ_i and θ_j can be arbitrarily close. When this happens, the corresponding nodes $x_i = \cos(\theta_i)$ and $x_j = \cos(\theta_j)$ become nearly identical, creating a *near-singular* interpolation problem.

In contrast, deterministic Chebyshev nodes have angles that are exactly $\frac{\pi}{N}$ apart. This *guaranteed minimum separation* ensures that no two nodes can cluster together.

The mathematical insight is profound: **the success of Chebyshev nodes comes not from the arcsine marginal distribution, but from the deterministic structure that guarantees minimum node separation.**

4.9.5 The Lesson

This étude teaches a crucial lesson about interpolation theory:

1. **Marginal distribution is not sufficient:** Having the “right” probability distribution for individual nodes does not guarantee good collective behavior.
2. **Structure matters:** The deterministic, equispaced placement of angles in true Chebyshev nodes provides essential guarantees that random sampling cannot replicate.
3. **Minimum separation is key:** The $O(\frac{1}{N})$ minimum spacing between Chebyshev angles translates to guaranteed node separation, preventing the catastrophic clustering that plagues random approaches.

This finding reinforces a fundamental principle: in numerical analysis, *structured* methods often outperform *random* methods precisely because structure provides guarantees that randomness cannot.

The code generating Figure 20 is available in:

- `codes/python/ch04/lebesgue_random_chebyshev.py`
- `codes/matlab/ch04/lebesgue_random_chebyshev.m`
- `codes/julia/ch04/lebesgue_random_chebyshev.jl`

4.10 Practical Guidelines and Outlook

4.10.1 When to Use Which Grid

Based on the theory developed in this chapter, we can state clear guidelines:

1. **Always prefer Chebyshev points** for polynomial interpolation on a finite interval. The $O(\ln N)$ growth of the Lebesgue constant ensures stability.

2. **Avoid equispaced nodes** for high-degree polynomial interpolation. The exponential growth of Λ_N makes this a losing proposition.
3. **Consider the singularity structure** of the function being approximated. The location of complex singularities determines the convergence rate.

4.10.2 Mapping to General Intervals

Chebyshev points are defined on $[-1, 1]$, but can be mapped to any interval $[a, b]$ via the linear transformation:

$$x_{\text{phys}} = \frac{a+b}{2} + \frac{b-a}{2}x_{\text{ref}},$$

where $x_{\text{ref}} \in [-1, 1]$ is the reference coordinate.

4.10.3 Preview of Differentiation Matrices

The Chebyshev points introduced in this chapter will play a central role in the differentiation matrices we develop in subsequent chapters. The clustering of nodes near the boundaries, far from being a peculiarity, is precisely what enables accurate spectral differentiation. The connection between node distribution, interpolation accuracy, and differentiation stability is one of the beautiful unifying themes of spectral methods.

In the next chapter, we will see how to convert function values at Chebyshev points into accurate approximations of derivatives, building on the geometric insights developed here.

4.11 A non-exhaustive literature overview

The theory of polynomial interpolation and node geometry presented in this chapter draws from a rich intellectual tradition spanning two centuries, from foundational existence theorems to modern algorithmic refinements. This overview traces the principal threads.

The theoretical bedrock. The genesis of polynomial approximation theory lies in the work of Karl Weierstrass [130], who established that continuous functions on a closed interval can be uniformly approximated by polynomials to arbitrary precision. This existence result was subsequently given constructive proofs by several mathematicians in rapid succession: Picard [100] using integral operators, Lebesgue [76] through geometric methods, Fejér [42] via Cesàro summation of Fourier series, Landau [75] through convolution with rational kernels, Jackson [67] using trigonometric approximation, and most elegantly by Bernstein [15] through the probabilistic construction that now bears his name. However, the negative result of Faber [41] demonstrated that no single triangular array of interpolation nodes can guarantee convergence for *all* continuous functions, shifting the focus from universal convergence to convergence for smooth functions on well-chosen grids.

The Runge phenomenon and its quantification. The pathology of equispaced interpolation was discovered by Runge [103], who showed that the seemingly innocuous function $(1 + 25x^2)^{-1}$ defies polynomial interpolation on uniform grids. The mechanism of failure was quantified through the Lebesgue constant: Turetskii [124] derived the asymptotic formula $\Lambda_N^{\text{eq}} \sim 2^{N+1}/(eN \ln N)$ for equispaced nodes, establishing that error amplification grows exponentially. In contrast, Brutman [22] proved that the Lebesgue constant for Chebyshev nodes grows only logarithmically, and Vértési [126] showed that this logarithmic rate is asymptotically optimal among all possible node distributions. Detailed numerical experiments by Luttmann and Rivlin [81] corroborated these theoretical predic-

tions and revealed fine structural properties of the Lebesgue function. A comprehensive modern survey of Lebesgue constants across different node distributions is provided by Smith [111].

The potential-theoretic framework. The deep explanation for why certain node distributions succeed lies in logarithmic potential theory, developed comprehensively by Saff and Totik [104]. The key insight is that the convergence of polynomial interpolation is governed by equipotential curves in the complex plane: the equilibrium measure $\mu_{\text{eq}}(x) = 1/(\pi\sqrt{1-x^2})$ generates Bernstein ellipses as its level curves, and any node distribution that asymptotically follows this density (such as Chebyshev or Legendre nodes) will guarantee convergence for functions analytic inside the corresponding ellipse. This framework unifies the disparate observations about node geometry into a single elegant theory.

Algorithmic realization. The translation of these theoretical insights into stable numerical algorithms required further innovation. Berrut and Trefethen [16] revived and popularized the barycentric form of Lagrange interpolation, which reduces evaluation cost to $O(N)$ operations while avoiding the catastrophic cancellation inherent in the classical Lagrange formula. Higham [64] provided a rigorous forward error analysis proving that the barycentric formula is numerically stable whenever the Lebesgue constant is moderate, thereby confirming that the theoretical stability of Chebyshev nodes translates faithfully into finite-precision arithmetic. These algorithmic developments underpin the practical implementations used throughout this book; see [120] and [21] for comprehensive treatments.

Multi-dimensional frontiers. While one-dimensional interpolation theory is well established, the extension to higher dimensions remains an active area of research. On triangular domains, the Fekete points (which maximize the Vandermonde determinant) provide the natural generalization of Chebyshev nodes and are now standard in high-order spectral element methods [63]. On quadrilateral domains, a breakthrough was achieved by Bos, Caliari, De Marchi, Vianello, and Xu [17, 18], who discovered the *Padua points*: the first explicitly known set of unisolvent nodes on the square $[-1, 1]^2$ with a Lebesgue constant growing only as $O(\ln^2 N)$. These points are generated by the self-intersections of a Lissajous curve and offer a practical alternative to tensor-product grids for bivariate interpolation.

4.12 Summary

This chapter has established that the choice of interpolation nodes is decisive for spectral accuracy:

1. **The Runge phenomenon:** High-degree polynomial interpolation on equispaced nodes diverges near the boundaries, with errors growing exponentially as N increases.
2. **Potential theory explanation:** The convergence rate of polynomial interpolation is governed by the location of the function's singularities in the complex plane and by the potential-theoretic properties of the node distribution.
3. **Chebyshev points:** Nodes distributed as $x_j = \cos(j\pi/N)$ cluster near the boundaries and yield a Lebesgue constant that grows only as $O(\ln N)$, ensuring stable, spectrally accurate interpolation.
4. **Lebesgue constant:** This quantity measures the worst-case amplification of interpolation error. Its slow logarithmic growth for Chebyshev nodes, versus exponential growth for equispaced nodes, is the geometric foundation of spectral methods.
5. **Barycentric interpolation:** The barycentric formula provides a numerically stable $O(N)$ algorithm for evaluating the interpolant at any point, avoiding the explicit construction of the Lagrange polynomials.
6. **Structure over randomness:** Random nodes with the correct marginal distribution perform catastrophically worse than deterministic Chebyshev nodes. The guaranteed minimum separation of deterministic grids, not the density profile alone, is what prevents ill-conditioning.

These geometric insights motivate the differentiation matrices developed in the next chapter, where function values at Chebyshev points are converted into accurate spectral approximations of derivatives.



CHAPTER 5

Differentiation Matrices

In the previous chapter, we mastered the art of polynomial interpolation — constructing polynomials that pass exactly through a set of data points. We discovered that the choice of nodes determines whether interpolation succeeds or fails, with Chebyshev points emerging as the optimal choice for non-periodic problems. Now we take the next logical step: having represented a function as an interpolating polynomial, how do we *differentiate* it? The answer leads us to one of the most elegant structures in numerical analysis: the differentiation matrix.

The remarkable insight of pseudospectral methods is that differentiation can be accomplished by a single matrix-vector multiplication. Given function values $\mathbf{u} = (u_0, u_1, \dots, u_N)^\top$ at the grid points, we can approximate the derivative values $\mathbf{u}' = (u'_0, u'_1, \dots, u'_N)^\top$ as

$$\mathbf{u}' \approx D\mathbf{u}, \quad (1)$$

where D is the *differentiation matrix*. This matrix encapsulates the entire differentiation process: interpolate, differentiate, evaluate.

This chapter develops the theory and practice of constructing differentiation matrices. We begin with familiar finite difference approximations, viewing them through the lens of matrix algebra. We then take the crucial limiting step: what happens when the stencil extends to include *all* grid points? The answer reveals that spectral methods are not a separate species from finite differences, but rather their natural culmination—the limiting case as the stencil width grows without bound.

5.1 From Interpolation to Differentiation

5.1.1 The Matrix Perspective

Let us begin with a fundamental observation. Given $N + 1$ distinct nodes $\{x_0, x_1, \dots, x_N\}$ and function values $\{u_0, u_1, \dots, u_N\}$, the unique interpolating polynomial of degree at most N can be written using the Lagrange formula:

$$p_N(x) = \sum_{j=0}^N u_j L_j(x), \quad (2)$$

where $L_j(x)$ are the Lagrange basis polynomials from Section 4.5.

To approximate the derivative of u at the nodes, we simply differentiate the interpolating polynomial:

$$p'_N(x) = \sum_{j=0}^N u_j L'_j(x). \quad (3)$$

Evaluating this at node x_i yields:

$$p'_N(x_i) = \sum_{j=0}^N u_j L'_j(x_i). \quad (4)$$

This is a linear combination of the function values! The coefficients $L'_j(x_i)$ depend only on the node locations, not on the function values. We can therefore write Equation 4 in matrix form:

$$\mathbf{u}' \approx D\mathbf{u}, \quad \text{where} \quad D_{ij} = L'_j(x_i). \quad (5)$$

The entry D_{ij} represents the weight given to the function value at x_j when approximating the derivative at x_i . The systematic construction and practical implementation of such matrices was the subject of the influential MATLAB differentiation matrix suite by Weideman and Reddy [128].

5.1.2 The Barycentric Formula for Differentiation Matrix Entries

For general (non-equispaced) nodes, the differentiation matrix entries can be computed using the barycentric weights w_j introduced in Section 4.5:

$$D_{ij} = \frac{w_j/w_i}{x_i - x_j} \quad \text{for } i \neq j. \quad (6)$$

This formula, whose remarkable stability properties were established by Berrut and Trefethen [16], provides a unified framework for computing differentiation matrix entries on any node distribution.

The diagonal entries are determined by the important *consistency condition*: the derivative of a constant function is zero. Since D applied to the constant vector $(1, 1, \dots, 1)^\top$ must yield zero, we have

$$\sum_{j=0}^N D_{ij} = 0 \quad \Rightarrow \quad D_{ii} = -\sum_{j \neq i} D_{ij}. \quad (7)$$

This row-sum property provides a convenient way to compute the diagonal entries. In fact, Baltensperger and Berrut [8] showed that computing diagonal entries via this “negative sum trick” is not merely convenient but essential for numerical stability, reducing the rounding error norm from $O(N^4)$ to $O(N^2)$. Baltensperger and Trummer [9] further quantified the importance of this technique in their paper “Spectral Differencing with a Twist,” which we revisit in Section 7.

5.2 Finite Difference Matrices

5.2.1 The Local Approach

Before tackling the full spectral differentiation matrix, let us review the familiar territory of finite differences. The classical approach approximates the derivative using only *nearby* function values—a local stencil.

The simplest example is the *second-order central difference*:

$$u'(x_i) \approx \frac{u_{i+1} - u_{i-1}}{2h}, \quad (8)$$

where h is the grid spacing. This formula is “second-order accurate” because the error is $O(h^2)$ for smooth functions, as can be verified by Taylor expansion.

For higher accuracy, we can include more neighbors. The *fourth-order central difference* uses five points:

$$u'(x_i) \approx \frac{-u_{i+2} + 8u_{i+1} - 8u_{i-1} + u_{i-2}}{12h}. \quad (9)$$

The *sixth-order central difference* extends to seven points:

$$u'(x_i) \approx \frac{u_{i+3} - 9u_{i+2} + 45u_{i+1} - 45u_{i-1} + 9u_{i-2} - u_{i-3}}{60h}. \quad (10)$$

Figure 21 illustrates these stencils schematically. Each formula uses only the function values at nodes within its stencil. The derivative at x_i depends only on nearby neighbors, not on distant points.




Figure 21: Finite difference stencils in one dimension, progressing from local to global. The derivative at the central node x_i (red) is approximated using only the highlighted nodes (blue) within each stencil. From top to bottom: 3-point stencil (2nd order), 5-point stencil (4th order), 7-point stencil (6th order), and spectral method (all nodes). As the stencil widens, accuracy improves. The spectral method represents the limiting case where every node contributes to the derivative approximation, yielding exponential rather than algebraic convergence.

The code generating Figure 21 is available in:

- codes/python/ch05/fd_stencil_schematic.py
- codes/matlab/ch05/fd_stencil_schematic.m
- codes/julia/ch05/fd_stencil_schematic.jl

5.2.2 Matrix View of Finite Differences

These formulas can be assembled into differentiation matrices. For the second-order scheme Equation 8, assuming periodic boundary conditions on N equispaced points with spacing $h = 2\pi/N$, the matrix is *tridiagonal* (plus corner entries for periodicity):

$$D^{(2)} = \frac{1}{2h} \begin{pmatrix} 0 & 1 & 0 & \cdots & 0 & -1 \\ -1 & 0 & 1 & \cdots & 0 & 0 \\ 0 & -1 & 0 & \cdots & 0 & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 1 & 0 & 0 & \cdots & -1 & 0 \end{pmatrix}. \quad (11)$$

The fourth-order scheme Equation 9 produces a *pentadiagonal* matrix with bandwidth 5, and the sixth-order scheme yields a *heptadiagonal* matrix with bandwidth 7.

Figure 22 illustrates this progression. As the order of accuracy increases, the stencil widens and the matrix bandwidth grows.

Finite Difference Matrix Sparsity: From Local to Global




Figure 22: Sparsity patterns of differentiation matrices for $N = 20$ grid points. Left: second-order finite differences (tridiagonal, bandwidth 3). Center: fourth-order finite differences (pentadiagonal, bandwidth 5). Right: spectral method (dense, bandwidth N). The progression illustrates the key insight: spectral methods are the limiting case of finite differences as the stencil width extends to the full domain.

The code generating Figure 22 is available in:

- codes/python/ch05/fd_matrix_bandwidth.py
- codes/matlab/ch05/fd_matrix_bandwidth.m
- codes/julia/ch05/fd_matrix_bandwidth.jl

5.2.3 The Limit Question

This observation leads to a fundamental question: *if widening the stencil improves accuracy, what is the ultimate limit?*

The answer is a stencil that uses *all* N grid points. The resulting matrix is *dense*—every point influences the derivative at every other point. This is precisely the spectral differentiation matrix.

From this perspective, spectral methods are not a separate species from finite differences. They are the natural endpoint of the progression: as we increase the order of accuracy by including more neighbors, we eventually include all points. The sparse banded matrix becomes dense, and algebraic convergence transforms into spectral (exponential) convergence.

This unifying viewpoint, emphasized by Fornberg in his influential textbook [47] and first demonstrated in his 1987 *Geophysics* paper [45], provides both conceptual clarity and practical guidance. It suggests that finite differences and spectral methods lie on a continuum, with the choice of stencil width being a tunable parameter balancing accuracy against computational cost.

5.3 The Periodic Spectral Differentiation Matrix

5.3.1 Periodic Problems and Equispaced Nodes

A remarkable simplification occurs for *periodic* problems. On a domain like $[0, 2\pi)$ with periodic boundary conditions, the natural interpolation basis consists of trigonometric polynomials (sines and cosines) rather than algebraic polynomials.

For periodic problems, the optimal node distribution is *equispaced*:

$$x_j = \frac{2\pi j}{N}, \quad j = 0, 1, \dots, N-1. \quad (12)$$

This is in stark contrast to the non-periodic case studied in Section 4.4, where equispaced nodes lead to the Runge phenomenon. The difference lies in the boundary: periodic functions have no endpoints, so there is no need for the clustering that Chebyshev nodes provide. The periodicity itself regularizes the problem.

5.3.2 Derivation via the Fourier Basis

The spectral differentiation matrix for periodic problems can be derived by considering the trigonometric interpolant and differentiating it analytically. For N equispaced points (with N even), the interpolating trigonometric polynomial is

$$p(x) = \sum_{j=0}^{N-1} u_j \varphi_j(x), \quad (13)$$

where $\varphi_j(x)$ is the *periodic cardinal function* (or discrete Dirichlet kernel) satisfying $\varphi_j(x_k) = \delta_{jk}$.

The periodic cardinal function can be written as a sum of complex exponentials:

$$\varphi_j(x) = \frac{1}{N} \sum_{k=-N/2+1}^{N/2} e^{ik(x-x_j)}. \quad (14)$$

This sum can be evaluated in closed form. Writing $\theta = (x - x_j)/2$ and using the geometric series, we obtain:

$$\varphi_j(x) = \frac{\sin(N\theta)}{N \sin(\theta)} = \frac{\sin(N(x - x_j)/2)}{N \sin((x - x_j)/2)}. \quad (15)$$

Figure 23 visualizes these periodic cardinal functions for $N = 16$ equispaced nodes. Each function peaks at value 1 at its corresponding node x_j and vanishes at all other nodes, satisfying the cardinal property $\varphi_j(x_k) = \delta_{jk}$. Unlike Lagrange basis polynomials on non-periodic domains, which can exhibit unbounded oscillations (recall Section 4.5), these periodic cardinal functions remain bounded. The damped oscillations away from the peak reflect the $\sin(x)/x$ -like structure of the discrete Dirichlet kernel.




Figure 23: Periodic cardinal functions $\varphi_j(x)$ for $N = 16$ equispaced nodes on $[0, 2\pi)$. Each function peaks at value 1 at its corresponding node x_j and vanishes at all other nodes, satisfying the cardinal property $\varphi_j(x_k) = \delta_{jk}$. The oscillations decay smoothly away from each peak, in contrast to the boundary-amplified oscillations seen in Lagrange basis polynomials for non-periodic interpolation.

The following Python code computes the periodic cardinal function Equation 15:

```

1  import numpy as np
2
3  def periodic_cardinal(x, x_j, N):
4      """
5      Compute the periodic cardinal function phi_j(x) centered at x_j.
6
7      Parameters:
8          x : array - Points at which to evaluate
9          x_j : float - Center point (node location)
10         N : int - Number of grid points
11
12     Returns:
13         phi : array - Values of phi_j at x
14     """
15     theta = (x - x_j) / 2.0
16     phi = np.zeros_like(x, dtype=float)
17
18     # Handle singularity at theta = 0 using L'Hopital's rule
19     small = np.abs(np.sin(theta)) < 1e-14
20     phi[~small] = np.sin(N * theta[~small]) / (N * np.sin(theta[~small]))
21     phi[small] = 1.0 # Limit as theta -> 0
22
23     return phi

```

Python

The equivalent MATLAB implementation:

```

1  function phi = periodic_cardinal(x, x_j, N)

```

Matlab

```

2      % Compute the periodic cardinal function phi_j(x) centered at x_j.
3      % Input:
4      %   x   - points at which to evaluate
5      %   x_j - center point (node location)
6      %   N   - number of grid points
7      % Output:
8      %   phi - values of phi_j at x
9
10     theta = (x - x_j) / 2.0;
11     phi = zeros(size(x));
12
13     % Handle singularity at theta = 0
14     small = abs(sin(theta)) < 1e-14;
15     phi(~small) = sin(N * theta(~small)) ./ (N * sin(theta(~small)));
16     phi(small) = 1.0; % Limit as theta -> 0
17 end

```

The Julia implementation:

```

1  function periodic_cardinal(x, x_j, N)
2      theta = (x .- x_j) ./ 2.0
3      phi = zeros(length(x))
4
5      # Handle singularity at theta = 0 using L'Hopital's rule
6      small = abs.(sin.(theta)) .< 1e-14
7      phi[.!small] = sin.(N .* theta[.!small]) ./ (N .* sin.(theta[.!small]))
8      phi[small] .= 1.0 # Limit as theta -> 0
9
10     return phi
11 end

```

The code implementing these functions and generating Figure 23 is available in:

- codes/python/ch05/periodic_cardinal_functions.py
- codes/matlab/ch05/periodic_cardinal_functions.m
- codes/julia/ch05/periodic_cardinal_functions.jl

To find the differentiation matrix, we need to compute $\varphi'_j(x_m)$ for all m . Let $\xi = (x - x_j)/2$ for brevity. Then Equation 15 becomes $\varphi_j = \sin(N\xi)/(N \sin \xi)$. Applying the quotient rule:

$$\frac{d\varphi_j}{d\xi} = \frac{N \cos(N\xi) \sin \xi - \sin(N\xi) \cos \xi}{N \sin^2 \xi}. \quad (16)$$

Since $d\xi/dx = 1/2$, we obtain:

$$\varphi'_j(x) = \frac{1}{2} \cdot \frac{\cos(N\xi)}{\sin \xi} - \frac{1}{2} \cdot \frac{\sin(N\xi) \cos \xi}{N \sin^2 \xi}. \quad (17)$$

We evaluate this at the grid points $x = x_m$.

Case 1: Diagonal entries ($m = j$).

When $x = x_j$, we have $\xi = 0$, so both numerator and denominator of Equation 17 vanish. Applying L'Hôpital's rule (twice) yields $\varphi'_j(x_j) = 0$. Geometrically, the cardinal function is symmetric about its peak at x_j , so its derivative must vanish there.

Case 2: Off-diagonal entries ($m \neq j$).

When $x = x_m$ with $m \neq j$, we have $\xi = \pi(m - j)/N$. At these points:

- $\sin(N\xi) = \sin(\pi(m - j)) = 0$ (the second term in the numerator vanishes),
- $\cos(N\xi) = \cos(\pi(m - j)) = (-1)^{m-j}$.

Substituting into Equation 17:

$$\varphi'_j(x_m) = \frac{1}{2} \frac{(-1)^{m-j}}{\tan(\pi(m - j)/N)} = \frac{(-1)^{m-j}}{2} \cot \frac{(m - j)\pi}{N}. \quad (18)$$

Since $D_{mk} = \varphi'_k(x_m)$, the spectral differentiation matrix entries are:

$$D_{jk} = \begin{cases} \frac{1}{2}(-1)^{j-k} \cot \left(\frac{(j - k)\pi}{N} \right) & \text{if } j \neq k \\ 0 & \text{if } j = k. \end{cases} \quad (19)$$

5.3.3 Properties of the Spectral Matrix

The matrix defined by Equation 19 has several remarkable properties:

1. **Skew-symmetry:** $D^\top = -D$. This reflects the fact that differentiation is an antisymmetric operator in appropriate function spaces. Recent work by Iserles [66] has highlighted the importance of this property for structure-preserving integration, where skew-symmetric differentiation matrices ensure conservation of physical invariants such as energy.
2. **Zero diagonal:** $D_{jj} = 0$. The derivative approximation at any point depends only on neighboring values, not on the value at the point itself.
3. **Toeplitz structure:** D_{jk} depends only on the difference $j - k$. This means the matrix has constant diagonals—a consequence of the translation-invariance of differentiation on periodic domains.
4. **Circulant:** Due to the periodic boundary conditions, the matrix is actually *circulant*: each row is a cyclic shift of the previous row. Circulant matrices can be diagonalized by the discrete Fourier transform [55], enabling $O(N \log N)$ matrix-vector products via the FFT [30].
5. **Dense:** Unlike finite difference matrices, every off-diagonal entry is nonzero. This is the price we pay for spectral accuracy.
6. **Exactness for trigonometric polynomials:** If $u(x)$ is a trigonometric polynomial of degree at most $N/2$, then Du gives the *exact* derivative values.

Figure 24 visualizes these properties for $N = 16$.




Figure 24: Structure of the periodic spectral differentiation matrix for $N = 16$. Left: heatmap showing the matrix entries. The Toeplitz (constant diagonal) structure is evident, as is the skew-symmetry ($D^T = -D$) with positive values (red) appearing where negative values (blue) appear in the transpose. Right: the first row entries $D_{0,k}$ (blue dots) plotted against the column offset $k = j - i$. The dashed red curve shows the continuous cotangent function $\frac{1}{2} \cdot (-1)^k \cot(k\pi/N)$ from Equation 19, confirming that the discrete matrix entries lie exactly on this curve. The antisymmetric pattern ($D_{0,k} = -D_{0,-k}$) reflects the skew-symmetry of D .

The code generating Figure 24 is available in:

- codes/python/ch05/spectral_matrix_structure.py
- codes/matlab/ch05/spectral_matrix_structure.m
- codes/julia/ch05/spectral_matrix_structure.jl

5.3.4 Python Implementation

The following Python code constructs the periodic spectral differentiation matrix:

```

1  import numpy as np
2
3  def spectral_diff_periodic(N):
4      """
5      Construct the periodic spectral differentiation matrix.
6
7      Parameters:
8          N : int - Number of grid points (should be even)
9
10     Returns:
11         D : ndarray (N, N) - Differentiation matrix
12         x : ndarray (N,) - Grid points on  $[0, 2\pi)$ 
13     """
14     h = 2 * np.pi / N
15     x = h * np.arange(N)
16     D = np.zeros((N, N))
17

```

Python

```

18     for i in range(N):
19         for j in range(N):
20             if i != j:
21                 D[i, j] = 0.5 * ((-1) ** (i - j)) / np.tan((i - j) * np.pi / N)
22
23     return D, x

```

The equivalent MATLAB implementation:

```

1  function [D, x] = spectral_diff_periodic(N)
2      h = 2 * pi / N;
3      x = h * (0:N-1)';
4      D = zeros(N, N);
5
6      for i = 1:N
7          for j = 1:N
8              if i ~= j
9                  D(i, j) = 0.5 * ((-1)^(i-j)) / tan((i-j) * pi / N);
10             end
11         end
12     end
13 end

```

 Matlab

The Julia implementation:

```

1  function spectral_diff_periodic(N)
2      h = 2π / N
3      x = h .* (0:N-1)
4      D = zeros(N, N)
5
6      for i in 1:N
7          for j in 1:N
8              if i != j
9                  D[i, j] = 0.5 * ((-1)^(i - j)) / tan((i - j) * π / N)
10             end
11         end
12     end
13
14     return D, collect(x)
15 end

```

 Julia

The code implementing these algorithms is available in:

- codes/python/ch05/spectral_matrix_periodic.py
- codes/matlab/ch05/spectral_matrix_periodic.m
- codes/julia/ch05/spectral_matrix_structure.jl

5.3.5 A Practical Demonstration

Let us put our spectral differentiation matrix to work. Consider the smooth periodic function

$$u(x) = e^{\sin^2 x}, \quad (20)$$

which is *not* a trigonometric polynomial—its Fourier series has infinitely many terms. The derivatives are:

$$\begin{aligned} u'(x) &= \sin(2x)e^{\sin^2 x}, \\ u''(x) &= [\sin^2(2x) + 2\cos(2x)]e^{\sin^2 x}. \end{aligned} \quad (21)$$

With $N = 64$ grid points, the spectral method computes both derivatives to near machine precision! Figure 25 shows the results: the numerical values (markers) lie exactly on the exact curves (solid lines), with maximum errors around 10^{-14} .




Figure 25: Spectral differentiation of $u(x) = e^{\sin^2 x}$ using $N = 64$ grid points. Left: the function u (navy) and its first derivative u' exact (coral) vs. spectral (teal circles). Right: second derivative u'' exact (purple) vs. spectral (teal squares). The maximum errors, displayed in each panel, are near machine precision ($\approx 10^{-14}$). This remarkable accuracy with relatively few points is the hallmark of spectral methods for smooth periodic functions.

Table 3 quantifies the convergence rate. The errors decrease dramatically—by roughly four orders of magnitude each time N doubles. This is the signature of *spectral convergence*: for analytic functions, the error decreases faster than any polynomial in $1/N$.

N	Error in u'	Error in u''
8	6.96×10^{-2}	2.00×10^0
16	4.33×10^{-4}	3.67×10^{-2}
32	1.12×10^{-9}	3.26×10^{-7}
64	2.84×10^{-14}	4.99×10^{-13}

Table 3: Convergence of spectral differentiation for $u(x) = e^{\sin^2 x}$. The maximum errors $\|u' - Du\|_\infty$ and $\|u'' - D^2u\|_\infty$ decrease exponentially as N increases, reaching machine precision by $N = 64$.

The code for this demonstration is available in:

- codes/python/ch05/spectral_derivatives_demo.py
- codes/matlab/ch05/spectral_derivatives_demo.m
- codes/julia/ch05/spectral_derivatives_demo.jl

5.4 Fornberg's Recursive Algorithm

5.4.1 Motivation: Robustness for Arbitrary Grids

The explicit formula Equation 19 is elegant but “brittle”—it applies only to equispaced periodic grids. What if we need differentiation weights for non-equispaced nodes? Or for non-periodic problems? Or for higher-order derivatives?

The direct approach would be to solve the Vandermonde system that arises from polynomial interpolation. However, Vandermonde matrices are notoriously ill-conditioned, especially for large N . We need a more robust algorithm.

In 1988, Bengt Fornberg published a remarkable recursive algorithm [46] that computes finite difference weights for *any* node distribution and *any* derivative order. The algorithm is:

- Numerically stable even for large stencils
- Efficient: $O(MN^2)$ for all derivatives up to order M on N nodes
- Elegant: the recursion has a beautiful mathematical structure

5.4.2 The Recursive Algorithm

The key insight is that we can build up the weights incrementally. Let $\delta_m^{(k,n)}$ denote the weight for node x_k when approximating the m -th derivative using nodes $x_0, x_1, \dots, x_n$ (a stencil of $n + 1$ points), evaluated at some point ξ .

The algorithm proceeds level by level:

- **Level 0** ($n = 0$): A single node. For interpolation ($m = 0$), the weight is simply 1.
- **Level 1** ($n = 1$): Two nodes. Weights are computed from level 0.
- **Level n** : Weights for $n + 1$ nodes are computed from level $n - 1$.

The recursion formulas are:

For $k < n$ (weights for nodes already in the previous stencil):

$$\delta_m^{(k,n)} = \frac{(\xi - x_n)\delta_m^{(k,n-1)} - m\delta_{m-1}^{(k,n-1)}}{x_n - x_k}. \quad (22)$$

For $k = n$ (weight for the newly added node):

$$\delta_m^{(n,n)} = \frac{c_n}{c_{n-1}} \left(m\delta_{m-1}^{(n-1,n-1)} - (\xi - x_{n-1})\delta_m^{(n-1,n-1)} \right), \quad (23)$$


where $c_n = \prod_{\ell=0}^{n-1} (x_n - x_\ell)$.

The base case is $\delta_0^{(0,0)} = 1$, and we define $\delta_m^{(k,n)} = 0$ when $m < 0$ or $m > n$.

5.4.3 The Stencil Pyramid

Figure 26 visualizes the recursive structure. Each level of the pyramid contains weights for one stencil size. Arrows show the data dependencies: weights at level n depend only on weights from level $n - 1$.

Fornberg's Recursive Algorithm: The Stencil Pyramid



Each box represents a weight $\delta_i^{(k,m)}$ for node k in a stencil of $m+1$ nodes. Arrows show data dependencies: weights at level m depend on weights from level $m-1$.

Figure 26: Fornberg's recursive algorithm visualized as a "stencil pyramid." Each box represents a weight $\delta_m^{(k,n)}$ for node index k in a stencil of $n+1$ nodes. The algorithm builds from top (single node) to bottom (full stencil). Green arrows indicate updates to existing node weights via Equation 22; red arrows show computation of new node weights via Equation 23. The recursive structure ensures numerical stability.

The code generating Figure 26 is available in:

- codes/python/ch05/stencil_pyramid.py
- codes/matlab/ch05/stencil_pyramid.m
- codes/julia/ch05/stencil_pyramid.jl

The pyramid structure explains why the algorithm is numerically stable. Each weight is computed as a simple combination of weights from the previous level, avoiding the accumulation of rounding errors that plagues direct methods.

5.4.4 Python Implementation

The following code implements Fornberg's algorithm, translated from the MATLAB version by Dr. Toby Driscoll [34]:

```
1 import numpy as np
2
3 def fdweights(xi, x, m):
4     """
5     Compute finite difference weights using Fornberg's algorithm.
6
7     Parameters:
8         xi : float - Point where derivative is approximated
9         x  : array - Node positions
10        m   : int - Derivative order (0=interpolation, 1=first, etc.)
```

Python

```

11
12     Returns:
13         w : array - Weights for the m-th derivative approximation
14         """
15     x = np.asarray(x, dtype=float)
16     n = len(x) - 1
17     w = np.zeros(n + 1)
18     x_shifted = x - xi # Translate evaluation point to origin
19
20     for k in range(n + 1):
21         w[k] = _weight(x_shifted, m, n, k)
22
23     return w
24
25 def _weight(x, m, j, k):
26     """Recursive weight computation (evaluation point at 0)."""
27     if m < 0 or m > j:
28         return 0.0
29     elif m == 0 and j == 0:
30         return 1.0
31     else:
32         if k < j:
33             c = (x[j] * _weight(x, m, j-1, k)
34                  - m * _weight(x, m-1, j-1, k)) / (x[j] - x[k])
35         else:
36             beta = np.prod(x[j-1] - x[:j-1]) / np.prod(x[j] - x[:j])
37             c = beta * (m * _weight(x, m-1, j-1, j-1)
38                        - x[j-1] * _weight(x, m, j-1, j-1))
39     return c

```

The equivalent MATLAB implementation, due to Dr. Toby Driscoll [34]:

```

1  function w = fdweights(xi, x, m)
2      % Compute finite difference weights using Fornberg's algorithm.
3      % Input:
4      %   xi - evaluation point for the derivative
5      %   x  - node positions (vector)
6      %   m  - derivative order (0=interpolation, 1=first, etc.)
7      % Output:
8      %   w  - weights for the m-th derivative approximation
9
10     p = length(x) - 1;
11     w = zeros(size(x));
12     x = x - xi; % Translate evaluation point to origin
13
14     for k = 0:p
15         w(k+1) = weight(x, m, p, k);

```

Matlab

```

16     end
17 end
18
19 function c = weight(x, m, j, k)
20     % Recursive weight computation (evaluation point at 0).
21     if (m < 0) || (m > j)
22         c = 0;
23     elseif (m == 0) && (j == 0)
24         c = 1;
25     else
26         if k < j
27             c = (x(j+1) * weight(x, m, j-1, k) ...
28                 - m * weight(x, m-1, j-1, k)) / (x(j+1) - x(k+1));
29         else
30             beta = prod(x(j) - x(1:j-1)) / prod(x(j+1) - x(1:j));
31             c = beta * (m * weight(x, m-1, j-1, j-1) ...
32                 - x(j) * weight(x, m, j-1, j-1));
33         end
34     end
35 end

```

The Julia implementation:

```

1  function fdweights(xi, x, m)
2      n = length(x) - 1
3      x_shifted = Float64.(x) .- xi
4
5      w = zeros(n + 1)
6      for k in 0:n
7          w[k+1] = _weight(x_shifted, m, n, k)
8      end
9      return w
10 end
11
12 function _weight(x, m, j, k)
13     (m < 0 || m > j) && return 0.0
14     (m == 0 && j == 0) && return 1.0
15
16     if k < j
17         c = (x[j+1] * _weight(x, m, j-1, k) -
18             m * _weight(x, m-1, j-1, k)) / (x[j+1] - x[k+1])
19     else
20         beta = j >= 2 ? prod(x[j] - x[i+1] for i in 0:j-2) /
21             prod(x[j+1] - x[i+1] for i in 0:j-1) :
22             j == 1 ? 1.0 / (x[2] - x[1]) : 1.0
23         c = beta * (m * _weight(x, m-1, j-1, j-1) -
24             x[j] * _weight(x, m, j-1, j-1))

```

 Julia

```

25     end
26     return c
27 end

```

This algorithm is the universal tool for computing differentiation matrix entries on any grid.

The code implementing this algorithm is available in:

- codes/python/ch05/fdweights.py
- codes/matlab/ch05/fdweights.m
- codes/julia/ch05/fdweights.jl

5.5 Computational Étude: The Rational Trigonometric Test

We now conduct a computational étude that reveals the dramatic difference between finite difference and spectral accuracy. The goal is not merely to verify theoretical convergence rates, but to develop intuition for *why* spectral methods achieve such remarkable precision.

5.5.1 Choosing a Test Function

The choice of test function is crucial. We need a function that is smooth and periodic (so that both finite difference and spectral methods apply), yet not so simple that it masks the differences between methods. A pure trigonometric function like $\sin(kx)$ would be inappropriate: since $\sin(kx)$ is an eigenfunction of both the continuous and discrete differentiation operators, spectral methods would give exact results trivially, revealing nothing about their convergence behavior.

Instead, we choose the *rational trigonometric function*:

$$u(x) = \frac{1}{2 + \sin(x)} \quad \text{on } [0, 2\pi). \quad (24)$$

This function satisfies our requirements admirably. It is infinitely differentiable (analytic) on the entire real line, and clearly periodic with period 2π . Most importantly, it has an *infinite* Fourier series—unlike $\sin(x)$ or finite trigonometric polynomials, its spectral content extends to all frequencies, ensuring there are no “lucky cancellations” in our numerical differentiation.

The exact derivative, computed by the quotient rule, is:

$$u'(x) = -\frac{\cos(x)}{(2 + \sin(x))^2}. \quad (25)$$

There is a deeper reason for choosing this particular function. Although $u(x)$ is smooth on the real line, it has singularities in the *complex plane*. The denominator $2 + \sin(x)$ vanishes when $\sin(x) = -2$, which has no real solutions but does have complex solutions at $x = -\pi/2 \pm i \cdot \operatorname{arcsinh}(2)$. The distance from the real axis to these nearest singularities is $d = \operatorname{arcsinh}(2) \approx 1.44$. As we shall see, this distance controls the convergence rate of spectral methods—a beautiful connection to potential theory that we explored in Section 4.3.

5.5.2 The Experiment

Our experimental design is straightforward. For a sequence of grid sizes $N = 4, 6, 8, 10, \dots, 64$, we construct four differentiation matrices: three finite difference matrices of orders 2, 4, and 6 (using Fornberg’s algorithm from the previous section), and the periodic spectral differentiation matrix given by Equation 19.

For each matrix D , we sample the test function at the N equispaced grid points to form the vector $\mathbf{u} = (u(x_0), u(x_1), \dots, u(x_{N-1}))^\top$, compute the numerical derivative $D\mathbf{u}$, and compare it to the exact derivative values $\mathbf{u}'_{\text{exact}} = (u'(x_0), u'(x_1), \dots, u'(x_{N-1}))^\top$. The error is measured in the maximum norm:

$$\varepsilon_N = \|\mathbf{u}'_{\text{exact}} - D\mathbf{u}\|_\infty = \max_{0 \leq j \leq N-1} |u'(x_j) - (D\mathbf{u})_j|. \quad (26)$$

5.5.3 Results and Interpretation

Figure 27 displays the results on a semi-logarithmic plot. The visual contrast between finite difference and spectral methods is striking and reveals a fundamental distinction in their convergence behavior.




Figure 27: Differentiation error comparison: finite differences versus spectral method. The test function is $u(x) = 1/(2 + \sin(x))$ on the periodic domain $[0, 2\pi)$. Finite difference methods (FD2, FD4, FD6) exhibit algebraic convergence with the expected rates $O(N^{-2})$, $O(N^{-4})$, $O(N^{-6})$. The spectral method achieves geometric (exponential) convergence, reaching near machine precision around $N = 50$. Dashed lines show theoretical convergence rates.

The finite difference methods exhibit *algebraic* (polynomial) convergence. On a log-log plot, their error curves would appear as straight lines; on our semi-log plot, they curve downward with decreasing slope. The second-order method (FD2) shows error decreasing as $O(N^{-2})$, which is equivalent to $O(h^2)$ since $h = 2\pi/N$. The fourth-order method (FD4) improves to $O(N^{-4})$, and the sixth-order method (FD6) achieves $O(N^{-6})$. These rates match the theoretical predictions from Taylor series analysis: a p -th order finite difference scheme has truncation error $O(h^p)$.

The spectral method, in contrast, exhibits *geometric* (exponential) convergence. Its error curve appears as a straight line on the semi-log plot, indicating that $\log \varepsilon_N$ decreases linearly with N . Mathematically, $\varepsilon_N = O(c^{-N})$ for some constant $c > 1$. The method reaches near machine precision ($\approx 10^{-14}$) at around $N = 50$ with roughly 50 grid points, we have computed the derivative to nearly the full precision available in double-precision arithmetic!

To appreciate what this means in practice, consider trying to achieve 14-digit accuracy with finite differences. For the second-order method, we would need to solve $CN^{-2} \approx 10^{-14}$. Even with a modest constant $C \approx 1$, this requires $N \approx 10^7 \approx 10$ million grid points. For a problem in three dimensions, this would mean 10^{21} unknowns, which is utterly impractical. The spectral method achieves the same accuracy with only 50 points.

5.5.4 Discussion: Why Does Spectral Win?

The difference between algebraic and spectral convergence is fundamental:

Algebraic convergence ($O(N^{-p})$): Doubling N reduces the error by a factor of 2^p . This is good, but the improvement is polynomial.

Spectral convergence ($O(c^{-N})$): Doubling N *squares* the error (roughly). This exponential improvement is why spectral methods can achieve machine precision with modest grid sizes.

The source of spectral convergence is the *analyticity* of the function being approximated. For the test function Equation 24, the nearest singularities in the complex plane are at distance approximately $\operatorname{arcsinh}(2) \approx 1.44$ from the real axis. Potential theory (cf. Section 4.3) tells us that the convergence rate is controlled by this distance: the interpolation error decreases like ρ^{-N} where $\rho = e^d$ and d is the distance to the nearest singularity. A comprehensive treatment of the connection between convergence rates and Bernstein ellipses in the complex plane is given by Trefethen [121]. The eigenvalue analysis of Weideman and Trefethen [129] further illuminates why spectral differentiation matrices achieve such high accuracy, showing that their spectral properties closely mirror those of the continuous differential operator.

The code generating Figure 27 is available in:

- `codes/python/ch05/convergence_comparison.py`
- `codes/matlab/ch05/convergence_comparison.m`
- `codes/julia/ch05/convergence_comparison.jl`

5.6 Higher-Order Derivatives

5.6.1 Second Derivatives: Squaring vs. Direct Construction

For second derivatives, we have two options:

1. **Matrix squaring:** $D^{(2)} = D \cdot D$, where D is the first-derivative matrix
2. **Direct construction:** Build $D^{(2)}$ directly using second-derivative weights

Matrix squaring is simpler and often sufficient. The resulting matrix $(D \cdot D)_{ij}$ represents the combined effect of differentiating twice. For smooth functions, this gives accurate second derivatives.

For the periodic spectral case, the second-derivative matrix has a closed form:

$$D_{jk}^{(2)} = \begin{cases} -\frac{1}{2} \frac{(-1)^{j-k}}{\sin^2((j-k)\pi/N)} & \text{if } j \neq k \\ -\frac{\pi^2(N^2+2)}{3(2\pi)^2} & \text{if } j = k. \end{cases} \quad (27)$$

This closed form was analyzed by Weideman and Trefethen [129], who studied the eigenvalues of second-order spectral differentiation matrices and established that their spectrum scales as $O(N^4)$,

providing the theoretical justification for using matrix squaring or explicit formulas. However, matrix squaring D^2 is often accurate enough and more convenient.

5.6.2 A Demonstration: Higher-Order Derivatives

Let us put higher-order spectral differentiation to the test. Consider the smooth periodic function

$$u(x) = e^{-\sin(2x)}, \quad (28)$$

which has higher frequency content than our earlier examples. Its derivatives become increasingly complex:

$$\begin{aligned} u'(x) &= -2 \cos(2x) e^{-\sin(2x)}, \\ u''(x) &= 4[\sin(2x) + \cos^2(2x)] e^{-\sin(2x)}. \end{aligned} \quad (29)$$

Figure 28 shows spectral approximations of the first four derivatives using $N = 32$ grid points. The numerical values (markers) align precisely with the exact curves (solid lines), with errors displayed in each panel. Notice how the error grows with derivative order—a fundamental limitation of numerical differentiation.




Figure 28: Higher-order spectral derivatives of $u(x) = e^{-\sin(2x)}$ using $N = 32$ grid points. Each panel compares the exact derivative (navy curve) with the spectral approximation (colored panel markers). The maximum errors, displayed in each panel, grow with derivative order—from $\approx 10^{-7}$ for the first derivative to $\approx 10^{-2}$ for the fourth. This error amplification is intrinsic to numerical differentiation.

Table 4 quantifies the convergence behavior. For each derivative order, the spectral method achieves exponential convergence as N increases. However, the errors at any fixed N grow significantly with derivative order, reflecting the ill-conditioning of higher-order differentiation.

N	Error u'	Error u''	Error u'''	Error u''''
8	3.50×10^{-1}	6.17×10^0	9.40×10^0	1.55×10^2
16	8.64×10^{-3}	3.92×10^{-1}	6.51×10^{-1}	2.82×10^1
32	3.52×10^{-7}	5.26×10^{-5}	9.44×10^{-5}	1.39×10^{-2}
64	2.42×10^{-14}	1.18×10^{-12}	1.45×10^{-11}	6.14×10^{-10}

Table 4: Convergence of spectral differentiation for $u(x) = e^{-\sin(2x)}$ at different derivative orders. Each column shows the maximum error $\|u^{(m)} - D^m u\|_\infty$ for the m -th derivative. While spectral convergence is achieved for all orders, higher derivatives require more grid points to reach a given accuracy.

Figure 29 illustrates the matrix squaring approach for second derivatives. The left panel shows the structure of $D^2 = D \cdot D$, which inherits the Toeplitz property from D . The center panel confirms that the eigenvalues of D^2 are $-k^2$ for $k = -N/2 + 1, \dots, N/2$, matching the eigenvalues of the continuous operator d^2/dx^2 acting on Fourier modes e^{ikx} . The right panel demonstrates the accuracy of the second derivative approximation.




Figure 29: Matrix squaring for second derivatives with $N = 16$. Left: structure of $D^2 = D \cdot D$, showing the Toeplitz pattern inherited from D . Center: eigenvalues of D^2 (circles) compared to the theoretical values $-k^2$ (squares), confirming spectral accuracy. Right: second derivative of $u(x) = e^{-\sin(2x)}$ computed via $D^2 u$, showing excellent agreement with the exact solution.

The following Python code computes higher-order derivatives via matrix powers:

```

1  import numpy as np
2
3  def higher_order_derivative(D, u, order):
4      """
5      Compute the m-th derivative using matrix powers.
6
7      Parameters:

```

Python

```

8      D      : ndarray (N, N) - First-derivative matrix
9      u      : ndarray (N,) - Function values at grid points
10     order : int - Derivative order (1, 2, 3, ...)
11
12     Returns:
13         u_m   : ndarray (N,) - m-th derivative values
14     """
15     D_m = np.linalg.matrix_power(D, order)
16     return D_m @ u

```

The equivalent MATLAB implementation:

```

1  function u_m = higher_order_derivative(D, u, order)
2      % Compute the m-th derivative using matrix powers.
3      % Input:
4      %   D      - first-derivative matrix (N x N)
5      %   u      - function values at grid points (N x 1)
6      %   order - derivative order (1, 2, 3, ...)
7      % Output:
8      %   u_m   - m-th derivative values
9
10     D_m = D^order;
11     u_m = D_m * u;
12 end

```

 Matlab

The Julia implementation:

```

1  using LinearAlgebra
2
3  function higher_order_derivative(D, u, order)
4      D_m = D^order
5      return D_m * u
6  end

```

 Julia

The code generating Figure 28 and Figure 29 is available in:

- codes/python/ch05/higher_order_derivatives.py
- codes/matlab/ch05/higher_order_derivatives.m
- codes/julia/ch05/higher_order_derivatives.jl

5.6.3 Fornberg's Algorithm for Higher Derivatives

A key advantage of Fornberg's algorithm is that it handles any derivative order m with no additional complexity. The same recursive structure computes interpolation weights ($m = 0$), first-derivative weights ($m = 1$), second-derivative weights ($m = 2$), and so on.

This generality is valuable when solving PDEs that involve mixed derivatives or high-order terms.

5.7 Computational Étude: The Quantum Harmonic Oscillator

We close the periodic spectral methods portion of this chapter with an example that demonstrates spectral accuracy for a physically important eigenvalue problem. Boyd [21] provides an extensive discussion of spectral methods for problems on unbounded domains, including the use of Hermite functions and domain truncation strategies. Recent work by Ma *et al.* [82] has applied Jacobi–Galerkin spectral methods to nonlinear fractional Schrödinger equations, demonstrating that the quantum oscillator remains a relevant testbed for cutting-edge research. The *quantum harmonic oscillator* is described by the time-independent Schrödinger equation:

$$-u'' + x^2u = \lambda u, \quad x \in \mathbb{R}. \quad (30)$$

This equation arises throughout physics: in quantum mechanics (the harmonic potential), in vibration analysis (normal modes), and in probability theory (Hermite functions).

5.7.1 Exact Solution

The eigenvalues of Equation 30 are well known:

$$\lambda_n = 2n + 1, \quad n = 0, 1, 2, \dots \quad (31)$$

The corresponding eigenfunctions are the *Hermite functions*:

$$u_n(x) = H_n(x)e^{-x^2/2}, \quad (32)$$

where H_n is the n -th Hermite polynomial. These functions decay like $e^{-x^2/2}$ as $|x| \rightarrow \infty$, which is faster than any polynomial. In fact, the Hermite functions are *entire* functions (analytic throughout $\mathbb{C}$), so spectral methods should achieve super-geometric convergence.

5.7.2 Numerical Approach: Periodic Spectral Method

Although the problem is posed on the infinite line $\mathbb{R}$, the rapid decay of the eigenfunctions allows us to truncate to a finite interval $[-L, L]$ for sufficiently large L . The key observation is that the eigenfunctions, being entire and rapidly decaying, are effectively periodic on $[-L, L]$ for large L . This makes the periodic spectral method from Section 5.3 an ideal tool.

The numerical scheme proceeds as follows:

1. Set up N equispaced grid points x_j on $[-L, L]$.
2. Construct the periodic second-derivative matrix $D_N^{(2)}$ rescaled by $(\pi/L)^2$.
3. Form the operator matrix $A = -D_N^{(2)} + \text{diag}(x_1^2, \dots, x_N^2)$.
4. Compute eigenvalues of A using standard linear algebra.

Note that no boundary conditions need to be imposed explicitly: the periodicity of the method automatically handles the (effectively zero) boundary values.

The core algorithm is remarkably simple. In Python:

```
1 def harmonic_oscillator_eigenvalues(N, L):
2     """Solve -u'' + x^2u = λu on [-L, L] using periodic spectral method."""
3     h = 2 * np.pi / N
4     x = h * np.arange(N)
5     x = L * (x - np.pi) / np.pi # Map to [-L, L]
6
7     # Second derivative matrix (rescaled)
```

 Python

```

8     D2 = build_periodic_D2(N) * (np.pi / L)**2
9
10    # Potential matrix
11    V = np.diag(x**2)
12
13    # Solve eigenvalue problem
14    eigenvalues = np.linalg.eigvalsh(-D2 + V)
15    return np.sort(eigenvalues)

```

The equivalent MATLAB implementation:

```

1  function eigenvalues = harmonic_oscillator(N, L)
2      h = 2*pi/N; x = h*(0:N-1)'; x = L*(x-pi)/pi;
3
4      % Periodic second derivative matrix (rescaled)
5      column = [-pi^2/(3*h^2)-1/6, ...
6                -.5*(-1).^(1:N-1)./sin(h*(1:N-1)/2).^2];
7      D2 = (pi/L)^2 * toeplitz(column);
8
9      eigenvalues = sort(eig(-D2 + diag(x.^2)));
10 end

```

Matlab

The Julia implementation:

```

1  function harmonic_oscillator_eigenvalues(N, L)
2      # Solve -u'' + x^2u = λu on [-L, L] using periodic spectral method.
3      h = 2π / N
4      x = L * (h * collect(0:N-1) .- π) / π # Map to [-L, L]
5
6      # Second derivative matrix (rescaled Toeplitz)
7      D2 = periodic_d2_matrix(N) * (π / L)^2
8
9      # Potential matrix
10     V = Diagonal(x .^ 2)
11
12     # Solve eigenvalue problem
13     eigenvalues = eigvals(Symmetric(-D2 + V))
14     return sort(eigenvalues)
15 end

```

Julia

5.7.3 Results

Figure 30 shows the remarkable performance of the periodic spectral method. The left panel displays the first five eigenfunctions computed with $N = 64$ equispaced points on $[-8, 8]$, along with the exact Hermite functions. The agreement is visually perfect.




Figure 30: The quantum harmonic oscillator solved with the periodic spectral method. Left: First five eigenfunctions computed with $N = 64$ equispaced points on $[-8, 8]$ (dots) compared to exact Hermite functions (lines). Right: Eigenvalue error versus N for the first four eigenvalues. The error decreases faster than any exponential, reaching machine precision around $N = 36$.

The right panel shows the eigenvalue convergence. For $N = 36$ and $L = 8$, the first four eigenvalues are computed to approximately 13-digit accuracy:

n	Computed λ_n	Error
0	0.999 999 999 999 96	4×10^{-14}
1	3.000 000 000 000 03	3×10^{-14}
2	4.999 999 999 999 97	3×10^{-14}
3	6.999 999 999 999 99	1×10^{-14}

Table 5: Computed eigenvalues of the harmonic oscillator with the periodic spectral method ($N = 36$, $L = 8$). The exact values are $\lambda_n = 2n + 1$.

This is spectral accuracy in action. With just 36 equispaced grid points, we have computed eigenvalues to essentially machine precision. The periodic method is particularly well suited here because the eigenfunctions decay rapidly, making the truncated domain effectively periodic. In Section 8, we will revisit this problem using Chebyshev spectral methods and compare the two approaches.

The code generating Figure 30 is available in:

- codes/python/ch05/harmonic_oscillator.py
- codes/matlab/ch05/harmonic_oscillator.m
- codes/julia/ch05/harmonic_oscillator.jl

5.8 Looking Ahead: The Non-Periodic Case

The periodic spectral matrix Equation 19 relies crucially on the periodicity of the problem. For *non-periodic* problems on finite intervals, such as boundary value problems with Dirichlet or Neumann conditions, we need a different approach.

The solution, which we develop in the next chapter, is to use *Chebyshev differentiation matrices*. These combine the Chebyshev nodes from Section 4.4 with differentiation matrix techniques similar to those presented here. The resulting matrices are dense but not Toeplitz, reflecting the non-uniform node distribution.

The key formulas are similar in spirit to Equation 19 but more complex due to the clustering of Chebyshev nodes near the boundaries. The Chebyshev differentiation matrix will become our primary tool for spectral solutions of differential equations on bounded domains.

5.9 A non-exhaustive literature overview

The conceptual genesis of spectral differentiation matrices lies in the realization that pseudospectral methods are not a distinct species of numerical approximation but rather the natural culmination of finite difference schemes. As the stencil width of a finite difference operator increases to encompass the entire domain, the discretization matrix transitions from a sparse, banded structure to a dense, global operator. This “limit” perspective, which provides the intuitive bedrock for understanding spectral accuracy, was forcefully articulated by Fornberg [45] in his seminal *Geophysics* paper, where he demonstrated that for the elastic wave equation, the pseudospectral method — viewed as the limit of high-order finite differences — achieves nondispersive wave propagation that low-order schemes fail to capture, even at high resolutions. The theoretical underpinnings of this transition trace back to the foundational work of Kreiss and Olinger [73], who formalized the comparison between accurate difference methods and Fourier techniques for hyperbolic equations. Their work established that for smooth solutions, the phase error of spectral methods decays exponentially, whereas finite difference errors decay only algebraically. In the context of periodic domains, the differentiation matrix assumes a highly structured form. As discussed by Gottlieb and Orszag [55], the periodic differentiation matrix is circulant, a property that allows its diagonalization via the Discrete Fourier Transform. This connection bridges the gap between the matrix-vector multiplication view ($O(N^2)$) and the Fast Fourier Transform view ($O(N \log N)$). Recent work by Ibrahim and Lawrence [65] on spectral rectangular collocation highlights how modern spectral formulations still rely on the fundamental error analysis established in these early comparative studies.

While the theoretical properties of global differentiation were understood by the 1970s, the practical construction of these matrices for arbitrary, non-uniform grids remained a significant computational bottleneck. The naive approach — inverting the Vandermonde matrix — is famously ill-conditioned, with condition numbers growing exponentially with N . The solution appeared in 1988, when Fornberg [46] published his recursive algorithm in *Mathematics of Computation*. This algorithm bypasses the Vandermonde system entirely; instead, it utilizes the recursive properties of Lagrange polynomial weights to generate finite difference coefficients of *any* order for *any* distribution of grid points. The algorithm’s stability and $O(N^2)$ complexity for generating full differentiation matrices revolutionized the implementation of spectral methods on non-standard intervals and adaptive grids. The significance of Fornberg’s contribution cannot be overstated: it decoupled the generation of differentiation stencils from the geometric constraints of the grid. As noted in recent reviews of high-order methods, such as the work on spectral stability by Hangelbroek *et al.* [60], Fornberg’s

algorithm remains the “gold standard” for weight generation in mesh-free and radial basis function (RBF) contexts, extending well beyond polynomial spectral methods. In the realm of software, Driscoll [34] provided a widely cited MATLAB implementation that made the recursive generation of weights accessible to a broad community of applied mathematicians. Contemporary applications continue to leverage this algorithmic foundation. For example, Cantero [26] utilized Fornberg’s weights to approximate moving point loads on bridges, demonstrating that high-order finite difference schemes derived via this recursion are crucial for processing signal data in structural engineering.

As spectral methods matured, practitioners encountered a subtle but pervasive instability: the sensitivity of differentiation matrices to rounding errors. On Chebyshev and Legendre grids, the entries of the differentiation matrix D grow as $O(N^2)$, and small errors in their computation can be amplified, leading to a degradation of accuracy that scales poorly with grid size. The seminal analysis of this phenomenon was provided by Bayliss, Class, and Matkowsky [14]. They identified that the primary source of error arises when the row sums of the computed differentiation matrix do not vanish. Theoretically, the derivative of a constant function must be zero, implying that $\sum_j D_{ij} = 0$; however, due to floating-point cancellations, this condition is rarely met exactly by explicit formulas. They proposed the “negative sum trick”: computing the off-diagonal entries explicitly and then setting the diagonal entry D_{ii} to be the negative sum of the off-diagonal elements in that row. This heuristic was rigorously analyzed and validated by Baltensperger and Berrut [8], and later expanded upon by Baltensperger and Trummer [9], who demonstrated that enforcing the row-sum condition reduces the norm of the rounding error from $O(N^4)$ to $O(N^2)$. This reduction is critical; without it, the maximum achievable accuracy for second-derivative problems on typical grids would be severely compromised. The “negative sum trick” is now a standard feature of every robust spectral code. Trefethen and Trummer [122] further explored the eigenvalue sensitivity of these matrices, showing that rounding errors could displace the eigenvalues of D significantly, potentially triggering instability in time-stepping schemes. In recent years, the stability of differentiation matrices has been re-examined in the context of preconditioning. Javeed *et al.* [68] developed a preconditioner for spectral collocation that approximates integration as a convolution, yielding an inverse operator that can be applied in $O(N \log N)$ operations — a development that addresses the ill-conditioning of spectral matrices for second derivatives.

The turn of the millennium marked the standardization of spectral software. In 2000, Weideman and Reddy [128] published “A MATLAB Differentiation Matrix Suite” in the *ACM Transactions on Mathematical Software*. This paper did for spectral methods what LINPACK did for dense linear algebra: it provided a robust, modular, and concise implementation of differentiation matrices for Chebyshev, Hermite, Laguerre, and Fourier domains. The Weideman–Reddy suite codified the best practices of the era — including the negative sum trick and the use of trigonometric identities for periodic grids — into a set of functions that were both pedagogically clear and computationally efficient. Their implementation leveraged MATLAB’s vectorization capabilities, replacing nested loops with outer products to construct matrices in a single logical step. This succinct coding style influenced a generation of spectral codes, including the Python and Julia implementations found in this textbook. Parallel to this software standardization, the **barycentric interpolation formula** experienced a profound renaissance. Although the formula dates back to Jacobi (1825), its numerical properties were largely overlooked until Higham [64] proved its backward stability. This theoretical result was popularized by Berrut and Trefethen [16] in their landmark *SIAM Review* paper, “Barycentric Lagrange Interpolation.” They demonstrated that the barycentric form is not only computationally efficient ($O(N)$ for evaluation) but also remarkably stable against rounding errors, provided the weights are computed correctly. In the context of differentiation matrices, the barycentric weights provide a unified formula for matrix entries that applies to any node distribution, further generalizing the tools available to the spectral practitioner. Trefethen’s *Spectral Methods in MATLAB* [120] remains the most accessible introduction to the pseudospectral approach, presenting the key ideas through compact, readable

programs. Recent advancements have extended the barycentric framework to **rational functions**. Driscoll and Zhou [36] presented stable algorithms for derivative evaluation in barycentric rational representations, expanding the scope of spectral methods to a broader class of function spaces and allowing for “spectral” accuracy even in the presence of complex-plane singularities that would stall polynomial convergence.

In the contemporary era (2020–2026), research in spectral differentiation has bifurcated into three exciting directions: overcoming the dense matrix limitation, adapting to complex geometries, and integrating with specific physical models. The traditional weakness of spectral methods has been the dense nature of the differentiation matrices, which implies $O(N^2)$ storage and $O(N^3)$ solve times for linear systems. A paradigm shift occurred with the introduction of **Ultraspherical Spectral Methods** by Olver and Townsend [92]. They observed that while differentiation is dense in the Chebyshev basis (T_n), it becomes a sparse, banded operation if one maps the range to the Ultraspherical (Gegenbauer) polynomials ($C_n^{(\lambda)}$). This insight transforms spectral differentiation from a dense linear algebra problem into a sparse band-solver problem, allowing differential equations to be solved in $O(N)$ operations. Xu [132] extended the ultraspherical method to nonlinear ODEs using an inexact Newton–GMRES framework, demonstrating that the sparse structure can be leveraged even when the operator is linearized locally. Furthermore, Olver and Townsend [93] have generalized these ideas to triangles and disks, utilizing hierarchies of orthogonal polynomials to maintain sparsity in two dimensions. Handling boundary conditions in spectral collocation has traditionally involved “matrix stripping” — removing rows associated with boundary points. However, this can lead to ill-conditioning. Driscoll and Hale [35] introduced **Rectangular Spectral Collocation**, a technique that formulates the problem using an overdetermined system, decoupling the resolution of the bulk equation from the boundary constraints and significantly improving the condition number of the resulting linear systems. The versatility of the differentiation matrix has also led to its adoption in fractional calculus. Abdelhamid *et al.* [1] have developed fractional differentiation matrices that approximate non-local operators like the Riemann–Liouville and Caputo derivatives, retaining spectral accuracy for functions with fractional regularity. Finally, the **quantum harmonic oscillator** — the focus of Section 5.7 — remains an active area of spectral research. While the classic solution uses Hermite functions, recent studies (e.g., Ma *et al.* [82]) have applied Jacobi–Galerkin spectral methods to nonlinear fractional Schrödinger equations, achieving spectral convergence for ground states and excited states in complex potentials. The ability of spectral methods to capture the exponential decay of wavefunctions makes them the premier tool for high-precision quantum dynamics.

5.10 Summary

This chapter has established differentiation matrices as the computational heart of pseudospectral methods:

1. **Matrix representation:** Differentiation of interpolating polynomials can be expressed as matrix-vector multiplication: $u' \approx Du$.
2. **Finite differences as sparse approximations:** Classical FD methods correspond to sparse (banded) differentiation matrices. Higher-order FD methods use wider stencils and denser matrices.
3. **Spectral methods as the limit:** When the stencil extends to all grid points, we obtain the dense spectral differentiation matrix. This limiting viewpoint unifies finite difference and spectral approaches.

4. **Periodic spectral matrix:** For periodic problems on equispaced grids, the matrix entries are given by the cotangent formula Equation 19.
5. **Fornberg's algorithm:** A stable, efficient recursive algorithm computes differentiation weights for arbitrary node distributions and derivative orders.
6. **Spectral accuracy:** For smooth functions, spectral methods achieve exponentially fast convergence, vastly outperforming any fixed-order finite difference scheme.

The differentiation matrix is our passport to spectral solutions of differential equations. In the chapters ahead, we will use these matrices to solve boundary value problems, initial value problems, and eventually the partial differential equations that motivated our journey.



CHAPTER 6

Smoothness and Spectral Accuracy

In the previous chapter, we witnessed spectral methods achieving machine precision with remarkably few grid points. The test function $u(x) = 1/(2 + \sin(x))$ was differentiated to fourteen-digit accuracy using only fifty nodes, while finite difference methods of any fixed order would require millions of points to achieve comparable precision. But *why*? What is the source of this extraordinary accuracy?

The answer lies in a beautiful chain of reasoning that connects the *smoothness* of a function to the *accuracy* of its numerical differentiation. This chapter develops the theoretical foundations that explain the spectacular performance of spectral methods. The argument proceeds in two steps:

1. **Smoothness implies rapid decay:** A smooth function has little energy at high wavenumbers, so its Fourier coefficients decay rapidly.
2. **Rapid decay implies small aliasing error:** When we sample a function on a discrete grid, high frequencies “fold” onto low frequencies through a phenomenon called aliasing. If the high-frequency coefficients are negligible, this aliasing causes negligible error.

These two steps, made precise by the theorems in this chapter, constitute the fundamental explanation for spectral accuracy. Understanding them provides insight not only into *why* spectral methods work, but also into *when* they work: the rate of convergence is controlled by the smoothness of the function being approximated.

6.1 The Two Steps of Accuracy

6.1.1 The Heuristic Picture

Consider a smooth, periodic function $f(x)$ on $[0, 2\pi)$. Its Fourier series representation is

$$f(x) = \sum_{k=-\infty}^{\infty} \hat{f}_k e^{ikx}, \quad (33)$$

where the Fourier coefficients are

$$\hat{f}_k = \frac{1}{2\pi} \int_0^{2\pi} f(x) e^{-ikx} dx. \quad (34)$$

The wavenumber k corresponds to oscillations with frequency $|k|$ per period. A smooth function, by definition, changes gradually; it has no abrupt jumps or corners. Such a function cannot have significant energy at high frequencies, for high-frequency waves oscillate rapidly. Therefore, the coefficients $\hat{f}_k$ must decay as $|k| \rightarrow \infty$.

Now suppose we sample f at N equispaced points, producing the grid function $v_j = f(x_j)$ where $x_j = 2\pi j/N$. The discrete Fourier transform of this grid function involves only N coefficients, corresponding to wavenumbers $|k| \leq N/2$. What happened to the higher wavenumbers?

The answer is *aliasing*. The discrete grid cannot distinguish between a wave with wavenumber k and a wave with wavenumber $k + N$, since $e^{ikx_j} = e^{i(k+N)x_j}$ for all grid points x_j . High frequencies masquerade as low frequencies; they “fold” into the resolved spectral range. The discrete Fourier coefficient $\tilde{f}_k$ is not equal to the continuous coefficient $\hat{f}_k$, but rather to the sum of all coefficients that alias to k :

$$\tilde{f}_k = \sum_{j=-\infty}^{\infty} \hat{f}_{k+jN}. \quad (35)$$

This is the *aliasing formula*, and it reveals the key insight: if $\hat{f}_k$ decays rapidly, the aliased contributions $\hat{f}_{k \pm N}, \hat{f}_{k \pm 2N}, \dots$ are negligible. The discrete coefficients $\tilde{f}_k$ are then excellent approximations to the continuous coefficients $\hat{f}_k$, and spectral methods achieve their remarkable accuracy.

6.1.2 Roadmap of This Chapter

We shall make this heuristic argument precise through four theorems:

- **Theorem 1** classifies functions by their smoothness and establishes the corresponding decay rates of their Fourier coefficients.
- **Theorem 2** (the aliasing formula) quantifies how high frequencies contaminate low frequencies upon discretization.
- **Theorems 3 and 4** combine these results to bound the errors in spectral interpolation and differentiation.

The following sections make this heuristic argument precise through four theorems that quantify the connection between smoothness and spectral accuracy.

6.1.3 The Density of Smooth Functions

Before diving into the theorems, we address a natural concern: the analysis that follows focuses on smooth functions, but what if our function of interest is merely continuous, or has limited regularity?

The answer lies in a classical approximation theorem. Smooth functions are *dense* in the space of continuous functions: for any continuous function f on a compact interval and any $\varepsilon > 0$, there exists a smooth function g such that $\|f - g\|_\infty < \varepsilon$. For periodic functions on $[0, 2\pi]$, trigonometric polynomials provide such approximations; this is the content of the Weierstrass approximation theorem [130]. The far-reaching generalization to arbitrary compact Hausdorff spaces, proved by Stone [112], confirmed that polynomial density is a fundamental structural property of continuous function spaces.

The construction is elegant: convolve f with a smooth “bump function” (a mollifier) of small support. The resulting function inherits the smoothness of the mollifier while approximating f arbitrarily well. This technique, known as *mollification*, is fundamental in analysis and partial differential equations.

The implication for spectral methods is profound. Suppose we wish to approximate a merely continuous function f . We can:

1. Approximate f by a smooth function g with $\|f - g\|_\infty < \varepsilon/2$;
2. Apply spectral methods to g , which (by Theorem 1) has rapidly decaying Fourier coefficients;
3. Achieve total error less than ε with sufficiently many grid points.

This justifies our focus on smooth functions: spectral methods provide a *universal approximation framework* for continuous functions, approaching any target through smooth intermediaries. The smoothness assumption is not a limitation but rather a pathway to understanding.

6.2 The Decay of Fourier Coefficients

6.2.1 Smoothness Classes

The following theorem, due in various parts to mathematicians from Riemann to Paley and Wiener [96], establishes the fundamental connection between function smoothness and spectral decay. We state it for functions on the real line $\mathbb{R}$; similar results hold for periodic functions.

Theorem 1 (Smoothness and Fourier decay). Let $u \in L^2(\mathbb{R})$ have Fourier transform $\hat{u}$.

(a) If u has $p - 1$ continuous derivatives in $L^2(\mathbb{R})$ for some $p \geq 0$ and a p -th derivative of bounded variation, then

$$\hat{u}_k = O(|k|^{-p-1}) \quad \text{as } |k| \rightarrow \infty. \quad (36)$$

(b) If u has infinitely many continuous derivatives in $L^2(\mathbb{R})$, then

$$\hat{u}_k = O(|k|^{-m}) \quad \text{as } |k| \rightarrow \infty \quad (37)$$

for every $m \geq 0$.

(c) If there exists $a > 0$ such that u can be extended to an analytic function in the complex strip $|\operatorname{Im}(z)| < a$ with bounded L^2 norm along horizontal lines, then

$$\hat{u}_k = O(e^{-a|k|}) \quad \text{as } |k| \rightarrow \infty. \quad (38)$$

(d) If u is entire (analytic throughout $\mathbb{C}$) and satisfies $|u(z)| = o(e^{a|z|})$ as $|z| \rightarrow \infty$ for some $a > 0$, then $\hat{u}$ has compact support: $\hat{u}_k = 0$ for all $|k| > a$.

Parts (c) and (d) are known as the Paley–Wiener theorems [96]. The theorem establishes a hierarchy of smoothness classes, each with its characteristic decay rate:

Smoothness Class	Decay Rate	Example
Bounded variation	$O(k ^{-1})$	Step function
p derivatives, p -th in BV	$O(k ^{-p-1})$	B-splines
C^∞ (infinitely smooth)	$O(k ^{-m})$ for all m	Bump functions
Analytic in strip	$O(e^{-a k })$	$1/(1+x^2)$
Entire	Faster than any exponential	e^{-x^2}
Band-limited	Compact support	$\operatorname{sinc}(x)$

Table 6: The smoothness hierarchy and corresponding Fourier decay rates. More smoothness implies faster decay.

6.2.2 Intuition via Integration by Parts

The algebraic decay rates in parts (a) and (b) can be understood through integration by parts. For a smooth periodic function f with period 2π , we have

$$\hat{f}_k = \frac{1}{2\pi} \int_0^{2\pi} f(x) e^{-ikx} dx = \frac{1}{ik} \cdot \frac{1}{2\pi} \int_0^{2\pi} f'(x) e^{-ikx} dx = \frac{1}{ik} \hat{f}'_k. \quad (39)$$

Each integration by parts gains a factor of k^{-1} in the decay rate. If f has p derivatives, we can integrate by parts p times:

$$\hat{f}_k = \frac{1}{(ik)^p} \widehat{f^{(p)}}_k. \quad (40)$$

If $f^{(p)}$ has bounded variation (so its Fourier coefficients are $O(k^{-1})$), then $\hat{f}_k = O(k^{-p-1})$. This integration-by-parts argument is developed in detail by Boyd [21].

6.2.3 Intuition via Complex Analysis

The exponential decay in part (c) arises from complex analysis. If $f(z)$ is analytic in a strip $|\operatorname{Im}(z)| < a$, we can shift the contour of integration in the Fourier integral:

$$\hat{f}_k = \int_{-\infty}^{\infty} f(x) e^{-ikx} dx = \int_{-\infty}^{\infty} f(x + i\sigma) e^{-ik(x+i\sigma)} dx = e^{k\sigma} \int_{-\infty}^{\infty} f(x + i\sigma) e^{-ikx} dx \quad (41)$$

for any $|\sigma| < a$. Taking $\sigma = a - \varepsilon$ for small $\varepsilon > 0$ (and $k > 0$) gives $|\hat{f}_k| \leq C e^{-ak}$. The distance from the real axis to the nearest singularity controls the decay rate.

6.2.4 Computational Demonstration

Figure 31 illustrates the decay hierarchy for three representative functions on $[0, 2\pi]$:

1. $f_1(x) = |\sin(x)|^3$: This function is C^2 but not C^3 (the third derivative has jump discontinuities). By Theorem 1(a), its Fourier coefficients decay as $O(k^{-4})$.
2. $f_2(x) = 1/(1 + \sin^2(x/2))$: This function is smooth on the real line but has poles in the complex plane. It is analytic in a strip of finite width, so by Theorem 1(c), its coefficients decay geometrically: $O(c^{-k})$ for some $c > 1$.
3. $f_3(x) = \exp(\sin(x))$: This function is entire (analytic throughout $\mathbb{C}$). Its coefficients decay faster than any exponential.




Figure 31: The decay hierarchy of Fourier coefficients. Functions with greater smoothness exhibit faster decay: algebraic ($O(k^{-4})$) for finite regularity, geometric ($O(e^{-\alpha k})$) for strip-analytic functions, and super-geometric for entire functions. The plot shows $|f_k|$ versus wavenumber k on a semi-log scale.

The following Python code computes the Fourier coefficients:

```
1 def fourier_coefficients(f, N):
2     """Compute Fourier coefficients via FFT."""
3     x = np.linspace(0, 2*np.pi, N, endpoint=False)
4     f_hat = np.fft.fft(f(x)) / N
5     return np.abs(f_hat[:N//2])
```

Python

The equivalent MATLAB implementation:

```
1 function f_hat = fourier_coefficients(f, N)
2     x = linspace(0, 2*pi, N+1);
3     x = x(1:end-1); % Remove endpoint for periodicity
4     f_hat = abs(fft(f(x))) / N;
5     f_hat = f_hat(1:N/2+1);
6 end
```

Matlab

The Julia implementation:

```
1 function compute_fourier_coefficients(f, N)
2     x = range(0, 2pi, length=N+1)[1:N]
3     f_hat = fft(f.(x)) / N
4     f_hat_abs = abs.(f_hat[1:N÷2+1])
5     k = 0:N÷2
6     return collect(k), f_hat_abs
7 end
```

Julia

The code generating Figure 31 is available in:

- `codes/python/ch06/fourier_decay.py`
- `codes/matlab/ch06/fourier_decay.m`
- `codes/julia/ch06/fourier_decay.jl`

6.3 The Aliasing Phenomenon

6.3.1 The Poisson Summation Formula

When we sample a continuous function at N equispaced points, we lose information about frequencies beyond the *Nyquist frequency* $N/2$, a concept formalized by Shannon [107] in the context of signal processing. This lost information does not simply disappear; it contaminates the lower frequencies through aliasing.

Theorem 2 (Aliasing formula). Let $u \in L^2(\mathbb{R})$ have a first derivative of bounded variation, and let v be the grid function on $h\mathbb{Z}$ defined by $v_j = u(x_j)$. Then for all $k \in [-\pi/h, \pi/h]$,

$$\hat{v}_k = \sum_{j=-\infty}^{\infty} \hat{u}_{k+2\pi j/h}. \quad (42)$$

For a periodic function sampled at N equispaced points with spacing $h = 2\pi/N$, this becomes:

$$\tilde{f}_k = \sum_{j=-\infty}^{\infty} \hat{f}_{k+jN}. \quad (43)$$

6.3.2 Spectral Folding

The aliasing formula has a vivid geometric interpretation: the frequency axis “folds” onto itself at the Nyquist frequency. All frequencies that differ by multiples of N become indistinguishable on the discrete grid.

Consider a wave e^{ikx} sampled at $x_j = 2\pi j/N$. At the grid points,

$$e^{ikx_j} = e^{2\pi i k j / N} = e^{2\pi i (k+N) j / N} = e^{i(k+N)x_j}. \quad (44)$$

The discrete samples cannot distinguish wavenumber k from wavenumber $k + N$.

Figure 32 demonstrates this phenomenon in three panels:

- **Left:** A function with high-frequency content sampled coarsely. The interpolant through the samples misses the true oscillations.
- **Center:** The frequency folding diagram, showing how wavenumbers outside $[-N/2, N/2]$ map into this range.
- **Right:** Comparison of the true spectrum (many points) with the aliased spectrum (few points).




Figure 32: The aliasing phenomenon. Left: A function $f(x) = \sin(x) + 0.5 \sin(10x)$ sampled at $N = 12$ points. The high-frequency component $\sin(10x)$ aliases to $\sin(-2x)$, producing a misleading interpolant. Center: Frequency folding diagram showing how wavenumbers $k = 10$ and $k = -10$ fold into the resolved range $[-6, 6]$. Right: True versus aliased spectrum for $f(x) = 1/(1.5 + \cos(x))$.

6.3.3 Why Smooth Functions Are Immune

For smooth functions, aliasing is harmless. If $\hat{f}_k$ decays rapidly, the aliased contributions $\hat{f}_{k \pm N}, \hat{f}_{k \pm 2N}, \dots$ are negligible compared to $\hat{f}_k$. The discrete coefficients $\tilde{f}_k$ are then excellent approximations to the continuous coefficients $\hat{f}_k$.

Quantitatively, the aliasing error is bounded by the tail of the Fourier series:

$$|\tilde{f}_k - \hat{f}_k| \leq \sum_{j \neq 0} |\hat{f}_{k+jN}| \leq 2 \sum_{m=N/2}^{\infty} |\hat{f}_m|. \quad (45)$$

If $\hat{f}_k = O(k^{-p-1})$, this sum is $O(N^{-p})$. If $\hat{f}_k = O(e^{-ak})$, the sum is $O(e^{-aN/2})$, which is exponentially small. The rigorous stability analysis showing that these small aliasing errors do not accumulate to destroy the solution in well-posed problems was provided by Tadmor [116] and Pasciak [97].

The code generating Figure 32 is available in:

- codes/python/ch06/aliasing_demo.py
- codes/matlab/ch06/aliasing_demo.m
- codes/julia/ch06/aliasing_demo.jl

6.4 Accuracy of Spectral Differentiation

6.4.1 From Fourier Coefficients to Differentiation Error

We now connect the decay of Fourier coefficients to the accuracy of spectral differentiation. The key observation is that differentiation in Fourier space is multiplication by ik :

$$\hat{f}'_k = ik \hat{f}_k. \quad (46)$$

For the discrete approximation, we have $\hat{f}'_{\text{discrete}} = ik \tilde{f}_k$. The error in the discrete derivative is therefore controlled by the aliasing error in the Fourier coefficients, amplified by the factor k .

Theorem 3 (Effect of discretization on the Fourier transform). Let $u \in L^2(\mathbb{R})$ have a first derivative of bounded variation, and let v be the grid function on $h\mathbb{Z}$ defined by $v_j = u(x_j)$. Then:

(a) If u has $p - 1$ continuous derivatives with a p -th derivative of bounded variation, then

$$|\hat{v}_k - \hat{u}_k| = O(h^{p+1}) \quad \text{as } h \rightarrow 0. \quad (47)$$

(b) If u has infinitely many continuous derivatives, then

$$|\hat{v}_k - \hat{u}_k| = O(h^m) \quad \text{as } h \rightarrow 0 \quad (48)$$

for every $m \geq 0$.

(c) If u is analytic in a strip $|\operatorname{Im}(z)| < a$, then

$$|\hat{v}_k - \hat{u}_k| = O(e^{-\pi(a-\varepsilon)/h}) \quad \text{as } h \rightarrow 0 \quad (49)$$

for every $\varepsilon > 0$.

These bounds were established rigorously by Gottlieb and Orszag [55]; the extension to arbitrary grids was developed by Carpenter and Gottlieb [29]. The corresponding result for differentiation error is:

Theorem 4 (Accuracy of spectral differentiation). Let $u \in L^2(\mathbb{R})$ have a ν -th derivative ($\nu \geq 1$) of bounded variation, and let w be the ν -th spectral derivative of u on the grid $h\mathbb{Z}$. Then:

(a) If u has $p - 1$ continuous derivatives with a p -th derivative of bounded variation ($p \geq \nu + 1$), then

$$|w_j - u^{(\nu)}(x_j)| = O(h^{p-\nu}) \quad \text{as } h \rightarrow 0. \quad (50)$$

(b) If u has infinitely many continuous derivatives, then

$$|w_j - u^{(\nu)}(x_j)| = O(h^m) \quad \text{as } h \rightarrow 0 \quad (51)$$

for every $m \geq 0$.

(c) If u is analytic in a strip $|\operatorname{Im}(z)| < a$, then

$$|w_j - u^{(\nu)}(x_j)| = O(e^{-\pi(a-\varepsilon)/h}) \quad \text{as } h \rightarrow 0 \quad (52)$$

for every $\varepsilon > 0$.

(d) If u is band-limited with $\hat{u}_k = 0$ for $|k| > \pi/h$, then

$$w_j = u^{(\nu)}(x_j) \quad (53)$$

exactly.

Remark on the case $\nu = 0$. Theorem 4 requires $\nu \geq 1$ because the case $\nu = 0$ is trivial: the “zeroth spectral derivative” is simply the sampled function values $w_j = v_j = u(x_j)$, so the pointwise error $|w_j - u(x_j)| = 0$ exactly. The non-trivial question for $\nu = 0$ concerns *interpolation error*: how well does the trigonometric interpolant $p_N(x)$ approximate $u(x)$ at points *between* grid points? This is controlled by Theorem 3, which bounds the Fourier coefficient error. The distinction is that Theorem 3 measures error in Fourier space, while Theorem 4 measures differentiation error in physical space at grid points.

6.4.2 Connection to Chapter 5

These theorems explain the convergence behavior observed in Section 5.5. The test function $u(x) = 1/(2 + \sin(x))$ is analytic on the real line, with nearest singularities in the complex plane at $x = -\pi/2 \pm i \cdot \operatorname{arcsinh}(2)$. The distance to the real axis is $a = \operatorname{arcsinh}(2) \approx 1.44$. By Theorem 4(c), the spectral differentiation error decays as $O(e^{-a^N})$ for N grid points on $[0, 2\pi)$. This geometric interpretation of convergence — controlled by the distance to the nearest complex singularity — is developed with particular clarity by Trefethen and Weideman [123].

The finite difference methods in Chapter 5 achieve only algebraic convergence ($O(N^{-2})$, $O(N^{-4})$, $O(N^{-6})$) because they use only local information, while spectral methods exploit global smoothness to achieve exponential convergence.

6.4.3 Computational Verification

Figure 33 demonstrates the convergence rates predicted by Theorem 4 for our three test functions:

- $|\sin(x)|^3$: Algebraic convergence $O(N^{-3})$ because the function has limited smoothness.
- $1/(1 + \sin^2(x/2))$: Geometric convergence $O(c^{-N})$ because the function is analytic in a strip.
- $\exp(\sin(x))$: Super-geometric convergence because the function is entire.




Figure 33: Spectral differentiation convergence rates for three functions of increasing smoothness. The error is measured in the maximum norm $\|f' - Df\|_\infty$. Functions with greater smoothness achieve faster convergence, exactly as predicted by Theorem 4.

The following Python code computes the differentiation error:

```
1 def spectral_diff_error(f, f_deriv, N):
2     """Compute max error in spectral differentiation."""
3     D, x = spectral_diff_periodic(N)
4     error = np.max(np.abs(D @ f(x) - f_deriv(x)))
5     return error
```

Python

The equivalent MATLAB implementation:


```
1 function error = spectral_diff_error(f, f_deriv, N)
2     [D, x] = spectral_diff_periodic(N);
3     error = max(abs(D * f(x) - f_deriv(x)));
```

Matlab

```
4 end
```

The Julia implementation:

```
1 function spectral_diff_error(f, f_deriv, N)
2     D, x = spectral_diff_periodic(N)
3     error = maximum(abs.(D * f.(x) - f_deriv.(x)))
4     return error
5 end
```



The code generating Figure 33 is available in:

- codes/python/ch06/convergence_rates.py
- codes/matlab/ch06/convergence_rates.m
- codes/julia/ch06/convergence_rates.jl

6.5 A non-exhaustive literature overview

The theoretical framework of spectral accuracy rests on the deep interplay between complex analysis, approximation theory, and signal processing. The fundamental connection between a function’s smoothness and the decay of its expansion coefficients (Theorem 1) has its origins in the golden age of classical analysis. Weierstrass [130] established the density of polynomials in continuous function spaces, an existential result that was made constructive and quantitative by Bernstein [15] and Jackson [67]. Their “converse theorems” provided the first rigorous link between the rate of polynomial approximation and the differential properties of the target function.

For analytic functions, the theory transcends the real line. The celebrated **Paley–Wiener Theorems**, published in the seminal colloquium *Fourier Transforms in the Complex Domain* [96], characterized the Fourier transforms of square-integrable functions supported on a half-line or interval. Their work allows us to interpret “spectral accuracy” geometrically: the convergence rate $O(e^{-\alpha N})$ is strictly determined by the width α of the strip of analyticity in the complex plane. This geometric viewpoint has been elaborated with great clarity and pedagogical force by Trefethen [123], who demonstrated that the error of the trapezoidal rule (and by extension, the periodic spectral method) is dominated by the poles of the integrand in the complex plane. Boyd [21] provides a comprehensive catalog of these decay rates for various basis functions and singularities.

The phenomenon of aliasing, formalized in Theorem 2, relies on the **Poisson Summation Formula**, a result dating back to Poisson [101] but finding its modern numerical significance through the **Shannon–Nyquist Sampling Theorem** [107]. While Shannon focused on the exact reconstruction of band-limited signals, the implications of aliasing for the *stability* of numerical PDE solvers were rigorously analyzed by Orszag [94] and Gottlieb and Orszag [55]. Their monograph remains the definitive reference for the error analysis of Galerkin and Collocation methods, proving that aliasing errors, while present, do not destroy stability for smooth flows if properly managed.

When smoothness breaks down — as in the case of shock waves or contact discontinuities — the “spectral promise” of exponential convergence is lost, manifesting as the Gibbs phenomenon. Tadmor [116] introduced the **Spectral Viscosity Method (SVM)**, demonstrating that high-order accuracy could be recovered in a weak sense even for nonlinear conservation laws by adding a spectrally localized dissipation.

In the contemporary era (2020–2025), research has focused on pushing spectral methods beyond the “smoothness barrier.” Wang [127] has derived sharp pointwise error estimates for functions with algebraic singularities, revealing how convergence deterioration localizes near endpoints versus

interior singularities. In the realm of data science, Calder and Trillos [24] have extended spectral convergence guarantees to graph Laplacians on random data clouds, establishing the discrete-to-continuum limits necessary for spectral clustering and manifold learning. Finally, the application of spectral methods to non-local operators has surged, with Zhang and Wang [135] and Bao *et al.* [10] proving optimal error bounds for spectral discretizations of fractional Schrödinger equations, demonstrating that the spectral approach remains vital for the most modern classes of differential equations.

6.6 Summary

This chapter has established the theoretical foundations for spectral accuracy:

1. **Smoothness implies spectral decay:** The Fourier coefficients of a function decay at a rate determined by its smoothness. Algebraic decay for functions with finite regularity; exponential decay for analytic functions; super-exponential decay for entire functions.
2. **Aliasing is controlled by decay:** When a function is sampled on a discrete grid, high frequencies fold onto low frequencies. If the high-frequency coefficients are small, this aliasing introduces negligible error.
3. **Spectral methods exploit smoothness:** By using global information (all grid points), spectral methods capture the smoothness of a function and achieve convergence rates that match the decay of its Fourier coefficients.
4. **Exponential convergence for analytic functions:** For functions that are analytic in a strip containing the real axis, spectral methods achieve exponential convergence. The rate is controlled by the distance to the nearest singularity in the complex plane.

The theorems in this chapter provide not only convergence rates but also insight into *when* spectral methods will excel (smooth functions) and when they may struggle (functions with limited regularity). In the next chapter, we extend these ideas to non-periodic problems on bounded domains, where Chebyshev methods take center stage.



CHAPTER 7

Chebyshev Differentiation Matrices

Having developed the theory of differentiation matrices for periodic problems in the previous chapter, we now face a new challenge: what happens when the domain is *bounded*? Many problems in science and engineering (e.g. heat conduction, wave propagation, fluid dynamics) are posed on finite intervals with boundary conditions at the endpoints. For such problems, the elegant trigonometric framework of Fourier methods must give way to something new. The key insight of this chapter is that polynomial interpolation, when done correctly, provides the foundation for spectral methods on bounded domains. The “correct” approach, as we discovered in Section 4, requires carefully chosen interpolation nodes. Equispaced points lead to the Runge phenomenon; Chebyshev points do not. This chapter builds on that foundation to construct the *Chebyshev differentiation matrix*—the non-periodic analog of the Fourier differentiation matrix. The construction follows the framework developed by Gottlieb and Orszag [55] and presented in detail by Trefethen [120].

7.1 The Non-Periodic Challenge

7.1.1 From Periodic to Bounded Domains

In the previous chapter, we exploited the periodicity of the domain $[0, 2\pi)$ to construct differentiation matrices using equispaced nodes and trigonometric interpolation. The resulting matrices had beautiful structure: circulant, with entries determined by a simple cotangent formula.

For problems on a bounded interval like $[-1, 1]$, we face several new difficulties:

1. **No periodicity:** The function values at the endpoints are independent, not related by periodicity.
2. **The Runge phenomenon:** Equispaced nodes lead to wild oscillations near the boundaries, as we demonstrated dramatically in Section 4.2.
3. **Boundary conditions:** Physical problems typically impose conditions at $x = \pm 1$, which must be incorporated into the differentiation process.

The solution to these challenges comes from our study of interpolation theory: the Chebyshev-Gauss-Lobatto points

$$x_j = \cos\left(\frac{j\pi}{N}\right), \quad j = 0, 1, \dots, N, \quad (54)$$

cluster near the boundaries, counteracting the Runge phenomenon and enabling spectral accuracy.

7.1.2 Grid Comparison: Equispaced vs. Chebyshev

Figure 34 illustrates the fundamental difference between equispaced and Chebyshev grids. The equispaced grid distributes points uniformly across $[-1, 1]$, while the Chebyshev grid clusters points near the boundaries according to the projection from a circle.

Equispaced vs. Chebyshev-Gauss-Lobatto Grids




Figure 34: Comparison of equispaced and Chebyshev-Gauss-Lobatto grids for $N = 16$ intervals. Top: equispaced points are distributed uniformly. Middle: Chebyshev points cluster near the boundaries. Bottom left: the circle projection interpretation: Chebyshev points are the projections of equally-spaced points on a semicircle. Bottom right: comparison of grid spacing near the left boundary, showing the $O(N^{-2})$ clustering of Chebyshev points.

The code generating Figure 34 is available in:

- codes/python/ch07/cheb_grid_comparison.py
- codes/matlab/ch07/cheb_grid_comparison.m
- codes/julia/ch07/cheb_grid_comparison.jl

The boundary clustering is not a minor detail, it is essential for stability. Near the boundaries, where polynomial interpolation tends to oscillate wildly, the Chebyshev grid provides many closely-spaced points to control the behavior. Near the center, where interpolation is naturally well-behaved, fewer points suffice.

The spacing of Chebyshev points near the boundary is $O(N^{-2})$, compared to $O(N^{-1})$ for equispaced points. This denser boundary clustering has important implications for time-stepping in differential equations, as we shall see in later chapters. In particular, this $O(N^{-2})$ clustering imposes a severe time-stepping restriction: explicit schemes for advection require $\Delta t \sim O(N^{-2})$, and for diffusion $\Delta t \sim O(N^{-4})$ — the so-called CFL bottleneck of Chebyshev methods [120].

7.2 The Chebyshev Differentiation Matrix

7.2.1 Differentiation via Interpolation

The construction of the Chebyshev differentiation matrix follows the same principle as in the periodic case: interpolate, then differentiate.

Given function values $\mathbf{v} = (v_0, v_1, \dots, v_N)^\top$ at the Chebyshev points Equation 54, we first construct the unique interpolating polynomial $p(x)$ of degree at most N satisfying $p(x_j) = v_j$. The derivative approximation at the grid points is then

$$\mathbf{w} = D_N \mathbf{v}, \quad \text{where} \quad w_i = p'(x_i). \quad (55)$$

The matrix D_N is the *Chebyshev differentiation matrix*. Its entries can be computed from the Lagrange basis polynomials:

$$(D_N)_{ij} = L'_j(x_i), \quad (56)$$

where L_j is the Lagrange interpolating polynomial centered at x_j .

7.2.2 Explicit Formulas

The entries of the Chebyshev differentiation matrix can be written explicitly. Define the weights

$$c_j = \begin{cases} 2 & \text{if } j = 0 \text{ or } j = N \\ 1 & \text{otherwise.} \end{cases} \quad (57)$$

These formulas are derived in full by Gottlieb and Orszag [55], Trefethen [120] [Chapter 6], and Boyd [21]. The matrix entries are:

Off-diagonal entries ($i \neq j$):

$$(D_N)_{ij} = \frac{c_i}{c_j} \frac{(-1)^{i+j}}{x_i - x_j}. \quad (58)$$

Diagonal entries (interior nodes, $j = 1, \dots, N-1$):

$$(D_N)_{jj} = -\frac{x_j}{2(1 - x_j^2)}. \quad (59)$$

Corner entries:

$$(D_N)_{00} = \frac{2N^2 + 1}{6}, \quad (D_N)_{NN} = -\frac{2N^2 + 1}{6}. \quad (60)$$

7.2.3 The Negative Sum Trick

While the formulas Equation 59 and Equation 60 give exact expressions for the diagonal entries, direct evaluation can be numerically unstable. A more robust approach uses the *negative sum trick*: since the derivative of a constant function must be zero, each row of D_N must sum to zero:

$$(D_N)_{jj} = -\sum_{k \neq j} (D_N)_{jk}. \quad (61)$$

This ensures that $D_N \mathbf{1} = \mathbf{0}$ to machine precision, where $\mathbf{1}$ is the vector of all ones. The numerical importance of this device was analyzed by Baltensperger and Berrut [8], who showed that naive evaluation of the diagonal formulas can lose several digits of accuracy; the comprehensive study by Weideman and Reddy [128] confirmed the negative sum trick as a mandatory ingredient of robust spectral codes.

The following code implements the Chebyshev differentiation matrix:

```
1 def cheb_matrix(N):
```

 Python

```


2  """Chebyshev differentiation matrix."""
3  if N == 0:
4      return np.array([[0.0]]), np.array([1.0])
5
6  # Chebyshev-Gauss-Lobatto points
7  x = np.cos(np.pi * np.arange(N + 1) / N)
8
9  # Weights: c_0 = c_N = 2, others = 1
10 c = np.ones(N + 1)
11 c[0], c[N] = 2.0, 2.0
12
13 # Off-diagonal entries
14 X = np.tile(x, (N + 1, 1))
15 dX = X - X.T
16 C = np.outer(c, 1.0 / c)
17 sign = np.outer((-1.0)**np.arange(N + 1), (-1.0)**np.arange(N + 1))
18
19 with np.errstate(divide='ignore'):
20     D = C * sign / (-dX)
21
22 # Negative sum trick for diagonal
23 np.fill_diagonal(D, 0.0)
24 D[np.diag_indices(N + 1)] = -np.sum(D, axis=1)
25
26 return D, x

```

```

1  function [D, x] = cheb_matrix(N)
2  % Chebyshev differentiation matrix.
3      if N == 0
4          D = 0; x = 1; return
5      end
6
7  % Chebyshev-Gauss-Lobatto points
8  x = cos(pi * (0:N)' / N);
9
10 % Weights: c_0 = c_N = 2, others = 1
11 c = ones(N+1, 1);
12 c(1) = 2; c(N+1) = 2;
13
14 % Off-diagonal entries
15 X = repmat(x, 1, N+1);
16 dX = X - X';
17 C = c * (1 ./ c');
18 sign = (-1).^((0:N)' + (0:N));
19
20 D = C .* sign ./ (-dX);

```



```

21
22     % Negative sum trick for diagonal
23     D(1:N+2:end) = 0;
24     D(1:N+2:end) = -sum(D, 2);
25 end


```

The Julia implementation:

```

1  function cheb_matrix(N)
2      N == 0 && return (zeros(1, 1), [1.0])
3
4      # Chebyshev-Gauss-Lobatto points
5      x = [cos(π * j / N) for j in 0:N]
6
7      # Weights: c_0 = c_N = 2, others = 1
8      c = ones(N + 1)
9      c[1] = 2.0; c[N+1] = 2.0
10
11     # Off-diagonal entries
12     X = repeat(x', N + 1, 1)
13     dX = X .- X'
14     C = c * (1.0 ./ c')
15     sign_mat = ((-1.0) .^ (0:N)) * ((-1.0) .^ (0:N))'
16
17     D = C .* sign_mat ./ (-dX .+ I(N + 1))
18     D = D .- Diagonal(diag(D))
19
20     # Negative sum trick for diagonal
21     for i in 1:N+1
22         D[i, i] = -sum(D[i, :])
23     end
24     return D, x
25 end

```



7.3 Small- N Examples

7.3.1 Hand Calculations

To develop intuition, let us compute the smallest Chebyshev matrices by hand.

For $N = 1$, we have two nodes: $x_0 = 1$ and $x_1 = -1$. The only polynomial passing through two points is a line, and its derivative is constant:

$$D_1 = \begin{pmatrix} \frac{1}{2} & -\frac{1}{2} \\ \frac{1}{2} & -\frac{1}{2} \end{pmatrix}. \quad (62)$$

For $N = 2$, we have three nodes: $x_0 = 1$, $x_1 = 0$, and $x_2 = -1$. The middle row of D_2 is particularly illuminating:

$$D_2 = \begin{pmatrix} \frac{3}{2} & -2 & \frac{1}{2} \\ \frac{1}{2} & 0 & -\frac{1}{2} \\ -\frac{1}{2} & 2 & -\frac{3}{2} \end{pmatrix}. \quad (63)$$

The middle row $(\frac{1}{2}, 0, -\frac{1}{2})$ is exactly the *centered finite difference* formula! This reveals a beautiful connection: at interior points where the grid happens to be locally symmetric, the spectral method reduces to the familiar finite difference formula.

7.4 Matrix Structure and Properties

7.4.1 The Dense Matrix

Figure 35 visualizes the structure of the Chebyshev differentiation matrix for $N = 16$.




Figure 35: Structure of the Chebyshev differentiation matrix for $N = 16$. Left: heatmap showing the matrix entries, with red indicating positive values and blue indicating negative. The large corner entries $(D_N)_{00}$ and $(D_N)_{NN}$ are visible. Right: row profiles showing boundary row (red) and interior row (green). The boundary row has large entries reflecting the $O(N^2)$ corner values.

The code generating Figure 35 and Figure 36 is available in:

- codes/python/ch07/cheb_matrix_structure.py
- codes/matlab/ch07/cheb_matrix_structure.m
- codes/julia/ch07/cheb_matrix_structure.jl

The row profile plot in Figure 35 shows both boundary and interior rows together, but the vastly different scales make the interior row structure difficult to discern. Figure 36 isolates the interior row to reveal its characteristic structure.




Figure 36: Interior row profile of the Chebyshev differentiation matrix ($i = 9$, $N = 16$). Unlike the boundary row with entries of order $O(N^2)$, interior rows have much smaller magnitudes. The largest entry in absolute value is the diagonal element $D_{9,9}$, with oscillating signs typical of differentiation stencils.

Several features are visible in Figure 36. The largest entry in magnitude is the diagonal element $(D_N)_{ii}$, which for interior nodes is relatively small (recall that diagonal entries vanish for truly centered points, while off-center interior points have $O(1)$ diagonal entries). The off-diagonal entries exhibit alternating signs, reminiscent of finite difference stencils, with the magnitude decaying as we move away from the diagonal. Notably, the entries near the row ends (columns $j = 0$ and $j = N$) remain significant, reflecting the global nature of spectral differentiation: even distant boundary points contribute to the derivative at interior locations.

Unlike the sparse banded matrices of finite difference methods, the Chebyshev differentiation matrix is *dense*: every entry is generally nonzero. This is the price we pay for spectral accuracy—information from every grid point contributes to the derivative at every other point.

7.4.2 Cardinal Functions

The columns of D_N have a natural interpretation: column j contains the derivatives of the j th Lagrange cardinal function evaluated at all the grid points. Figure 37 illustrates this connection.




Figure 37: Chebyshev cardinal functions (Lagrange basis polynomials). Left: several cardinal functions for $N = 10$, each peaking at value 1 at its corresponding node and vanishing at all others. Right: a single cardinal function with tangent lines at the grid points—the slopes of these tangent lines are precisely the entries in the corresponding column of the differentiation matrix.

The code generating Figure 37 is available in:

- codes/python/ch07/cheb_cardinal.py
- codes/matlab/ch07/cheb_cardinal.m
- codes/julia/ch07/cheb_cardinal.jl

The barycentric form of Lagrange interpolation, highlighted by Berrut and Trefethen [16], provides an alternative route to computing these derivatives that decouples the node distribution from the algorithmic structure.

7.5 Demonstration: The Witch of Agnesi

7.5.1 A Smooth Test Function

To demonstrate spectral differentiation in action, we use the *Witch of Agnesi*:

$$u(x) = \frac{1}{1 + 4x^2}, \quad (64)$$

with exact derivative

$$u'(x) = \frac{-8x}{(1 + 4x^2)^2}. \quad (65)$$

This function is smooth and analytic on $[-1, 1]$, with poles at $x = \pm i/2$ in the complex plane. The distance from $[-1, 1]$ to the nearest singularity determines the rate of exponential convergence. This geometric picture is a direct application of the Bernstein ellipse theory developed in Section 6; see also Trefethen [120] [Chapter 8] for a thorough treatment.

Figure 38 shows the function and its spectral derivative approximation for $N = 10$ and $N = 20$ grid points.


Chebyshev Differentiation of $u(x) = 1/(1 + 4x^2)$ (Witch of Agnesi)

Figure 38: Chebyshev spectral differentiation of the Witch of Agnesi $u(x) = 1/(1 + 4x^2)$. Top row: function values at Chebyshev points. Bottom row: comparison of exact and spectral derivatives, with maximum errors indicated. The error decreases exponentially with N .

The code generating Figure 38 is available in:

- codes/python/ch07/cheb_diff_demo.py
- codes/matlab/ch07/cheb_diff_demo.m
- codes/julia/ch07/cheb_diff_demo.jl

The exponential convergence is evident from Table 7, which shows the maximum differentiation error for various values of N .

N	Max error	N	Max error
4	8.5×10^{-1}	20	1.9×10^{-3}
6	4.8×10^{-1}	24	3.3×10^{-4}
8	2.4×10^{-1}	28	5.6×10^{-5}
10	1.2×10^{-1}	32	9.4×10^{-6}
12	5.3×10^{-2}	36	1.5×10^{-6}
14	2.4×10^{-2}	40	2.5×10^{-7}
16	1.0×10^{-2}	50	5.8×10^{-9}

Table 7: Maximum differentiation error for the Witch of Agnesi $u(x) = 1/(1 + 4x^2)$. The error decreases exponentially with N .

The following code demonstrates spectral differentiation:

```

1  def differentiate_witch(N):
2      """Differentiate u = 1/(1+4x^2) using Chebyshev spectral method."""
3      D, x = cheb_matrix(N)
4
5      # Function and exact derivative
6      u = 1.0 / (1.0 + 4.0 * x**2)
7      du_exact = -8.0 * x / (1.0 + 4.0 * x**2)**2
8
9      # Spectral derivative
10     du_spectral = D @ u
11
12     # Maximum error
13     max_error = np.max(np.abs(du_spectral - du_exact))
14     return x, du_spectral, du_exact, max_error

```

 Python

```

1  function [x, du_spectral, du_exact, max_error] =
2      differentiate_witch(N)
3      % Differentiate u = 1/(1+4x^2) using Chebyshev spectral method.
4
5      [D, x] = cheb_matrix(N);
6
7      % Function and exact derivative
8      u = 1 ./ (1 + 4*x.^2);
9      du_exact = -8*x ./ (1 + 4*x.^2).^2;
10
11     % Spectral derivative
12     du_spectral = D * u;
13
14     % Maximum error
15     max_error = max(abs(du_spectral - du_exact));
16 end

```

 Matlab

The Julia implementation:

```

1  function differentiate_witch(N)
2      # Differentiate u = 1/(1+4x^2) using Chebyshev spectral method.
3      D, x = cheb_matrix(N)
4
5      # Function and exact derivative
6      u = @. 1.0 / (1.0 + 4.0 * x^2)
7      du_exact = @. -8.0 * x / (1.0 + 4.0 * x^2)^2
8
9      # Spectral derivative
10     du_spectral = D * u
11
12     # Maximum error
13     max_error = maximum(abs.(du_spectral .- du_exact))
14     return x, du_spectral, du_exact, max_error
15 end

```

 Julia

7.6 Spectral Convergence

7.6.1 Four Functions of Increasing Smoothness

The rate of spectral convergence depends critically on the smoothness of the function being differentiated. The theoretical framework developed in Section 6 applies equally to Chebyshev methods on bounded domains: Theorem 1 establishes that smoother functions have more rapidly decaying spectral coefficients, and Theorem 4 translates this decay into bounds on the differentiation error. The convergence rates we report below measure how fast the maximum error $\|D_N v - u'(x)\|_\infty$ decreases as N increases.

To illustrate these principles, we examine four test functions with different regularity:

1. $|x|^{5/2}$: This function has fractional Hölder regularity $s = 5/2$. The first two derivatives,

$$(|x|^{5/2})' = \frac{5}{2}|x|^{3/2} \operatorname{sgn}(x), \quad (|x|^{5/2})'' = \frac{15}{4}|x|^{1/2}, \quad (66)$$

are continuous, but the third derivative $(15/8)|x|^{-1/2} \operatorname{sgn}(x)$ is unbounded at $x = 0$. By the generalization of Theorem 1 to non-integer regularity, the Chebyshev coefficients decay as $O(n^{-s-1}) = O(n^{-7/2})$, and by Theorem 4 the differentiation error is $O(N^{-s}) = O(N^{-5/2}) = O(N^{-2.5})$. The fractional exponent in the function directly determines the convergence rate.

2. $e^{-1/(1-x^2)}$: The “bump function” is infinitely differentiable (C^∞) on $[-1, 1]$, vanishing together with all its derivatives at $x = \pm 1$. However, it is not analytic at the endpoints: no Taylor series converges there. By Theorem 1(b), C^∞ functions have coefficients that decay faster than any algebraic rate, $O(n^{-m})$ for all m . The differentiation error thus converges superalgebraically (faster than any power of N) but not exponentially.
3. $\tanh(5x)$: This function is real-analytic on $[-1, 1]$, but $\tanh(z)$ has poles at $z = \pm i\pi/2$ in the complex plane. Scaling by 5 moves the nearest poles to $z = \pm i\pi/10$. By Theorem 1(c), analyticity in a strip of half-width $a = \pi/10$ yields Chebyshev coefficients decaying as $O(e^{-an})$. For the

Chebyshev ellipse with parameter ρ , the convergence rate is $O(\rho^{-N})$ where $\rho \approx 1 + a = 1 + \pi/10 \approx 1.31$.

4. x^8 : A polynomial of degree 8 is its own interpolant for $N \geq 8$. This is the bounded-domain analog of band-limited functions: since x^8 has only finitely many nonzero Chebyshev coefficients (up to T_8), the spectral derivative is *exact* once N is large enough to resolve all terms. This corresponds to Theorem 4(d) for the periodic case.

Figure 39 displays the maximum differentiation error versus N for these four functions.




Figure 39: Convergence of the spectral differentiation error $\|D_N v - u'(x)\|_\infty$ as a function of N for four test functions. Top left: $|x|^{5/2}$ shows algebraic convergence $O(N^{-2.5})$ due to limited smoothness (Hölder regularity $5/2$). Top right: the bump function $e^{-1/(1-x^2)}$ achieves superalgebraic (faster than any power) but not exponential convergence, consistent with C^∞ but non-analytic behavior. Bottom left: $\tanh(5x)$ demonstrates exponential convergence until machine precision, as expected for analytic functions. Bottom right: the polynomial x^8 is differentiated exactly for $N \geq 8$.

The code generating Figure 39 is available in:

- `codes/python/ch07/cheb_convergence.py`
- `codes/matlab/ch07/cheb_convergence.m`
- `codes/julia/ch07/cheb_convergence.jl`

The message is clear: spectral methods achieve their promised exponential convergence only for analytic functions. For less smooth functions, convergence is still rapid but algebraic, with the rate determined by the degree of smoothness. A comprehensive catalogue of convergence rates for different

classes of functions, including the interplay between algebraic and geometric decay, is given by Boyd [21].

7.7 A non-exhaustive literature overview

The intellectual roots of Chebyshev spectral differentiation reach back to Fourier’s 1822 treatise on heat conduction [48], which demonstrated that functions — even those with discontinuities — could be represented as infinite sums of trigonometric modes. While Fourier series are naturally suited to periodic domains, their direct application to bounded intervals destroys uniform convergence at the boundaries. The resolution came with the work of Cornelius Lanczos in the 1930s [74], who recognized that Chebyshev polynomials — defined as $T_n(x) = \cos(n \arccos x)$ — inherit the favourable convergence properties of Fourier series while providing a robust basis for non-periodic functions on $[-1, 1]$. This insight remained largely theoretical until the rediscovery of the Fast Fourier Transform by Cooley and Tukey [30] made the transformation between physical and spectral space computationally feasible in $O(N \log N)$ operations. The subsequent formalization by Gottlieb and Orszag [55] established the mathematical rigour that underpins modern practice.

The transition from theory to reliable computation required several algorithmic refinements. Baltensperger and Berrut [8] provided a careful analysis of rounding errors in the pseudospectral differentiation matrix, demonstrating that the naive evaluation of diagonal entries can lose several digits of accuracy and establishing the negative sum trick as essential practice. Weideman and Reddy [128] consolidated these insights into a widely-used MATLAB differentiation matrix suite that made spectral methods accessible to a broad community of practitioners. Berrut and Trefethen [16] highlighted the extraordinary stability of the barycentric form of Lagrange interpolation, providing an alternative algorithmic pathway that decouples the choice of nodes from the structure of the differentiation algorithm. Trefethen’s *Spectral Methods in MATLAB* [120] remains the most accessible introduction to the pseudospectral approach, presenting the key ideas through compact, readable programs.

Several comprehensive textbooks have codified the theory for practitioners. The foundational monograph by Gottlieb and Orszag [55] provides the rigorous mathematical framework, including convergence theorems and stability analysis. Boyd [21] offers an encyclopaedic treatment of Chebyshev and Fourier methods with an emphasis on practical algorithms and an extensive catalogue of convergence rates. Fornberg [47] gives a unified perspective on pseudospectral and finite-difference methods, while Canuto *et al.* [27] develop the functional-analytic foundations in full generality.

In recent years, Chebyshev spectral methods have found striking applications beyond their traditional home in fluid dynamics. Batic, Dutykh, and Beek [13] employed Chebyshev differentiation matrices to compute the quasinormal mode frequencies of noncommutative geometry-inspired wormholes, a problem in general relativity where the high accuracy of spectral methods is essential for resolving the small imaginary parts of the complex eigenfrequencies. Melia *et al.* [86] developed hardware-accelerated hierarchical Poincaré–Steklov solvers that exploit the structure of spectral discretizations on modern GPU architectures, achieving significant speedups for large-scale elliptic problems.

Contemporary research is extending Chebyshev differentiation in several new directions. Zayernouri *et al.* [134] have developed spectral and spectral element methods for fractional differential equations, exploiting the global nature of Chebyshev polynomials to handle the non-local operators that arise in anomalous diffusion and memory effects. Tasnin [118] has explored the integration of Chebyshev spectral methods with physics-informed neural networks, embedding spectral layers into deep learning architectures to overcome the spectral bias that limits standard networks in resolving high-frequency features. These developments, together with emerging work on quantum algorithms

for spectral linear systems, suggest that Chebyshev differentiation matrices will remain central tools in computational science for the foreseeable future.

7.8 Summary

This chapter has developed the Chebyshev differentiation matrix for spectral methods on bounded domains. Key points include:

- **Chebyshev points** Equation 54 provide the optimal node distribution, clustering near boundaries to avoid the Runge phenomenon.
- **The negative sum trick** Equation 61 ensures numerical stability by computing diagonal entries from row sums.
- **Matrix structure:** The Chebyshev differentiation matrix is dense, with $O(N^2)$ corner entries.
- **Spectral convergence** depends on smoothness: exponential for analytic functions, algebraic for functions with limited regularity.

In the next chapter, we will use these differentiation matrices to solve boundary value problems—the natural next step in applying spectral methods to differential equations.



CHAPTER 8

Boundary Value Problems

With the Chebyshev differentiation matrix in hand, we are ready to tackle one of the most important classes of problems in applied mathematics: boundary value problems (BVPs). Unlike initial value problems, where we march forward in time from given initial conditions, BVPs impose constraints at multiple locations, typically at the boundaries of the domain. This spatial coupling makes BVPs inherently global, and spectral methods are ideally suited to exploit this structure.

This chapter demonstrates how to transform differential equations into linear algebra problems. The differentiation matrix D_N from Section 7 becomes the workhorse, and its square D_N^2 handles second-order equations. Imposing boundary conditions requires a simple “matrix surgery”, removing rows and columns corresponding to boundary points. The result is a systematic approach that handles linear, variable-coefficient, and even nonlinear problems with remarkable ease.

8.1 Second Derivatives and Matrix Squaring

8.1.1 The Need for D^2

Most BVPs in physics involve second-order derivatives: the heat equation, the wave equation, the Poisson equation, and many others all feature u_{xx} . These classical equations are treated extensively in the foundational spectral methods texts by Gottlieb and Orszag [55] and Canuto *et al.* [27]. To apply spectral collocation, we need a second derivative matrix.

Two approaches present themselves:

1. **Direct formulas:** Derive explicit expressions for $(D^2)_{ij}$ analogous to the first derivative formulas in Section 7. These exist but are complex. The definitive treatment is the MATLAB Differentiation Matrix Suite of Weideman and Reddy [128], which provides algorithmically stable recurrence relations for arbitrary-order derivatives on Chebyshev and Hermite grids.
2. **Matrix squaring:** Simply compute $D^2 = D \times D$. This is $O(N^3)$ but entirely adequate for spectral N -values (typically $N < 200$).

We adopt the second approach for its simplicity. Trefethen [120] advocates this pedagogical choice, noting that optimised BLAS routines make the $O(N^3)$ matrix multiplication extremely efficient for the moderate N typical of spectral methods. Baltensperger and Berrut [8] showed that for $N < 200$, the conditioning of the spectral operator ($O(N^4)$ for second derivatives) dominates the error budget long before the method of matrix construction becomes the limiting factor. The product D_N^2 gives us the second derivative matrix: if $\mathbf{v}$ contains function values at the Chebyshev points, then $D_N^2 \mathbf{v}$ approximates the second derivative values.

8.2 Imposing Boundary Conditions

8.2.1 Dirichlet Conditions: Matrix Stripping

The most common boundary conditions specify the function values at the endpoints:

$$u(-1) = \alpha, \quad u(1) = \beta. \quad (67)$$

These are *Dirichlet conditions*. Boyd [21] classifies boundary conditions into “numerical” (imposed through the discrete operator) and “behavioural” (satisfied by the basis functions automatically); matrix stripping belongs to the numerical category.

With Chebyshev points ordered as $x_0 = 1, x_1, \dots, x_{N-1}, x_N = -1$, the boundary conditions fix $v_0 = \beta$ and $v_N = \alpha$. The differential equation need only be enforced at the *interior* points $x_1, \dots, x_{N-1}$.

The implementation is straightforward:

- **Remove** the first and last rows of the equation (we don’t need the ODE at boundary points).
- **Remove** the first and last columns (the boundary values are known, not unknowns).

The result is an $(N - 1) \times (N - 1)$ interior system. Denoting the entries of D_N^2 by $(D_N^2)_{ij}$ with $i, j = 0, 1, \dots, N$, we define

$$\tilde{D}_{ij}^2 = (D_N^2)_{ij}, \quad i, j = 1, \dots, N - 1. \quad (68)$$

For homogeneous conditions ($\alpha = \beta = 0$), the boundary terms vanish and we simply solve $\tilde{D}^2 \mathbf{u}_{\text{int}} = \mathbf{f}_{\text{int}}$.

8.2.2 Neumann and Robin Conditions: Boundary Bordering

Matrix stripping works well for Dirichlet conditions, but *Neumann* conditions ($u'(\text{boundary}) = g$) or *Robin* conditions ($\alpha u + \beta u' = \gamma$) require a different approach. Since the boundary condition involves the derivative, we cannot simply remove the boundary unknowns.

The *boundary bordering* technique keeps the full $(N + 1) \times (N + 1)$ system and *replaces* boundary rows with the appropriate conditions. This technique, codified by Boyd [21], provides a unified framework for imposing Dirichlet, Neumann, Robin, and even nonlinear boundary conditions within the same linear algebra structure. Specifically:

- **Neumann** $u'(x_0) = g$: replace row 0 of D^2 with row 0 of D_N (the first derivative matrix), and set the corresponding right-hand side entry to g .
- **Robin** $\alpha u(x_0) + \beta u'(x_0) = \gamma$: replace row 0 with $\alpha \mathbf{e}_0^\top + \beta D_N[0, :]$, where $\mathbf{e}_0$ is the first unit vector.
- **Dirichlet** $u(x_0) = \alpha$: replace row 0 with $\mathbf{e}_0^\top$ and set the right-hand side to α .

This approach handles any combination of boundary conditions uniformly. We demonstrate it in Section 8.5 below.

8.3 Linear BVP: The Poisson Equation

8.3.1 A Model Problem

Consider the one-dimensional Poisson equation:

$$u_{xx} = \sin(\pi x) + 2 \cos(2\pi x), \quad x \in (-1, 1), \quad u(\pm 1) = 0. \quad (69)$$

This has the exact solution

$$u(x) = -\frac{\sin(\pi x)}{\pi^2} + \frac{1 - \cos(2\pi x)}{2\pi^2}. \quad (70)$$

8.3.2 Spectral Solution

The solution procedure is direct:

1. Construct D_N and compute D_N^2 .
2. Extract the interior submatrix $\tilde{D}^2$ (rows and columns $1, \dots, N-1$ of D_N^2).
3. Evaluate the right-hand side at interior points.
4. Solve the linear system $\tilde{D}^2 \mathbf{u}_{\text{int}} = \mathbf{f}_{\text{int}}$.
5. Embed the result in the full vector with boundary values.

Figure 40 shows the solution and demonstrates spectral convergence.




Figure 40: Solution of the 1D Poisson equation Equation 69. Left: numerical solution (circles) compared with exact solution (line) for $N = 16$. Right: exponential convergence of the maximum error, reaching machine precision by $N = 24$.

Table 8 quantifies the convergence rate. The error decreases by roughly two orders of magnitude for each increment of two in N , the hallmark of spectral (exponential) convergence. The theoretical basis for this exponential convergence rate was established by Gottlieb and Orszag [55], who proved that for analytic functions the spectral approximation error decays faster than any polynomial in $1/N$; the monograph by Canuto *et al.* [27] provides a comprehensive modern treatment. Machine precision is reached by $N = 24$, after which further refinement cannot improve the result due to floating-point arithmetic limitations.

N	Max Error
4	5.78×10^{-1}
6	4.44×10^{-2}
8	3.01×10^{-3}
10	1.33×10^{-4}
12	7.55×10^{-6}
16	7.84×10^{-9}
20	3.62×10^{-12}
24	6.07×10^{-16}
32	3.75×10^{-16}
48	7.91×10^{-16}

Table 8: Convergence of the Chebyshev spectral method for the Poisson equation Equation 69. The error decreases exponentially with N until machine precision ($\approx 2.2 \times 10^{-16}$) is reached at $N = 24$.

The implementation is remarkably concise:

```

1 def solve_poisson(N, f):
2     """Solve  $u_{xx} = f(x)$  with  $u(\pm 1) = 0$ ."""
3     D, x = cheb_matrix(N)
4     D2 = D @ D
5
6     # Extract interior system
7     D2_int = D2[1:N, 1:N]
8     f_int = f(x[1:N])
9
10    # Solve
11    u_int = np.linalg.solve(D2_int, f_int)
12
13    # Assemble full solution
14    u = np.zeros(N + 1)
15    u[1:N] = u_int
16    return x, u

```

 Python

```

1 function [x, u] = solve_poisson(N, f)
2 % Solve  $u_{xx} = f(x)$  with  $u(\pm 1) = 0$ .
3     [D, x] = cheb_matrix(N);
4     D2 = D * D;
5
6     % Extract interior system
7     D2_int = D2(2:N, 2:N);
8     f_int = f(x(2:N));
9

```

 Matlab

```

10     % Solve
11     u_int = D2_int \ f_int;
12
13     % Assemble full solution
14     u = zeros(N+1, 1);
15     u(2:N) = u_int;
16 end

```

The Julia implementation:

```

1  function solve_poisson(N, f)
2      # Solve  $u_{xx} = f(x)$  with  $u(\pm 1) = 0$ .
3      D, x = cheb_matrix(N)
4      D2 = D * D
5
6      # Extract interior system
7      D2_int = D2[2:N, 2:N]
8      f_int = f.(x[2:N])
9
10     # Solve and assemble
11     u_int = D2_int \ f_int
12     u = zeros(N + 1)
13     u[2:N] = u_int
14     return x, u
15 end

```

 Julia

The code generating Figure 40 is available in:

- codes/python/ch08/bvp_linear.py
- codes/matlab/ch08/bvp_linear.m
- codes/julia/ch08/bvp_linear.jl

8.4 Variable Coefficient Problems

8.4.1 The Airy-Type Equation

Variable coefficients pose no additional difficulty for spectral methods. Consider:

$$u_{xx} - (1 + x^2)u = 1, \quad x \in (-1, 1), \quad u(\pm 1) = 0. \quad (71)$$

The variable coefficient $(1 + x^2)$ becomes a diagonal matrix. The discretized operator is:

$$L = D_N^2 - \text{diag}(1 + x^2). \quad (72)$$

After extracting the interior system, we solve $\tilde{L}u_{\text{int}} = \mathbf{1}_{\text{int}}$.

Figure 41 compares this solution with the constant-coefficient case.

Variable Coefficient BVP: $u_{xx} - (1 + x^2)u = 1$, $u(\pm 1) = 0$




Figure 41: Variable coefficient BVP Equation 71. Left: comparison of solutions for variable coefficient $(1 + x^2)$ and constant coefficient (1). The variable coefficient reduces the solution amplitude, especially near the boundaries where $1 + x^2$ is largest. Right: verification of the constant-coefficient case against its exact solution.

The implementation requires only a minor modification to the Poisson solver:

```
1 def solve_variable_coeff(N, coeff_func):
2     """Solve  $u_{xx} - c(x)u = 1$  with  $u(\pm 1) = 0$ ."""
3     D, x = cheb_matrix(N)
4     D2 = D @ D
5
6     # Build operator  $L = D^2 - \text{diag}(c(x))$ 
7     c = coeff_func(x)
8     L = D2 - np.diag(c)
9
10    # Extract interior system
11    L_int = L[1:N, 1:N]
12    rhs = np.ones(N - 1)
13
14    # Solve and assemble
15    u_int = np.linalg.solve(L_int, rhs)
16    u = np.zeros(N + 1)
17    u[1:N] = u_int
18    return x, u
```

Python

```
1 function [x, u] = solve_variable_coeff(N, coeff_func)
2 % Solve  $u_{xx} - c(x)u = 1$  with  $u(\pm 1) = 0$ .
3 [D, x] = cheb_matrix(N);
4 D2 = D * D;
5
6 % Build operator  $L = D^2 - \text{diag}(c(x))$ 
7 c = coeff_func(x);
8 L = D2 - diag(c);
9
10 % Extract interior system
```

Matlab

```

11     L_int = L(2:N, 2:N);
12     rhs = ones(N-1, 1);
13
14     % Solve and assemble
15     u_int = L_int \ rhs;
16     u = zeros(N+1, 1);
17     u(2:N) = u_int;
18 end

```

The Julia implementation:

```

1  function solve_variable_coeff(N, coeff_func)
2      # Solve  $u_{xx} - c(x)u = 1$  with  $u(\pm 1) = 0$ .
3      D, x = cheb_matrix(N)
4      D2 = D * D
5
6      # Build operator  $L = D^2 - \text{diag}(c(x))$ 
7      c = coeff_func.(x)
8      L = D2 - Diagonal(c)
9
10     # Extract interior system
11     L_int = L[2:N, 2:N]
12     rhs = ones(N - 1)
13
14     # Solve and assemble
15     u_int = L_int \ rhs
16     u = zeros(N + 1)
17     u[2:N] = u_int
18     return x, u
19 end

```

 Julia

The code generating Figure 41 is available in:

- codes/python/ch08/bvp_variable_coeff.py
- codes/matlab/ch08/bvp_variable_coeff.m
- codes/julia/ch08/bvp_variable_coeff.jl

8.5 Mixed Boundary Conditions

8.5.1 Boundary Bordering in Practice

Not all problems come with Dirichlet conditions at both endpoints. Many physical situations involve *Neumann* conditions (prescribed flux or insulation) or a mix of condition types. Consider the steady-state heat equation with a fixed temperature at one end and an insulated boundary at the other:

$$u_{xx} = -\cos(\pi x/2), \quad x \in (-1, 1), \quad u(-1) = 0, \quad u'(1) = 0. \quad (73)$$

The Dirichlet condition $u(-1) = 0$ fixes the temperature at the left boundary, while the Neumann condition $u'(1) = 0$ models perfect insulation at the right boundary (zero heat flux).

The exact solution is

$$u(x) = \frac{4}{\pi^2} \cos(\pi x/2) + \frac{2}{\pi}(x+1), \quad (74)$$

which one can verify satisfies the ODE and both boundary conditions.

Following the boundary bordering approach from Section 8.2.2, we keep the full $(N+1) \times (N+1)$ system and replace boundary rows:

1. Start with $L = D_N^2$ and $\mathbf{f} = f(x)$.
2. **Row 0** (Neumann at $x_0 = 1$): replace $L[0, :] = D_N[0, :]$ and set $f_0 = 0$.
3. **Row N** (Dirichlet at $x_N = -1$): replace $L[N, :] = \mathbf{e}_N^\top$ and set $f_N = 0$.
4. Solve $L\mathbf{u} = \mathbf{f}$.

In Python:

```
1 def solve_mixed_bc(N):
2     """Solve u_xx = f(x) with u(-1) = 0, u'(1) = 0."""
3     D2, D, x = cheb_second_derivative_matrix(N)
4
5     L = D2.copy()
6     rhs = rhs_function(x)
7
8     # Neumann at x[0] = 1: replace with derivative row
9     L[0, :] = D[0, :]
10    rhs[0] = 0.0
11
12    # Dirichlet at x[N] = -1: replace with identity row
13    L[N, :] = 0.0
14    L[N, N] = 1.0
15    rhs[N] = 0.0
16
17    return x, np.linalg.solve(L, rhs)
```

 Python

The equivalent MATLAB implementation:

```
1 function [x, u] = solve_mixed_bc(N, rhs_func)
2     [D, x] = cheb_matrix(N);
3     D2 = D * D;
4
5     L = D2;
6     rhs = rhs_func(x);
7
8     % Neumann at x(1) = 1: replace with derivative row
9     L(1, :) = D(1, :);
10    rhs(1) = 0;
11
12    % Dirichlet at x(N+1) = -1: replace with identity row
13    L(N+1, :) = 0;
14    L(N+1, N+1) = 1;
15    rhs(N+1) = 0;
16
```

 Matlab


```

17     u = L \ rhs;
18 end

```

And in Julia:

```

1  function solve_mixed_bc(N)
2      D2, D, x = cheb_second_derivative_matrix(N)
3
4      L = copy(D2)
5      rhs = rhs_function.(x)
6
7      # Neumann at x[1] = 1: replace with derivative row
8      L[1, :] = D[1, :]
9      rhs[1] = 0.0
10
11     # Dirichlet at x[N+1] = -1: replace with identity row
12     L[N+1, :] .= 0.0
13     L[N+1, N+1] = 1.0
14     rhs[N+1] = 0.0
15
16     return x, L \ rhs
17 end

```

Julia

Figure 42 shows the solution and convergence behavior.




Figure 42: Mixed boundary conditions Equation 73. Left: the solution rises from $u(-1) = 0$ (Dirichlet) to a maximum at $x = 1$ where the slope is zero (Neumann, insulated). Right: spectral convergence reaches approximately 10^{-14} by $N = 16$, after which round-off effects from the large entries in D_N near the boundary cause a slight accuracy plateau.

N	Max Error
4	1.46×10^{-2}
6	2.04×10^{-4}
8	1.70×10^{-6}
10	9.40×10^{-9}
12	3.69×10^{-11}
16	8.06×10^{-14}
20	1.25×10^{-13}
24	3.11×10^{-14}
32	3.79×10^{-13}
48	1.80×10^{-12}

Table 9: Convergence for the mixed BC problem Equation 73. The error decreases spectrally until $N \approx 16$, reaching $\approx 10^{-14}$. Unlike the pure Dirichlet case (Table 8), which reaches machine precision at 10^{-16} , the Neumann condition enforced through the first derivative matrix D_N introduces a slightly higher noise floor. This is because the entries of D_N at the boundary grow as $\mathcal{O}(N^2)$, amplifying round-off errors.

The code generating Figure 42 is available in:

- `codes/python/ch08/bvp_mixed_bc.py`
- `codes/matlab/ch08/bvp_mixed_bc.m`
- `codes/julia/ch08/bvp_mixed_bc.jl`

8.6 Nonlinear BVP: The Bratu Equation

8.6.1 A Classic Nonlinear Problem

The Bratu equation models combustion and thermal explosion. The equation was originally derived by Frank-Kamenetskii [49] in the context of thermal ignition theory, where the exponential nonlinearity models the Arrhenius temperature dependence of reaction rates:

$$u_{xx} + \lambda e^u = 0, \quad x \in (-1, 1), \quad u(\pm 1) = 0. \quad (75)$$

This equation exhibits a *turning point* phenomenon: solutions exist only for $\lambda \leq \lambda_c$, where $\lambda_c \approx 0.878$ for the domain $[-1, 1]$. Boyd [20] produced the definitive spectral analysis of this bifurcation, using arc-length continuation to resolve the turning point singularity and compute λ_c to high precision. Above this critical value, no solution exists. For $\lambda = 0.5$ (well below the critical value), a unique solution exists.

The nonlinearity e^u prevents direct linear algebra. Instead, we use *Newton iteration*: linearize, solve, update, repeat.

8.6.2 Newton Iteration

To solve a nonlinear equation, we cannot simply invert a matrix. Instead, we employ an *iterative* strategy: start from an initial guess, and progressively refine it until the solution is accurate enough.

Newton's method is the most widely used approach for nonlinear equations — a comprehensive treatment, including convergence theory and affine invariance, can be found in Deuffhard [32]. The key idea is *linearization*: at each iteration, we approximate the nonlinear problem by a linear one, solve that linear problem exactly, and use the result to improve our current approximation. Concretely, suppose we have a current approximation $\mathbf{u}^{(k)}$. We seek a correction $\delta\mathbf{u}$ such that $\mathbf{u}^{(k)} + \delta\mathbf{u}$ satisfies the equation more closely. Expanding the residual to first order:

$$\mathbf{F}(\mathbf{u}^{(k)} + \delta\mathbf{u}) \approx \mathbf{F}(\mathbf{u}^{(k)}) + J(\mathbf{u}^{(k)})\delta\mathbf{u} = \mathbf{0}, \quad (76)$$

where $J = \partial\mathbf{F}/\partial\mathbf{u}$ is the *Jacobian matrix*. Setting this to zero gives the Newton step:

$$J(\mathbf{u}^{(k)})\delta\mathbf{u} = -\mathbf{F}(\mathbf{u}^{(k)}). \quad (77)$$

For our Bratu problem, the discrete residual is $\mathbf{F}(\mathbf{u}) = \tilde{D}^2\mathbf{u} + \lambda e^{\mathbf{u}}$, and the Jacobian is:

$$J(\mathbf{u}) = \tilde{D}^2 + \lambda \operatorname{diag}(e^{\mathbf{u}}). \quad (78)$$

Note that the Jacobian changes at every iteration because $\operatorname{diag}(e^{\mathbf{u}})$ depends on the current solution.

The complete algorithm is:

1. Choose an initial guess $\mathbf{u}^{(0)}$ (we use $\mathbf{u}^{(0)} = \mathbf{0}$).
2. For $k = 0, 1, 2, \dots$:
 1. Compute the residual $\mathbf{F}^{(k)} = \tilde{D}^2\mathbf{u}^{(k)} + \lambda e^{\mathbf{u}^{(k)}}$.
 2. Compute the Jacobian $J^{(k)} = \tilde{D}^2 + \lambda \operatorname{diag}(e^{\mathbf{u}^{(k)}})$.
 3. Solve the linear system $J^{(k)}\delta\mathbf{u} = -\mathbf{F}^{(k)}$ for the Newton correction $\delta\mathbf{u}$.
 4. Update: $\mathbf{u}^{(k+1)} = \mathbf{u}^{(k)} + \delta\mathbf{u}$.
 5. If the stopping criterion is satisfied, terminate.

Stopping criterion. A critical practical detail is deciding when to stop iterating. Common choices include:

- *Correction norm*: $\|\delta\mathbf{u}\|_{\infty} < \text{tol}$, i.e., the Newton step becomes negligibly small.
- *Residual norm*: $\|\mathbf{F}^{(k)}\|_{\infty} < \text{tol}$, i.e., the equation is satisfied to the desired accuracy.
- *Relative correction*: $\|\delta\mathbf{u}\|_{\infty}/\|\mathbf{u}^{(k)}\|_{\infty} < \text{tol}$, which accounts for the magnitude of the solution.

In our implementation, we use the correction norm $\|\delta\mathbf{u}\|_{\infty} < 10^{-10}$ as the stopping criterion, combined with a maximum iteration count of 50 as a safety guard against non-convergence. This is a natural choice: when the correction becomes smaller than 10^{-10} , the solution has stabilized to at least 10 significant digits.

Convergence rate. Newton's method enjoys *quadratic convergence* when sufficiently close to a solution: the error at each step is roughly the square of the error at the previous step. This means that once the iteration begins to converge, the number of correct digits approximately doubles with each step. In practice, for $\lambda = 0.5$, Newton's method converges in about 5 to 8 iterations starting from $\mathbf{u}^{(0)} = \mathbf{0}$, as shown in Figure 43.

In the contemporary era, interest has shifted to fractional-order generalisations of the Bratu equation. Salama *et al.* [105] have extended spectral collocation using shifted Lucas polynomials to solve fractional Bratu equations, modelling non-local memory effects in combustion. Gheorghiu [53] revisited the classical problem to benchmark the accuracy of modern Newton–Kantorovich solvers, achieving spectral accuracy with as few as 8–16 grid points.

Figure 43 shows the solution and convergence behavior.


Nonlinear BVP: Bratu Equation with $\lambda = 0.5$ 

Figure 43: Nonlinear BVP: the Bratu equation Equation 75 with $\lambda = 0.5$. Left: solutions for different values of N . Right: convergence history comparing Newton iteration (quadratic convergence) with fixed-point iteration (linear convergence). Newton typically converges in 5–8 iterations.

The Newton iteration is straightforward to implement:

```

1 def solve_bratu_newton(N, lam=0.5, tol=1e-10, max_iter=50):
2     """Solve  $u_{xx} + \lambda \exp(u) = 0$  with  $u(\pm 1) = 0$  using Newton iteration."""
3     D, x = cheb_matrix(N)
4     D2 = D @ D
5     D2_int = D2[1:N, 1:N]
6
7     u = np.zeros(N - 1) # Initial guess
8
9     for k in range(max_iter):
10         exp_u = np.exp(u)
11         # Residual:  $F = D^2 u + \lambda \exp(u)$ 
12         F = D2_int @ u + lam * exp_u
13         # Jacobian:  $J = D^2 + \lambda \text{diag}(\exp(u))$ 
14         J = D2_int + lam * np.diag(exp_u)
15         # Newton step
16         delta_u = np.linalg.solve(J, -F)
17         u = u + delta_u
18
19         if np.max(np.abs(delta_u)) < tol:
20             break
21
22     # Assemble full solution
23     u_full = np.zeros(N + 1)
24     u_full[1:N] = u
25     return x, u_full, k + 1

```

Python

```

1 function [x, u_full, iterations] = solve_bratu_newton(N, lam, tol,
2 max_iter)
3 % Solve  $u_{xx} + \lambda \exp(u) = 0$  with  $u(\pm 1) = 0$  using Newton iteration.

```

Matlab

```

3     if nargin < 4, max_iter = 50; end
4     if nargin < 3, tol = 1e-10; end
5     if nargin < 2, lam = 0.5; end
6
7     [D, x] = cheb_matrix(N);
8     D2 = D * D;
9     D2_int = D2(2:N, 2:N);
10
11    u = zeros(N-1, 1); % Initial guess
12
13    for k = 1:max_iter
14        exp_u = exp(u);
15        % Residual: F = D^2 u + λ*exp(u)
16        F = D2_int * u + lam * exp_u;
17        % Jacobian: J = D^2 + λ*diag(exp(u))
18        J = D2_int + lam * diag(exp_u);
19        % Newton step
20        delta_u = J \ (-F);
21        u = u + delta_u;
22
23        if max(abs(delta_u)) < tol
24            break
25        end
26    end
27
28    % Assemble full solution
29    u_full = zeros(N+1, 1);
30    u_full(2:N) = u;
31    iterations = k;
32 end


```

The Julia implementation:

```

1  function solve_bratu_newton(N; lam=0.5, tol=1e-10, max_iter=50)
2      # Solve  $u_{xx} + \lambda \exp(u) = 0$  with  $u(\pm 1) = 0$  using Newton iteration.
3      D, x = cheb_matrix(N)
4      D2 = D * D
5      D2_int = D2[2:N, 2:N]
6
7      u = zeros(N - 1) # Initial guess
8
9      for k in 1:max_iter
10         exp_u = exp.(u)
11         # Residual and Jacobian
12         F = D2_int * u + lam * exp_u
13         J = D2_int + lam * Diagonal(exp_u)
14         # Newton step

```



```

15     delta_u = J \ (-F)
16     u .+= delta_u
17
18     maximum(abs.(delta_u)) < tol && break
19 end
20
21 # Assemble full solution
22 u_full = zeros(N + 1)
23 u_full[2:N] = u
24 return x, u_full
25 end

```

The code generating Figure 43 is available in:

- codes/python/ch08/bvp_nonlinear.py
- codes/matlab/ch08/bvp_nonlinear.m
- codes/julia/ch08/bvp_nonlinear.jl

8.7 Eigenvalue Problems

8.7.1 Resolution Limits

The eigenvalue problem for the Laplacian,

$$u_{xx} = \lambda u, \quad x \in (-1, 1), \quad u(\pm 1) = 0, \quad (79)$$

has exact eigenvalues $\lambda_n = -(n\pi/2)^2$ and eigenfunctions $u_n(x) = \sin(n\pi(x+1)/2)$.

Spectral methods compute these eigenvalues with spectral accuracy, *but only for the low modes*. High-frequency modes require many points per wavelength (ppw) for accurate representation. The rule of thumb is:

$$\text{Need } \geq \pi \text{ points per wavelength for spectral accuracy.} \quad (80)$$

This resolution criterion was established in the foundational work of Kreiss and Oliger [73], who demonstrated the superior efficiency of spectral methods over finite differences for wave propagation. Recent dispersion analysis by Gkanos *et al.* [54] confirms that while spectral methods minimise phase lag significantly better than finite differences, the “ π points” rule is an asymptotic stability limit rather than an accuracy guarantee for large domains.

Figure 44 demonstrates this resolution limit.

Eigenvalue Problem: $u_{xx} = \lambda u$, $u(\pm 1) = 0$ ($N = 36$)




Figure 44: Eigenvalue problem Equation 79 for $N = 36$. Six eigenmodes are shown with their points per wavelength (ppw) and eigenvalue errors. Low modes (high ppw) are resolved to near machine precision. As ppw drops below $\pi \approx 3.14$, accuracy degrades rapidly. This is not a bug but a fundamental resolution limit.

Table 10 quantifies this resolution effect. For the first few modes, the points per wavelength is large and the spectral method delivers near machine precision. As the mode number increases and the ppw drops toward $\pi \approx 3.14$, accuracy degrades sharply. Beyond this threshold, the computed eigenvalues bear little resemblance to the exact values.

n	ppw	λ_n (exact)	λ_n (numerical)	Absolute Error	Relative Error
1	72.00	-2.4674	-2.4674	5.68×10^{-14}	2.30×10^{-14}
5	14.40	-61.6850	-61.6850	2.77×10^{-13}	4.49×10^{-15}
10	7.20	-246.7401	-246.7401	5.56×10^{-11}	2.25×10^{-13}
15	4.80	-555.1652	-555.1653	1.97×10^{-5}	3.56×10^{-8}
20	3.60	-986.9604	-988.0301	1.07×10^0	1.08×10^{-3}
23	3.13	-1305.2552	-1270.1437	3.51×10^1	2.69×10^{-2}
25	2.88	-1542.1257	-1567.3682	2.52×10^1	1.64×10^{-2}
30	2.40	-2220.6610	-5860.9103	3.64×10^3	1.64×10^0
35	2.06	-3022.5663	-79 900.13	7.69×10^4	2.54×10^1

Table 10: Eigenvalue accuracy for $N = 36$. The points per wavelength (ppw) equals $2N/n$. For modes with $\text{ppw} \geq \pi$, both absolute and relative errors are near machine precision. The transition at $\text{ppw} \approx \pi$ is remarkably sharp: mode 23 ($\text{ppw} = 3.13 \approx \pi$) already shows $\approx 3\%$ relative error, and beyond this threshold accuracy collapses completely.

The eigenvalue computation uses standard linear algebra:

```

1  def compute_laplacian_eigenvalues(N):
2      """Compute eigenvalues of  $u_{xx} = \lambda u$  with  $u(\pm 1) = 0$ ."""
3      D, x = cheb_matrix(N)
4      D2 = D @ D
5      D2_int = D2[1:N, 1:N]
6
7      # Compute eigenvalues and eigenvectors
8      eigenvalues, eigenvectors = np.linalg.eig(D2_int)
9
10     # Sort by magnitude (most negative first)
11     idx = np.argsort(eigenvalues)
12     eigenvalues = eigenvalues[idx]
13     eigenvectors = eigenvectors[:, idx]
14
15     # Exact eigenvalues:  $\lambda_n = -(n\pi/2)^2$ 
16     n = np.arange(1, N)
17     exact_eigenvalues = -(n * np.pi / 2)**2
18
19     return eigenvalues, eigenvectors, exact_eigenvalues

```

 Python

```

1  function [eigenvalues, eigenvectors, exact_eigenvalues] =
2  compute_laplacian_eigenvalues(N)
3      % Compute eigenvalues of  $u_{xx} = \lambda u$  with  $u(\pm 1) = 0$ .
4      [D, x] = cheb_matrix(N);
5      D2 = D * D;
6      D2_int = D2(2:N, 2:N);
7
8      % Compute eigenvalues and eigenvectors
9      [eigenvectors, Lambda] = eig(D2_int);
10     eigenvalues = diag(Lambda);
11
12     % Sort by magnitude (most negative first)
13     [eigenvalues, idx] = sort(eigenvalues);
14     eigenvectors = eigenvectors(:, idx);
15
16     % Exact eigenvalues:  $\lambda_n = -(n\pi/2)^2$ 
17     n = (1:N-1)';
18     exact_eigenvalues = -(n * pi / 2).^2;
19 end

```

 Matlab

The Julia implementation:

```

1  function compute_laplacian_eigenvalues(N)
2      # Compute eigenvalues of  $u_{xx} = \lambda u$  with  $u(\pm 1) = 0$ .
3      D, x = cheb_matrix(N)
4      D2 = D * D
5      D2_int = D2[2:N, 2:N]
6

```

 Julia


```

7      # Compute eigenvalues and eigenvectors
8      eig = eigen(D2_int)
9      idx = sortperm(real.(eig.values))
10     eigenvalues = real.(eig.values[idx])
11     eigenvectors = real.(eig.vectors[:, idx])
12
13     # Exact eigenvalues:  $\lambda_n = -(\pi/2)^2$ 
14     exact_eigenvalues = [-(n *  $\pi$  / 2)^2 for n in 1:N-1]
15     return eigenvalues, eigenvectors, exact_eigenvalues
16 end

```

The code generating Figure 44 is available in:

- codes/python/ch08/bvp_eigenvalue.py
- codes/matlab/ch08/bvp_eigenvalue.m
- codes/julia/ch08/bvp_eigenvalue.jl

8.8 Two-Dimensional Problems

8.8.1 Tensor Products

For problems on a rectangle $[-1, 1]^2$, we use *tensor product* grids: Chebyshev points in both x and y directions. The total number of grid points is $(N + 1)^2$.

Remark (Padua points). Tensor product grids are not the only option for bivariate polynomial interpolation on the square. The *Padua points*, discovered by De Marchi, Caliari, and Vianello [83], are the first known example (and to date the only one) of a unisolvent point set on $[-1, 1]^2$ with *minimal growth* of the Lebesgue constant, proven to be $\mathcal{O}(\log^2 n)$ by Bos, Caliari, De Marchi, Vianello, and Xu [17]. For total polynomial degree n , Padua points require only $(n + 1)(n + 2)/2$ points, roughly half the $(n + 1)^2$ needed by the tensor product grid. The points admit an elegant geometric construction: they lie exactly on the self-intersections and boundary contacts of a generating Lissajous curve on $[-1, 1]^2$, which gives rise to four distinct families obtained by successive 90-degree rotations [18]. An efficient implementation is available as Algorithm 886 [25]. While we use tensor product grids throughout this chapter for their simplicity and natural compatibility with the Kronecker product structure of differential operators, Padua points offer a more economical alternative for pure interpolation and approximation problems.

The 2D Laplacian operator is built using *Kronecker products*:

$$L = I \otimes D^2 + D^2 \otimes I, \quad (81)$$

where I is the identity matrix and $\otimes$ denotes the Kronecker product.

Figure 45 shows the tensor product grid structure.




Figure 45: Chebyshev tensor product grid for $N = 16$. Points cluster near all four boundaries, providing the boundary resolution needed for spectral accuracy in two dimensions.

8.8.2 2D Poisson Problem

Consider the 2D Poisson equation:

$$u_{xx} + u_{yy} = -2\pi^2 \sin(\pi x) \sin(\pi y), \quad u = 0 \text{ on boundary.} \quad (82)$$

The exact solution is $u(x, y) = \sin(\pi x) \sin(\pi y)$.

Figure 46 shows the solution.

$$\text{2D Poisson: } \nabla^2 u = -2\pi^2 \sin(\pi x) \sin(\pi y)$$




Figure 46: Solution of the 2D Poisson equation Equation 82. Left: 3D surface plot of the solution. Right: contour plot with Chebyshev grid points overlaid. For $N = 16$, the maximum error is approximately 10^{-12} .

The Kronecker product formulation leads to concise code:

```

1  def solve_poisson_2d(N, f):
2      """Solve  $\nabla^2 u = f$  on  $[-1,1]^2$  with  $u = 0$  on boundary."""
3      D, x = cheb_matrix(N)
4      D2 = D @ D
5      D2_int = D2[1:N, 1:N]
6      x_int = x[1:N]
7
8      # Build 2D Laplacian:  $L = I \otimes D^2 + D^2 \otimes I$ 
9      I = np.eye(N - 1)
10     L = np.kron(I, D2_int) + np.kron(D2_int, I)
11
12     # Right-hand side on interior grid
13     X, Y = np.meshgrid(x_int, x_int)
14     F = f(X, Y).flatten(order='F')
15
16     # Solve and reshape
17     u_vec = np.linalg.solve(L, F)
18     U_int = u_vec.reshape((N-1, N-1), order='F')
19
20     # Embed in full grid with zero boundary
21     U = np.zeros((N+1, N+1))
22     U[1:N, 1:N] = U_int
23     return np.meshgrid(x, x), U

```

 Python

```

1  function [grids, U] = solve_poisson_2d(N, f)
2      % Solve  $\nabla^2 u = f$  on  $[-1,1]^2$  with  $u = 0$  on boundary.
3      [D, x] = cheb_matrix(N);
4      D2 = D * D;
5      D2_int = D2(2:N, 2:N);
6      x_int = x(2:N);
7
8      % Build 2D Laplacian:  $L = I \otimes D^2 + D^2 \otimes I$ 
9      I = eye(N - 1);
10     L = kron(I, D2_int) + kron(D2_int, I);
11
12     % Right-hand side on interior grid
13     [X, Y] = meshgrid(x_int, x_int);
14     F = f(X, Y);
15
16     % Solve and reshape
17     u_vec = L \ F(:);
18     U_int = reshape(u_vec, N-1, N-1);
19
20     % Embed in full grid with zero boundary
21     U = zeros(N+1, N+1);
22     U(2:N, 2:N) = U_int;

```

 Matlab

```

23     [grids{1}, grids{2}] = meshgrid(x, x);
24 end

```

The Julia implementation:

```

1  function solve_poisson_2d(N, f)
2      # Solve  $\nabla^2 u = f$  on  $[-1,1]^2$  with  $u = 0$  on boundary.
3      D, x = cheb_matrix(N)
4      D2 = D * D
5      D2_int = D2[2:N, 2:N]
6      x_int = x[2:N]
7      n_int = N - 1
8
9      # Build 2D Laplacian:  $L = I \otimes D^2 + D^2 \otimes I$ 
10     Imat = I(n_int)
11     L = kron(Imat, D2_int) + kron(D2_int, Imat)
12
13     # Right-hand side on interior grid
14     F = [f(x_int[j], x_int[i]) for i in 1:n_int, j in 1:n_int]
15
16     # Solve and reshape
17     u_vec = L \ vec(F)
18     U_int = reshape(u_vec, n_int, n_int)
19
20     # Embed in full grid with zero boundary
21     U = zeros(N + 1, N + 1)
22     U[2:N, 2:N] = U_int
23     return x, U
24 end

```

The sparsity pattern of the 2D Laplacian operator, shown in Figure 47, reveals the Kronecker product structure.




Figure 47: Sparsity pattern of the 2D Laplacian matrix for $N = 16$. The $(N - 1)^2 \times (N - 1)^2$ matrix shows the characteristic block structure arising from the Kronecker product formulation Equation 81.

Despite the visual impression of sparsity in Figure 47, a closer look reveals an important structural property: the Chebyshev second derivative matrix D^2 is *fully dense* (every collocation point couples to every other). The apparent sparsity arises entirely from the Kronecker product structure. The term $I \otimes D^2$ contributes $(N - 1)$ dense diagonal blocks of size $(N - 1) \times (N - 1)$, while $D^2 \otimes I$ scatters $(N - 1) \times (N - 1)$ identity-sized blocks across the off-diagonal positions. By inclusion-exclusion (subtracting the $(N - 1)^2$ diagonal entries counted twice), the exact number of nonzero entries is

$$\text{nnz}(L) = (N - 1)^2(2(N - 1) - 1) = (N - 1)^2(2N - 3), \quad (83)$$

so that each row contains exactly $2(N - 1) - 1$ nonzeros.

Table 11 quantifies how the sparsity evolves with N .

N	Matrix Size	Nonzeros	Density (%)	nnz/row
8	49×49	637	26.5	13
12	121×121	2541	17.4	21
16	225×225	6525	12.9	29
20	361×361	13 357	10.2	37
24	529×529	23 805	8.5	45
32	961×961	58 621	6.3	61

Table 11: Sparsity statistics for the 2D Laplacian matrix $L = I \otimes D^2 + D^2 \otimes I$. While the density decreases as $\mathcal{O}(1/N)$, each row has $2(N - 1) - 1$ nonzeros, growing linearly with N . For comparison, a standard second-order finite difference stencil on the same grid has exactly 5 nonzeros per row regardless of N . This denser coupling is the price paid for spectral (exponential) convergence.

The code generating Figure 46 is available in:

- codes/python/ch08/bvp_2d_poisson.py
- codes/matlab/ch08/bvp_2d_poisson.m
- codes/julia/ch08/bvp_2d_poisson.jl

8.9 The Helmholtz Equation

8.9.1 Near-Resonance Behavior

The Helmholtz equation models wave phenomena:

$$u_{xx} + u_{yy} + k^2 u = f(x, y), \quad u = 0 \text{ on boundary.} \quad (84)$$

When k^2 approaches an eigenvalue of the Laplacian, the system becomes *nearly resonant* and the solution amplitude grows dramatically. The near-singularity of the Helmholtz operator near resonance is a manifestation of the *pollution effect* analysed by Babuska and Sauter [6], who showed that maintaining accuracy at high wavenumbers requires $N \sim k^{3/2}$ or even $N \sim k^2$ rather than the naïve $N \approx k/\pi$. Moiola and Spence [89] further clarified the sign-indefinite nature of the Helmholtz operator, explaining why standard iterative solvers stall near resonance.

For a localized Gaussian forcing $f(x, y) = e^{-20[(x-0.3)^2 + (y+0.4)^2]}$ and $k = 7$, we are near resonance with the $(2, 4)$ mode (theoretical $k \approx 7.02$).

Figure 48 shows the characteristic modal structure of the near-resonant solution.




Figure 48: Helmholtz equation Equation 84 with $k = 7$, near resonance with the $(2, 4)$ eigenmode ($k_{2,4} \approx 7.02$). Left: 3D surface showing the wave-like structure. Right: contour plot with forcing location marked. The solution exhibits the characteristic pattern of the $(2, 4)$ eigenmode.

The Helmholtz solver modifies the 2D Laplacian by adding the $k^2 I$ term:

```
1 def solve_helmholtz(N, k, f):
2     """Solve  $\nabla^2 u + k^2 u = f$  on  $[-1, 1]^2$  with  $u = 0$  on boundary."""
3     D, x = cheb_matrix(N)
```

 Python

```

4     D2 = D @ D
5     D2_int = D2[1:N, 1:N]
6     x_int = x[1:N]
7
8     # Build Helmholtz operator:  $L = I \otimes D^2 + D^2 \otimes I + k^2 I$ 
9     I = np.eye(N - 1)
10    n_int = (N - 1)**2
11    L = np.kron(I, D2_int) + np.kron(D2_int, I) + k**2 * np.eye(n_int)
12
13    # Right-hand side
14    X, Y = np.meshgrid(x_int, x_int)
15    F = f(X, Y).flatten(order='F')
16
17    # Solve
18    u_vec = np.linalg.solve(L, F)
19    U_int = u_vec.reshape((N-1, N-1), order='F')
20
21    U = np.zeros((N+1, N+1))
22    U[1:N, 1:N] = U_int
23    return np.meshgrid(x, x), U

```

```

1  function [grids, U] = solve_helmholtz(N, k, f)
2  % Solve  $\nabla^2 u + k^2 u = f$  on  $[-1,1]^2$  with  $u = 0$  on boundary.
3      [D, x] = cheb_matrix(N);
4      D2 = D * D;
5      D2_int = D2(2:N, 2:N);
6      x_int = x(2:N);
7
8      % Build Helmholtz operator:  $L = I \otimes D^2 + D^2 \otimes I + k^2 I$ 
9      I = eye(N - 1);
10     n_int = (N - 1)^2;
11     L = kron(I, D2_int) + kron(D2_int, I) + k^2 * eye(n_int);
12
13     % Right-hand side
14     [X, Y] = meshgrid(x_int, x_int);
15     F = f(X, Y);
16
17     % Solve
18     u_vec = L \ F(:);
19     U_int = reshape(u_vec, N-1, N-1);
20
21     U = zeros(N+1, N+1);
22     U(2:N, 2:N) = U_int;
23     [grids{1}, grids{2}] = meshgrid(x, x);
24 end

```

 Matlab

The Julia implementation:

```

1  function solve_helmholtz(N, k, f)
2      # Solve  $\nabla^2 u + k^2 u = f$  on  $[-1,1]^2$  with  $u = 0$  on boundary.
3      D, x = cheb_matrix(N)
4      D2 = D * D
5      D2_int = D2[2:N, 2:N]
6      x_int = x[2:N]
7      n_int = N - 1
8
9      # Build Helmholtz operator:  $L = I \otimes D^2 + D^2 \otimes I + k^2 I$ 
10     Imat = I(n_int)
11     L = kron(Imat, D2_int) + kron(D2_int, Imat) + k^2 * I(n_int^2)
12
13     # Right-hand side
14     F = [f(x_int[j], x_int[i]) for i in 1:n_int, j in 1:n_int]
15
16     # Solve
17     u_vec = L \ vec(F)
18     U_int = reshape(u_vec, n_int, n_int)
19
20     U = zeros(N + 1, N + 1)
21     U[2:N, 2:N] = U_int
22     return x, U
23 end

```

The code generating Figure 48 is available in:

- codes/python/ch08/bvp_helmholtz.py
- codes/matlab/ch08/bvp_helmholtz.m
- codes/julia/ch08/bvp_helmholtz.jl

8.10 Computational Étude: The Quantum Harmonic Oscillator Revisited

In Section 5.7, we solved the quantum harmonic oscillator eigenvalue problem

$$-u'' + x^2 u = \lambda u, \quad x \in \mathbb{R} \quad (85)$$

using the periodic (Fourier) spectral method. We now revisit this same problem using the Chebyshev spectral methods developed in this chapter, and compare the two approaches. The comparison of Fourier (equispaced) and Chebyshev (clustered) grids for problems on unbounded domains is a recurring theme in spectral methods; Boyd [21] discusses the interplay between domain truncation and basis choice, noting that when eigenfunctions decay rapidly, Fourier grids can outperform Chebyshev grids despite the latter's theoretical optimality on bounded domains.

8.10.1 Chebyshev Approach

Recall that the exact eigenvalues are $\lambda_n = 2n + 1$ and the eigenfunctions are Hermite functions that decay like $e^{-x^2/2}$. We truncate to $[-L, L]$ with homogeneous Dirichlet conditions $u(\pm L) = 0$, which are automatically satisfied by the true eigenfunctions for large L .

The Chebyshev method proceeds as follows:

1. Compute the Chebyshev–Gauss–Lobatto points $t_j = \cos(j\pi/N)$ on $[-1, 1]$ and the differentiation matrix D_N .
2. Map to the physical domain: $x_j = Lt_j$, giving points on $[-L, L]$.
3. Form the second-derivative matrix $D^{(2)} = D_N^2/L^2$ (the factor $1/L^2$ accounts for the chain rule).
4. Strip boundary rows and columns (matrix stripping, as in Section 8.2) to obtain the interior operator.
5. Assemble $A = -D_{\text{int}}^{(2)} + \text{diag}(x_1^2, \dots, x_{N-1}^2)$.
6. Compute eigenvalues of A .

In Python, using the Chebyshev matrix from Section 7.2:

```
1 def solve_chebyshev(N, L):
2     """Solve -u'' + x^2u = λu on [-L, L] using Chebyshev."""
3     D, t = cheb_matrix(N)
4     x = L * t                # Map to [-L, L]
5     D2 = (D @ D) / L**2      # Second derivative
6
7     # Strip boundaries (matrix stripping)
8     D2_int = D2[1:N, 1:N]
9     x_int = x[1:N]
10
11     A = -D2_int + np.diag(x_int**2)
12     eigenvalues, eigenvectors = np.linalg.eig(A)
13     idx = np.argsort(np.real(eigenvalues))
14     return np.real(eigenvalues[idx]), np.real(eigenvectors[:, idx]), x
```

 Python

The equivalent MATLAB implementation:

```
1 function [eigenvalues, eigvecs, x] = solve_chebyshev(N, L)
2     [D, t] = cheb_matrix(N);
3     x = L * t;                % Map to [-L, L]
4     D2 = (D * D) / L^2;      % Second derivative
5
6     % Strip boundaries
7     D2_int = D2(2:N, 2:N);
8     x_int = x(2:N);
9
10    A = -D2_int + diag(x_int.^2);
11    [eigvecs, E] = eig(A);
12    eigenvalues = sort(diag(E));
13 end
```

 Matlab

And in Julia:

```
1 function solve_chebyshev(N, L)
2     D, t = cheb_matrix(N)
```

 Julia

```

3      x = L * t                                # Map to [-L, L]
4      D2 = (D * D) / L^2                       # Second derivative
5
6      # Strip boundaries
7      D2_int = D2[2:N, 2:N]
8      x_int = x[2:N]
9
10     A = -D2_int + Diagonal(x_int.^2)
11     eig_result = eigen(Matrix(A))
12     return sort(real.(eig_result.values)), eig_result.vectors, x
13 end

```

8.10.2 Results

Figure 49 shows the Chebyshev method in action. The left panel displays the first five eigenfunctions computed with $N = 64$ Chebyshev points on $[-8, 8]$, overlaid with the exact Hermite functions. The agreement is visually perfect.




Figure 49: The quantum harmonic oscillator solved with Chebyshev spectral methods. Left: First five eigenfunctions computed with $N = 64$ Chebyshev–Gauss–Lobatto points on $[-8, 8]$ (dots) compared to exact Hermite functions (lines). Right: Eigenvalue convergence comparison between the Fourier and Chebyshev methods for λ_0 and λ_1 . Both methods achieve spectral convergence, but the Fourier approach reaches machine precision with fewer points.

8.10.3 Fourier vs Chebyshev: A Comparison

The right panel of Figure 49 compares the convergence of the Fourier and Chebyshev methods on the same problem with $L = 8$. Table 12 presents the eigenvalue errors for both approaches.

N	Fourier λ_0	Chebyshev λ_0	Fourier λ_1	Chebyshev λ_1
12	2.19×10^{-2}	3.39×10^{-1}	1.72×10^{-1}	1.96×10^0
18	3.00×10^{-5}	3.87×10^{-2}	6.44×10^{-4}	2.54×10^{-1}
24	2.37×10^{-9}	1.20×10^{-3}	9.84×10^{-8}	1.49×10^{-2}
30	1.77×10^{-14}	1.44×10^{-5}	7.54×10^{-13}	2.95×10^{-4}
36	6.00×10^{-15}	7.38×10^{-8}	6.66×10^{-15}	2.17×10^{-6}
42	2.11×10^{-14}	1.75×10^{-10}	2.80×10^{-14}	6.82×10^{-9}
48	1.72×10^{-14}	2.01×10^{-13}	1.02×10^{-14}	1.01×10^{-11}

Table 12: Eigenvalue errors for the quantum harmonic oscillator: Fourier vs Chebyshev methods with $L = 8$. Both methods achieve machine precision, but the Fourier method converges faster for this particular problem.

The comparison reveals an interesting phenomenon: the Fourier method converges *faster* than the Chebyshev method for this problem, despite being designed for periodic functions. This can be understood by considering the nature of the eigenfunctions and the grid point distributions:

- **Eigenfunctions decay rapidly.** The Hermite functions are effectively zero beyond $|x| \approx 5$ for the lowest modes. On $[-8, 8]$, the solution is essentially zero near the boundaries, making it *effectively periodic*.
- **Equispaced points sample uniformly.** The Fourier method places points uniformly across $[-L, L]$, providing resolution where the eigenfunctions are non-negligible.
- **Chebyshev points cluster near boundaries.** The Chebyshev–Gauss–Lobatto points accumulate near $x = \pm L$, precisely where the eigenfunctions are vanishingly small. This “wastes” resolution on regions where there is nothing to resolve.

The Chebyshev method remains more general: it handles any non-periodic boundary value problem, including those where the solution has significant structure near the boundaries. The Fourier method exploits the specific structure of this problem. This comparison illustrates that the best numerical method depends on the problem at hand; matching the method to the solution structure can yield significant efficiency gains.

The code generating Figure 49 is available in:

- `codes/python/ch08/harmonic_oscillator.py`
- `codes/matlab/ch08/harmonic_oscillator.m`
- `codes/julia/ch08/harmonic_oscillator.jl`

8.11 A non-exhaustive literature overview

The application of spectral methods to boundary value problems is a narrative of overcoming the rigidity of the Fast Fourier Transform to handle non-periodic constraints. The foundational shift from the integral-based Galerkin formulations of the early 20th century — pioneered by Galerkin [51] and formalised by Lanczos [74] — to the differential point-wise approach used in this chapter occurred in the 1970s. The seminal works of Kreiss and Oliger [73] and Orszag [94] established the viability of the collocation (pseudospectral) method, demonstrating that satisfying the differential equation

exactly at grid points could yield exponential convergence without the computational burden of convolution sums. The construction of the differentiation matrix, the algebraic heart of this chapter, was systematised by Weideman and Reddy [128] in their MATLAB Differentiation Matrix Suite, which provided not only the explicit algorithms for Chebyshev and Hermite matrices but also the rigorous error analysis justifying the “negative sum trick” to mitigate round-off errors. The debate between computing second derivatives via matrix squaring (D^2) versus explicit recurrence formulas was settled for practical purposes by Baltensperger and Berrut [8], who showed that for the moderate resolutions ($N < 200$) typical of spectral methods, the operational simplicity of matrix squaring outweighs theoretical stability concerns.

The technique of boundary bordering — replacing rows of the differentiation matrix with boundary constraints — was codified by Boyd [21] in his influential text on Chebyshev and Fourier spectral methods. While alternative methods like basis recombination (constructing basis functions that inherently satisfy the boundary conditions) offer better conditioning for high-order problems, as demonstrated by Shen [108], boundary bordering remains the standard for general-purpose solvers due to its flexibility in handling Robin and mixed conditions. The chapter’s “matrix surgery” approach — stripping boundary rows and columns for Dirichlet problems — follows the pedagogical tradition established by Gottlieb and Orszag [55] and refined by Trefethen [120], who emphasised the elegance of reducing differential equations to small, dense linear systems solvable by direct methods.

For nonlinear problems, the Bratu equation serves as the historical benchmark. Originally derived by Frank-Kamenetskii [49] in the context of thermal ignition theory, this equation was given its definitive spectral treatment by Boyd [20], who used Chebyshev spectral methods with arc-length continuation to resolve the turning point singularity and compute the critical parameter $\lambda_c \approx 3.513830719$ to high precision. In the contemporary era (2020–2026), the focus has shifted to fractional-order generalisations. Salama *et al.* [105] have extended spectral collocation using shifted Lucas polynomials to solve fractional Bratu equations, modelling non-local memory effects in combustion. Gheorghiu [53] revisited the classical problem to benchmark the accuracy of modern Newton–Kantorovich solvers, reconfirming the method’s unrivalled efficiency for smooth nonlinearities. A comprehensive treatment of Newton methods for nonlinear PDEs, including convergence theory and affine invariance, can be found in the monograph by Deuflhard [32].

In the realm of wave propagation, the Helmholtz equation reveals the subtle limitations of spectral accuracy. The “ π points per wavelength” rule is necessary but not sufficient for high-frequency problems due to the *pollution effect*, a phenomenon rigorously analysed by Babuska and Sauter [6]. Recent work by Gkanos *et al.* [54] utilises dispersion analysis to quantify these errors in complex acoustical environments, while Galkowski and Spence [52] have applied semiclassical analysis to the high-frequency Helmholtz equation, providing theoretical upper bounds on the grid densities required to resolve scattering from non-convex obstacles. Moiola and Spence [89] clarified the sign-indefinite nature of the Helmholtz operator, explaining why standard iterative solvers stall near resonance and motivating the development of complex-shifted Laplacian preconditioners.

For two-dimensional problems, the chapter’s tensor product approach using Kronecker products is efficient ($O(N^3)$) but restricted to rectangular geometries. The mention of Padua points connects to a major breakthrough in approximation theory: De Marchi, Caliari, and Vianello [83] discovered these as the first explicit unisolvent point set on the square with optimal Lebesgue constant growth ($O(\log^2 N)$), later proven rigorously by Bos *et al.* [17]. The elegant geometric construction via Lissajous curves [18] and the efficient Algorithm 886 [25] make Padua points a practical alternative to tensor product grids when interpolation accuracy, rather than differential operator structure, is the primary concern.

Finally, the field is currently witnessing a paradigm shift towards quantum algorithms for spectral methods. Liu *et al.* [80] have recently proposed a Quantum Spectral Method for non-periodic BVPs, encoding Chebyshev basis coefficients into quantum states to achieve polylogarithmic computational

complexity ($O(\text{poly}(\log N))$) for solving the linear systems. This suggests that the “matrix surgery” described in this chapter — currently an $O(N^3)$ operation on classical silicon — could one day become an $O(\text{poly}(\log N))$ operation on a quantum processor. The treatment of non-local boundary conditions, as explored by Mustapha *et al.* [90] for fractional differential equations, further extends the boundaries of what spectral methods can address, requiring operational matrices derived from shifted Jacobi or Lucas polynomials to handle integral constraints rather than local derivative conditions.

8.12 Summary

This chapter has demonstrated the power and simplicity of spectral collocation for boundary value problems:

- **Matrix stripping** handles Dirichlet boundary conditions by removing boundary rows and columns.
- **Variable coefficients** become diagonal matrices—no additional complexity.
- **Nonlinear problems** yield to Newton iteration, with the Jacobian easily computed from the differentiation matrix.
- **Eigenvalue problems** reveal the resolution limits of spectral methods: $\geq \pi$ points per wavelength are needed.
- **Tensor products** extend the method to higher dimensions via Kronecker products.

The key insight throughout is that spectral methods transform differential equations into *dense linear algebra*. The matrices are small (because spectral accuracy requires few points), and the resulting systems can be solved directly. This directness, no iteration and no mesh refinement, is the hallmark of spectral methods for smooth problems.



CHAPTER 9

Physical and Fourier Space on Grids

Every function tells two stories. In *physical space*, we see it as a curve: values $u(x)$ plotted against position x . In *Fourier space*, we see it decomposed into waves: amplitude coefficients $\hat{u}(k)$ plotted against wavenumber k . These two representations are equivalent; each contains complete information about the function. The art of spectral methods lies in moving fluently between these perspectives, exploiting whichever view makes a problem simpler [120].

This chapter develops the mathematical machinery connecting physical and Fourier space across three increasingly practical settings:

1. **Continuous functions on $\mathbb{R}$:** The classical Fourier transform, where both x and k range over the entire real line.
2. **Functions on an infinite grid $h\mathbb{Z}$:** The semidiscrete Fourier transform, where sampling at spacing h restricts wavenumbers to a bounded interval.
3. **Functions on a periodic grid:** The discrete Fourier transform (DFT), computable via the Fast Fourier Transform (FFT), which is the workhorse of practical spectral computation.

Along the way, we encounter the phenomenon of *aliasing*, which explains why high frequencies masquerade as low frequencies on discrete grids, and the *sinc function*, which provides the bridge between discrete samples and continuous band-limited interpolants. These concepts form the foundation for everything that follows in spectral methods.

9.1 Motivation: Two Views of the Same Function

Consider a smooth function such as $u(x) = e^{-x^2} \cos(3x)$, a Gaussian-modulated cosine. Figure 50 shows this function from both perspectives.




Figure 50: Two views of the same function $u(x) = e^{-x^2} \cos(3x)$. Left: *Physical space representation* showing the oscillating, localized wavepacket. Right: *Fourier space representation* showing the magnitude of the Fourier transform on a logarithmic scale. The rapid decay of $|\hat{u}(k)|$ reflects the smoothness of $u(x)$.

In physical space (left panel), we see a wavepacket: oscillations modulated by a Gaussian envelope. The function is smooth, localized, and decays rapidly as $|x| \rightarrow \infty$.

In Fourier space (right panel), we see the same information differently. The Fourier transform $\hat{u}(k)$ shows two peaks near $k = \pm 3$, corresponding to the carrier frequency $\cos(3x)$. The transform decays rapidly as $|k| \rightarrow \infty$, a hallmark of smooth functions [38, 120].

The following Python code approximates the Fourier transform using the FFT on a large computational domain:


```

1  import numpy as np
2
3  def compute_spectrum(u_func, L=20.0, N=1024):
4      """Approximate Fourier transform via FFT on [-L/2, L/2]."""
5      x = np.linspace(-L/2, L/2, N, endpoint=False)
6      dx = x[1] - x[0]
7      u = u_func(x)
8
9      # Approximate continuous FT via FFT
10     k = 2 * np.pi * np.fft.fftfreq(N, d=dx)
11     u_hat = dx * np.fft.fft(u)
12     return k, u_hat

```

 Python

The equivalent MATLAB implementation:

```

1  function [k, u_hat] = compute_spectrum(u_func, L, N)
2      % Approximate Fourier transform via FFT on [-L/2, L/2]
3      x = linspace(-L/2, L/2, N+1); x(end) = [];
4      dx = x(2) - x(1);
5      u = u_func(x);
6
7      % Wavenumber vector (MATLAB ordering)
8      k = 2*pi * [0:N/2-1, -N/2:-1] / (N*dx);
9      u_hat = dx * fft(u);
10 end

```

 Matlab

The Julia implementation:

```

1  using FFTW
2
3  function compute_spectrum(u_func; L=20.0, N=1024)
4      x = range(-L/2, stop=L/2, length=N+1)[1:N]
5      dx = x[2] - x[1]
6      u = u_func.(x)
7
8      # Approximate continuous FT via FFT
9      k = 2π .* fftfreq(N, 1/dx)
10     u_hat = dx .* fft(u)
11     return collect(x), u, k, u_hat
12 end

```

 Julia

The code generating Figure 50 is available in:

- codes/python/ch09/two_views_function.py
- codes/matlab/ch09/two_views_function.m
- codes/julia/ch09/two_views_function.jl

This visual preview illustrates the central theme: *smooth functions have rapidly decaying Fourier transforms*. This connection between physical smoothness and spectral decay is what makes spectral methods so powerful.

9.2 The Continuous Fourier Transform on $\mathbb{R}$

9.2.1 Definitions and Conventions

The *Fourier transform* of a function $u(x)$ defined on $\mathbb{R}$ is

$$\hat{u}(k) = \int_{-\infty}^{\infty} e^{-ikx} u(x) dx, \quad k \in \mathbb{R}. \quad (86)$$

The quantity $\hat{u}(k)$ represents the *amplitude density* of u at wavenumber k . This definition goes back to Fourier's foundational treatise on heat conduction [48]. The process of decomposing a function into its constituent waves is called *Fourier analysis*.

Conversely, we can reconstruct u from $\hat{u}$ by the *inverse Fourier transform*:

$$u(x) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{ikx} \hat{u}(k) dk, \quad x \in \mathbb{R}. \quad (87)$$

This reconstruction, building u from a superposition of complex exponentials e^{ikx} , is called *Fourier synthesis*.

The variable x is the *physical variable* (position), and k is the *Fourier variable* (wavenumber). For a wave e^{ikx} , the wavenumber k determines the spatial frequency: the number of oscillations per unit length is $|k| / (2\pi)$.

9.2.2 Fundamental Properties

The Fourier transform satisfies several important properties that make it a powerful tool for analysis. Table 13 summarizes the key relationships.

Property	Physical Space	Fourier Space
Linearity	$au + bv$	$a\hat{u} + b\hat{v}$
Translation	$u(x - x_0)$	$e^{-ikx_0} \hat{u}(k)$
Modulation	$e^{ik_0x} u(x)$	$\hat{u}(k - k_0)$
Differentiation	$u'(x)$	$ik\hat{u}(k)$
Convolution	$u * v$	$\hat{u} \cdot \hat{v}$

Table 13: *Fundamental properties of the Fourier transform. Operations in physical space correspond to simple algebraic operations in Fourier space.*

The *differentiation property* is particularly important for spectral methods: differentiation in physical space becomes multiplication by ik in Fourier space. This simple observation underlies the efficiency of spectral methods for differential equations [48, 120].

9.2.3 Physical and Fourier Variables: A Summary

Table 14 summarizes the relationship between physical and Fourier variables across the three settings we develop in this chapter. The key insight is that discreteness in one domain implies boundedness in the other.

Setting	Physical Variable	Fourier Variable
Continuous functions	$x \in \mathbb{R}$	$k \in \mathbb{R}$
Infinite grid $h\mathbb{Z}$	$x_j = jh$	$k \in [-\pi/h, \pi/h]$
Periodic grid, N points	$x_j \in [0, 2\pi)$	integer $k \in \mathbb{Z}$

Table 14: Physical versus Fourier variables for three settings. Discreteness in physical space implies boundedness in Fourier space; periodicity in physical space implies discreteness in Fourier space.

9.3 Sampling on an Infinite Grid: The Semidiscrete Fourier Transform

9.3.1 The Infinite Grid $h\mathbb{Z}$

We now consider sampling a continuous function at equally spaced points. The *infinite grid* $h\mathbb{Z}$ consists of points

$$x_j = jh, \quad j \in \mathbb{Z}, \quad (88)$$

where $h > 0$ is the *grid spacing* or *mesh width*.




Figure 51: The infinite grid $h\mathbb{Z}$ with grid points $x_j = jh$ for $j \in \mathbb{Z}$.

A grid function v assigns values v_j to each grid point x_j . We want to define a Fourier transform for such grid functions.

9.3.2 Aliasing: Why Wavenumbers Are Bounded

The crucial observation is that on the discrete grid, not all wavenumbers are distinguishable. Consider two complex exponentials $f(x) = e^{ik_1x}$ and $g(x) = e^{ik_2x}$. On the continuum $\mathbb{R}$, these are different functions whenever $k_1 \neq k_2$. But when restricted to the grid $h\mathbb{Z}$:

$$f_j = e^{ik_1x_j} = e^{ik_1jh}, \quad g_j = e^{ik_2x_j} = e^{ik_2jh}. \quad (89)$$

If $k_1 - k_2$ is an integer multiple of $2\pi/h$, then $f_j = g_j$ for all j !

This phenomenon is called *aliasing*: waves with wavenumbers differing by multiples of $2\pi/h$ are indistinguishable on the grid. They are *aliases* of each other [69, 107].

Key Insight: On a grid with spacing h , wavenumbers differing by $2\pi/h$ are indistinguishable:

$$e^{ikx_j} = e^{i(k+2\pi m/h)x_j} \quad \text{for all } j \in \mathbb{Z}, m \in \mathbb{Z}. \quad (90)$$

Consequently, it suffices to measure wavenumbers in an interval of length $2\pi/h$. By convention, we choose the symmetric interval $[-\pi/h, \pi/h]$, called the *Nyquist interval* [69, 107].

9.3.3 The Semidiscrete Fourier Transform

For a grid function v defined on $h\mathbb{Z}$, the *semidiscrete Fourier transform* is

$$\hat{v}(k) = h \sum_{j=-\infty}^{\infty} e^{-ikx_j} v_j, \quad k \in [-\pi/h, \pi/h]. \quad (91)$$

The inverse transform reconstructs the grid values:

$$v_j = \frac{1}{2\pi} \int_{-\pi/h}^{\pi/h} e^{ikx_j} \hat{v}(k) dk, \quad j \in \mathbb{Z}. \quad (92)$$

Note the duality: the forward transform Equation 91 is a sum (discrete in physical space), while the inverse Equation 92 is an integral (continuous in Fourier space, but bounded).

These formulas approximate the continuous Fourier transform and its inverse: Equation 91 is a trapezoid rule approximation to Equation 86, and Equation 92 truncates the integration domain of Equation 87. As $h \rightarrow 0$, both pairs of formulas converge.

9.3.4 Computational Étude 1: Aliasing of $\sin(\pi x)$ and $\sin(9\pi x)$

Figure 52 demonstrates aliasing concretely. Consider the two functions $\sin(\pi x)$ and $\sin(9\pi x)$ sampled on the grid $h = 1/4$ (i.e., $1/4\mathbb{Z}$). Despite being completely different continuous functions, they produce *identical* samples at every grid point!




Figure 52: Aliasing on the grid $h = 1/4$. The functions $\sin(\pi x)$ and $\sin(9\pi x)$ are distinct continuous functions (solid and dashed curves), but they coincide at every grid point (black dots). The wavenumber 9π aliases to π because $9\pi = \pi + 2\pi \cdot 4 = \pi + 2\pi/h$.

The aliasing relation is straightforward to verify. The Nyquist interval for $h = 1/4$ is $[-4\pi, 4\pi]$. The wavenumber 9π lies outside this interval. To find its alias, we compute:

$$9\pi = \pi + 8\pi = \pi + 2\pi \cdot \frac{1}{h} = \pi + \frac{2\pi}{1/4}. \quad (93)$$

Thus 9π aliases to π : both waves oscillate identically at the grid points.

The following Python code demonstrates this aliasing:

```
1 import numpy as np
2
3 h = 0.25
4 x_grid = np.arange(-2, 2 + h, h)
```

 Python

```

5
6 # Both functions give identical samples!
7 f_samples = np.sin(np.pi * x_grid)
8 g_samples = np.sin(9 * np.pi * x_grid)
9
10 print(f"Max difference: {np.max(np.abs(f_samples - g_samples)):.2e}")
11 # Output: Max difference: 1.11e-16 (machine precision)

```

The equivalent MATLAB implementation:

```

1 h = 0.25;
2 x_grid = -2:h:2;
3
4 % Both functions give identical samples
5 f_samples = sin(pi * x_grid);
6 g_samples = sin(9 * pi * x_grid);
7
8 fprintf('Max difference: %.2e\n', max(abs(f_samples - g_samples)));
9 % Output: Max difference: 1.11e-16

```

 Matlab

The Julia implementation:

```

1 h = 0.25
2 x_grid = collect(-2:h:2)
3
4 # Both functions give identical samples!
5 f_samples = sin.(pi .* x_grid)
6 g_samples = sin.(9pi .* x_grid)
7
8 println("Max difference: $(maximum(abs.(f_samples .- g_samples)))")
9 # Output: Max difference: 1.1102230246251565e-16

```

 Julia

General aliasing formula: Given a wavenumber $\kappa \in \mathbb{R}$, its alias $k \in [-\pi/h, \pi/h]$ is determined by

$$\kappa = k + \frac{2\pi m}{h} \quad (94)$$

for some integer m . Explicitly, $k = \kappa - \frac{2\pi}{h} \text{round}(\frac{\kappa h}{2\pi})$ [120].

The code generating Figure 52 is available in:

- codes/python/ch09/aliasing_demo.py
- codes/matlab/ch09/aliasing_demo.m
- codes/julia/ch09/aliasing_demo.jl

9.4 Band-Limited Interpolation and the Sinc Kernel

9.4.1 From Discrete Samples to Continuous Functions

Given samples v_j on the grid $h\mathbb{Z}$, how do we construct a continuous function that interpolates these values? The semidiscrete Fourier transform provides a canonical answer: construct the unique *band-limited interpolant* whose Fourier transform is supported in the Nyquist interval $[-\pi/h, \pi/h]$. The

mathematical foundations of this construction trace back to Whittaker [131] and have a rich history documented by Jerri [69] and Meijering [85].

9.4.2 The Discrete Delta Function

The simplest grid function is the *Kronecker delta*:

$$\delta_j = \begin{cases} 1 & \text{if } j = 0 \\ 0 & \text{if } j \neq 0. \end{cases} \quad (95)$$

Its semidiscrete Fourier transform is constant:

$$\hat{\delta}(k) = h \sum_{j=-\infty}^{\infty} e^{-ikx_j} \delta_j = h \cdot e^0 = h, \quad k \in [-\pi/h, \pi/h]. \quad (96)$$

9.4.3 The Sinc Function

Applying the inverse transform Equation 92 to $\hat{\delta}(k) = h$ yields the band-limited interpolant of δ :

$$p(x) = \frac{h}{2\pi} \int_{-\pi/h}^{\pi/h} e^{ikx} dk = \frac{\sin(\pi x/h)}{\pi x/h}. \quad (97)$$

This famous function is called the *sinc function*:

$$S_h(x) = \frac{\sin(\pi x/h)}{\pi x/h}. \quad (98)$$

By convention, $S_h(0) = 1$ (the limit as $x \rightarrow 0$).

Sir Edmund Whittaker [131] called S_1 “a function of royal blood in the family of entire functions, whose distinguished properties separate it from its bourgeois brethren” [84, 120].

9.4.4 Sinc Interpolation

The sinc function is the fundamental interpolation kernel for band-limited functions. Any grid function v can be written as

$$v_j = \sum_{m=-\infty}^{\infty} v_m \delta_{j-m}, \quad (99)$$

where δ_{j-m} equals 1 if $j = m$ and 0 otherwise. Since band-limited interpolation is linear and translation-invariant, the interpolant of v is:

$$p(x) = \sum_{m=-\infty}^{\infty} v_m S_h(x - x_m). \quad (100)$$

This is the *Whittaker–Shannon interpolation formula*, the foundation of the sampling theorem.

Sampling Theorem (Whittaker–Shannon–Nyquist): A band-limited function u with $\hat{u}(k) = 0$ for $|k| > \pi/h$ is completely determined by its samples $\{u(x_j)\}$ on $h\mathbb{Z}$. The reconstruction is given by Equation 100.

For a comprehensive tutorial on the history and various extensions of this theorem, see the survey by Jerri [69], which traces the result from its origins with Whittaker and Nyquist to Shannon’s [107] information-theoretic formulation.

9.4.5 Computational Étude 2: Sinc Interpolation of Three Signals

Figure 53 shows the band-limited interpolants of three grid functions: a discrete delta, a discrete square wave, and a discrete triangular (hat) function.




Figure 53: Band-limited interpolation of three grid functions ($h = 1$). Top: The Kronecker delta δ_j and its interpolant, the sinc function $S_h(x)$. Middle: A discrete square wave and its interpolant, exhibiting the Gibbs phenomenon near discontinuities. Bottom: A discrete triangular (hat) function and its interpolant, showing moderate oscillations.

Several observations:

- **Top panel:** The sinc function is smooth and analytic, decaying as $|x|^{-1}$.
- **Middle panel:** The square wave interpolant exhibits *Gibbs phenomenon*: oscillations near discontinuities that do not diminish as $h \rightarrow 0$ [38, 120]. Band-limited interpolation cannot approximate discontinuous functions well.
- **Bottom panel:** The hat function is continuous but not differentiable at its corners. Its interpolant is smoother than the square wave case but still oscillatory.

The core of the sinc interpolation algorithm is simple:

```

1 def sinc_kernel(z):
2     """Compute sinc(z) = sin(pi*z)/(pi*z), with sinc(0) = 1."""
3     out = np.ones_like(z, dtype=float)
4     mask = z != 0
5     out[mask] = np.sin(np.pi * z[mask]) / (np.pi * z[mask])
6     return out
7

```

 Python

```

8  def sinc_interpolate(x_grid, v, x_fine, h=1.0):
9      """Band-limited interpolation via sinc functions."""
10     p = np.zeros_like(x_fine)
11     for xi, vi in zip(x_grid, v):
12         p += vi * sinc_kernel((x_fine - xi) / h)
13     return p

```

The equivalent MATLAB implementation:

```

1  function p = sinc_interpolate(x_grid, v, x_fine, h)
2      % Band-limited interpolation via sinc functions
3      p = zeros(size(x_fine));
4      for i = 1:length(x_grid)
5          z = (x_fine - x_grid(i)) / h;
6          s = ones(size(z));
7          mask = z ~= 0;
8          s(mask) = sin(pi*z(mask)) ./ (pi*z(mask));
9          p = p + v(i) * s;
10     end
11 end

```

 Matlab

The Julia implementation:

```

1  function sinc_kernel(z)
2      out = ones(length(z))
3      for i in eachindex(z)
4          z[i] != 0 && (out[i] = sin(π * z[i]) / (π * z[i]))
5      end
6      return out
7  end
8
9  function sinc_interpolate(x_grid, v, x_fine; h=1.0)
10     p = zeros(length(x_fine))
11     for (xi, vi) in zip(x_grid, v)
12         p .+= vi .* sinc_kernel((x_fine .- xi) ./ h)
13     end
14     return p
15 end

```

 Julia

The code generating Figure 53 is available in:

- codes/python/ch09/sinc_interpolation.py
- codes/matlab/ch09/sinc_interpolation.m
- codes/julia/ch09/sinc_interpolation.jl

9.5 Periodic Grids: The Discrete Fourier Transform and FFT

9.5.1 The Periodic Setting

We now turn to the practical setting: a bounded, periodic domain. Consider the interval $[0, 2\pi)$ with periodic boundary conditions. We place N equispaced grid points:

$$x_j = \frac{2\pi j}{N}, \quad j = 0, 1, \dots, N-1. \quad (101)$$

The grid spacing is $h = 2\pi/N$, which gives the fundamental relation:

$$\frac{\pi}{h} = \frac{N}{2}. \quad (102)$$

This equation appears repeatedly in periodic spectral methods.

A *periodic grid function* v satisfies $v_{j+N} = v_j$ for all j . Equivalently, we can view it as one period of length N extracted from an infinite periodic sequence.

9.5.2 Fourier Space for Periodic Functions

Two constraints shape the Fourier space for periodic grids:

1. **Periodicity in physical space:** The function must have period 2π , so only waves e^{ikx} with integer wavenumbers k have the correct period.
2. **Discreteness in physical space:** Aliasing restricts distinguishable wavenumbers to an interval of length $2\pi/h = N$.

Combining these constraints:

Physical space: discrete, bounded $\Rightarrow x_j \in \{0, h, 2h, \dots, (N-1)h\}$
 Fourier space: bounded, discrete $\Rightarrow k \in \{-N/2 + 1, \dots, N/2\}$

9.5.3 The Discrete Fourier Transform

For N even, the *discrete Fourier transform (DFT)* is

$$\hat{v}_k = h \sum_{j=0}^{N-1} e^{-ikx_j} v_j, \quad k = -N/2 + 1, \dots, N/2. \quad (103)$$

The *inverse DFT* is

$$v_j = \frac{1}{2\pi} \sum_{k=-N/2+1}^{N/2} e^{ikx_j} \hat{v}_k, \quad j = 0, 1, \dots, N-1. \quad (104)$$

These are exact inverses: no approximation is involved. The DFT maps N function values to N Fourier coefficients and back.

9.5.4 The Nyquist Mode and Symmetry

The wavenumber $k = N/2$ requires special treatment. The wave $e^{i(N/2)x}$ equals $\cos(Nx/2)$ on the grid (a “sawtooth” pattern alternating ± 1). For differentiation, we symmetrize the highest mode by setting $\hat{v}_{-N/2} = \hat{v}_{N/2}$, giving:

$$v_j = \frac{1}{2\pi} \sum_{k=-N/2}^{N/2} \{\}' e^{ikx_j} \hat{v}_k, \quad (105)$$

where the prime indicates that the $k = \pm N/2$ terms are multiplied by $1/2$.

9.5.5 The Fast Fourier Transform

The DFT can be computed naively in $O(N^2)$ operations, but the *Fast Fourier Transform (FFT)* reduces this to $O(N \log N)$. Discovered by Cooley and Tukey [30] in 1965 (though Gauss had the idea in 1805, as documented by Heideman, Johnson, and Burrus [62]), the FFT revolutionized scientific computing.

The key insight is that a DFT of size N can be decomposed into two DFTs of size $N/2$: one for even-indexed entries, one for odd. This divide-and-conquer approach yields the $O(N \log N)$ complexity.

9.5.6 FFT Indexing Conventions

MATLAB and Python store FFT output with different conventions than our mathematical formulas. Table 15 clarifies the correspondence.

MATLAB (1-based)	Python (0-based)	Math mode k	Description
1	0	0	Constant (mean)
2 to $N/2$	1 to $N/2 - 1$	1 to $N/2 - 1$	Positive modes
$N/2 + 1$	$N/2$	$N/2$	Nyquist mode
$N/2 + 2$ to N	$N/2 + 1$ to $N - 1$	$-N/2 + 1$ to -1	Negative modes

Table 15: FFT indexing conventions for N even. The wavenumber vector in Python is `np.fft.fftfreq(N)*N`, and in MATLAB it is `[0:N/2, -N/2+1:-1]`.

9.5.7 Spectral Differentiation via FFT

To differentiate a periodic grid function spectrally [50, 120]:

1. Compute $\hat{v} = \text{FFT}(v)$
2. Multiply: $\hat{w}_k = ik\hat{v}_k$ for $k \neq N/2$, and $\hat{w}_{N/2} = 0$
3. Compute $w = \text{IFFT}(\hat{w})$

The following Python code implements this:

```
1 def spectral_derivative(v):
2     """Compute spectral derivative of periodic function."""
3     N = len(v)
4     v_hat = np.fft.fft(v)
5     k = np.fft.fftfreq(N) * N # Wavenumbers: 0,1,...,N/2,-N/2+1,...,-1
6     k = 1j * k
7     k[N//2] = 0 # Zero out Nyquist mode for odd derivatives
8     w_hat = k * v_hat
9     return np.real(np.fft.ifft(w_hat))
```

 Python

The equivalent MATLAB implementation:

```

1 function w = spectral_derivative(v)
2     N = length(v);
3     v_hat = fft(v);
4     k = [0:N/2-1, 0, -N/2+1:-1]; % Wavenumbers with zeroed Nyquist
5     w_hat = 1i * k .* v_hat;
6     w = real(ifft(w_hat));
7 end

```

 Matlab

The Julia implementation:

```

1 using FFTW
2
3 function spectral_derivative(v)
4     N = length(v)
5     v_hat = fft(v)
6     k = [0:N÷2-1; 0; -N÷2+1:-1] # Wavenumbers with zeroed Nyquist
7     w_hat = 1im .* k .* v_hat
8     return real.(ifft(w_hat))
9 end

```

 Julia

9.6 Aliasing and Spectra on Periodic Grids

9.6.1 Aliasing in FFT Computations

When we compute the FFT of a sampled function, any frequency content above the Nyquist frequency $N/2$ folds back into the resolved range. This is the periodic analog of the aliasing we saw in Section 9.3. The foundational work on mitigating this error in nonlinear computations is due to Orszag [95], who introduced the famous “2/3 rule” for dealiasing; see also Bowman and Roberts [19] for efficient modern implementations.

9.6.2 Computational Étude 3: Aliasing in FFT of High-Frequency Data

Consider $u(x) = \sin(17x)$ sampled at $N = 32$ points. Since $17 > N/2 = 16$, this frequency is above the Nyquist limit and will alias.




Figure 54: Aliasing in the FFT of $\sin(17x)$ with $N = 32$ points. The true frequency $k = 17$ lies above the Nyquist limit $N/2 = 16$. The FFT shows a spike at $k = -15$ because $17 = -15 + 32$: the frequency “wraps around” the Nyquist boundary.

The alias is computed as: $17 = -15 + 32$, so wavenumber 17 appears as wavenumber -15 in the FFT output.

```
1 N = 32
2 x = 2 * np.pi * np.arange(N) / N
3 u = np.sin(17 * x)
4
5 u_hat = np.fft.fft(u) / N
6 k = np.fft.fftshift(np.fft.fftfreq(N) * N)
7
8 # The spike appears at k = -15, not k = 17
```

Python

The equivalent MATLAB implementation:

```
1 N = 32;
2 x = 2*pi*(0:N-1)/N;
3 u = sin(17*x);
4
5 u_hat = fft(u) / N;
6 k = fftshift([0:N/2-1, -N/2:-1]);
7
8 % The spike appears at k = -15, not k = 17
```

Matlab

The Julia implementation:

```
1 N = 32
2 x = 2π * (0:N-1) / N
3 u = sin.(17 .* x)
4
5 u_hat = fft(u) / N
6 k = fftshift(fftfreq(N) .* N)
7
8 # The spike appears at k = -15, not k = 17
```

Julia

The code generating Figure 54 is available in:

- `codes/python/ch09/fft_aliasing.py`
- `codes/matlab/ch09/fft_aliasing.m`
- `codes/julia/ch09/fft_aliasing.jl`

The concept of aliasing has found renewed significance beyond numerical analysis. In the field of deep learning, aliasing in neural radiance fields (NeRF) has been identified as a critical bottleneck; see Mildenhall *et al.* [88] for a comprehensive treatment. The classical spectral perspective on aliasing developed by Trefethen [120] thus remains fundamental across a wide range of modern applications.

9.6.3 What Your Eye Aliases

Aliasing occurs not just in computation but in perception. Execute the following in Python:

```
1 import matplotlib.pyplot as plt
2 plt.plot(np.sin(np.arange(1, 3001)), '.')
```

 Python

or in MATLAB:

```
1 plot(sin(1:3000), '.')
```

 Matlab

or in Julia:

```
1 using CairoMakie
2 scatter(sin.(1:3000), markersize=2)
```

 Julia

The result appears to oscillate slowly, even though $\sin(n)$ for integer n oscillates rapidly. Your visual system, sampling the plotted points, aliases the high frequency to a low-frequency pattern!

9.7 Spectra and Smoothness

9.7.1 The Smoothness-Decay Connection

One of the most important results in Fourier analysis is the connection between smoothness in physical space and decay in Fourier space:

- **Discontinuous functions:** Fourier coefficients decay as $O(|k|^{-1})$
- **p continuous derivatives:** Decay as $O(|k|^{-p-1})$
- **Infinitely smooth (C^∞):** Faster than any polynomial
- **Analytic functions:** Exponential decay $O(e^{-a|k|})$ [21, 121]

This hierarchy is a classical result in harmonic analysis; see, e.g., Trefethen [120], Fourier [48], and the general theory of (p, q) -Fourier coefficient decay by Edmunds, Gurka, and Lang [37]. It explains why spectral methods are so accurate for smooth problems: smooth functions have negligible high-frequency content, so the aliasing error from discretization is tiny.

9.7.2 Computational Demonstration

Figure 55 shows the spectra of three periodic functions with different smoothness:

1. $\exp(\sin x)$: Analytic (infinitely differentiable and more), with exponential spectral decay
2. Periodic hat function: C^0 but not C^1 at corners, with algebraic decay $\approx |k|^{-2}$
3. Square wave: Jump discontinuities, with slow decay $\approx |k|^{-1}$




Figure 55: Fourier spectra of three functions with different smoothness. Left: $\exp(\sin x)$ (analytic) shows exponential decay on a semilog plot. Center: Periodic hat function (C^0) shows algebraic decay $\approx |k|^{-2}$ on a log-log plot. Right: Square wave (discontinuous) shows slow decay $\approx |k|^{-1}$.

Function	Regularity	Decay of $ \hat{f}_k $
$\exp(\sin x)$	Analytic	Faster than any power
Hat function	C^0 , piecewise C^1	$\approx k ^{-2}$
Square wave	Jump discontinuities	$\approx k ^{-1}$

Table 16: Smoothness classes and their characteristic Fourier decay rates.

The code generating Figure 55 is available in:

- codes/python/ch09/smoothness_spectra.py
- codes/matlab/ch09/smoothness_spectra.m
- codes/julia/ch09/smoothness_spectra.jl

9.8 Periodic Band-Limited Interpolation

9.8.1 The Periodic Sinc Function

Just as the sinc function is the band-limited interpolant of the Kronecker delta on $h\mathbb{Z}$, the *periodic sinc* is the band-limited interpolant of the periodic delta on the N -point periodic grid.

The periodic delta is

$$\delta_j = \begin{cases} 1 & \text{if } j \equiv 0 \pmod{N} \\ 0 & \text{otherwise.} \end{cases} \quad (106)$$

Its DFT satisfies $\hat{\delta}_k = h$ for all k . The inverse DFT gives the periodic sinc:

$$S_N(x) = \frac{\sin(\pi x/h)}{(2\pi/h) \tan(x/2)}, \quad (107)$$

with $h = 2\pi/N$. This derivation is presented in detail by Trefethen [120] and follows naturally from the discrete orthogonality of complex exponentials on the periodic grid [55].

For small x , the periodic sinc behaves like the nonperiodic sinc: $S_N(x) \approx \sin(\pi x/h)/(\pi x/h)$. The difference appears for larger $|x|$, where the periodic sinc has period 2π .

9.8.2 Computational Étude 4: Zero-Padding Interpolation via FFT

A practical way to perform band-limited interpolation on a periodic grid is *zero-padding* in Fourier space:

1. Given N samples v_j , compute the DFT $\hat{v}_k$.
2. Create a larger array of size $M = qN$ (say $q = 4$), placing $\hat{v}_k$ in the low-frequency positions and zeros in the high-frequency positions.
3. Compute the inverse DFT to get M interpolated values.

This is equivalent to evaluating the trigonometric interpolant at M points, but using FFTs makes it $O(M \log M)$ rather than $O(MN)$ [50, 120].




Figure 56: Band-limited interpolation via zero-padding. Black dots show $N = 32$ samples of $\exp(\sin x)$. The solid curve shows the interpolant obtained by zero-padding the FFT to $M = 128$ points. The interpolant passes exactly through the original samples and provides a smooth reconstruction between them.

```

1  def zero_pad_interpolate(v, q=4):
2      """Interpolate by zero-padding in Fourier space."""
3      N = len(v)
4      M = q * N
5
6      v_hat = np.fft.fft(v)
7      v_hat_padded = np.zeros(M, dtype=complex)
8
9      # Copy low frequencies (0 to N/2-1 and -N/2+1 to -1)
10     v_hat_padded[:N//2] = v_hat[:N//2]
11     v_hat_padded[-N//2+1:] = v_hat[-N//2+1:]
12     # Handle Nyquist: split between +N/2 and -N/2
13     v_hat_padded[N//2] = v_hat[N//2] / 2
14     v_hat_padded[-N//2] = v_hat[N//2] / 2
15
16     return np.real(np.fft.ifft(v_hat_padded)) * q

```


Python

The equivalent MATLAB implementation:

```

1  function v_fine = zero_pad_interpolate(v, q)
2      N = length(v);
3      M = q * N;
4
5      v_hat = fft(v);
6      v_hat_padded = zeros(1, M);
7
8      % Copy low frequencies
9      v_hat_padded(1:N/2) = v_hat(1:N/2);
10     v_hat_padded(end-N/2+2:end) = v_hat(N/2+2:end);
11     % Handle Nyquist mode
12     v_hat_padded(N/2+1) = v_hat(N/2+1) / 2;
13     v_hat_padded(M-N/2+1) = v_hat(N/2+1) / 2;
14
15     v_fine = real(ifft(v_hat_padded)) * q;
16 end

```




The Julia implementation:

```

1  using FFTW
2
3  function zero_pad_interpolate(v; q=4)
4      N = length(v)
5      M = q * N
6      v_hat = fft(v)
7      v_hat_padded = zeros{ComplexF64, M}
8
9      # Copy low frequencies
10     v_hat_padded[1:N÷2] = v_hat[1:N÷2]
11     v_hat_padded[end-N÷2+2:end] = v_hat[end-N÷2+2:end]
12     # Handle Nyquist mode: split between +N/2 and -N/2
13     v_hat_padded[N÷2+1] = v_hat[N÷2+1] / 2
14     v_hat_padded[M-N÷2+1] = v_hat[N÷2+1] / 2
15
16     return real.(ifft(v_hat_padded)) .* q
17 end

```



The code generating Figure 56 is available in:

- codes/python/ch09/zero_padding_interpolation.py
- codes/matlab/ch09/zero_padding_interpolation.m
- codes/julia/ch09/zero_padding_interpolation.jl

9.9 A non-exhaustive literature overview

The mathematical machinery codified in this chapter — the Fourier transform, sampling, aliasing, and the FFT — is the product of three centuries of analytical progress. While the name “Fast Fourier Trans-

form” is inextricably linked to the 1965 publication by Cooley and Tukey [30], the intellectual history runs far deeper. Historical analysis of Carl Friedrich Gauss’s *Nachlass* (unpublished notes) reveals that in 1805, while seeking to interpolate the orbits of asteroids Pallas and Juno from limited observational data, Gauss derived an algorithm functionally equivalent to the FFT [62]. Gauss’s method, designed for hand calculation, recognized that trigonometric sums could be split based on the prime factors of the number of data points, effectively reducing the arithmetic load from $O(N^2)$ to $O(N \log N)$ long before the “Big O” notation existed. It is a striking historical irony that this method remained largely dormant, overshadowed by Fourier’s analytical work on heat conduction published in 1822 [48]. Fourier’s insistence that arbitrary functions — even those with discontinuities — could be represented by trigonometric series was revolutionary and controversial, sparking decades of debate regarding convergence that eventually birthed the field of harmonic analysis. The *discrete* algorithmic efficiency discovered by Gauss was lost to the broader scientific community until the digital age necessitated it. The 1965 Cooley–Tukey paper was the catalyst for the modern spectral revolution, not because it was the first discovery, but because it appeared at the precise moment when digital computers became capable of executing these transforms at scale.

The transition from continuous signals to discrete samples, treated in Section 9.3 and Section 9.4, rests on the Whittaker–Shannon–Kotelnikov (WKS) sampling theorem. While often attributed simply to Shannon (1949) or Nyquist (1928) in engineering contexts, the mathematical rigorousness of the theorem traces back to the interpolation theory of the early 20th century. E. T. Whittaker’s seminal 1915 paper [131] is of particular importance to the narrative of this textbook. Whittaker investigated the properties of the “cardinal function” — what we now call the sinc function expansion — and established that it provides the unique consistent interpolation for band-limited functions. Whittaker famously described this function as “a function of royal blood in the family of entire functions, whose distinguished properties separate it from its bourgeois brethren” [84]. This evocative phrasing highlights the exceptional nature of the sinc kernel: it is an entire function (analytic everywhere in the complex plane) that perfectly reconstructs band-limited data, satisfying the “spectral promise” of exponential accuracy. The genealogy extends even further back. Historical scholarship indicates that Augustin-Louis Cauchy was aware of the mechanisms of band-limited sampling as early as 1841, and Émile Borel stated the essential features of the reconstruction theorem in 1897 [85]. Shannon’s contribution in 1949 [107] was to transplant these mathematical identities into the fertile soil of information theory, proving that the sampling limit (the Nyquist rate) was a fundamental bound on information capacity, not just a mathematical curiosity. For a comprehensive tutorial review, see Jerri [69]. For the student of spectral methods, this distinction is crucial: the sampling theorem validates the use of discrete grids to solve continuous PDEs, guaranteeing that no information is lost as long as the solution remains sufficiently smooth (band-limited).

Aliasing, introduced in Section 9.6 as the phenomenon where high frequencies “masquerade” as low frequencies, is often presented solely as a source of error to be eliminated. The literature, however, presents a more nuanced view, evolving from strict mitigation strategies in the 1970s to sophisticated exploitation in modern machine learning and optics. The foundational work on mitigating this error was established by Orszag [95], who introduced the famous “2/3 rule” (often called the 3/2 rule in terms of dealiasing padding), which proves that aliasing can be completely eliminated for quadratic nonlinearities by padding the spectrum with zeros to extend the grid by a factor of 3/2 before performing the nonlinear multiplication in physical space. This technique remains the gold standard in Direct Numerical Simulation (DNS) of turbulence, where preserving the fidelity of the energy cascade is paramount. Bowman and Roberts [19] developed efficient dealiased convolution algorithms that avoid the overhead of explicit zero-padding. In the period 2024–2026, the literature reflects a paradigm shift where aliasing is viewed through new lenses. In the field of deep learning, specifically in Physics-Informed Neural Networks (PINNs) and Neural Operators (like the Fourier Neural Operator, FNO), aliasing has emerged as a critical bottleneck [88]. Standard activation functions introduce infinite

spectral content, leading to aliasing errors that degrade the convergence of neural PDE solvers. This has led to the development of “aliasing-free” architectures and spectral activation functions designed to respect the Nyquist limit of the underlying grid [4], effectively reinventing Orszag’s principles for the age of AI. Furthermore, in the domain of optics and nanophotonics, recent research describes “anti-aliasing metasurfaces” [72]. Here, aliasing is not a computational artifact but a physical diffraction phenomenon. By developing multidimensional sampling theories that transcend the traditional Nyquist limit, researchers have created flat optical devices that suppress diffraction noise, enabling ultra-compact high-resolution imaging systems. This illustrates a profound “third-order insight”: the mathematical concept of aliasing, once confined to time-series analysis, now governs the design of physical materials and neural architectures.

This chapter primarily focuses on uniform grids, which allow for the use of the standard FFT. However, the “Spectral Promise” is increasingly being applied to problems where data cannot be sampled uniformly — from MRI scans to astrophysical observations. The Non-Uniform Fast Fourier Transform (NUFFT) has become a critical area of algorithmic research, exploding in importance between 2023 and 2026. The NUFFT generalizes the FFT to off-grid data points, usually by convolving the non-uniform data onto a regular grid using a carefully chosen kernel (like the Kaiser–Bessel or the “exponential of semicircle” kernel). Barnett, Magland, and af Klinteberg [12] developed a high-performance parallel NUFFT library (FINUFFT) based on this exponential of semicircle kernel, achieving order-of-magnitude speedups on modern hardware. In radio astronomy, the detection of the faint 21-cm hydrogen signal from the Epoch of Reionization requires processing visibility data from massive interferometer arrays; Cox *et al.* [31] demonstrated that high-performance NUFFT algorithms are essential for simulating these visibilities at the scale required for the Square Kilometer Array (SKA). Parallel to the NUFFT, the Sparse Fast Fourier Transform (sFFT) breaks the $O(N \log N)$ barrier by exploiting the observation that many real-world signals have only K non-zero coefficients in the frequency domain, where $K \ll N$. The sFFT algorithms compute the transform in sub-linear time, often $O(K \log N)$ or even $O(K)$ [61]. This literature suggests a divergence in spectral methods: one branch (NUFFT) focuses on geometric flexibility, while the other (sFFT) focuses on algorithmic efficiency for massive, sparse datasets.

Perhaps the most visionary development in the recent literature (2025–2026) is the intersection of Fourier analysis with quantum computing. While the Quantum Fourier Transform (QFT) has been a staple of quantum algorithms (e.g., Shor’s algorithm) for decades, recent work has generalized Fourier analysis to higher-order structures to characterize quantum complexity. A series of breakthrough papers in 2025, including work published in *PNAS*, introduces “Quantum Higher-Order Fourier Analysis” (q-HOFA) [23]. This theoretical framework generalizes classical higher-order Fourier analysis (used in additive combinatorics) to the quantum realm. The researchers define “quantum uniformity norms” (analogous to Gowers norms) that characterize the “Clifford hierarchy” — a classification of quantum gates based on their complexity and fault-tolerance properties. This development is profound for the future of spectral methods. It suggests that the tools of Fourier analysis are not limited to classical scalar functions but can be lifted to operate on operators in Hilbert space. While classical spectral methods scale as $O(N \log N)$, quantum spectral methods scale polylogarithmically with the dimension of the vector space, potentially offering exponential speedups for specific classes of smooth, high-dimensional PDEs.

Despite these advances, the core theoretical engine of spectral methods remains the connection between smoothness in physical space and decay in Fourier space. This relationship, formalized in Section 9.7, guarantees that for analytic functions, the Fourier coefficients decay exponentially ($O(e^{-a|k|})$), leading to the famed “spectral accuracy.” The recent literature reaffirms and refines this principle. Work in *SIAM Journal on Numerical Analysis* and *Mathematics of Computation* (2024–2025) has derived sharp pointwise error estimates for functions with algebraic singularities, showing exactly how convergence deteriorates near non-smooth points [78]. Furthermore, the theory of “Fourier

extensions” or “Fourier continuation” has matured, providing robust methods to map non-periodic functions onto periodic domains while preserving spectral convergence rates, effectively bypassing the Gibbs phenomenon for non-periodic BVPs [136]. The comprehensive treatments by Trefethen [121] and Boyd [21] provide the rigorous justification for the computational heuristics used in this textbook. They explain *why* the method works so well for the test cases in this and subsequent chapters and provide the diagnostic tools to understand failure modes when smoothness is lost.

9.10 Summary

This chapter has developed the mathematical framework connecting physical and Fourier space across three settings:

1. **Continuous Fourier transform:** Functions on $\mathbb{R}$ have Fourier transforms on $\mathbb{R}$. Differentiation becomes multiplication by ik .
2. **Semidiscrete Fourier transform:** Sampling at spacing h restricts wavenumbers to the Nyquist interval $[-\pi/h, \pi/h]$. Wavenumbers differing by $2\pi/h$ are aliases of each other.
3. **Discrete Fourier transform:** On a periodic grid with N points, both physical and Fourier variables are discrete. The FFT computes the DFT in $O(N \log N)$ operations.

Key concepts:

- **Aliasing:** High frequencies masquerade as low frequencies on discrete grids
- **Sinc interpolation:** Band-limited reconstruction from samples via the Whittaker–Shannon formula
- **Smoothness implies decay:** Smooth functions have rapidly decaying Fourier coefficients, making aliasing negligible

These concepts form the foundation for spectral methods. In the chapters that follow, we apply them to solve differential equations, exploiting the connection between smoothness and spectral accuracy.

9.10.1 Quick Reference

Setting	Variables	Forward Transform	Inverse Transform
Continuous	$x, k \in \mathbb{R}$	$\hat{u} = \int e^{-ikx} u \, dx$	$u = \frac{1}{2\pi} \int e^{ikx} \hat{u} \, dk$
Semidiscrete	$x_j = jh, \quad k \in [-\frac{\pi}{h}, \frac{\pi}{h}]$	$\hat{v} = h \sum_j e^{-ikx_j} v_j$	$v_j = \frac{1}{2\pi} \int e^{ikx_j} \hat{v} \, dk$
Discrete (DFT)	$x_j = \frac{2\pi j}{N}, k \in \mathbb{Z}$	$\hat{v}_k = h \sum_j e^{-ikx_j} v_j$	$v_j = \frac{1}{2\pi} \sum_k e^{ikx_j} \hat{v}_k$

Table 17: Quick reference for Fourier transforms in three settings.

9.10.2 FFT Idioms

Python (NumPy):

```
1 import numpy as np
2 k = np.fft.fftfreq(N) * N      # Wavenumbers: 0, 1, ..., N/2, -N/2+1, ..., -1
```

 Python

```

3 u_hat = np.fft.fft(u)           # Forward FFT
4 u = np.fft.ifft(u_hat)         # Inverse FFT
5 k_shifted = np.fft.fftshift(k) # Reorder to -N/2,...,N/2-1 for plotting

```

MATLAB:

```

1 k = [0:N/2-1, -N/2:-1];        % Wavenumbers
2 u_hat = fft(u);                % Forward FFT
3 u = ifft(u_hat);               % Inverse FFT
4 k_shifted = fftshift(k);        % Reorder for plotting

```

 Matlab
Julia (FFTW.jl):

```

1 using FFTW
2 k = fftfreq(N) .* N             # Wavenumbers: 0,1,...,N/2,-N/2+1,...,-1
3 u_hat = fft(u)                  # Forward FFT
4 u = ifft(u_hat)                 # Inverse FFT
5 k_shifted = fftshift(k)         # Reorder for plotting

```

 Julia



CHAPTER 10

Spectral PDE Solvers with Chebyshev and Fourier Grids

The machinery of spectral differentiation, developed carefully over the preceding chapters, now stands ready for its primary purpose: solving partial differential equations. In this chapter, we bring together Chebyshev and Fourier methods to tackle the three fundamental classes of PDEs: hyperbolic equations (waves), parabolic equations (diffusion), and elliptic equations (equilibrium states). The unifying theme is the *method of lines*: discretise space with spectral accuracy, then advance in time (or solve directly for steady states) using appropriate numerical schemes.

By the end of this chapter, you should be able to:

1. Build Chebyshev differentiation operators via FFT and apply them in 1D and 2D.
2. Discretise in space and time the main PDE classes: wave equations, heat equations, and elliptic problems (Laplace, Poisson, Helmholtz).
3. Understand CFL-type stability limits for explicit schemes, and why implicit schemes are often preferred for parabolic and elliptic problems.
4. Read space–time plots and contour plots, and relate numerical behaviour to physical intuition.
5. Implement spectral PDE solvers for vibrating strings, drum membranes, heat diffusion, and electrostatics.

10.1 Chebyshev–Fourier Recap

Before diving into PDEs, let us briefly recall the geometric connection between Chebyshev and Fourier methods that makes FFT-accelerated Chebyshev differentiation possible. This connection, developed in Section 7 and Section 9, is the foundation for everything that follows.

10.1.1 The Circle-to-Interval Map

The Chebyshev–Gauss–Lobatto points on $[-1, 1]$ are

$$x_j = \cos\left(j\frac{\pi}{N}\right), \quad j = 0, 1, \dots, N. \quad (108)$$

These $N + 1$ points arise naturally as projections: if we place $N + 1$ equally-spaced angles

$$\theta_j = j\frac{\pi}{N}, \quad j = 0, 1, \dots, N, \quad (109)$$

on the upper semicircle and project them onto the x -axis, we obtain exactly the Chebyshev points. The map connecting the angle and physical variables is

$$x = \cos \theta. \quad (110)$$

This geometric interpretation explains the characteristic clustering of Chebyshev points near the boundaries $x = \pm 1$: points near $\theta = 0$ and $\theta = \pi$ project to $x \approx \pm 1$, but the horizontal projection compresses them together. Near the centre ($\theta \approx \frac{\pi}{2}$), the points spread out more evenly.

10.1.2 Chebyshev Polynomials as Wrapped Cosines

The Chebyshev polynomial of degree n is defined by the remarkable formula

$$T_n(x) = \cos(n\theta), \quad \text{where } x = \cos \theta. \quad (111)$$

This identity reveals that Chebyshev polynomials are simply cosine waves “in disguise”: the n th Chebyshev polynomial oscillates exactly as $\cos(n\theta)$ would, but viewed through the nonlinear lens of $x = \cos \theta$.

The practical consequence is profound: data $v_j = f(x_j)$ sampled at Chebyshev points corresponds to an *even, 2π -periodic function* in the θ variable:

$$F(\theta) = f(\cos \theta). \quad (112)$$

Since $F(\theta) = F(-\theta)$ (evenness) and $F(\theta + 2\pi) = F(\theta)$ (periodicity), we can apply Fourier methods, specifically the FFT, to work with Chebyshev expansions efficiently.




Figure 57: *The Chebyshev–Fourier connection. Left: The unit semicircle with equally-spaced angles $\theta_j = j \frac{\pi}{N}$ projecting to Chebyshev points $x_j = \cos \theta_j$ on the interval $[-1, 1]$. The clustering near $x = \pm 1$ is evident. Right: A Chebyshev polynomial $T_n(x)$ can be visualised as a cosine wave $\cos(n\theta)$ “wrapped around a cylinder and viewed from the side.”*

The code generating Figure 57 is available in:

- `codes/python/ch10/cheb_fourier_geometry.py`
- `codes/matlab/ch10/cheb_fourier_geometry.m`
- `codes/julia/ch10/cheb_fourier_geometry.jl`

10.2 Chebyshev Differentiation via FFT

Having the geometric picture, we now build the computational machinery for FFT-based Chebyshev differentiation. Given function values v_j at the Chebyshev points Equation 108, our goal is to compute accurate approximations to the derivatives $u'(x_j)$ and $u''(x_j)$.

10.2.1 The “ x to θ to Fourier” Pipeline

The algorithm proceeds through a sequence of transformations. The efficiency of this differentiation relies on the $O(N \log N)$ Fast Fourier Transform algorithm introduced by Cooley and Tukey [30]. The specific implementation using the “extended vector” approach to enforce symmetries is a staple of modern spectral software, comprehensively detailed by Trefethen [120] and Boyd [21].

1. **Extend to an even vector:** Given values $v_0, v_1, \dots, v_N$ at Chebyshev points, form the extended vector V of length $2N$ by reflecting:

$$V_j = v_j \quad (j = 0, \dots, N), \quad V_{2N-j} = v_j \quad (j = 1, \dots, N-1). \quad (113)$$

This extended vector represents an even function sampled at $2N$ equally-spaced angles in $[0, 2\pi)$.

2. **FFT in θ -space:** Compute the discrete Fourier transform $\hat{V}_k$ using the FFT.

3. **Differentiate in Fourier space:** Multiply by ik (for the first derivative) or $(ik)^2$ (for the second derivative). For odd-order derivatives, set the Nyquist mode $\hat{V}_N$ to zero to avoid aliasing artefacts.
4. **Inverse FFT:** Transform back to obtain the derivative $W_j = Q'(\theta_j)$ in θ -space.
5. **Convert from θ to x :** Apply the chain rule to transform derivatives with respect to θ into derivatives with respect to x . Since $x = \cos \theta$, we have $d\frac{x}{d}\theta = -\sin \theta = -\sqrt{1-x^2}$, giving

$$q'(x) = \frac{Q'(\theta)}{dx/d\theta} = -\frac{Q'(\theta)}{\sqrt{1-x^2}}. \quad (114)$$

For the second derivative, the relation is more involved:

$$q''(x) = \frac{x}{(1-x^2)^{3/2}} Q'(\theta) + \frac{1}{1-x^2} Q''(\theta). \quad (115)$$

6. **Boundary values:** At $x = \pm 1$, the formulas Equation 114 and Equation 115 involve division by zero. Special formulas, derived via l'Hôpital's rule, handle the endpoints:

$$q'(1) = \sum_{n=0}^N n^2 a_n, \quad q'(-1) = \sum_{n=0}^N (-1)^{n+1} n^2 a_n, \quad (116)$$

where a_n are the Chebyshev coefficients.

10.2.2 Implementation: The *chebfft* Function

The following Python code implements the first derivative via FFT on the Chebyshev grid:

```

1  def chebfft(v):
2      """Chebyshev first derivative via FFT."""
3      N = len(v) - 1
4      if N == 0:
5          return np.array([0.0])
6
7      x = np.cos(np.pi * np.arange(N + 1) / N)
8
9      # Extend to even periodic function in theta
10     V = np.concatenate([v, v[-2:0:-1]])
11
12     # FFT and multiply by ik
13     U = np.fft.fft(V).real
14     k = np.concatenate([np.arange(N), [0], np.arange(1-N, 0)])
15     W = np.fft.ifft(1j * k * np.fft.fft(V)).real
16
17     # Transform from theta to x derivative
18     w = np.zeros(N + 1)
19     w[1:N] = -W[1:N] / np.sqrt(1 - x[1:N]**2)
20
21     # Boundary values via summation formulas
22     ii = np.arange(N)
23     w[0] = np.sum(ii**2 * U[ii]) / N + 0.5 * N * U[N]
24     w[N] = np.sum((-1)**(ii+1) * ii**2 * U[ii]) / N + \
25         0.5 * (-1)**(N+1) * N * U[N]
```

 Python

```

26
27     return w

```

The equivalent MATLAB implementation:

```

1  function w = chebfft(v)
2  % CHEBFFT Chebyshev first derivative via FFT.
3      N = length(v) - 1;
4      if N == 0, w = 0; return, end
5
6      x = cos((0:N)' * pi / N);
7      V = [v(:); flipud(v(2:N))];
8      U = real(fft(V));
9
10     ii = 0:N-1;
11     W = real(ifft(1i * [ii, 0, 1-N:-1]' .* fft(V)));
12
13     w = zeros(N+1, 1);
14     w(2:N) = -W(2:N) ./ sqrt(1 - x(2:N).^2);
15     w(1) = sum(ii'.^2 .* U(ii'+1)) / N + 0.5*N*U(N+1);
16     w(N+1) = sum((-1).^(ii+1)' .* ii'.^2 .* U(ii'+1)) / N + ...
17             0.5*(-1)^(N+1)*N*U(N+1);
18 end

```

 Matlab

The Julia implementation:

```

1  function chebfft(v)
2      N = length(v) - 1
3      N == 0 && return [0.0]
4
5      x = [cos(π * j / N) for j in 0:N]
6
7      # Extend to even periodic function in theta
8      V = [v; v[N:-1:2]]
9
10     # FFT and multiply by ik
11     U = real.(fft(V))
12     k = [collect(0:N-1); 0; collect(1-N:-1)]
13     W = real.(ifft(1im .* k .* fft(V)))
14
15     # Transform from theta to x derivative
16     w = zeros(N + 1)
17     w[2:N] = -W[2:N] ./ sqrt.(1.0 .- x[2:N].^2)
18
19     # Boundary values via summation formulas
20     ii = 0:N-1
21     w[1] = sum(ii.^2 .* U[ii .+ 1]) / N + 0.5 * N * U[N+1]
22     w[N+1] = sum((-1.0).^(ii .+ 1) .* ii.^2 .* U[ii .+ 1]) / N +
23             0.5 * (-1.0)^(N + 1) * N * U[N+1]

```

 Julia

```

24     return w
25 end

```

The code generating the figures in this section is available in:

- codes/python/ch10/chebfft.py
- codes/matlab/ch10/chebfft.m
- codes/julia/ch10/chebfft.jl

10.2.3 Accuracy Demonstration

To verify spectral accuracy, consider the test function

$$f(x) = e^x \cos(4x), \quad (117)$$

with exact derivative

$$f'(x) = e^x (\cos(4x) - 4 \sin(4x)). \quad (118)$$

Figure 58 shows the function, its spectral derivative, and the convergence of the maximum error as N increases.




Figure 58: Chebyshev differentiation via FFT applied to $f(x) = e^x \cos(4x)$. Left: The function $f(x)$ on $[-1, 1]$ with Chebyshev nodes marked. Centre: Comparison of exact and FFT-computed derivative for $N = 16$. Right: Maximum error versus N on a semilog scale, showing the characteristic straight-line descent of spectral convergence until machine precision is reached around $N = 24$.

The code generating Figure 58 is available in:

- codes/python/ch10/chebfft_accuracy.py
- codes/matlab/ch10/chebfft_accuracy.m
- codes/julia/ch10/chebfft_accuracy.jl

The straight-line descent in the error plot confirms exponential (spectral) convergence: each additional grid point reduces the error by a roughly constant factor. This is the hallmark of spectral methods applied to smooth functions.

10.3 Method of Lines and Time Stepping

Before tackling specific PDEs, we establish the general framework connecting spatial discretisation to time evolution. This *method of lines* approach treats space and time separately: first discretise in space to obtain a system of ordinary differential equations (ODEs), then apply a suitable time integrator. The

Method of Lines was formally systematised by Schiesser [106], providing a framework that decouples spatial accuracy (spectral) from temporal stability (A-stability of ODE solvers). The connection between this separation and the stiffness inherent in Chebyshev spectral derivatives has been explored in the comprehensive review by Hamdi *et al.* [59].

10.3.1 Semidiscretisation

Consider a PDE of the form

$$u_t = \mathcal{L}u, \quad (119)$$

where $\mathcal{L}$ is a spatial differential operator (e.g., $\mathcal{L} = \kappa \frac{\partial^2}{\partial x^2}$ for the heat equation). After discretising in space with $N + 1$ grid points, we obtain the *semidiscrete* system

$$\dot{\mathbf{u}} = \mathbf{L}\mathbf{u}, \quad (120)$$

where $\mathbf{u}(t) = (u_0(t), u_1(t), \dots, u_{N(t)}(t))^T$ is the vector of grid values and $\mathbf{L}$ is the matrix (or matrix-free operator) representing the discretised spatial operator.

For second-order in time PDEs like the wave equation $u_{tt} = c^2 u_{xx}$, the semidiscrete form is

$$\ddot{\mathbf{u}} = \mathbf{L}\mathbf{u}, \quad (121)$$

which can be rewritten as a first-order system by introducing the velocity $\mathbf{v} = \dot{\mathbf{u}}$.

10.3.2 The Spatial Operator

The operator $\mathbf{L}$ can be implemented in two ways:

1. **Matrix form:** Build the full Chebyshev differentiation matrix D_N (or D_N^2 for second derivatives) as in Section 7, then apply it via matrix-vector multiplication.
2. **Matrix-free form:** Use the FFT-based `chebfft` function, avoiding explicit matrix storage. For the second derivative, simply apply `chebfft` twice.

For 2D problems on tensor-product grids, the Laplacian takes the form

$$\mathbf{L} = D_2 \otimes I + I \otimes D_2, \quad (122)$$

where $\otimes$ denotes the Kronecker product. In practice, we compute $D_2 U + U D_2^T$ using matrix-matrix multiplication, exploiting the separable structure.

10.3.3 Time Integrators by PDE Type

Different PDE types call for different time-stepping strategies. Table 18 summarises the typical choices.

PDE Type	Semidiscrete Form	Typical Schemes	Stability Notes
Wave	$\ddot{\mathbf{u}} = \mathbf{L}\mathbf{u}$	Leapfrog, Newmark	Explicit; $\Delta t \sim N^{-2}$
Diffusion	$\dot{\mathbf{u}} = \mathbf{L}\mathbf{u}$	Backward Euler, CN	Implicit preferred
Advection	$\dot{\mathbf{u}} = -c\mathbf{D}\mathbf{u}$	Leapfrog, RK3	Explicit; $\Delta t \sim N^{-2}$

Table 18: Model time integrators for different PDE types. CN denotes Crank–Nicolson.

10.3.4 The Chebyshev CFL Condition

A crucial point for explicit time stepping: on Chebyshev grids, the time step must scale as

$$\Delta t = O(N^{-2}), \quad (123)$$

not $O(N^{-1})$ as for equispaced finite differences. This more restrictive stability limit arises from the $O(N^{-2})$ clustering of Chebyshev points near the boundaries. The explicit stability limit for Chebyshev spectral methods was rigorously analysed by Gottlieb and Turkel [56], who demonstrated that the clustering of nodes necessitates time steps scaling as the square of the reciprocal of the number of modes. Trefethen and Trummer [122] further illuminated this by analysing the spectra of Chebyshev differentiation matrices, revealing that the eigenvalues of the first derivative operator grow as $O(N^2)$, creating the severe stability constraint known as the “Chebyshev penalty.”

Physically, information propagates across the smallest grid spacing in each time step. Since the minimum spacing near the boundaries is $O(N^{-2})$, the CFL condition requires correspondingly small time steps. This is the price we pay for the superior spatial accuracy of spectral methods.

Implicit methods (backward Euler, Crank–Nicolson) avoid this restriction entirely, making them attractive for diffusion-dominated problems where time accuracy requirements are less stringent.

10.4 One-Dimensional Wave Equation

We begin our tour of spectral PDE solvers with the classic vibrating string: the one-dimensional wave equation with fixed endpoints.

10.4.1 Problem Formulation

The *wave equation* in one spatial dimension is

$$u_{tt} = c^2 u_{xx}, \quad -1 < x < 1, \quad t > 0, \quad (124)$$

where $c > 0$ is the wave speed. The *Dirichlet boundary conditions*

$$u(-1, t) = u(1, t) = 0 \quad (125)$$

model a string fixed at both ends. We also need initial conditions:

$$u(x, 0) = u_0(x), \quad u_t(x, 0) = v_0(x). \quad (126)$$

For our demonstration, we choose a half-sine initial displacement with zero initial velocity:

$$u(x, 0) = \sin\left(\frac{\pi}{2}(1+x)\right), \quad u_t(x, 0) = 0. \quad (127)$$

This initial profile satisfies the boundary conditions and has an explicit eigenfunction expansion, allowing comparison with the exact solution.

10.4.2 Spatial Discretisation

We discretise space using $N + 1$ Chebyshev points Equation 108. The second spatial derivative is approximated using either:

- The Chebyshev differentiation matrix $D_N^2 = D_N \times D_N$, or
- Two applications of `chebfft`: $u_{xx} \approx \text{chebfft}(\text{chebfft}(u))$.

To enforce the Dirichlet boundary conditions, we work with the interior nodes $j = 1, 2, \dots, N - 1$ and set $u_0 = u_N = 0$ after each time step.

10.4.3 Time Stepping: Leapfrog with Taylor Start

For the second-order-in-time wave equation, the *leapfrog* (Störmer–Verlet) scheme is a natural choice. The seminal comparison of spectral and finite difference methods for hyperbolic equations by Kreiss and Oliger [73] demonstrated that spectral methods require significantly fewer grid points per wave-

length ($N \approx \pi$) compared to finite differences ($N \approx 10\text{--}20$), making them ideal for long-range wave propagation. The leapfrog scheme is:

$$\mathbf{u}^{n+1} = 2\mathbf{u}^n - \mathbf{u}^{n-1} + (c\Delta t)^2 D_2 \mathbf{u}^n. \quad (128)$$

This scheme requires values at two previous time levels. For the first step ($n = 0$), we use a Taylor expansion:

$$\mathbf{u}^1 = \mathbf{u}^0 + \Delta t \dot{\mathbf{u}}^0 + \frac{(\Delta t)^2}{2} \ddot{\mathbf{u}}^0 = \mathbf{u}^0 + \frac{(c\Delta t)^2}{2} D_2 \mathbf{u}^0, \quad (129)$$

where we used $\dot{\mathbf{u}}^0 = \mathbf{v}_0 = 0$ (zero initial velocity).

For stability, we choose

$$\Delta t = \frac{\alpha}{N^2}, \quad (130)$$

with $\alpha \approx 4$ to 8. This $O(N^{-2})$ scaling reflects the Chebyshev CFL condition.

10.4.4 Implementation

The following Python code solves the 1D wave equation:

```
1 def waveld_cheb(N=80, c=1.0, tmax=6.0, alpha=4.0):
2     """Solve 1D wave equation on Chebyshev grid."""
3     # Chebyshev grid
4     x = np.cos(np.pi * np.arange(N + 1) / N)
5
6     # Initial data: half-sine
7     u = np.sin(0.5 * np.pi * (1 + x))
8
9     # Time stepping
10    dt = alpha / N**2
11    nsteps = int(np.ceil(tmax / dt))
12
13    # Taylor start (zero initial velocity)
14    u_xx = chebfft(chebfft(u))
15    u_xx[0] = u_xx[N] = 0 # enforce BCs
16    u_prev = u - 0.5 * (c * dt)**2 * u_xx
17
18    for n in range(nsteps):
19        u_xx = chebfft(chebfft(u))
20        u_xx[0] = u_xx[N] = 0
21        u_new = 2*u - u_prev + (c * dt)**2 * u_xx
22        u_prev, u = u, u_new
23
24    return x, u
```

 Python

The equivalent MATLAB implementation:

```
1 function [x, u] = waveld_cheb(N, c, tmax, alpha)
2 % Solve 1D wave equation on Chebyshev grid.
3     x = cos(pi * (0:N)' / N);
4
5     % Initial data: half-sine
6     u = sin(0.5 * pi * (1 + x));
```

 Matlab

```

7
8     dt = alpha / N^2;
9     nsteps = ceil(tmax / dt);
10
11     % Taylor start
12     u_xx = chebfft(chebfft(u));
13     u_xx(1) = 0; u_xx(N+1) = 0;
14     u_prev = u - 0.5 * (c*dt)^2 * u_xx;
15
16     for n = 1:nsteps
17         u_xx = chebfft(chebfft(u));
18         u_xx(1) = 0; u_xx(N+1) = 0;
19         u_new = 2*u - u_prev + (c*dt)^2 * u_xx;
20         u_prev = u; u = u_new;
21     end
22 end

```

The Julia implementation:

```

1  function waveld_cheb(; N=80, c=1.0, tmax=6.0, alpha=4.0)
2      x = [cos(pi * j / N) for j in 0:N]
3
4      # Initial data: half-sine
5      u = sin.(0.5 * pi .* (1.0 .+ x))
6
7      dt = alpha / N^2
8      nsteps = Int(ceil(tmax / dt))
9
10     # Taylor start (zero initial velocity)
11     u_xx = chebfft(chebfft(u))
12     u_xx[1] = 0.0; u_xx[N+1] = 0.0
13     u_prev = u .- 0.5 * (c * dt)^2 .* u_xx
14
15     for n in 1:nsteps
16         u_xx = chebfft(chebfft(u))
17         u_xx[1] = 0.0; u_xx[N+1] = 0.0
18         u_new = 2.0 .* u .- u_prev .+ (c * dt)^2 .* u_xx
19         u_prev = u; u = u_new
20     end
21     return x, u
22 end

```

 Julia

10.4.5 Results and Physical Interpretation

Figure 59 shows the solution as a waterfall plot in space and time.

1D Wave Equation: Vibrating String




Figure 59: Solution of the 1D wave equation Equation 124 with initial half-sine profile. The waterfall plot shows $u(x, t)$ for $t \in [0, 6]$. The initial shape splits into two counter-propagating waves, which reflect at the fixed boundaries ($x = \pm 1$) with a sign change. The waves pass through each other and interfere, creating the characteristic standing wave pattern.

The physics is clearly visible: the initial displacement splits into left- and right-moving pulses (d'Alembert's solution). Upon reaching the boundaries, each pulse reflects with a sign change (required by the Dirichlet condition). As the waves pass through each other, they superpose linearly, creating the complex interference patterns visible in the waterfall plot.

The code generating Figure 59 is available in:

- `codes/python/ch10/wave1d_cheb.py`
- `codes/matlab/ch10/wave1d_cheb.m`
- `codes/julia/ch10/wave1d_cheb.jl`

10.5 Two-Dimensional Wave Equation

We now extend the wave equation to two spatial dimensions: the vibrating drum membrane.

10.5.1 Problem Formulation

The 2D wave equation on the square $[-1, 1]^2$ is

$$u_{tt} = c^2(u_{xx} + u_{yy}), \quad -1 < x, y < 1, \quad t > 0, \quad (131)$$

with homogeneous Dirichlet boundary conditions $u = 0$ on $\partial\Omega$ (the membrane is fixed along its entire boundary).

For initial data, we choose an offset Gaussian bump:

$$u(x, y, 0) = \exp(-30[(x - 0.2)^2 + (y + 0.3)^2]), \quad u_t(x, y, 0) = 0. \quad (132)$$

The asymmetric position $(0.2, -0.3)$ breaks the symmetry of the domain, producing richer interference patterns.

10.5.2 Spatial Discretisation: Tensor Products

We use a *tensor product grid*: Chebyshev points in both x and y directions:

$$x_i = \cos\left(i\frac{\pi}{N}\right), \quad y_j = \cos\left(j\frac{\pi}{N}\right), \quad i, j = 0, 1, \dots, N. \quad (133)$$

The solution is represented as an $(N + 1) \times (N + 1)$ matrix U with entries $U_{ij} \approx u(x_i, y_j)$.

The key insight is that the 2D Laplacian separates into 1D operators:

$$\Delta u \approx D_2 U + U D_2^\top, \quad (134)$$

where $D_2 = D_N^2$ is the 1D second-derivative matrix. The first term applies D_2 to each row (differentiation in x); the second applies $D_2^\top$ to each column (differentiation in y).

This tensor-product structure is computationally efficient: applying Equation 134 costs $O(N^3)$ operations (two matrix-matrix products), compared to $O(N^4)$ for a naive implementation using the full $(N + 1)^2 \times (N + 1)^2$ Laplacian matrix.

10.5.3 Time Stepping

Leapfrog extends naturally to 2D:

$$U^{n+1} = 2U^n - U^{n-1} + (c\Delta t)^2(D_2 U^n + U^n D_2^\top). \quad (135)$$

After each time step, we enforce the boundary conditions by setting the boundary rows and columns of U to zero.

For stability, we take $\Delta t = \frac{\alpha}{N^2}$ with $\alpha \approx 3$ to 6.

10.5.4 Implementation

```

1  def wave2d_cheb(N=32, c=1.0, tmax=1.0, alpha=3.0):
2      """Solve 2D wave equation on Chebyshev tensor grid."""
3      # Grid
4      x = np.cos(np.pi * np.arange(N + 1) / N)
5      xx, yy = np.meshgrid(x, x)
6
7      # Build D2 matrix
8      D, _ = cheb_matrix(N)
9      D2 = D @ D
10
11     # Initial data: offset Gaussian
12     U = np.exp(-30 * ((xx - 0.2)**2 + (yy + 0.3)**2))
13
14     dt = alpha / N**2
15     nsteps = int(np.ceil(tmax / dt))
16
17     # Taylor start

```

 Python

```

18 Lap_U = D2 @ U + U @ D2.T
19 U_prev = U - 0.5 * (c * dt)**2 * Lap_U
20
21 for n in range(nsteps):
22     Lap_U = D2 @ U + U @ D2.T
23     U_new = 2*U - U_prev + (c * dt)**2 * Lap_U
24
25     # Enforce boundary conditions
26     U_new[0, :] = U_new[-1, :] = 0
27     U_new[:, 0] = U_new[:, -1] = 0
28
29     U_prev, U = U, U_new
30
31 return xx, yy, U

```

The equivalent MATLAB implementation:

```

1 function [xx, yy, U] = wave2d_cheb(N, c, tmax, alpha)
2 % Solve 2D wave equation on Chebyshev tensor grid.
3 x = cos(pi * (0:N)' / N);
4 [xx, yy] = meshgrid(x, x);
5
6 [D, ~] = cheb_matrix(N);
7 D2 = D * D;
8
9 % Initial data
10 U = exp(-30 * ((xx - 0.2).^2 + (yy + 0.3).^2));
11
12 dt = alpha / N^2;
13 nsteps = ceil(tmax / dt);
14
15 Lap_U = D2 * U + U * D2';
16 U_prev = U - 0.5 * (c*dt)^2 * Lap_U;
17
18 for n = 1:nsteps
19     Lap_U = D2 * U + U * D2';
20     U_new = 2*U - U_prev + (c*dt)^2 * Lap_U;
21     U_new(1,:) = 0; U_new(end,:) = 0;
22     U_new(:,1) = 0; U_new(:,end) = 0;
23     U_prev = U; U = U_new;
24 end
25 end

```

Matlab

The Julia implementation:

```

1 function wave2d_cheb(; N=32, c=1.0, tmax=1.0, alpha=3.0)
2     x = [cos(pi * j / N) for j in 0:N]
3     xx = [x[i] for i in 1:N+1, j in 1:N+1]
4     yy = [x[j] for i in 1:N+1, j in 1:N+1]

```

Julia

```

5
6     D, _ = cheb_matrix(N)
7     D2 = D * D
8
9     # Initial data: offset Gaussian
10    U = exp.(-30.0 .* ((xx .- 0.2).^2 .+ (yy .+ 0.3).^2))
11
12    dt = alpha / N^2
13    nsteps = Int(ceil(tmax / dt))
14
15    # Taylor start
16    Lap_U = D2 * U + U * D2'
17    U_prev = U .- 0.5 * (c * dt)^2 .* Lap_U
18
19    for n in 1:nsteps
20        Lap_U = D2 * U + U * D2'
21        U_new = 2.0 .* U .- U_prev .+ (c * dt)^2 .* Lap_U
22        U_new[1, :] .= 0.0;   U_new[end, :] .= 0.0
23        U_new[:, 1] .= 0.0;   U_new[:, end] .= 0.0
24        U_prev = U; U = U_new
25    end
26    return xx, yy, U
27 end

```

10.5.5 Results

Figure 60 shows four snapshots of the drum membrane at different times.

2D Wave Equation: Vibrating Membrane




Figure 60: Solution of the 2D wave equation Equation 131 showing a vibrating drum membrane. Four snapshots at times $t = 0, 0.33, 0.67$, and 1.0 . The initial Gaussian bump spreads into a circular wavefront, which reflects off the fixed boundaries and creates complex interference patterns. The tensor-product Chebyshev grid clusters points near all four edges.

The initial Gaussian bump spreads radially as a circular wavefront. Upon reaching the boundaries, the wave reflects and interferes with itself, creating intricate patterns. Note that despite using only $N = 24$ points in each direction (576 total grid points), the solution captures the wave dynamics accurately, a testament to spectral accuracy.

The code generating Figure 60 is available in:

- `codes/python/ch10/wave2d_cheb.py`
- `codes/matlab/ch10/wave2d_cheb.m`
- `codes/julia/ch10/wave2d_cheb.jl`

10.6 One-Dimensional Heat Equation

We now turn to parabolic equations, where diffusion rather than wave propagation dominates the physics.

10.6.1 Problem Formulation

The *heat equation* (or diffusion equation) in one dimension is

$$u_t = \kappa u_{xx}, \quad -1 < x < 1, \quad t > 0, \quad (136)$$

where $\kappa > 0$ is the thermal diffusivity. With homogeneous Dirichlet conditions $u(\pm 1, t) = 0$, this models a rod with endpoints held at zero temperature.

For initial data, we choose a two-bump profile:

$$u(x, 0) = \exp(-20(x + 0.5)^2) + 0.5 \exp(-30(x - 0.4)^2). \quad (137)$$

This allows us to observe how two separate hot spots diffuse and eventually merge.

10.6.2 Time Stepping: Crank–Nicolson

Unlike the wave equation, the heat equation is *dissipative*: solutions decay over time as heat diffuses away. Explicit methods face severe stability restrictions ($\Delta t \leq O(N^{-4})$ for spectral methods!), making implicit schemes essential. The Implicit-Explicit (IMEX) strategy, formalized by Ascher, Ruuth, and Wetton [5], treats the stiff linear diffusion implicitly while keeping nonlinear terms explicit. For highly stiff systems with diagonal linear operators, Kassam and Trefethen [70] demonstrated the superiority of Exponential Time Differencing (ETD) methods, which analytically incorporate the stiff linear part.

The *Crank–Nicolson* scheme averages explicit and implicit Euler:

$$\frac{\mathbf{u}^{n+1} - \mathbf{u}^n}{\Delta t} = \frac{\kappa}{2} (D_{2,i} \mathbf{u}^{n+1} + D_{2,i} \mathbf{u}^n), \quad (138)$$

where $D_{2,i}$ is the interior block of D_2 (rows and columns 1 through $N - 1$).

Rearranging gives the linear system:

$$\left(I - \frac{\kappa \Delta t}{2} D_{2,i} \right) \mathbf{u}^{n+1} = \left(I + \frac{\kappa \Delta t}{2} D_{2,i} \right) \mathbf{u}^n. \quad (139)$$

Crank–Nicolson is:

- **Unconditionally stable**: no CFL restriction on Δt
- **Second-order accurate** in time: $O((\Delta t)^2)$ local truncation error
- **Implicit**: requires solving a linear system at each step

The linear system is dense (because $D_{2,i}$ is dense), but for moderate N , direct solution via LU factorisation is efficient. For very large N , iterative methods become preferable.

10.6.3 Implementation

```

1 def heat1d_cheb(N=64, kappa=1.0, tmax=1.0, dt=0.01):
2     """Solve 1D heat equation with Crank-Nicolson."""
3     D, x = cheb_matrix(N)
4     D2 = D @ D
5
6     # Interior block
7     ii = slice(1, N)
8     D2i = D2[np.ix_(np.arange(1, N), np.arange(1, N))]
9     I = np.eye(N - 1)
10
11    # Crank-Nicolson matrices
12    A = I - 0.5 * kappa * dt * D2i
13    B = I + 0.5 * kappa * dt * D2i

```

 Python

```

14
15     # Initial condition (interior points only)
16     u_full = np.exp(-20*(x + 0.5)**2) + 0.5*np.exp(-30*(x - 0.4)**2)
17     u = u_full[ii]
18
19     nsteps = int(np.ceil(tmax / dt))
20     for n in range(nsteps):
21         u = np.linalg.solve(A, B @ u)
22
23     # Reconstruct full solution with BCs
24     u_full = np.zeros(N + 1)
25     u_full[ii] = u
26     return x, u_full

```

The equivalent MATLAB implementation:

```

1  function [x, u_full] = heat1d_cheb(N, kappa, tmax, dt)
2  % Solve 1D heat equation with Crank-Nicolson.
3      [D, x] = cheb_matrix(N);
4      D2 = D * D;
5
6      % Interior block
7      D2i = D2(2:N, 2:N);
8      I = eye(N - 1);
9
10     A = I - 0.5 * kappa * dt * D2i;
11     B = I + 0.5 * kappa * dt * D2i;
12
13     u_full = exp(-20*(x + 0.5).^2) + 0.5*exp(-30*(x - 0.4).^2);
14     u = u_full(2:N);
15
16     nsteps = ceil(tmax / dt);
17     for n = 1:nsteps
18         u = A \ (B * u);
19     end
20
21     u_full = zeros(N+1, 1);
22     u_full(2:N) = u;
23 end

```

Matlab

The Julia implementation:

```

1  function heat1d_cheb(; N=64, kappa=1.0, tmax=1.0, dt=0.01)
2      D, x = cheb_matrix(N)
3      D2 = D * D
4
5      # Interior block
6      ii = 2:N
7      D2i = D2[ii, ii]

```

Julia

```

8   Ie = Matrix{Float64}(I, N - 1, N - 1)
9
10  # Crank-Nicolson matrices
11  A = Ie - 0.5 * kappa * dt * D2i
12  B = Ie + 0.5 * kappa * dt * D2i
13
14  # Initial condition (interior points only)
15  u_full = exp.(-20.0 .* (x .+ 0.5).^2) .+
16           0.5 .* exp.(-30.0 .* (x .- 0.4).^2)
17  u = u_full[ii]
18
19  nsteps = Int(ceil(tmax / dt))
20  for n in 1:nsteps
21      u = A \ (B * u)
22  end
23
24  # Reconstruct full solution with BCs
25  u_full = zeros(N + 1)
26  u_full[ii] = u
27  return x, u_full
28 end

```

10.6.4 Results

Figure 61 shows the temperature evolution as a space–time surface and as four snapshots.




Figure 61: Solution of the 1D heat equation Equation 136 with two-bump initial data. Main panel: Space–time surface showing diffusion and decay. Insets: Four snapshots at $t = 0, 0.1, 0.3$, and 1.0 . The two initial hot spots spread, merge, and eventually decay towards the zero steady state required by the boundary conditions.

The physics of diffusion is evident: sharp features smooth out, amplitudes decrease, and the solution approaches the steady state $u \equiv 0$. The two initial bumps spread and eventually merge into a single, broader profile before decaying completely.

The code generating Figure 61 is available in:

- codes/python/ch10/heat1d_cheb.py
- codes/matlab/ch10/heat1d_cheb.m
- codes/julia/ch10/heat1d_cheb.jl

10.7 Two-Dimensional Heat Equation

The extension to 2D follows the same pattern as the wave equation: tensor-product grids and Kronecker structure.

10.7.1 Problem Formulation

The 2D heat equation on the square $[-1, 1]^2$ is

$$u_t = u_{xx} + u_{yy}, \quad -1 < x, y < 1, \quad t > 0, \quad (140)$$

with homogeneous Dirichlet conditions $u = 0$ on $\partial\Omega$.

Initial data:

$$u(x, y, 0) = \exp(-25[(x + 0.3)^2 + (y - 0.1)^2]). \quad (141)$$

10.7.2 Time Stepping: Backward Euler with Kronecker Structure

The semidiscrete system is

$$U_t = D_2 U + U D_2^\top. \quad (142)$$

For implicit time stepping (e.g., backward Euler), we must solve

$$(I - \Delta t \mathbf{L}) \text{vec}(U^{n+1}) = \text{vec}(U^n), \quad (143)$$

where $\mathbf{L} = D_{2,i} \otimes I + I \otimes D_{2,i}$ is the *Kronecker sum* of the interior blocks.

The Kronecker sum $\mathbf{L}$ is an $(N-1)^2 \times (N-1)^2$ matrix. For moderate N (say, $N \leq 64$), we can form it explicitly and solve with a direct method. For larger N , iterative methods (conjugate gradient, GMRES) or specialised tensor solvers are more efficient.

10.7.3 Implementation

```

1  def heat2d_cheb(N=32, tmax=0.5, dt=0.01):
2      """Solve 2D heat equation with backward Euler."""
3      D, x = cheb_matrix(N)
4      D2 = D @ D
5
6      xx, yy = np.meshgrid(x, x)
7
8      # Interior indices
9      ii = np.arange(1, N)
10     D2i = D2[np.ix_(ii, ii)]

```

 Python


```

11     I = np.eye(N - 1)
12
13     # Kronecker sum Laplacian
14     L = np.kron(D2i, I) + np.kron(I, D2i)
15     A = np.eye((N-1)**2) - dt * L
16
17     # Initial condition
18     U_full = np.exp(-25 * ((xx + 0.3)**2 + (yy - 0.1)**2))
19     u = U_full[np.ix_(ii, ii)].flatten()
20
21     nsteps = int(np.ceil(tmax / dt))
22     for n in range(nsteps):
23         u = np.linalg.solve(A, u)
24
25     # Reconstruct full solution
26     U_full = np.zeros((N+1, N+1))
27     U_full[np.ix_(ii, ii)] = u.reshape(N-1, N-1)
28     return xx, yy, U_full

```

The equivalent MATLAB implementation:

```

1  function [xx, yy, U] = heat2d_cheb(N, tmax, dt)
2      [D, x] = cheb_matrix(N);
3      D2 = D * D;
4      [xx, yy] = meshgrid(x, x);
5
6      % Interior block and Kronecker sum
7      D2i = D2(2:N, 2:N);
8      I = eye(N - 1);
9      L = kron(D2i, I) + kron(I, D2i);
10     A = eye((N-1)^2) - dt * L;
11
12     % Initial condition
13     U = exp(-25 * ((xx+0.3).^2 + (yy-0.1).^2));
14     U(1,:) = 0; U(N+1,:) = 0;
15     U(:,1) = 0; U(:,N+1) = 0;
16
17     nsteps = ceil(tmax / dt);
18     for n = 1:nsteps
19         u = reshape(U(2:N, 2:N), [], 1);
20         U = zeros(N+1, N+1);
21         U(2:N, 2:N) = reshape(A \ u, N-1, N-1);
22     end
23 end

```

Matlab

The Julia implementation:

```

1  function heat2d_cheb(; N=32, tmax=0.5, dt=0.01)
2      D, x = cheb_matrix(N)

```

Julia

```

3      D2 = D * D
4      xx = [x[i] for i in 1:N+1, j in 1:N+1]
5      yy = [x[j] for i in 1:N+1, j in 1:N+1]
6
7      # Interior block and Kronecker sum
8      D2i = D2[2:N, 2:N]
9      Imat = Matrix{Float64}(I, N-1, N-1)
10     L = kron(D2i, Imat) + kron(Imat, D2i)
11     A = Matrix{Float64}(I, (N-1)^2, (N-1)^2) - dt * L
12
13     # Initial condition
14     U = exp.(-25 .* ((xx .+ 0.3).^2 .+ (yy .- 0.1).^2))
15     u = vec(U[2:N, 2:N])
16
17     nsteps = ceil(Int, tmax / dt)
18     for n in 1:nsteps
19         u = A \ u
20     end
21
22     U = zeros(N+1, N+1)
23     U[2:N, 2:N] = reshape(u, N-1, N-1)
24     return xx, yy, U
25 end

```

10.7.4 Results

Figure 62 shows four snapshots of the diffusing hot spot.

2D Heat Equation: Diffusion on a Square




Figure 62: Solution of the 2D heat equation Equation 140. Four snapshots at times $t = 0, 0.05, 0.15$, and 0.5 . The initial hot spot diffuses isotropically, spreading and decaying towards the zero steady state.

The code generating Figure 62 is available in:

- `codes/python/ch10/heat2d_cheb.py`
- `codes/matlab/ch10/heat2d_cheb.m`
- `codes/julia/ch10/heat2d_cheb.jl`

10.8 Two-Dimensional Poisson Equation

We now turn to *elliptic* equations, which describe equilibrium (steady-state) phenomena rather than time evolution.

10.8.1 Problem Formulation

The *Poisson equation* on the square $[-1, 1]^2$ is

$$-\Delta u(x, y) = f(x, y), \quad -1 < x, y < 1, \quad (144)$$

with Dirichlet boundary conditions $u = g$ on $\partial\Omega$.

Physically, Equation 144 describes:

- **Electrostatics:** u is the electric potential, f is the charge density
- **Steady heat conduction:** u is the temperature, f is the heat source
- **Membrane deflection:** u is the vertical displacement, f is the applied load

10.8.2 Manufactured Solution

To test our solver, we use the *method of manufactured solutions*: choose an exact solution, compute the corresponding right-hand side f , and verify that the numerical solution matches.

We take

$$u(x, y) = (1 - x^2)(1 - y^2) \cos\left(\frac{\pi x}{2}\right) \cos\left(\frac{\pi y}{2}\right). \quad (145)$$

The factors $(1 - x^2)$ and $(1 - y^2)$ ensure that $u = 0$ on the boundary (homogeneous Dirichlet conditions). The right-hand side $f = -\Delta u$ can be computed symbolically or numerically by applying the discrete Laplacian to the exact solution values.

10.8.3 Discrete System

The discrete Poisson equation is a linear system:

$$Au = f, \quad (146)$$

where:

- u contains the $(N - 1)^2$ interior unknowns, stacked column-by-column
- f is the corresponding right-hand side
- $A = D_{2,i} \otimes I + I \otimes D_{2,i}$ is the Kronecker sum of interior second-derivative blocks

The efficiency of spectral elliptic solvers relies on exploiting the structure of the operator. Haidvogel and Zang [58] introduced the matrix diagonalization method, which decouples the multidimensional system into a sequence of independent one-dimensional problems, reducing the complexity from $O(N^4)$ to $O(N^3)$. Shen [108] later developed spectral-Galerkin methods using basis functions that satisfy the boundary conditions automatically, leading to sparse matrices and optimal $O(N^2)$ complexity.

10.8.4 Implementation

```

1  def poisson2d_cheb(N=32):
2      """Solve 2D Poisson equation with manufactured solution."""
3      D, x = cheb_matrix(N)
4      D2 = D @ D
5
6      xx, yy = np.meshgrid(x, x)
7
8      # Interior indices
9      ii = np.arange(1, N)
10     D2i = D2[np.ix_(ii, ii)]
11     I = np.eye(N - 1)
12
13     # Kronecker sum Laplacian

```

 Python

```

14     A = np.kron(D2i, I) + np.kron(I, D2i)
15
16     # Manufactured exact solution
17     U_exact = (1 - xx**2) * (1 - yy**2) * \
18               np.cos(np.pi*xx/2) * np.cos(np.pi*yy/2)
19
20     # RHS from discrete Laplacian of exact solution
21     Lap_U = D2 @ U_exact + U_exact @ D2.T
22     f = -Lap_U[np.ix_(ii, ii)]
23
24     # Solve
25     u_vec = np.linalg.solve(A, f.flatten())
26
27     # Reconstruct full solution
28     U = np.zeros((N+1, N+1))
29     U[np.ix_(ii, ii)] = u_vec.reshape(N-1, N-1)
30
31     return xx, yy, U, U_exact

```

The equivalent MATLAB implementation:

```

1  function [xx, yy, U, U_exact] = poisson2d_cheb(N)
2  % Solve 2D Poisson equation with manufactured solution.
3      [D, x] = cheb_matrix(N);
4      D2 = D * D;
5
6      [xx, yy] = meshgrid(x, x);
7
8      ii = 2:N;
9      D2i = D2(ii, ii);
10     I = eye(N - 1);
11
12     A = kron(D2i, I) + kron(I, D2i);
13
14     U_exact = (1 - xx.^2) .* (1 - yy.^2) .* ...
15              cos(pi*xx/2) .* cos(pi*yy/2);
16
17     Lap_U = D2 * U_exact + U_exact * D2';
18     f = -Lap_U(ii, ii);
19
20     u_vec = A \ f(:);
21
22     U = zeros(N+1, N+1);
23     U(ii, ii) = reshape(u_vec, N-1, N-1);
24 end

```

Matlab

The Julia implementation:

```

1  function poisson2d_cheb(; N=32)

```

Julia

```

2    D, x = cheb_matrix(N)
3    D2 = D * D
4
5    xx = [x[i] for i in 1:N+1, j in 1:N+1]
6    yy = [x[j] for i in 1:N+1, j in 1:N+1]
7
8    # Interior indices
9    ii = 2:N
10   D2i = D2[ii, ii]
11   Ie = Matrix{Float64}(I, N - 1, N - 1)
12
13   # Kronecker sum Laplacian
14   L = kron(D2i, Ie) + kron(Ie, D2i)
15
16   # Manufactured exact solution
17   U_exact = (1 .- xx.^2) .* (1 .- yy.^2) .*
18             cos.(pi .* xx ./ 2) .* cos.(pi .* yy ./ 2)
19
20   # RHS from discrete Laplacian of exact solution
21   Lap_U = D2 * U_exact + U_exact * D2'
22   f = -Lap_U[ii, ii]
23
24   # Solve -L * u = f
25   u_vec = (-L) \ vec(f)
26
27   # Reconstruct full solution
28   U = zeros(N + 1, N + 1)
29   U[ii, ii] = reshape(u_vec, N - 1, N - 1)
30   return xx, yy, U, U_exact
31 end

```

10.8.5 Results

Figure 63 compares the exact and numerical solutions.

2D Poisson Equation: Spectral Solution




Figure 63: Solution of the 2D Poisson equation Equation 144 with manufactured solution Equation 145. Left: Exact solution. Centre: Numerical solution (visually indistinguishable from exact). Right: Error on a logarithmic scale, showing machine-precision accuracy throughout the interior. The spectral method achieves essentially exact results for this smooth problem.

The error is at machine precision (roughly 10^{-14} to 10^{-15}), confirming spectral accuracy. This is the ideal case: a smooth manufactured solution where the only errors come from finite-precision arithmetic.

The code generating Figure 63 is available in:

- `codes/python/ch10/poisson2d_cheb.py`
- `codes/matlab/ch10/poisson2d_cheb.m`
- `codes/julia/ch10/poisson2d_cheb.jl`

10.8.6 Laplace Equation

The *Laplace equation* is the special case $f = 0$:

$$\Delta u = 0. \quad (147)$$

Solutions are called *harmonic functions*. Physically, this describes equilibrium with no sources: the electrostatic potential in a charge-free region, or the steady temperature in a plate with no internal heat sources.

For Laplace's equation, the solution is entirely determined by the boundary conditions. The code adapts trivially: set the interior right-hand side to zero and include any nonzero boundary data in the system.

10.8.7 Helmholtz Equation

The *Helmholtz equation* adds a “mass” term:

$$-\Delta u - k^2 u = f(x, y), \quad (148)$$

where $k > 0$ is the wavenumber. This equation arises in time-harmonic wave problems (acoustics, electromagnetics) and quantum mechanics.

The discrete system becomes

$$(-L - k^2 I)u = f. \quad (149)$$

As k increases, the matrix $-L - k^2 I$ can become ill-conditioned, especially when k^2 approaches an eigenvalue of $-L$. Near such *resonance* values, the problem becomes nearly singular, and iterative solvers may struggle. This is a fundamental physical phenomenon: near resonance, small forcing produces large response.

10.9 Additional Physics Études

This section presents three additional examples demonstrating the versatility of spectral methods. These études are optional but illustrate how the same FFT-based machinery applies across different physical contexts.

10.9.1 Variable-Coefficient Transport

Consider advection with spatially varying velocity:

$$u_t + c(x)u_x = 0, \quad 0 < x < 2\pi, \quad u \text{ periodic}, \quad (150)$$

with

$$c(x) = 0.3 + \sin^2(x - 1). \quad (151)$$

This models transport in a medium where the wave speed depends on position. Using a *periodic Fourier grid* with FFT-based differentiation and RK4 time stepping, we can track a pulse as it moves through the variable medium.




Figure 64: Solution of the variable-coefficient transport equation Equation 150. Left: space-time diagram showing the characteristic trajectory of a Gaussian pulse in a medium where $c(x) = 0.3 + \sin^2(x - 1)$. The pulse accelerates in fast regions and decelerates in slow regions. Right: waterfall view of the same solution with vertical offset proportional to time.

The code generating Figure 64 is available in:

- codes/python/ch10/transport_variable.py
- codes/matlab/ch10/transport_variable.m
- codes/julia/ch10/transport_variable.jl

10.9.2 Schrödinger Equation

The time-dependent Schrödinger equation with a harmonic potential is

$$iu_t = -u_{xx} + x^2u, \quad (152)$$

where $u(x, t)$ is the complex wavefunction and x^2 is the harmonic oscillator potential.

A natural approach is *Strang splitting*, an operator splitting strategy introduced by Strang [113] that has become the method of choice for dispersive equations. Taha and Ablowitz [117] established the split-step Fourier method as the robust standard for the nonlinear Schrödinger equation, while Bao,

Jin, and Markowich [11] provided rigorous error analysis for time-splitting spectral approximations in the semiclassical regime. The algorithm proceeds in three sub-steps:

1. **Half potential step** (in physical space): $u \rightarrow u \cdot e^{-ix^2\Delta t/2}$
2. **Full kinetic step** (in Fourier space): $\hat{u}_k \rightarrow \hat{u}_k \cdot e^{-ik^2\Delta t}$
3. **Half potential step**: $u \rightarrow u \cdot e^{-ix^2\Delta t/2}$

This splitting is second-order accurate in time and preserves the L^2 norm of the wavefunction (unitarity). The kinetic step uses exactly the same FFT machinery as our other spectral methods.




Figure 65: Solution of the Schrödinger equation Equation 152 with harmonic potential using Strang splitting. Left: space-time evolution of the probability density $|u(x,t)|^2$ showing the oscillatory dynamics of a Gaussian wavepacket in the harmonic trap. Right: waterfall plot of $|u(x,t)|^2$ at selected times.

The code generating Figure 65 is available in:

- codes/python/ch10/schrodinger.py
- codes/matlab/ch10/schrodinger.m
- codes/julia/ch10/schrodinger.jl

10.9.3 Nonlinear Reaction-Diffusion

The Allen–Cahn equation combines diffusion with a bistable nonlinearity:

$$u_t = u_{xx} + u(1 - u^2). \quad (153)$$

The nonlinear term $u(1 - u^2)$ has stable equilibria at $u = \pm 1$ and an unstable equilibrium at $u = 0$. Solutions tend to separate into regions where $u \approx 1$ and $u \approx -1$, with thin transition layers (fronts) between them. The Allen–Cahn equation, originally formulated by Allen and Cahn [3] to model antiphase boundary motion in crystalline solids, demands numerical schemes that respect the thermodynamic structure of the system. Shen and Yang [109] introduced unconditionally energy-stable spectral schemes that preserve the dissipation of the Ginzburg–Landau free energy, building on the convex splitting framework of Eyre [40].

An IMEX (implicit-explicit) scheme treats the stiff linear diffusion implicitly and the nonlinear reaction explicitly:

$$(I - \Delta t D_{2,i}) \mathbf{u}^{n+1} = \mathbf{u}^n + \Delta t (\mathbf{u}^n - (\mathbf{u}^n)^3), \quad (154)$$

where the cubic is applied componentwise.




Figure 66: Solution of the Allen–Cahn equation Equation 153 using IMEX time stepping. Left: space-time evolution showing phase separation into regions where $u \approx \pm 1$. Right: snapshots at selected times showing the formation and coarsening of fronts between the two stable phases.

The code generating Figure 66 is available in:

- codes/python/ch10/allen_cahn.py
- codes/matlab/ch10/allen_cahn.m
- codes/julia/ch10/allen_cahn.jl

10.10 Efficiency: Matrices versus FFT

Throughout this chapter, we have presented both matrix-based and FFT-based approaches. Table 19 summarises the computational complexity.

Operation	Matrix-Based	FFT-Based
1D first derivative	$O(N^2)$	$O(N \log N)$
1D second derivative	$O(N^2)$	$O(N \log N)$
2D Laplacian (general)	$O(N^4)$	$O(N^2 \log N)$
2D Laplacian (tensor)	$O(N^3)$	$O(N^2 \log N)$

Table 19: Computational complexity of spectral differentiation by method. The matrix-based approach uses dense matrix-vector or matrix-matrix multiplication. The FFT-based approach uses the Chebyshev–Fourier connection. The tensor-product form $D_2U + UD_2^\top$ achieves $O(N^3)$ without explicit Kronecker products.

For small to moderate N (say, $N \leq 128$), the matrix approach is often simpler and fast enough. For larger N , the FFT approach becomes essential.

A further improvement, not explored here, uses the *Discrete Cosine Transform* (DCT) instead of the full FFT. Since the extended Chebyshev data is both real and even, the DCT can exploit both symmetries for an additional factor of 2–4 speedup. See the references for details.

10.11 A non-exhaustive literature overview

The development of spectral methods for partial differential equations is a rich narrative that intertwines approximation theory, linear algebra, and computational physics. What began as a technique for high-accuracy meteorological forecasting has evolved into a dominant paradigm for simulating turbulence, quantum dynamics, and general relativity. This overview highlights the seminal contributions that shaped the field and points toward contemporary advancements.

While the roots of spectral expansions can be traced to Fourier and Chebyshev, the computational era of spectral PDEs properly began in the early 1970s. The seminal work of Kreiss and Oliger [73] compared Fourier methods with finite difference schemes for hyperbolic equations, famously demonstrating that for a fixed error tolerance, spectral methods require significantly fewer grid points per wavelength ($N \approx \pi$) compared to finite differences ($N \approx 10\text{--}20$). This efficiency was critical for problems where memory was a bottleneck. Simultaneously, Orszag [94] applied spectral methods to the Orr–Sommerfeld stability equation, solving problems in fluid stability that were previously intractable. This work, along with his development of transform methods to evaluate convolution sums efficiently, marked the transition from “spectral Galerkin” methods (working entirely in coefficient space) to “pseudospectral” or “collocation” methods (working on physical grids), which we have utilised throughout this chapter. The theoretical bedrock for these methods was laid in the monograph by Gottlieb and Orszag [55], which provided the first rigorous error analysis, establishing the principle of “spectral accuracy” — that for analytic functions, the error decays exponentially with the number of modes.

A central challenge in spectral PDE solvers is the stiffness introduced by the high-order spatial discretisation. As derived in this chapter, the second-derivative operator on a Chebyshev grid scales as $O(N^4)$, imposing severe time-step restrictions for explicit schemes. Gottlieb and Turkel [56] and Trefethen and Trummer [122] analysed these stability limits, revealing that the eigenvalues of spectral differentiation matrices are extremely sensitive to boundary conditions and grid distribution. Trefethen’s analysis of the “pseudo-spectra” of these non-normal matrices provided deep insight into why spectral schemes could be sensitive to rounding errors and initial transients. To mitigate these restrictions, the community moved toward implicit and semi-implicit schemes. The Implicit-Explicit (IMEX) strategies formalised by Ascher, Ruuth, and Wetton [5] became standard for convection-diffusion problems, allowing the stiff diffusive terms to be treated implicitly while the nonlinear convective terms remain explicit. For highly stiff PDEs like the Kuramoto–Sivashinsky equation, Kassam and Trefethen [70] demonstrated the superiority of Exponential Time Differencing (ETD) and integrating factor methods, sparking a renewed interest in fourth-order time-stepping for stiff PDEs.

For equilibrium problems (Poisson, Helmholtz), the “matrix surgery” approach presented in this chapter is effective for moderate N . However, for large-scale simulations, efficiency becomes paramount. Haidvogel and Zang [58] introduced a matrix diagonalisation technique that decoupled the system into independent one-dimensional problems, reducing the complexity of two-dimensional Poisson solves from $O(N^4)$ to $O(N^3)$. This was refined by Shen [108], who constructed basis functions satisfying the boundary conditions automatically, leading to sparse matrices and $O(N^2)$ complexity. The tensor product structure emphasised in this chapter — exploiting Kronecker sums — is a modern realisation of these early insights. The modern canonical treatment of these algorithms is found in the texts by Shen, Tang, and Wang [110] and Boyd [21], which detail how to solve Poisson and Helmholtz equations in complex geometries using spectral-element and domain decomposition techniques.

The versatility of spectral methods is best seen in their application to specific physical systems, where preserving physical invariants is often as important as raw accuracy. For the Schrödinger equation, the preservation of unitarity (probability density) is critical in quantum mechanics. Taha and Ablowitz [117] compared various schemes for the nonlinear Schrödinger (NLS) equation, identifying

the split-step Fourier method as a robust performer. Bao, Jin, and Markowich [11] later provided a rigorous convergence analysis for time-splitting spectral approximations in the semiclassical regime, where highly oscillatory wavefunctions require the high resolution that only spectral methods can provide. For the Allen–Cahn and Cahn–Hilliard equations, energy stability is the primary concern. Eyre [40] proposed the convex splitting idea, which ensures unconditional gradient stability. Shen and Yang [109] extended these ideas to spectral discretisations, proving that high-order spectral accuracy can be combined with thermodynamic consistency to accurately model phase separation dynamics.

The field continues to evolve, increasingly intersecting with machine learning and high-performance computing. The integration of spectral methods with deep learning has led to Fourier Neural Operators (FNO). Li *et al.* [79] introduced the FNO architecture, and subsequent works such as Cao *et al.* [28] use spectral layers to learn PDE solution operators that generalise across different discretisations. These “Spectral-Refiner” models fine-tune FNO predictions using physical constraints, merging the data-driven power of AI with the rigorous accuracy of spectral methods. Moving beyond the method of lines, Kaur *et al.* [71] have proposed discretising time spectrally as well, treating the entire spacetime domain as a single tensor product grid, yielding exponential convergence in both space and time. Spectral methods are also uniquely suited for nonlocal operators (like the fractional Laplacian) which appear in anomalous diffusion models. Mustapha *et al.* [90] and Zhang and Wang [135] have developed fast spectral solvers for these equations, exploiting the fact that nonlocal convolution operators become diagonal multiplications in the spectral domain. Recent implementations, such as those described by Melia *et al.* [86], leverage graphical processing units (GPUs) to accelerate hierarchical spectral solvers, enabling the simulation of turbulent flows and large-scale elliptic problems at resolutions previously thought impossible. This brief overview underscores that spectral methods are not a static collection of algorithms but a dynamic field adapting to new computational architectures and physical models.

10.12 Summary

This chapter has developed spectral methods for the three fundamental classes of PDEs:

1. **Hyperbolic equations** (wave type):
 - Use explicit time stepping (leapfrog, Verlet)
 - Stability requires $\Delta t = O(N^{-2})$ on Chebyshev grids
 - The wave equation demonstrates propagation, reflection, and interference
2. **Parabolic equations** (diffusion type):
 - Use implicit time stepping (Crank–Nicolson, backward Euler)
 - Avoid severe CFL restrictions; take time steps dictated by accuracy
 - The heat equation demonstrates smoothing and decay
3. **Elliptic equations** (equilibrium type):
 - No time stepping; solve a linear system directly
 - Kronecker structure enables efficient tensor-product methods
 - Poisson and Helmholtz equations arise throughout physics

Key computational tools:

- **Chebyshev differentiation via FFT** (chebfft): The $x \rightarrow \theta \rightarrow$ Fourier pipeline transforms Chebyshev differentiation into FFT operations
- **Tensor products**: The 2D Laplacian $D_2 U + U D_2^\top$ exploits separability for efficiency

• **Kronecker sums:** For implicit solvers, $D_{2,i} \otimes I + I \otimes D_{2,i}$ gives the discrete Laplacian. Table 20 provides a quick reference for all PDEs treated in this chapter.

Section	PDE	Grid Type	Time Scheme	Key Operation
10.4	1D wave	Chebyshev	Leapfrog	chebfft twice
10.5	2D wave	Cheb. tensor	Leapfrog	$D_2 U + U D_2^\top$
10.6	1D heat	Chebyshev	Crank–Nicolson	Implicit solve
10.7	2D heat	Cheb. tensor	Backward Euler	Kronecker system
10.8	2D Poisson	Cheb. tensor	Direct solve	Kronecker sum
10.8	2D Helmholtz	Cheb. tensor	Direct solve	$L + k^2 I$
10.9	1D transport	Fourier	RK4	FFT derivative

Table 20: Summary of PDEs and methods presented in this chapter.

Bibliography

- [1] D. Abdelhamid, W. Albalawi, K. S. Nisar, and others, “Mixed Chebyshev and Legendre polynomials differentiation matrices for solving initial-boundary value problems”, *AIMS Mathematics*, vol. 8, no. 10, 2023, pp. 24609–24631.
- [2] Advanpix LLC., “Multiprecision Computing Toolbox for MATLAB”, 2022.
- [3] S. M. Allen, and J. W. Cahn, “A microscopic theory for antiphase boundary motion and its application to antiphase domain coarsening”, *Acta Metallurgica*, vol. 27, no. 6, 1979, pp. 1085–1095.
- [4] Anonymous, “A Dynamic Multiscale Anti-Aliasing Network for Time Series Forecasting”, *OpenReview*, 2025.
- [5] U. M. Ascher, S. J. Ruuth, and B. T. R. Wetton, “Implicit-Explicit methods for time-dependent partial differential equations”, *SIAM J. Numer. Anal.*, 1995, pp. 797–823.
- [6] I. Babuska, and S. A. Sauter, “Is the pollution effect of the FEM avoidable for the Helmholtz equation considering high wave numbers?”, *SIAM Journal on Numerical Analysis*, vol. 34, no. 6, 1997, pp. 2392–2423.
- [7] I. Babuška, B. A. Szabó, and I. N. Katz, “The p -version of the finite element method”, *SIAM Journal on Numerical Analysis*, vol. 18, no. 3, 1981, pp. 515–545.
- [8] R. Baltensperger, and J.-P. Berrut, “The errors in calculating the pseudospectral differentiation matrices for Chebyshev–Gauss–Lobatto points”, *Computers & Mathematics with Applications*, vol. 37, no. 1, 1999, pp. 41–56.
- [9] R. Baltensperger, and M. R. Trummer, “Spectral differencing with a twist”, *SIAM Journal on Scientific Computing*, vol. 24, no. 5, 2003, pp. 1465–1487.
- [10] W. Bao, Y. Feng, and C. Su, “Optimal error bounds on an exponential wave integrator Fourier spectral method for fractional nonlinear Schrödinger equations”, *Mathematics of Computation*, 2024.
- [11] W. Bao, S. Jin, and P. A. Markowich, “On time-splitting spectral approximations for the Schrödinger equation in the semiclassical regime”, *Journal of Computational Physics*, vol. 175, no. 2, 2002, pp. 487–524.
- [12] A. H. Barnett, J. Magland, and L. af Klinteberg, “A Parallel Nonuniform Fast Fourier Transform Library Based on an ‘Exponential of Semicircle’ Kernel”, *SIAM Journal on Scientific Computing*, vol. 41, no. 5, 2019, pp. C479–C504.
- [13] D. Batic, D. Dutykh, and J. Beek, “A Spectral Approach for Quasinormal Frequencies of Noncommutative Geometry-inspired Wormholes”, *Classical and Quantum Gravity*, 2025.
- [14] A. Bayliss, A. Class, and B. J. Matkowsky, “Roundoff error in computing derivatives using the Chebyshev differentiation matrix”, *Journal of Computational Physics*, vol. 116, no. 2, 1995, pp. 380–383.
- [15] S. N. Bernstein, “Démonstration du théorème de Weierstrass fondée sur le calcul des probabilités”, *Communications de la Société Mathématique de Kharkow*, vol. 13, 1912, pp. 1–2.
- [16] J.-P. Berrut, and L. N. Trefethen, “Barycentric Lagrange Interpolation”, *SIAM Review*, vol. 46, no. 3, 2004, pp. 501–517.
- [17] L. Bos, M. Caliari, S. D. Marchi, M. Vianello, and Y. Xu, “Bivariate Lagrange interpolation at the Padua points: the generating curve approach”, *Journal of Approximation Theory*, vol. 143, no. 1, 2006, pp. 15–25.
- [18] L. Bos, S. D. Marchi, M. Vianello, and Y. Xu, “Bivariate Lagrange interpolation at the Padua points: the ideal theory approach”, *Numerische Mathematik*, vol. 108, no. 1, 2007, pp. 43–57.
- [19] J. C. Bowman, and M. Roberts, “Efficient Dealiasing Convolutions without Padding”, *SIAM Journal on Scientific Computing*, vol. 33, no. 1, 2011, pp. 386–406.
- [20] J. P. Boyd, “An analytical and numerical study of the two-dimensional Bratu equation”, *Journal of Scientific Computing*, vol. 1, no. 2, 1986, pp. 183–206.
- [21] J. P. Boyd, *Chebyshev and Fourier Spectral Methods*. Dover Publications, New York, 2000, pp. 688.

- [22] L. Brutman, “On the Lebesgue Function for Polynomial Interpolation”, *SIAM Journal on Numerical Analysis*, vol. 15, no. 4, 1978, pp. 694–704.
- [23] K. Bu, W. Gu, and A. Jaffe, “Quantum higher-order Fourier analysis and the Clifford hierarchy”, *Proceedings of the National Academy of Sciences*, vol. 122, no. 45, 2025, pp. e2515667122.
- [24] J. Calder, and N. Garcia Trillos, “Improved spectral convergence rates for graph Laplacians on ϵ -graphs and k -NN graphs”, *Foundations of Data Science*, 2022.
- [25] M. Caliari, S. D. Marchi, and M. Vianello, “Algorithm 886: Padua2D—Lagrange Interpolation at Padua Points on Bivariate Domains”, *ACM Transactions on Mathematical Software*, vol. 35, no. 3, 2008, pp. 1–11.
- [26] D. Cantero, “Moving point loads on bridges using high-order finite difference weights”, 2021.
- [27] C. Canuto, M. Y. Hussaini, A. Quarteroni, and T. A. Zang, *Spectral Methods: Fundamentals in Single Domains*. Springer-Verlag Berlin Heidelberg, 2006, pp. 563.
- [28] S. Cao, F. Brarda, R. Li, and Y. Xi, “Spectral-Refiner: Accurate Fine-Tuning of Spatiotemporal Fourier Neural Operator for Turbulent Flows”, in *International Conference on Learning Representations (ICLR)*, 2025.
- [29] M. H. Carpenter, and D. Gottlieb, “Spectral Methods on Arbitrary Grids”, *Journal of Computational Physics*, vol. 129, no. 1, 1996, pp. 74–86.
- [30] J. W. Cooley, and J. W. Tukey, “An algorithm for the machine calculation of complex Fourier series”, *Mathematics of Computation*, vol. 19, no. 90, 1965, pp. 297.
- [31] T. A. Cox, S. G. Murray, A. R. Parsons, J. S. Dillon, and others, “fftvis: a non-uniform Fast Fourier Transform based interferometric visibility simulator”, *RAS Techniques and Instruments*, vol. 4, no. 1, 2025.
- [32] P. Deufhard, *Newton Methods for Nonlinear Problems: Affine Invariance and Adaptive Algorithms*. Springer Science & Business Media, 2011.
- [33] D. L. Donoho, A. Maleki, I. U. Rahman, M. Shahram, and V. Stodden, “Reproducible Research in Computational Harmonic Analysis”, *Computing in Science & Engineering*, vol. 11, no. 1, 2009, pp. 8–18.
- [34] T. A. Driscoll, “Finite difference weights MATLAB function”, 2007.
- [35] T. A. Driscoll, and N. Hale, “Rectangular spectral collocation”, *IMA Journal of Numerical Analysis*, vol. 36, no. 1, 2016, pp. 108–132.
- [36] T. A. Driscoll, and Y. Zhou, “Stable evaluation of derivatives for barycentric and continued fraction representations of rational functions”, *arXiv preprint arXiv:2601.10667*, 2026.
- [37] D. E. Edmunds, P. Gurka, and J. Lang, “Decay of (p, q) -Fourier coefficients”, *Proceedings of the Royal Society A*, vol. 470, no. 2170, 2014, pp. 20140221.
- [38] K. T. Elgindy, “Fourier–Gegenbauer Integral Galerkin Method for Solving the Advection–Diffusion Equation with Periodic Boundary Conditions”, *Computation*, vol. 13, no. 9, 2025, pp. 219.
- [39] L. Euler, “Sur la vibration des cordes”, *Mémoires de l'Académie des Sciences de Berlin*, vol. 4, 1748, pp. 69–85.
- [40] D. J. Eyre, “Unconditionally gradient stable time marching the Cahn–Hilliard equation”, in *MRS Proceedings*, 1998, pp. 39.
- [41] G. Faber, “Über die interpolatorische Darstellung stetiger Funktionen”, *Jahresbericht der Deutschen Mathematiker-Vereinigung*, vol. 23, 1914, pp. 192–210.
- [42] L. Fejér, “Sur les fonctions bornées et intégrables”, *Comptes Rendus Hebdomadaires des Séances de l'Académie des Sciences*, vol. 131, 1900, pp. 984–987.
- [43] B. A. Finlayson, *The Method of Weighted Residuals and Variational Principles*. Academic Press, 1972.
- [44] B. A. Finlayson, and L. E. Scriven, “The method of weighted residuals—a review”, *Applied Mechanics Reviews*, vol. 19, no. 9, 1966, pp. 735–748.

- [45] B. Fornberg, “The pseudospectral method: Comparisons with finite differences for the elastic wave equation”, *Geophysics*, vol. 52, no. 4, 1987, pp. 483–501.
- [46] B. Fornberg, “Generation of finite difference formulas on arbitrarily spaced grids”, *Mathematics of Computation*, vol. 51, no. 184, 1988, pp. 699–706.
- [47] B. Fornberg, *A practical guide to pseudospectral methods*. Cambridge University Press, 1996, pp. 242.
- [48] J. Fourier, *Théorie analytique de la chaleur*. Didot, Paris, 1822.
- [49] D. A. Frank-Kamenetskii, *Diffusion and Heat Exchange in Chemical Kinetics*. Princeton University Press, 1955.
- [50] M. Frigo, and S. G. Johnson, “The Design and Implementation of FFTW3”, *Proceedings of the IEEE*, vol. 93, no. 2, 2005, pp. 216–231.
- [51] B. G. Galerkin, “Rods and plates: Series occurring in various questions concerning the elastic equilibrium of rods and plates”, *Vestnik Inzhenerov i Tekhnikov*, vol. 19, 1915, pp. 897–908.
- [52] J. Galkowski, and E. A. Spence, “Numerical analysis of the high-frequency Helmholtz equation using semiclassical analysis”, *arXiv preprint arXiv:2511.15287*, 2025.
- [53] C. I. Gheorghiu, “Accurate spectral collocation solutions to some Bratu's type boundary value problems”, *arXiv preprint arXiv:2011.13212*, 2020.
- [54] K. Gkanos, and others, “Numerical wave propagation in complex acoustical environments using hybrid wave-based and geometrical acoustics methods”, *Journal of Sound and Vibration*, vol. 598, 2025, pp. 118821.
- [55] D. Gottlieb, and S. A. Orszag, *Numerical Analysis of Spectral Methods: Theory and Applications*. SIAM, 1977.
- [56] D. Gottlieb, and E. Turkel, “On time discretizations for spectral methods”, *Studies in Applied Mathematics*, vol. 63, no. 1, 1980, pp. 67–86.
- [57] Y. Gu, and J. Zhou, “An Efficient Spectral Method for Elliptic PDEs in Complex Domains with Circular Embedding”, *SIAM Journal on Scientific Computing*, 2021.
- [58] D. B. Haidvogel, and T. A. Zang, “The accurate solution of Poisson's equation by expansion in Chebyshev polynomials”, *Journal of Computational Physics*, vol. 30, no. 2, 1979, pp. 167–180.
- [59] S. Hamdi, W. Schiesser, and G. Griffiths, “Method of lines”, *Scholarpedia*, vol. 2, no. 7, 2007, pp. 2859.
- [60] T. Hangelbroek, and others, “Spectral stability and perturbation results for kernel differentiation matrices on the sphere”, *arXiv preprint arXiv:2311.06982*, 2023.
- [61] H. Hassanieh, P. Indyk, D. Katabi, and E. Price, “Nearly Optimal Sparse Fourier Transform”, in *Proceedings of the 44th Annual ACM Symposium on Theory of Computing (STOC)*, 2012, pp. 563–578.
- [62] M. T. Heideman, D. H. Johnson, and C. S. Burrus, “Gauss and the History of the Fast Fourier Transform”, *IEEE ASSP Magazine*, vol. 1, no. 4, 1984, pp. 14–21.
- [63] J. S. Hesthaven, and T. Warburton, *Nodal Discontinuous Galerkin Methods: Algorithms, Analysis, and Applications*. Springer, 2007.
- [64] N. J. Higham, “The numerical stability of barycentric Lagrange interpolation”, *IMA Journal of Numerical Analysis*, vol. 24, no. 4, 2004, pp. 547–556.
- [65] Ibrahim, and Lawrence, “Spectral rectangular collocation”, 2019.
- [66] A. Iserles, and others, “Stable spectral methods for time-dependent problems and the preservation of structure”, 2020.
- [67] D. Jackson, *Über die Genauigkeit der Annäherung stetiger Funktionen durch ganze rationale Funktionen gegebenen Grades und trigonometrische Summen gegebener Ordnung*, 1911.
- [68] A. Javeed, D. P. Kouri, D. Ridzal, and J. D. Steinman, “A preconditioner for spectral collocation”, *SIAM Journal on Scientific Computing*, vol. 47, no. 5, 2025, pp. A2828–A2850.
- [69] A. J. Jerri, “The Shannon sampling theorem — Its various extensions and applications: A tutorial review”, *Proceedings of the IEEE*, vol. 65, no. 11, 1977, pp. 1565–1596.

- [70] A.-K. Kassam, and L. N. Trefethen, “Fourth-order time-stepping for stiff PDEs”, *SIAM Journal on Scientific Computing*, vol. 26, no. 4, 2005, pp. 1214–1233.
- [71] A. Kaur, S. Lui, S. Nataj, and C. Liyanage, “Space-Time Spectral Methods for Linear and Nonlinear PDEs”, 2025, pp. 89–112.
- [72] S. Kim, J. Kim, K. Kim, M. Jeong, and J. Rho, “Anti-aliased metasurfaces beyond the Nyquist limit”, *Nature Communications*, vol. 16, 2025, pp. 411.
- [73] H.-O. Kreiss, and J. Oliger, “Comparison of accurate methods for the integration of hyperbolic equations”, *Tellus*, vol. 24, no. 3, 1972, pp. 199–215.
- [74] C. Lanczos, “Trigonometric Interpolation of Empirical and Analytical Functions”, *Journal of Mathematics and Physics*, vol. 17, 1938, pp. 123–199.
- [75] E. Landau, “Über die Approximation einer stetigen Funktion durch eine ganze rationale Funktion”, *Rendiconti del Circolo Matematico di Palermo*, vol. 25, 1908, pp. 337–345.
- [76] H. Lebesgue, “Sur l'approximation des fonctions”, *Bulletin des Sciences Mathématiques*, vol. 22, 1898, pp. 278–287.
- [77] R. J. LeVeque, I. M. Mitchell, and V. Stodden, “Reproducible research for scientific computing: Tools and strategies for changing the culture”, *Computing in Science & Engineering*, vol. 14, no. 4, 2012, pp. 13–17.
- [78] X. Li, Z. Zhao, and X. Jia, “Reconstruction of Piecewise Smooth Functions Based on Fourier Extension”, *Axioms*, vol. 14, no. 7, 2025, pp. 489.
- [79] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhatt, A. Stuart, and A. Anandkumar, “Fourier Neural Operator for Parametric Partial Differential Equations”, in *International Conference on Learning Representations (ICLR)*, 2021.
- [80] Y. Liu, and others, “A Quantum Spectral Method for Non-Periodic Boundary Value Problems”, *arXiv preprint arXiv:2511.11494*, 2025.
- [81] F. W. Luttman, and T. J. Rivlin, “Some numerical experiments in the theory of polynomial interpolation”, *IBM Journal of Research and Development*, vol. 9, 1965, pp. 187–191.
- [82] X. Ma, Y. Zhang, L. Chen, and X. Yang, “Efficient Jacobi–Galerkin spectral and second-order time discretization scheme for ground and first excited states of nonlinear fractional Schrödinger equation”, *Electronic Research Archive*, vol. 33, no. 9, 2025, pp. 5776–5793.
- [83] S. D. Marchi, M. Caliarì, and M. Vianello, “Bivariate polynomial interpolation at new nodal sets”, *Applied Mathematics and Computation*, vol. 165, no. 2, 2005, pp. 261–274.
- [84] J. McNamee, and F. Stenger, “Whittaker's Cardinal Function in Retrospect”, *Mathematics of Computation*, vol. 25, no. 113, 1971, pp. 141–154.
- [85] E. Meijering, “A Chronology of Interpolation: From Ancient Astronomy to Modern Signal and Image Processing”, *Proceedings of the IEEE*, vol. 90, no. 3, 2002, pp. 319–342.
- [86] O. Melia, D. Fortunato, and others, “Hardware-Accelerated Hierarchical Poincaré–Steklov Solvers”, *Journal of Computational Physics*, 2026.
- [87] B. Meuris, S. Qadeer, and P. Stinis, “Machine-learning-based spectral methods for partial differential equations”, *Scientific Reports*, vol. 13, 2023, pp. 1739.
- [88] B. Mildenhall, P. P. Srinivasan, M. Tancik, J. T. Barron, R. Ramamoorthi, and R. Ng, “NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis”, *Communications of the ACM*, vol. 65, no. 1, 2021, pp. 99–106.
- [89] A. Moiola, and E. A. Spence, “Is the Helmholtz equation really sign-indefinite?”, *SIAM Review*, vol. 56, no. 2, 2014, pp. 274–312.
- [90] I. Mustapha, B. Alali, and N. Albin, “Fourier spectral method for nonlocal equations on bounded domains”, *Advances in Computational Science and Engineering*, 2025.
- [91] Y. V. Nguéabou, and S. D. Olonijú, “Integrating Spectral Methods with Neural Network Architectures: A Review of Hybrid Approaches to Solving Differential Equation”, *Archives of Computational Methods in Engineering*, 2025.

- [92] S. Olver, and A. Townsend, “A fast and well-conditioned spectral method”, *SIAM Review*, vol. 55, no. 3, 2013, pp. 462–489.
- [93] S. Olver, R. M. Slevinsky, and A. Townsend, “Fast algorithms using orthogonal polynomials”, *Acta Numerica*, vol. 29, 2020, pp. 573–699.
- [94] S. A. Orszag, “Accurate solution of the Orr–Sommerfeld stability equation”, *Journal of Fluid Mechanics*, vol. 50, no. 4, 1971, pp. 689–703.
- [95] S. A. Orszag, “On the Elimination of Aliasing in Finite-Difference Schemes by Filtering High-Wavenumber Components”, *Journal of the Atmospheric Sciences*, vol. 28, 1971, pp. 1074–1075.
- [96] R. E. A. C. Paley, and N. Wiener, *Fourier Transforms in the Complex Domain*. American Mathematical Society, 1934.
- [97] J. E. Pasciak, “Spectral and Pseudospectral Methods for Advection Equations”, *Mathematics of Computation*, vol. 35, no. 152, 1980, pp. 1081–1092.
- [98] A. T. Patera, “A spectral element method for fluid dynamics: Laminar flow in a channel expansion”, *Journal of Computational Physics*, vol. 54, no. 3, 1984, pp. 468–488.
- [99] R. P. Peláez, P. Palacios-Alonso, and R. Delgado-Buscalioni, “Spectral solver for the oscillatory Stokes frequency-based equation in doubly periodic confined domains”, *Journal of Fluid Mechanics*, vol. 1010, 2025, pp. A57.
- [100] É. Picard, “Sur la représentation approchée des fonctions”, *Comptes Rendus*, vol. 112, 1891, pp. 185–186.
- [101] S. D. Poisson, “Suite du mémoire sur les intégrales définies et sur la sommation des séries”, *Journal de l'École Polytechnique*, vol. 12, 1823, pp. 404–509.
- [102] J. le R. d'Alembert, “Recherches sur la courbe que forme une corde tendue mise en vibration”, *Histoire de l'Académie Royale des Sciences et Belles Lettres de Berlin*, vol. 3, 1747, pp. 214–219.
- [103] C. Runge, “Über empirische Funktionen und die Interpolation zwischen äquidistanten Ordinaten”, *Zeitschrift für Mathematik und Physik*, vol. 46, 1901, pp. 224–243.
- [104] E. B. Saff, and V. Totik, *Logarithmic Potentials with External Fields*. Springer-Verlag, 1997.
- [105] M. H. Salama, and others, “Spectral Treatment of the Fractional Bratu Equation via Shifted Lucas Polynomials”, *Contemporary Mathematics*, vol. 6, no. 5, 2025.
- [106] W. E. Schiesser, *The Numerical Method of Lines: Integration of Partial Differential Equations*. Academic Press, 1991.
- [107] C. E. Shannon, “Communication in the Presence of Noise”, *Proceedings of the IRE*, vol. 37, no. 1, 1949, pp. 10–21.
- [108] J. Shen, “Efficient spectral-Galerkin method II. Direct solvers of second- and fourth-order equations using Chebyshev polynomials”, *SIAM Journal on Scientific Computing*, vol. 16, no. 1, 1995, pp. 74–87.
- [109] J. Shen, and X. Yang, “Numerical approximations of Allen–Cahn and Cahn–Hilliard equations”, *Discrete and Continuous Dynamical Systems – Series A*, vol. 28, no. 4, 2010, pp. 1669.
- [110] J. Shen, T. Tang, and L.-L. Wang, *Spectral Methods: Algorithms, Analysis and Applications*. Springer, 2011.
- [111] S. J. Smith, “Lebesgue constants in polynomial interpolation”, *Annales Mathematicae et Informaticae*, vol. 33, 2006, pp. 109–123.
- [112] M. H. Stone, “The Generalized Weierstrass Approximation Theorem”, *Mathematics Magazine*, vol. 21, no. 4–5, 1948, pp. 167–184, 237–254.
- [113] G. Strang, “On the construction and comparison of difference schemes”, *SIAM Journal on Numerical Analysis*, vol. 5, no. 3, 1968, pp. 506–517.
- [114] C. Sturm, “Mémoire sur les équations différentielles linéaires du second ordre”, *J. Math. Pures Appl.*, vol. 1, 1836, pp. 106–186.
- [115] C. Sturm, and J. Liouville, “Mémoire sur le développement des fonctions en séries...”, *Journal de Mathématiques Pures et Appliquées*, vol. 2, 1837, pp. 373–444.

- [116] E. Tadmor, “Convergence of Spectral Methods for Nonlinear Conservation Laws”, *SIAM Journal on Numerical Analysis*, vol. 26, no. 1, 1989, pp. 30–44.
- [117] T. R. Taha, and M. J. Ablowitz, “Analytical and numerical aspects of certain nonlinear evolution equations. III. Numerical, Korteweg-de Vries equation”, *J. Comput. Phys.*, vol. 55, no. 2, 1984, pp. 231–253.
- [118] F. Tasnin, “Revisiting Fourier and Chebyshev Spectral Methods with Physics-Informed Machine Learning”, in *IEEE Conference on Control Technology and Applications (CCTA)*, 2025.
- [119] A. N. Tikhonov, and A. A. Samarskii, *Equations of Mathematical Physics*. Dover Publications, Inc., 1963, pp. 785.
- [120] L. N. Trefethen, *Spectral methods in MatLab*. Society for Industrial, Applied Mathematics, Philadelphia, PA, USA, 2000, pp. 184.
- [121] L. N. Trefethen, *Approximation Theory and Approximation Practice*. SIAM, 2013.
- [122] L. N. Trefethen, and M. R. Trummer, “An instability phenomenon in spectral methods”, *SIAM Journal on Numerical Analysis*, vol. 24, no. 5, 1987, pp. 1008–1023.
- [123] L. N. Trefethen, and J. A. C. Weideman, “The Exponentially Convergent Trapezoidal Rule”, *SIAM Review*, vol. 56, no. 3, 2014, pp. 385–458.
- [124] A. H. Turetskii, “The bounding of polynomials defined at equally distributed points”, *Proc. Pedag. Inst. Vitebsk*, vol. 3, 1940, pp. 117–127.
- [125] J. V. Villadsen, and W. E. Stewart, “Solution of boundary-value problems by orthogonal collocation”, *Chemical Engineering Science*, vol. 22, no. 11, 1967, pp. 1483–1501.
- [126] P. Vértési, “Optimal Lebesgue constant for Lagrange interpolation”, *SIAM J. Numer. Anal.*, vol. 27, 1990, pp. 1322–1331.
- [127] H. Wang, “On the deterioration of convergence rate of spectral differentiations for functions with singularities”, *arXiv preprint*, 2025.
- [128] J. A. C. Weideman, and S. C. Reddy, “A MATLAB differentiation matrix suite”, *ACM Transactions on Mathematical Software*, vol. 26, no. 4, 2000, pp. 465–519.
- [129] J. A. C. Weideman, and L. N. Trefethen, “The eigenvalues of second-order spectral differentiation matrices”, *SIAM Journal on Numerical Analysis*, vol. 25, no. 6, 1988, pp. 1279–1298.
- [130] K. Weierstrass, “Über die analytische Darstellbarkeit sogenannter willkürlicher Functionen einer reellen Veränderlichen”, *Sitzungsberichte der Königlich Preussischen Akademie der Wissenschaften zu Berlin*, 1885, pp. 633–639/789–805.
- [131] E. T. Whittaker, “On the Functions which are represented by the Expansions of the Interpolation Theory”, *Proceedings of the Royal Society of Edinburgh*, vol. 35, 1915, pp. 181–194.
- [132] K. Xu, “Solving nonlinear ODEs with the ultraspherical spectral method”, *IMA Journal of Numerical Analysis*, vol. 44, no. 6, 2024, pp. 3749–3779.
- [133] G. Yao, Z. Wang, and Z. Wang, “A conforming multi-domain Legendre spectral method for solving diffusive-viscous wave equations in the exterior domain with separated star-shaped obstacles”, *IMA Journal of Numerical Analysis*, vol. 45, 2025, pp. 3235–3263.
- [134] M. Zayernouri, L.-L. Wang, and G. E. Karniadakis, *Spectral and Spectral Element Methods for Fractional Ordinary and Partial Differential Equations*. Cambridge University Press, 2024.
- [135] Z. Zhang, and H. Wang, “Optimal error bounds for spectral Galerkin methods applied to fractional diffusion equations”, *Mathematics of Computation*, 2025.
- [136] Z. Zhao, and Y. Wang, “A local Fourier extension method for function approximation”, *Inverse Problems and Imaging*, vol. 22, 2026, pp. 1–13.